


THINK LIKE Folang



A developer's guide to elegant,
idiomatic code

Kameswara Rao Mithipati



Preface

Why This Book Exists

Learning the syntax of a programming language tells you how to write an instruction. It does not automatically tell you how to look at a problem, find the right boundaries, or grow a small idea into a clear application.

Think Like FoLang is about that second part.

This book follows a small problem from its first plain-language description to a working FoLang application. It does not begin with a catalogue of language features. It begins with the problem, identifies what the program needs, and introduces a FoLang construct only when that construct helps express the solution more clearly.

The central question throughout the book is:

How should we think when designing a solution in FoLang?

The goal is not merely to produce code that compiles. The goal is to make the problem, the program's state, its behavior, and its boundaries easy to understand.

What This Book Teaches

The book develops a way of reasoning that can be used beyond the application we build together. It asks you to consider:

- What information belongs to the domain?
- Which behavior belongs with that information?
- Where should a boundary be introduced?
- Which details should remain visible, and which should remain internal?
- When is a simple function enough?
- When has the problem earned a stronger abstraction?
- How can the program grow without losing the shape of the original problem?

These questions come before the choice of a particular construct. FoLang offers several ways to express a solution, but this book does not introduce them all at once. Each choice appears in response to a concrete need in the application.

The reasoning matters as much as the final code. When we revise an earlier design, the book explains what pressure caused the change, what the new design improves, and what trade-off it introduces.

How the Book Proceeds

We start with one small problem described in ordinary language. We then build the smallest useful solution.

As the requirements grow, we return to that solution and ask whether its current shape still communicates the intent clearly. When it does, we keep it. When it does not, we introduce only the next idea needed to improve it.

The recurring rhythm is:

1. understand a requirement;
2. model the smallest clear solution;
3. write and run the code;
4. observe the next design pressure;
5. revise the model deliberately; and
6. verify that the application still expresses the original intent.

By the end, the result is a small but complete application. More importantly, you will have followed the decisions that shaped it. The intention is that you can apply the same reasoning when designing your own FoLang programs.

This is not an exhaustive tour of the language. FoLang contains mechanisms that this application may never need, and the book does not introduce a feature merely because it exists. The journey is complete for the problem we solve, while remaining intentionally incomplete as a catalogue of everything FoLang can express.

Who This Book Is For

This book is intended for readers who know FoLang's basic single-source constructs from *Hello, FoLang!* or have equivalent familiarity. You should be comfortable with basic programming ideas such as variables, conditions, loops, functions, and values.

Experience with another programming language is helpful but no particular language background is required. Where a familiar-looking idea has a different meaning in FoLang, the book explains the difference when it becomes relevant.

You do not need to understand compiler internals, advanced type theory, or the complete FoLang language model before beginning.

The Language Reference and This Book

The FoLang language reference is the normative definition of the language. It defines exact syntax, semantics, type rules, source structure, capability boundaries, diagnostics, and conformance requirements.

This book has a different responsibility. It explains how to use the language to reason about and solve a problem.

FoLang Language Reference

- > normative rules
- > exact language behavior
- > compiler-facing precision

Think Like FoLang

- > developer-facing reasoning
- > one problem-led application
- > design decisions and trade-offs
- > working examples

The examples in this book are written to follow the current language reference. If an explanation in this book ever conflicts with the normative reference, the reference governs.

FoLang is still evolving toward version 1.0. Syntax and semantics may be refined during that period, and examples in this book may be updated to remain aligned with the reference.

How to Use This Book

The chapters are intended to be read in order because each one continues the same application. You will get the most from the book by running the examples, changing them, and observing how those changes affect the design.

At each decision point, pause before reading the chosen solution. Decide what you would do and why. The purpose is not to guess the author's answer; it is to make your own reasoning visible so that you can compare it with the reasoning presented in the chapter.

Do not worry about memorizing every piece of syntax. Pay attention to why a construct is introduced, what responsibility it receives, and whether the resulting code makes the problem clearer.

A Note on Clarity

In this book, elegance means clarity rather than cleverness.

Sometimes clarity produces a short expression. Sometimes it requires an explicit declaration or a carefully chosen boundary. A solution is considered successful when another developer can understand what the program models, where its state lives, and why each part exists.

The standard used throughout the book is:

Does this design make the problem, its semantics, and our intent clearer?

When the answer is yes, the design is moving in the right direction.

Before We Begin

The first step is not choosing an abstraction. It is understanding the problem.

That is where the first chapter begins.

Appendix: Design Influences — What FoLang

Borrowed and Changed

FoLang was not designed in isolation. It draws from several programming traditions whose languages solved particular problems well: object-oriented languages, functional languages, systems languages, metaprogramming languages, and virtual-machine platforms.

Influence does not mean duplication. FoLang rarely adopts a mechanism without changing its boundaries, selection rules, or relationship to the rest of the language. The useful question is therefore not only:

|

Where did this idea come from?

It is also:

|

What did FoLang preserve, what did it change, and why does that matter to a developer?

This appendix is a map of those influences. It is explanatory rather than normative. When an exact syntax or semantic question arises, the FoLang language reference governs.

How to Read an Influence

The languages named here are reference points, not claims that one language was the sole historical origin of an idea. Many programming ideas developed across several communities at the same time or appeared in earlier research languages.

Each influence is described through four questions:

Question	Purpose
What problem did the earlier tradition address?	Establish the original motivation
What idea did FoLang retain?	Identify the recognizable connection
What did FoLang change?	Explain FoLang’s own boundary or rule
What should a FoLang developer remember?	Turn history into a practical mental model

The comparison is most useful when it helps a developer transfer existing knowledge without assuming that familiar-looking syntax has identical semantics.

Recurring FoLang Adaptation Principles

Across otherwise unrelated features, FoLang repeatedly makes the same kinds of adjustments.

Explicit selection over invisible search

When several behaviors could apply, FoLang generally asks the developer to make the intended source visible. Named typeclass instances, file-scoped extension activation, explicit effect resolutions, and explicit projected surfaces all follow this preference.

Powerful mechanisms with visible boundaries

FoLang retains advanced facilities such as multiple inheritance, custom operators, macros, continuations, native access, and dynamic runtime behavior. It does not allow those mechanisms to silently affect every compilation domain. Their placement and activation boundaries are part of their meaning.

One developer model, multiple backend strategies

The language specifies observable behavior rather than requiring a backend to copy one internal implementation. Copy-on-write, effects, execution models, runtime operations, and managed-object policies may be implemented differently by C, C++, LLVM, JVM, or another backend as long as FoLang programs observe the same semantics.

Introduce a distinction only when it carries meaning

FoLang separates concepts when combining them would hide an important decision: equality from identity, traits from mixins, projected libraries from packaged libraries, recoverable effects from fatal failures, and ordinary types from refinement or dependent types.

Control Flow and Expressions

Object-centered control — Smalltalk

Smalltalk demonstrated that control structures can be understood through objects and messages rather than requiring every control operation to be a special statement form.

FoLang retains the idea that conditions, loops, and iteration can participate in an object-centered expression model. It adapts that idea through standard associated control functions. The normative model states that FoLang control flow is implemented through dynamically scoped associated functions rather than by introducing unrelated capture rules for every block form.

The practical lesson is that a condition is not merely punctuation surrounding a branch. It is a value participating in an operation whose consequences are attached explicitly.

Pattern matching — ML, Haskell, Scala, and Rust

The algebraic and functional language traditions made pattern matching a primary way to reveal the shape of data.

FoLang uses pattern matching and developer-defined matchers to make structural alternatives explicit. It retains the idea that matching should describe the form being handled rather than hide the decision inside a chain of unrelated tests.

FoLang's custom matcher model also reflects comparison-strategy traditions such as Java's `Comparable` and `Comparator`, but it is not a direct copy of either Java facility. The influence is the separation of a value from a named strategy used to compare or interpret it.

Lambdas, comprehensions, and expression composition — functional traditions

Haskell, ML-family languages, Scala, Python, and many later languages helped make functions, transformations, and comprehensions ordinary parts of expression-oriented programming.

FoLang adopts lambdas, callbacks, function values, matching, and comprehension forms where they make data transformation clearer. It combines them with its own associated-function scope rules rather than importing another language's closure or capture semantics unchanged.

Continuations — Lisp and Scheme

Lisp and Scheme made continuations visible as a model of the remaining computation. Later work distinguished full continuations from several forms of delimited continuation.

FoLang exposes continuation families through `co.control`, including full and delimited forms. It keeps the boundary explicit: different continuation families state where capture stops and whether the continuation is exposed to application code or retained by the control runtime.

The lesson is to choose a continuation mechanism because the required control boundary is meaningful, not merely because a backend can implement suspension with fibers, stacks, or callbacks.

Deferred completion — Go and structured cleanup traditions

Go popularized a compact way to register completion work at the point where that work becomes necessary.

FoLang retains registration during forward execution and last-in-first-out completion. It defines `defer` alongside effects and ordinary returns so that cleanup behavior remains predictable across normal completion, `return_with_error`, and effect propagation.

Types and Abstraction

Typeclasses — Haskell

Haskell established typeclasses as contracts implemented independently for types.

FoLang retains the separation between a typeclass contract and a named instance. It removes implicit package-wide instance search: the caller names or activates the intended instance rather than asking the compiler to discover one silently.

This eliminates the instance-selection coherence problem created by implicit search, but it does not remove FoLang’s ownership and placement rules. An instance must still be declared with an owner permitted by the normative rules. When neither the typeclass nor the represented type is owned, the developer uses an owned `co.lang.newtype` wrapper.

The developer should remember:

A FoLang instance is an explicit named implementation, not an invisible global candidate.

Generics — C++, Java, and the wider parametric tradition

C++ and Java brought parameterized types and functions into mainstream object-oriented development, while ML and functional languages established a longer tradition of parametric polymorphism.

FoLang supports generic types and generic functions while keeping generic identity, bounds, variance, ownership, and placement in the language’s semantic model. The specification defines observable type behavior without requiring every backend to use C++ template instantiation, Java erasure, or any other particular runtime representation.

Modules and signatures — ML and OCaml

ML-family languages, particularly OCaml, established signatures and modules as a rich abstraction system over values, functions, and types.

FoLang retains the separation between a signature contract and a matching module. A FoLang module is a named singleton component and may be used through a compatible signature. Signatures may describe values, functions, associated type requirements, and fixed or manifest type components.

FoLang does not treat a module as a repeatedly constructed class or as a namespace containing arbitrary nested declarations. The developer should think of it as one named implementation component with one module identity and shared module state.

Structs and companion units — Go and companion-oriented languages

Go reinforced the value of keeping data declarations distinct from associated behavior. Other languages established companion namespaces and type-associated functions.

FoLang combines those ideas through struct companion units. A struct remains a data declaration, while eligible behavior belongs to a canonical same-package `<StructName>.comp.unit.fo1` file. The relationship is established by the filename-derived owner and package rules, not by informal naming convention.

Companion units distinguish receiverless functions, instance-associated value-receiver functions, type-associated functions, and operator functions. Embedding may promote accessible public instance-associated behavior, but it does not widen private access or turn all companion declarations into members of the embedding struct.

Classes, interfaces, traits, and mixins — object-oriented and composition traditions

Classes and interfaces reflect established object-oriented traditions from Smalltalk, C++, Java, C#, and related languages. Traits and mixins draw from composition models developed across languages such as Self, Scala, Rust, and others.

FoLang keeps these constructs separate because they answer different questions:

- a class defines constructable objects with identity and instance state;
- an interface defines a behavioral contract;
- a trait contributes behavior without state;
- a mixin may contribute both reusable state and behavior.

FoLang permits at most two direct class parents while allowing broader composition through interfaces, traits, and mixins. Concrete inherited behavior is not automatically exposed as part of the child API; relationship selection, override, implementation, and conflict rules remain explicit.

The design preserves reuse without treating every inherited member as an automatic public promise.

Extensions — C#, Kotlin, Swift, and related languages

Extension mechanisms allow behavior to be associated with a type without modifying the type's original declaration.

FoLang keeps this convenience but makes activation explicit and file-scoped through `@co.ddap.use`. Importing a package does not silently activate its extensions, and an existing call cannot change meaning merely because another package was imported.

The developer should distinguish availability from activation:

An extension may be present in the program without being active in the current source file.

Associated and opaque types — ML, Haskell, Rust, Swift, and Scala

Several language families separate a public type relationship from the representation chosen by an implementation.

FoLang signatures can require associated types or establish fixed/manifest type components. Opaque types hide representation behind a distinct public type identity. These mechanisms are not arbitrary nested type declarations; they are explicit contracts about type relationships and visibility.

Variants and absence — algebraic and nullable traditions

ML, Haskell, Scala, Rust, and many other languages use tagged alternatives to represent a closed set of possible values. Other mainstream languages allow a general null or empty state.

FoLang deliberately uses both ideas for different purposes. Variants communicate expected alternatives. Ordinary FoLang types also admit the universal none state, tested through `isNone()`. Refinement and dependent values form the exception: they cannot contain none and are protected by compile-time definite-initialization analysis.

The practical distinction is communicative as well as operational: a variant states that an alternative is part of the intended model, while none represents the absence available to ordinary values.

Refinement types — Liquid Haskell, F*, and refinement-type research

Refinement-type systems describe values of a base type that also satisfy a predicate.

FoLang retains the base-type-plus-predicate model but does not require a general theorem prover. Refinement validation is a normative runtime operation whenever a value enters the refinement type through assignment, initialization, argument passing, return, conversion, or another admissible transfer.

Definite-initialization analysis remains compile-time. Before a refinement value is proven assigned, reading it produces the `Uninitialized` diagnostic; once a value enters the type, its predicate is validated.

Dependent types — ATS, Idris, and dependent-type research

Dependent-type traditions allow values to participate in type identity, as in a vector whose length is part of its type.

FoLang adopts value-indexed types with deliberately restricted index forms. Indices must satisfy the normative resolution and evaluability rules; arithmetic expressions and arbitrary calls are not admitted merely to prove that two indices are mathematically equal.

This trades some proof power for predictable, decidable frontend checking. A developer uses a dependent type when the value index is essential to the program's contract, not merely to attach documentation to an ordinary value.

Effects, Errors, and Execution

Effects — condition systems, algebraic effects, and result-oriented APIs

Exception systems, Common Lisp condition systems, algebraic-effect research, and result types each separate failure signaling from some part of ordinary computation.

FoLang combines lessons from those traditions while keeping distinct channels:

- recoverable effects are represented by objects satisfying `co.lang.error`;
- `@co.dap.effects` advertises known effects without closing the effect set;
- the caller supplies invocation policy through `@co.dap.onEffect`;

- `return_with_error` converts an active effect into an ordinary error result;
- fatal failures remain outside the recoverable effect system and terminate the application.

Execution-model boundaries do not propagate recoverable effects into the invoking execution context. They use an implicit `return_with_error` policy and deliver the original error object through their required error-compatible result position.

The central FoLang adaptation is ownership of policy: the callable advertises what it may emit, but the ordinary caller decides what the invocation should do about it.

Concurrency and execution models — Go, Java, and structured execution traditions

Go, Java, and many runtime systems provide threads, tasks, schedulers, coroutines, or lightweight execution carriers.

FoLang exposes an execution-model distinction only when it changes observable language behavior. A backend may use an operating-system thread, task pool, virtual thread, fiber-like mechanism, event loop, or another implementation strategy, but the declared completion and error-boundary semantics must remain the same.

The developer chooses an execution model for its FoLang contract, not for an assumption about one backend's scheduler.

Metaprogramming and Language Extension

Macros — Lisp, Julia, and Nim

Lisp established quotation as a way to treat program structure as data. Julia and Nim demonstrated modern structured macros that operate on syntax rather than raw text.

FoLang adopts quotation, unquotation, composition, guards, and structured macro declarations while keeping their placement and compilation behavior explicit. Macros transform syntax under the language's macro rules; they are not an unrestricted textual preprocessor.

Templates — Nim and compile-time generation traditions

Template systems provide reusable compile-time structure without requiring the same semantics as ordinary runtime functions.

FoLang retains templates as a distinct metaprogramming mechanism. Their separate identity prevents compile-time generation from being mistaken for ordinary callable dispatch.

Decorators and annotations — Python and Java

Python popularized decorators as transformations or wrappers around declared behavior. Java made declarative annotations a mainstream metadata mechanism.

FoLang keeps decorators and annotations distinct. A decorator participates in behavioral wrapping or transformation according to its rules; an annotation describes or selects declaration semantics. Directives and pragmas remain separate compiler-recognized metadata categories rather than user-defined synonyms for annotations.

Operator overloading and new symbols — C++ and Raku

C++ made overloads of existing operators familiar in user-defined types. Raku demonstrated that a language can permit programs to introduce new operator spellings and parse properties.

FoLang separates these two capabilities:

- implementing an already known operator belongs to ordinary legal overload ownership;
- creating a new spelling, precedence, associativity, or fixity belongs only in `components/operators/component.fo1`.

The operator table is prepared before the permitted ordinary source domain is parsed. Custom syntax is confined to its owning application or projected application-library source domain and is not exported as an importable parser table.

The result preserves expressive custom notation without allowing imported libraries to silently reinterpret existing source text.

Systems, Runtime Boundaries, and Packages

Native representation — C and systems-language traditions

C remains the central reference point for pointers, addresses, ABI layout, calling conventions, and foreign interoperability.

FoLang exposes those facilities only inside the native capability domain and uses `co.lang.cstruct` for explicitly value-semantic ABI representation. Ordinary managed FoLang objects retain their own reference and object-graph semantics rather than becoming C representations merely because a backend is implemented in C or C++.

Dynamic runtime isolation — virtual-machine and plugin traditions

Java and other virtual-machine ecosystems demonstrated dynamic loading, runtime-created types, reflection, and managed execution domains.

FoLang places dynamic-runtime facilities behind a projected boundary. Runtime implementation types do not silently join the consuming application's compiled type or overload universe. The boundary makes dynamic behavior available without turning the whole application into an open runtime type graph.

Path-based packages — Python and hierarchical module traditions

Python and other hierarchical module systems made filesystem paths a natural source of package identity.

FoLang uses path-based packages but combines them with explicit unit merging, canonical primary filenames, package-family access, and separate projected or packaged library exposure models. A directory establishes package identity; it does not automatically make every declaration public or exportable.

Object Identity and State Policies

References and identity — object-oriented traditions

Object-oriented languages distinguish an object's identity from the names that refer to it, though many make value equality and reference equality easy to confuse.

FoLang makes the distinction explicit. Assignment of managed objects copies a reference, `==` compares values deeply, and `sameRef()` asks whether two references identify the same managed object. A `co.lang.cstruct` is the explicit value-semantic ABI exception.

Immutability, sharing, and copy-on-write — functional and systems traditions

Functional languages emphasize immutable values; concurrent systems provide shared synchronization; systems and collection implementations frequently use copy-on-write optimization.

FoLang presents Immutable, Shared, and CopyOnWrite as deep, mutually exclusive, permanent policies of a managed object graph rather than as unrelated parallel type families. CopyOnWrite has whole-graph isolation semantics, while a backend remains free to optimize physical storage as long as mutation cannot leak between the source and isolated logical states.

The developer chooses an observable state policy. The backend chooses a conforming implementation strategy.

Summary

FoLang area	Major influence	Characteristic FoLang adaptation
Object-centered control	Smalltalk	Dynamically scoped associated control functions
Pattern matching	ML, Haskell, Scala, Rust	Matching plus explicit developer-defined matcher strategies
Continuations	Lisp, Scheme	Explicit full and delimited control families
Deferred completion	Go	LIFO completion integrated with effects and returns
Typeclasses	Haskell	Named instances with no implicit package-wide search
Generics	C++, Java, parametric languages	Backend-neutral generic semantics and explicit bounds

Modules and signatures	ML, OCaml	Named singleton modules with signature and type-component contracts
Struct companions	Go and companion traditions	Canonical same-package companion ownership
Classes, traits, mixins	OOP and composition traditions	Bounded class inheritance and explicit behavior exposure
Extensions	C#, Kotlin, Swift	Explicit file-scoped activation
Variants and absence	Algebraic and nullable traditions	Expected alternatives separated from universal none
Refinement types	Liquid Haskell, F*	Runtime predicate validation plus definite initialization
Dependent types	ATS, Idris	Restricted, decidable value indices
Effects	Conditions, algebraic effects, result APIs	Definer advertises; caller owns ordinary invocation policy
Execution models	Go, Java, runtime systems	Observable contract independent of backend carrier
Macros and templates	Lisp, Julia, Nim	Structured compile-time facilities with explicit boundaries
Decorators and annotations	Python, Java	Behavioral transformation separated from declarative metadata
Operators	C++, Raku	Existing overloads separated from centrally registered new syntax
Native facilities	C	Capability-restricted ABI and memory access
Dynamic runtime	VM and plugin ecosystems	Projected isolation from the application's compiled type graph

Packages	Python and hierarchical modules	Path identity plus explicit visibility and exposure
Object policies	Functional and systems traditions	Deep graph-level Immutable, Shared, and CopyOnWrite semantics

Closing Perspective

FoLang's identity does not come from inventing every ingredient. It comes from the way those ingredients are selected, bounded, and made to work together.

The language repeatedly asks the developer to make consequential choices visible: which implementation is intended, which behavior is activated, which boundary is crossed, which error policy applies, and which state policy an object graph follows.

Knowing the influences helps when learning the syntax. Understanding the adaptations is what helps a developer think in FoLang.