



Folang Programming Language

Folang is a general-purpose programming language designed to be **expressive, consistent, and extensible**, merging functional fluency with object-centric abstractions.

Normative Status

This document is the authoritative and normative definition of the FoLang programming language.

All FoLang syntax, semantics, type-system rules, name-resolution rules, execution behavior, and other language requirements are defined by this document.

Every addition, correction, clarification, modification, or extension to the FoLang language must be recorded in this document before it becomes part of the language.

A compiler or other language implementation may use any internal architecture, parser, intermediate representation, runtime, or backend. However, to be identified as a conforming FoLang implementation, its externally observable behavior must adhere to the requirements defined in this document.

When an implementation conflicts with this specification, the specification governs unless the discrepancy is formally accepted and incorporated as a specification revision.

Implementation-specific extensions must be clearly identified as extensions and must not be represented as standard FoLang features unless they have been incorporated into this specification.

Design Overview




Figure: Frontend and Backend are independent components




Figure: Multiple backend implementations can coexist or be swapped without modifying the Frontend.

Folang follows a deliberately different approach from conventional programming language designs. The system is structured to ensure **clear separation of concerns, license isolation, and extensibility through well-defined integration boundaries**.

Compiler and Backend

1. Frontend

The Frontend is responsible for source-level analysis and semantic processing of FoLang programs.

Components

- Scanner / Lexer
- Parser
- AST / Parse Tree Generator
- Symbol Table Generator
- Semantic Analyzer

Implementation

- Implemented in **Go**
- Uses Go structures internally for AST, symbol-table, and semantic processing.
- The externally consumable frontend output is serialized according to the selected backend's interchange contract.
- The frontend always writes that artifact under the reserved project-root `build/` directory.
- The backend-supplied contract identifies the supported FoLang/plugin protocol version, HIR schema version, and wire format.
- The backend installs/provides this interchange-contract file in the same installation directory as the FoLang compiler executable.
- The frontend reads that installed backend contract to determine the interchange representation required by the selected backend.
- The `build/` directory is generated output, never a source/package/library discovery root.
- If `build/` does not exist when compilation begins, the compiler creates it. If it already exists, it may be empty or contain previous generated output.

License

- **GNU General Public License v3 (GPLv3)**
-

2. Backend

The Backend is responsible for transforming validated frontend output into executable artifacts.

The default backend must be downloaded or built separately. It is not bundled with the frontend binary.

Components

- Intermediate Representation (IR) Generator
- Native Binary Executable Generation

Implementation

A backend may be implemented in any language. It consumes the validated frontend artifact from the project-root `build/` directory. The artifact encoding is selected by the backend-supplied interchange contract.

During backend installation, that interchange-contract file is placed in the same installation directory as the FoLang compiler executable. The frontend reads the installed contract to determine the compatible FoLang/plugin protocol version, HIR schema version, and wire format required by the selected backend. The exact artifact basename and format-specific schema/version may therefore vary by backend contract, but the artifact location is always beneath `<project-root>/build/`.

Default / Reference Backend

The default FoLang backend is also the **reference backend implementation**. Its purpose is to provide an executable, inspectable example of how the FoLang specification can be implemented and to give backend implementers a concrete behavioural baseline for conformance testing.

The written FoLang specification remains normative. The reference backend demonstrates the required externally observable semantics, but its internal algorithms, allocation strategy, memory-management choices, data structures, optimization level, and performance characteristics are not themselves language requirements unless this specification explicitly says otherwise. If an implementation defect in the reference backend conflicts with the written specification, the written specification takes precedence.

A third-party backend may use a completely different runtime architecture or memory model and may optimize semantics differently, provided that FoLang programs observe behaviour conforming to this specification. Backend implementations may validate their behaviour against the reference backend and the FoLang conformance tests where applicable.

- Backend orchestration is implemented in **Go**
- Code generation target is **C++**
- Uses **Clang** or **GCC** to generate native binaries from generated C++ IR

License

Third-party backends may use their own licensing terms and implementation choices. The default backend uses the following license. Default backend is not part of the complete compiler binary and is separate; it must be downloaded or built separately.

- BSD 3-Clause License ***

Frontend Output Contract

Installed Backend Interchange Contract

The selected backend supplies this contract to tell the frontend which FoLang/plugin protocol, HIR schema, and wire format it accepts. During backend installation, the contract file is placed in the same installation directory as the FoLang compiler executable. The frontend reads it from that installation location. The default backend uses the same mechanism.

```
{
  "protocol": "folang-plugin/1.0",
  "hir_schema": "folang-hir/1",
  "wire": "protobuf",
}
```

The frontend has one fixed location contract and a backend-selected encoding contract:

```
<project-root>/build/
└─ <frontend-artifact> encoding/version selected by backend contract
```

Rules:

- the output location is always the reserved root-level build/ directory;
- the selected backend supplies the supported protocol version, HIR schema version, and wire format through its installed interchange-contract file;
- that contract file resides in the same installation directory as the FoLang compiler executable;
- wire="protobuf" in the example above means that particular backend contract requests Protocol Buffers; another supported backend contract may request a different compatible wire format/version;
- build/ is compiler-generated and excluded from all source, package, source-library, operator-source, and packaged-library discovery;
- if build/ is absent, the frontend creates it;
- if build/ already exists, being empty or non-empty is not a project-layout error;
- .fol source files placed under build/ are invalid project layout and are never compiled as source;
- backends consume the validated frontend artifact from build/.

Licensing Summary

Layer	Responsibility	Implementation	License
Frontend	Parsing and semantic analysis	Go	GPLv3
Backend (default)	IR processing and native code generation	Go (orchestration) + C++ (target)	BSD 3-Clause

The copyrightable material in the [FoLang Language Definition and Documentation](#), including its syntax, grammar, and semantic-rule descriptions, is licensed separately under [CC BY 4.0](#). ***

3. Capability Security Model

FoLang’s compiler ships with all language features compiled in but **systems and FFI features are disabled by default**. The compiler has no hardcoded keys — capability configuration happens entirely at install time. This moves authorization from source code (developer-controlled) to the compiler installation (organization-controlled).

Feature Tiers

Tier	Features	Default State
application	All standard language features, co.net, co.core, co.encoding, co.crypto, etc.	✔ Always enabled
system	Raw pointers, pointer arithmetic, @co.dap.native, co.native, co.sys.unsafe, MMIO, heap allocators, low-level platform/runtime implementation	🔒 Disabled — requires install-time configuration

<code>ffi</code>	<code>co.sys.ffi</code> , extern declarations/types, foreign calling conventions and linkage, <code>co.lang.void</code> pointers, C ABI	🔒 Disabled — requires install-time configuration
------------------	---	--

Disclaimer

- In the current alpha specification, the **example requirement** applies to core language syntax and semantic forms whose availability is defined by this specification. A syntax-level feature is considered implemented only when this document gives it a normative definition and at least one example establishing its intended use. Merely appearing in a language table or inventory does not by itself enable a grammar form, declaration kind, special expression form, or language-owned operator.
- This example requirement does **not** govern declared `@co.*` directives, annotations, pragmas, or decorators, and it does **not** determine which ordinary declarations exist inside the language-provided `co.*` package artifact. Those two cases follow the rules below.
- A documented but unimplemented core-language spelling remains reserved. The lexer recognizes the spelling as language-owned rather than treating it as an ordinary user identifier, and the parser rejects its use with an unsupported/unimplemented-feature diagnostic. A table entry by itself does not define a new grammar or declaration form.

Exceptions and Special Handling

`@co.*` Metadata Forms

- Declared built-in directives, annotations, pragmas, and decorators under `@co.*` are accepted by the generic metadata grammar whether or not this document contains an example for the particular metadata name or for each of its fields.
- The parser must collect the **complete metadata application**. This includes the qualified metadata name and every supplied positional argument, named argument, field, attribute, and associated expression. Parser acceptance and collection do not depend on an example being present, and the parser must not discard a field merely because that field is absent from an example or descriptive table.
- If no semantic implementation handles a collected `@co.*` form, the metadata form and all of its collected arguments remain syntactically valid but are silently ignored and have no semantic effect. If an implementation does handle the form, semantic resolution—not parsing—validates required fields, accepted field names, field types, values, defaults, target restrictions, and other form-specific rules. A malformed metadata argument list is still a syntax error under the generic metadata grammar.
- Examples therefore document the semantics and common use of `@co.*` metadata; they are not a prerequisite for parsing or collecting the metadata name or its fields.
- Existing `@co.*` metadata forms may gain additional supported fields or attributes in later revisions provided the generic metadata syntax remains unchanged. User-defined metadata outside the `@co.*` namespace is limited to **annotations and decorators**, which follow the ordinary declaration, import, and name-resolution rules defined for those external metadata forms. Directives and pragmas are language-internal metadata categories; FoLang provides no user declaration construct for creating them.

Language-Provided `co.*` Packages

- The `co.*` package tree is the FoLang standard package API. Although these packages are provided by the language distribution and are implicitly available, their declarations are not hardcoded language grammar merely because their names begin with `co`.
- The language distribution supplies the standard `co.*` package tree through a separate **export-kind `.foLenc` artifact**. The compiler/toolchain loads this language-provided artifact automatically and makes the `co` root available without `@co.ddap.import`. After the package contexts are loaded, their declarations participate in the same ordinary package, symbol, type, member, visibility, extension, and overload resolution model used for packages obtained from any other export-kind `.foLenc` artifact.
- Consequently, the example requirement above does not determine whether an ordinary declaration inside a `co.*` package exists. Availability of a package member is determined by the standard `co.*` export artifact for the applicable FoLang version. If that artifact does not provide the referenced declaration, ordinary name resolution reports a compile-time error.
- After version 1.0, the **package namespace structure** rooted at `co` is frozen. A standard `co.*` package or subpackage cannot be added, removed, renamed, reparented, or otherwise replaced by a different package path. This freeze applies to package and subpackage identity and hierarchy only; it does **not** freeze the declarations contained inside those packages.
- Existing `co.*` packages may continue to evolve as a versioned standard API. They may add or update unit-level functions, methods, types, classes, structs, modules, module instances, algorithms, data structures, and other ordinary declarations expressible with the already-defined FoLang language. For example, a later standard package artifact may add `co.core.RBTree` even though that type was not present in an earlier package inventory. Such an addition is package-API evolution, not a new grammar form or a change to the `co.*` package hierarchy.

A package API addition does not automatically create new special syntax. If a new API type requires a distinct language syntax form rather than ordinary existing declaration, construction, call, generic, operator, or member syntax, that syntax remains subject to the core-language rules of this specification.

Alpha Reserved-but-Unimplemented Policy and Post-1.0 Structural Freeze

The reserved-but-unimplemented category is a property of the pre-1.0 alpha core-language profile only. It is not a permanent part of FoLang. Before version 1.0, every syntax-level or structural language feature listed by this specification without a normative definition and at least one example establishing its purpose and use is resolved in exactly one of two ways:

- it gains a normative definition, at least one example, and a working implementation, and becomes part of the language; or
- it is removed from the core-language inventory entirely.

This rule excludes the declared `@co.*` metadata forms described above and ordinary declarations supplied by the standard `co.*` export artifact. Their parsing and package-member availability are governed by their respective rules rather than by the example requirement.

Nothing is carried into 1.0 as reserved-but-unsupported core syntax. At version 1.0, every entry that remains in a core-language syntax or structural-feature table is implemented rather than merely promised.

The alpha period exists for experimentation, consolidation, implementation, renaming, syntax changes, and removal. Beta and release-candidate versions may continue to harden the resulting design, but the **structural FoLang language surface** closes permanently at version 1.0.

After version 1.0, no later major or minor release introduces new core grammar forms, keywords, declaration kinds, operator spellings or fixities, or built-in `@co.*` metadata-form names. Existing structural constructs must retain the externally observable semantics defined by the 1.0 specification, except for corrections that restore already-stated intent.

The standard-package rule is narrower and separate: the `co.*` **package/subpackage hierarchy** is frozen, while declarations inside those existing packages may evolve through later language-provided `.folenc` versions. Adding `co.core.RBTree`, adding a unit-level function, or updating an existing standard-package API does not add a package path and therefore does not violate the package-tree freeze.

FoLang also remains extensible after 1.0 through extension mechanisms that are themselves already part of the 1.0 language, including third-party libraries, user-defined annotations and decorators, macros, custom operators, native and FFI integration, dynamic-runtime facilities, backend integration points, and other extension mechanisms explicitly defined by the 1.0 specification. These mechanisms may provide new capabilities to programs without changing the structural language surface.

Post-1.0 FoLang releases may therefore evolve through implementation improvements, standard-package API evolution, and external extensions within the frozen structural language. The frontend, backend, runtime, optimizer, intermediate representations, diagnostics, compilation strategy, execution performance, memory management, code generation, supported processor and accelerator architectures, and hardware-specific capabilities may all improve substantially. Such improvements may exploit new CPU, GPU, NPU, accelerator, or other hardware capabilities, but they must preserve the externally observable semantics defined by the FoLang specification.

A post-1.0 correction may fix an implementation/specification discrepancy or remove an internal specification contradiction when doing so restores the already-stated intent of an existing FoLang language feature. A correction does not introduce a new grammar form, keyword, declaration kind, operator grammar, or previously unavailable structural language capability. The permanent structural freeze must not be bypassed by classifying a language addition as a correction. Standard-package API evolution inside an existing `co.*` package is governed separately by the package rules above and is not a structural-language correction.

Quick Start

Hello World

```
// src/appl.fol – fixed application entry file, no annotation needed
co.out.println("Hello FoLang!");
```

Or with an alias for shorter form:

```
@co.ddap.alias(co.out, as="out")

out.println("Hello FoLang!");
```

Variables

```
// typed
name co.lang.string = "Rao";
age  co.lang.int    = 30;

// inferred from value - := errors if already declared
name := "Rao";
age  := 30;

// define, infer, and assign if not defined; otherwise assign a new value
name ?= "Kumar";
```

`co.lang.string` and `co.lang.int` are built-in data types. For more information, see [Builtin Data Types](#).

`=`, `:=`, and `?=` are built-in operators. For more information, see [Builtin Operators](#).

Constants and Immutability

FoLang separates two properties that are often confused. One is about *when* a value is known. The other is about *whether* it can change.

```
// compile-time constant - value known while compiling, substitutable
@co.dap.const SIZE co.lang.int = 1024;

// immutable binding - cannot be reassigned, value need not be known early
@co.dap.final startedAt co.lang.int = co.sys.now();
```

Annotation	Guarantees	Value known at compile time	Usable as an index
<code>@co.dap.const</code>	value is fixed and known while compiling	✓	✓
<code>@co.dap.final</code>	binding cannot be reassigned	✗ not required	✗

Every `@co.dap.const` is also immutable. The reverse does not hold: a `@co.dap.final` binding may be initialised from a function call, a file, or user input, so the compiler cannot substitute a literal for it.

`@co.dap.final` is the declaration-site form of the immutable object kind described in the object mutation policy. `makeImmutable` applies the same property to a value at run time.

Outside a declaration whose signature binds a dependent index parameter, only `@co.dap.const` may supply a named array size or dependent type index. A signature-bound index name is permitted symbolically even when its concrete value is supplied by the caller. See [Dependent Type Index Rules](#).

Single Source Application File

FoLang developers can create a complete executable program in one source file. A **single-source application** is an application whose entry file contains the complete program and does not depend on user package source files.

A single-source application file and an application entry file are the same fixed structural source, `src/app1.fo1`, and use the same entry-file grammar, context, and restrictions. A project is a single-source application when `src/app1.fo1` contains the complete program and there are no application package directories below `src/`. This section presents the allowed constructs, so a developer can start programming without first reading the complete specification.

For the single-source application layout, see [Project Layout](#).

Allowed Constructs

The application file may contain:

- built-in directives that are valid for an entry file
- package and library imports
- import aliases declared with `as=`
- file-local aliases for `co.*` paths declared with `@co.ddap.alias`
- type aliases and ADTs declared with `co.lang.type`
- parameterized `co.lang.type` constructors such as `Option(T) co.lang.type = Some(T) | None()`
- new types declared with `co.lang.newtype`
- opaque types declared with `co.lang.opaquetype`
- dependent-type aliases and dependent-type usages that do not declare an ordinary or type-level function
- refinement-type declarations

- subtype declarations
- supertype declarations
- non-capturing entry-local function-pattern groups
- capturing entry-local `let` function-pattern groups
- variable declarations, initialization, assignment, and mutation
- calls to built-in methods and functions
- calls to imported package and library APIs
- ordinary expressions
- comprehensions
- ternary expressions
- conditions and loops
- containment operations such as `contains`
- iterators such as `each`
- pattern matching and destructuring

```
a co.lang.int = 10;
b co.lang.int = 20;

co.out.println( a + b );
```

`co` is a reserved word in FoLang. For more information, see [Reserved Words](#).

`co` is the built-in root package in FoLang. For more information, see [Builtin Packages](#).

`a + b` is an expression in FoLang.

For the expression rules, see [Expressions](#).

Built-in and Imported Names

All `co.*` paths are always available.

A developer may use the complete built-in path:

```
co.out.println("Hello");
co.core.List.of(1, 2, 3);
```

A developer may optionally create a file-local alias:

```
@co.ddap.alias(co.out, as="out")
@co.ddap.alias(co.core.List, as="list")

out.println("Hello");
values := list.of(1, 2, 3);
```

Creating an alias does not hide the complete `co.*` name; both forms remain valid in that file.

Third-party packages, user packages, and libraries are not automatically available. They must be imported using `@co.ddap.import`. When `as=` is present, the imported API is accessed through that alias. When `as=` is omitted, the complete imported package or library path must be used.

```
@co.ddap.import(package="hr.employee", as="emp")
first := emp.EmployeeService.find(1001);

@co.ddap.import(package="finance.payroll")
second := finance.payroll.calculate(request);
```

A developer must import a package before using it.

`@co.ddap.import` and `@co.ddap.alias` are built-in directives. For more information, see [Built-in Directives](#).

For more information about FoLang imports, see [Import Details](#).

`println` is a built-in method on the `co.out` object.

Variable Kinds

FoLang supports several variable kinds for different purposes. Their availability is context-dependent; a developer cannot use every variable kind in every location. For the contexts in which each kind is supported, see [Variable Kind Support](#).

Simple Variable Declaration

```
someVar co.lang.int;
someString co.lang.string;
```

With Initialization

```
someBool co.lang.bool = co.const.true;
someInt co.lang.int = 42;
```

With Type Inference

```
someVal := "Hello, World!";
someNum := 3.14; // if not defined, define and initialize; else throws error
someR ?= "Kamesh"; // if not defined, define and initialize; else reassign
```

Alpha Character and String Literals

The alpha release accepts only the basic literal subset:

```
message := "hello";
letter := 'A';
```

A string is unprefixes, remains on one source line, and cannot contain a backslash or an embedded double quote. A character literal contains exactly one non-backslash character.

Encoding prefixes, escaped characters, raw strings, and universal character names are reserved for a later release. The scanner recognizes the complete spelling and reports an unsupported-feature error instead of tokenizing it as several unrelated expressions. For example, all of these are rejected:

```
x := "quoted: \"text\"";
x := '\n';
x := u8"hello";
x := R"(raw text)";
x := '\N{LATIN CAPITAL LETTER A}';
```

Pointer Declaration

```
somePtr    co.lang.int->>(*);
someDblPtr co.lang.int->(**);
someDeepPtr co.lang.int->(****);
```

The number of consecutive `*` characters is the pointer degree. Any positive degree is permitted; `*` and `**` above are common examples, not a maximum.

Array Declaration

```
someArray    co.lang.int->([5]); // single dimension
someDblArray co.lang.int->([2,3]); // multi dimension
someJaggedArray co.lang.int->([2][3]); // jagged
someVLAArray co.lang.int->([...]); // variable length
someZeroLA    co.lang.int->([0]); // zero length array
someZeroDimA  co.lang.int->([.]); // zero-dimensional array
```

Array Declaration with Initialization

```
someInitializedArray co.lang.int->([3]) = [1, 2, 3];
someInitializedArray1 co.lang.int->([ ]) = [1, 2, 3];
someInitializedDblArray co.lang.int->([,]) = [[1, 2], [3, 4]];
```


Reference Declaration

```
someRef      co.lang.int->(&);    // reference
someLValueRef co.lang.int->(&&);  // LValue reference
someHpRef     co.lang.int->(~);   // heap allocated reference
someAddr      co.lang.int->(@);   // address
someThunk     co.lang.int->(^);   // thunk
someSlice     co.lang.int->(:);   // slice
```

Range Declaration

```
// Typed range variable declaration
someRange co.lang.int->(..);

// Inferred range declarations
rangeI := 1 .. 10;    // [1, 10]   ExcludeStart=false, ExcludeEnd=false
rangeS := 0 <.. 5;    // (0, 5]    ExcludeStart=true,  ExcludeEnd=false
rangeL := 0 ..< 100;  // [0, 100)  ExcludeStart=false, ExcludeEnd=true
rangeB := 0 <..
```

Auto and Dynamic Variable Declaration

```
someAutoVar   co.lang.auto    = "Hello"; // type inferred from value; initialization required
someDynamicVar co.lang.dynamic;           // dynamic typing
```

Lazy

```
@co.dap.lazy
x = add(1, 2); // evaluating/using x invokes add(1, 2); until then the right-hand function is not invoked
```

Bind Variables

```
$[1-9][0-9]*
```

Bind variables are available in function-chaining expressions. At each chained call, `$1`, `$2`, `$3`, ... denote the return components of the **immediately preceding function invocation**, in declared return order. If the preceding function returns one value, only `$1` is available for that step; if it returns multiple values, `$1` through `$N` correspond to those return values. `$0` is not a bind variable.

```
dosomething(a co.lang.int, b co.lang.int)->(co.lang.int)
=>> somePack.someMethod(a)
=>> someOthPack.someOtherMeth($1, b);

// If previous() returns (A, B, C), the next chained call may consume
// those return components as $1, $2, and $3 respectively.
previous()
=>> consume($1, $2, $3);
```

Discard / Wildcard Variable

```
_
```

`_` is contextual rather than an ordinary identifier. In discard/wildcard positions it means “ignore this binding or match position.” It is also used as the filename-derived primary-declaration placeholder and as an unnamed type-parameter slot in type-constructor forms such as `F(_)`. In the predicate of a `co.lang.refinementType` declaration, `_` instead denotes the candidate value of the base type being tested. See [Refinement Types](#).

Comma and Grouping

```
// Comma
x co.lang.int = 10, y co.lang.string = "Hello", z co.lang.bool = co.const.true;

// Grouping
(x co.lang.int = 10, y co.lang.string = "Hello", z co.lang.bool = co.const.true);

// Tuple/grouping expression
pair := (x, y);
```

A comma inside a parenthesized expression or grouped declaration must introduce another expression or declarator. Therefore `(x,)`, `(x, y,)`, and `(x co.lang.int = 10,)` are invalid. Collection literals have their own rules and may permit a trailing comma where their production says so. Record patterns also require another field after every comma; `Employee{id: value,}` is invalid.

Fat Pointers

```
x co.lang.int->>(*, kind="", meta={});

co.lang.int->>(*, meta={});

y co.lang.int->>(*, meta={len:co.lang.usize, vtab:somepkg.VTable->(*)});

z co.lang.int->>(*kind=region, meta={});
```

```
Pointer
├─ base_type: T
├─ kind: <FatKind>
│   ├── thin
│   ├── slice
│   ├── relative
│   ├── trait
│   ├── buffer
│   ├── view
│   ├── opaque
│   ├── custom
│   ├── mem
│   ├── nullptr
│   ├── sptr
│   ├── uptr
│   ├── ptrdiff
│   ├── usize
│   ├── ssize
│   └─ (region) ← optional syntactic sugar
└─ meta:
    ├── region: heap | stack | global | numa(N) | mmio | constant | ...
    ├── len, cap, vtab, bits, endian, ...
```

Pointers for address manipulation

```
y co.lang.word->(repr=intptr);
z co.lang.word->(sign=unsigned, repr=uintptr);
p co.lang.word->(repr=ptrdiff);
n co.lang.word->(sign=unsigned, repr=usize);
m co.lang.word->(repr=isize);
o co.lang.void->(repr=nullptr);
```

Relative Pointers

```
z co.lang.int->>(*kind=relative, meta={});
```

Note Other than normal variable declaration rest have restrictions

Conditions, Loops and Iterators

Conditions

`then` is the one-shot conditional branch verb. Its argument may be a block or an ordinary value/expression. `otherwise(condition)` always introduces another Boolean condition; a conditionless `otherwise` form does not exist. `default(result)` is the optional terminal fallback and may likewise receive a block or value.

```

(boolean truth).then({
}).otherwise(boolean truth).then({
}).default({
});

//conditionEg1.unit.fol

_ co.lang.unit = {

    someFun()->()={

        x := co.const.true;

        x.then({

        }).default({

        });

    }

    someOtherFun()->()={

        (co.const.true).then({ //parenthesis mandatory

        }).default({

        });

    }

    someOtherFun1()->()={
        x co.lang.bool =co.const.true;
        (x).then({ // parentheses optional

        }).default({

        });

    }

    someOtherDiffFun1()->()={

        x co.lang.int=10;
        y:=30;

        (x > y).then({ // parentheses around (x > y) are mandatory

        }).otherwise(x < y).then({

        }).default({

        });

    }

}

```

Loops

`loop` is the repeated-execution verb and always has exactly one loop condition and one loop body. Unlike `then`, a `loop(...)` form cannot participate in an `otherwise(condition)` chain, and `default(...)` cannot follow a loop. When the loop condition is false, whether on the first test or after one or more iterations, the loop simply terminates.

To choose between different looping behaviours, first use a `then` / `otherwise(condition)` / `default` selection chain and place the required ordinary loop inside the selected block. The selection chain chooses a behaviour; each nested `loop(...)` then performs repetition using only its own condition.

```

(boolean truth).loop({
});

//loopsEg1.unit.fol
_ co.lang.unit = {
  someFun()->()={
    x := co.const.true;

    x.loop({
    });
  }

  someOtherFun()->()={
    (co.const.true).loop({ //parenthesis mandatory
    });
  }

  someOtherFun1()->()={
    x co.lang.bool =co.const.true;
    (x).loop({ // parentheses optional
    });
  }

  someOtherDiffFun1()->()={
    x co.lang.int=10;
    y:=30;

    (x > y).loop({ // parentheses around (x > y) are mandatory
    });
  }
}

```

```

//loopsEg2.unit.fol
_ co.lang.unit = {
  someFun()->()={
    x := co.const.true;
    v := 0;
    x.loop({
      (v == 10).then({
        this.break;
      });
      v += 1;
    });
  }

  someOtherFun()->()={
    v := 20;
    (co.const.true).loop({ //parenthesis mandatory
      (v == 30).then({
        this.continue;
      });
      v += 5;
    });
  }
}

```

Combining Conditions and Loops

A loop does not become a branch verb inside an `otherwise(condition)` chain and has no `default(...)` fallback of its own. When a program must choose between different looping behaviours, the selection is expressed separately with `then` / `otherwise(condition)` / `default`, and each selected block may contain its own ordinary single-condition loop.

```

(first condition).then({
  (first loop condition).loop({
    ...
  });
}).otherwise(second condition).then({
  (second loop condition).loop({
    ...
  });
}).default({
  ...
});

```

Ternary Operator

The ternary/value form uses the same `then` / `otherwise(condition)` / `default` selection vocabulary as block conditionals. Only the branch arguments differ: a value-producing chain supplies values or expressions instead of statement blocks.

```

s = (boolean truth).then(some var/value).default(some val/var);
s = (boolean truth).then(some var/val).otherwise(boolean truth).then(some var/val).default(some var/val);

//TernaryExample.unit.fol

_ co.lang.unit = {

  someFunction()->()={

    s := (co.const.true).then(20).default(10);
    y co.lang.int;
    z co.lang.int=20;
    y ?= (co.const.true).then(20).default(z);

  }

  someOtherFunction()->()={

    k co.lang.int=10;
    p co.lang.int =20;
    s ?= (k>10).then(30).otherwise(k<10).then(p).default(10);

  }

}

```

Parenthesizing the chain head

A conditional or ternary selection chain is a sequence of postfix method calls using `.then(...)`, `.otherwise(condition)`, and optional terminal `.default(...)`. A loop uses the same postfix-call style but is a single-condition repetition form: `.loop(...)` cannot be followed by `.otherwise(condition)` or `.default(...)`.

These calls always carry their own call parentheses, whatever the argument is. Nothing about the argument changes that.

`otherwise` always has the form `.otherwise(condition)` and therefore never acts as a terminal else marker. The terminal fallback is `.default(...)`; its argument may be a block or a value/expression according to the applicable form.

The only place a choice exists is the **head**: the subject before the first `.then(...)` or `.loop(...)` call. The head must already be a complete postfix expression that cannot absorb the following control verb.

Chain head	Parentheses	Why
<code>x</code>	optional	an identifier is already a complete postfix expression
<code>arr[0]</code> , <code>f()</code>	optional	index and call suffixes are postfix too
<code>co.const.true</code> , <code>myPkg.flag</code>	required	a qualified name would absorb <code>.then</code> as a further segment, giving <code>co.const.true.then</code>
<code>x > y</code>	required	<code>.then</code> binds tighter than <code>></code> , so <code>x > y.then({...})</code> groups as <code>x > (y.then({...}))</code>

A qualified name needs parentheses whether it names a literal or an identifier, so `myPkg.flag` is the same case as `co.const.true`.

```

x.then({ ... }); // head is an identifier
(co.const.true).then({ ... }); // head is a qualified name
(x > y).then({ ... }).otherwise(x < y).then({ ... }); // head is an expression
(k > 10).then(30).otherwise(k < 10).then(p).default(10);

```

In the last two lines the parentheses after `.otherwise` are its call parentheses, not an application of this rule.

Iterating Arrays / Lists / Maps / Ranges

`each` is itself the element-iteration operation. It does not produce a value that must subsequently be passed to `loop`, and `.loop(...)` must not follow an `each(...)` call.

The explicit-binding form takes the index/key binding, the value binding, and the per-element action directly. The action may be a block or any ordinary expression valid in that position. A callable may also be supplied directly as the sole callback argument; the callable receives the receiver's defined iteration tuple. A lambda is a callable callback and is permitted directly in this collection-operation context.

```

arr co.lang.int->([5]) = [6,7,8,9,10];

// block action
arr.each(idx, val, {
  co.out.print(idx);
  co.out.print(" :: ");
  co.out.println(val);
});

arr.each(_, val, {
  co.out.println(val);
});

// expression action – evaluated once for each element
arr.each(_, val, co.out.println(val));

// direct callable callback – receives the receiver's iteration tuple
arr.each(printItem);

// direct lambda callback
arr.each(|idx, val| => co.out.println(val));

```

The explicit-binding form establishes its bindings separately for every iteration and evaluates its action exactly once for that element. The action is therefore the body of `each`; `each` never participates in a `then`, `otherwise`, `default`, or `loop` chain. In the single-argument callback form, the argument must resolve to a callable compatible with the receiver's iteration tuple.

Array / List / Map / Range — Contains Element

```

arr co.lang.int->([5]) = [35,57,96,81,31];
k co.lang.int = 31;
arr.contains(k).then({
  co.out.println(k);
}).default({
  co.out.println("Not Found");
});

```

Comprehensions

```

k := (1 .. 10).filter(|x| => x % 2 == 0).map(|x| => x * x);

result := for (x <- List[1,2,3]).yield(x * 2); // List[2, 4, 6]
result := for (x <- Set(1,2,3)).yield(x * 2); // Set(2, 4, 6)
result := for (x <- Some(5)).yield(x * 2); // Some(10)
result := for (x <- fetchData()).yield(x.process()); // Future

ages := Map{"A":30,"B":40,"C":66,"E":88};
upper := for ((name, age) <- ages).yield(name.toUpperCase, age);

```

Pattern Matching

```
x co.lang.int = 10;

x.match.case(n: n > 10 => { n = n+100; "GT" }).case( n: n < 10 => "LT").default("EQ");
x.match.case(n: n > 10 => { n = n+100; "GT" }).case( n: n < 10 => "LT").case(_=>"EQ");
x.match(co.pattern.Type).case(co.lang.int => ...).case(co.lang.float => ...);
x.match(co.pattern.Value).case(0 => ...).case(1 => ...);
x.match(co.pattern.Instance).case(xx.CAT => ...).case(xx.DOG => ...).default("Animal");
x.match(co.pattern.Object).case(xx.Ball => "Ball").case(xx.CAT => "CAT").default("Unknown");
x.match(co.pattern.Shape).case(Point{x, y} => ...).default(...);

x.match(co.pattern.Any).case(co.lang.int => ...).case(co.lang.float => ...).case(0 => ...).default( ...);

x.match(PositiveEvenMatcher).case(0 => "Neither even nor odd").case(2 => "First Even Prime").default(...);
```

A match chain contains one or more `.case(...)` arms followed by at most one terminal `.default(...)` arm. `.otherwise` is not a match arm; it always introduces another Boolean condition in conditional `then` chains. Loops do not participate in `otherwise(condition)` chains. `default` is the common terminal fallback vocabulary for Boolean selection and match/case selection. A case or default result may itself be a ternary then expression, and that nested expression may consequently contain `.default(...)`.

For value dispatch, `match().case(...).default(...)` is also FoLang's generalized analogue of a traditional `switch / case / default` construct; match cases additionally support the pattern, type, object, and custom-matcher forms defined by this specification.

`_` is a special discard/wildcard variable. In a call it is permitted only as the first, index/key binding of the explicit receiver-qualified `.each` form, as in `items.each(_, value, { ... })`. Transparent grouping around the member callee, for example `(items.each)(_, value, { ... })`, does not change this rule. The value binding and the iteration-action argument of `each` cannot be discarded, and `contains(_)` and `containsVal(_)` are invalid because containment must compare a real value. Patterns and the filename-derived declaration-name form, type-constructor placeholder forms such as `F(_)`, and refinement-type predicates give `_` their own explicitly described meanings. In a refinement predicate, `_` denotes the candidate value being tested; it is not a general expression identifier.

`PositiveEvenMatcher` is a custom matcher. For more information about defining custom matchers, see [Custom Matcher](#).

Type Declarations

```
// Alias
x co.lang.type = co.lang.int;

// New
x co.lang.newtype = co.lang.int;

// Opaque
EmpIdType co.lang.opaquetype = co.lang.int;
DeptIdType co.lang.opaquetype = co.lang.int;

empId EmpIdType = 10;           // valid: base representation is accepted when constructing EmpIdType
deptId DeptIdType = 20;        // valid

empId = deptId;                // invalid: distinct opaque types are not assignment-compatible
empId2 EmpIdType = empId;      // valid: same opaque type

x co.lang.int = empId;         // invalid: an opaque value does not implicitly become its base type

// ADT (tagged union)
y co.lang.type = co.lang.int | co.lang.char;

// Proper subtypes of Employee; Employee itself is excluded.
// Employee is a class.
empSubType co.lang.subtype = somePackage.Employee;

// PermanentEmployee inherits Employee.
permanentEmp empSubType = PermanentEmployee{}; // valid: proper subtype

// ContractualEmployee inherits Employee.
contractualEmp empSubType = ContractualEmployee{}; // valid: proper subtype

dancerEmp empSubType = DashingDancer{}; // compiler error: not a subtype of Employee
baseEmp empSubType = Employee{}; // compiler error: Employee itself is excluded

// To accept Employee itself together with all proper subtypes:
empPlusType co.lang.type = empSubType | Employee;

// Proper supertypes of Toyota; Toyota itself is excluded.
superToyota co.lang.supertype = somePackage.Toyota;

// somePackage.Toyota extends somePackage.Car extends somePackage.FourWheeler extends somePackage.Vehicle

toyota superToyota = somePackage.Toyota{}; // compiler error: Toyota itself is excluded
car superToyota = somePackage.Car{}; // valid
four superToyota = somePackage.FourWheeler{}; // valid
vehicle superToyota = somePackage.Vehicle{}; // valid
truck superToyota = somePackage.Truck{}; // compiler error: not a supertype of Toyota

// To accept Toyota itself together with all proper supertypes:
superToyotaPlus co.lang.type = superToyota | somePackage.Toyota;

// Refinement type
positiveInt co.lang.refinementType = (co.lang.int).where(_ > 0);

percentage co.lang.refinementType =
  (co.lang.int).where(_ >= 0 && _ <= 100);

evenInt co.lang.refinementType =
  (co.lang.int).where(_ % 2 == 0);

nonEmptyString co.lang.refinementType =
  (co.lang.string).where(_.length > 0);

// Inside a refinement predicate, _ denotes the candidate value of the base type.
```

For the normative refinement-type rules, see [Refinement Types](#).

For `co.lang.opaquetype`, the declared base type supplies the representation accepted where opaque-type construction permits it, but the resulting opaque type has distinct type identity. Distinct opaque types are not assignment-compatible merely because they share the same base type, and an opaque value is not implicitly assignable back to its base type. Thus `x co.lang.int = empId;` is invalid when `empId` has type `EmpIdType`.

`co.lang.subtype` and `co.lang.supertype`

`co.lang.subtype` and `co.lang.supertype` define type sets for these two declaration kinds only. Their semantics do not define or replace inheritance, interface, generic, object-model, variance, or other assignability rules elsewhere in FoLang; those facilities retain their own independently defined semantics.

A declaration of the form:

```
TSub co.lang.subtype = BaseType;
```

defines a type whose admissible values have concrete types that are **proper transitive subtypes** of `BaseType`. `BaseType` itself is excluded. Direct and indirect subtypes are included, while unrelated types are excluded.

To admit `BaseType` itself together with those proper subtypes, form an explicit union:

```
TSubPlusBase co.lang.type = TSub | BaseType;
```

A declaration of the form:

```
TSuper co.lang.supertype = BaseType;
```

defines a type whose admissible values have concrete types that are **proper transitive supertypes** of `BaseType`. `BaseType` itself is excluded. Direct and indirect supertypes are included, while unrelated types are excluded.

To admit `BaseType` itself together with those proper supertypes, form an explicit union:

```
TSuperPlusBase co.lang.type = TSuper | BaseType;
```

Let Bindings

```
y co.lang.int = let({x = 10}).in({x + 1});
y co.lang.int = let({$ = 10}).in({$ + 1}); // $ refers to the value being defined

x co.lang.int = (x + 1).where(x = 10);
x co.lang.int = ($ + 1).where($ = 10);

offset := 100;

let adjust(0) = offset;
let adjust(n) = n + offset;
```

`$` is a special identifier usable in ordinary `let` binding expressions for recursive or self-referential expressions.

Ordinary `let` value-binding expressions remain available in language contexts that permit them, but they are forbidden directly in the application entry file. In the entry file, `let` is reserved exclusively for a named function-pattern group that captures at least one surrounding runtime binding. It cannot introduce an anonymous function, a general closure value, or a curried function.

Function Pattern

```
Option(T) co.lang.type= Some(T) | None();

f(Some(x)) => { this.return x + 1; }
f(None()) => { this.return 0; }

// desugars to:
f(v Option(co.lang.int))->(co.lang.int) = {
  this.return v.match()
    .case(x: Some(x) => x + 1)
    .case(_: None() => 0);
}
```

`=>` introduces a bare function-pattern clause. `=>>` is the distinct function-delegation operator, while `=>>>` is the closure/curry expression; neither introduces a function pattern.

Function-pattern groups are permitted in the application entry file as restricted entry-local dispatch helpers. A bare group cannot capture surrounding runtime variables. A `let` function-pattern group must capture at least one already initialized entry-file runtime binding and is the only entry-file construct that permits such capture. Neither form permits ordinary function declarations, anonymous functions, general closure values, currying, partial application, or escape as a function value.

For more information about `let` and function patterns, see [Let and Function Patterns](#).

A single-source application file is useful for testing FoLang and becoming familiar with the language. Real-world applications normally contain more than a single source file. They use abstraction, encapsulation, inheritance, polymorphism, and other language features, and they often depend on external packages and libraries to establish clear boundaries. At minimum, such applications require the following structural features:

1. [package source files](#) under [packages](#)
2. [Entry File](#)
3. [Libraries](#) and [Library surface files](#)
4. [imports](#)

FoLang supports many features for developing enterprise applications. The following list should be read together with the language [intent](#) and [FoLang Philosophy — Uniform Object Model](#).

Complete feature list:

1. [Packages](#)
2. [UDT](#)
3. [Functions](#)
4. [Units](#)
5. [Imports](#)
6. [Macros](#)
7. [Templates](#)
8. [Annotations and Decorators](#)
9. [Type Classes](#)
10. [Types](#)
11. [Generics](#)
12. [Matchers](#)
13. [Lambdas](#)
14. [Execution Models and Control Abstractions](#)
15. [Extensions](#)
16. [Native code](#)
17. [Indexers](#)
18. [Refinement Types](#)
19. [Dependent Types and Type-Level Functions](#)
20. [Dynamic Runtime](#)
21. [Local/Nested Types and Functions](#)
22. [Libraries](#)
23. [Export Projects](#)
24. [Operators](#)
25. [Forward / Extern Declarations](#)
26. [Labels and Named Blocks](#)
27. [Reflections](#)
28. [Comprehensions](#)

In FoLang, file-backed primary declarations use their own `<Name>.fo1` files. Package functions and non-UDT type declarations are grouped in any number of `*.unit.fo1` files, while struct-associated behavior is placed in `<StructName>.comp.unit.fo1`. These are all [package source files](#). The following sections begin with packages before moving to UDTs and functions.

Packages

Package Identity

FoLang projects use four standardized project-root domains: `src/`, `src1ib/`, `lib/`, and `build/`. These directory names are compiler-defined filesystem domains, not packages, and never contribute namespace components.

Only `src/` is mandatory before compilation. `src1ib/` and `lib/` are optional and must be omitted when unused; if either is physically present, it must contain valid content for that domain and cannot be empty. `build/` is compiler-managed generated output: it may be absent before compilation, in which case the frontend creates it, and if already present it may be empty or non-empty without causing a layout error.

Ordinary project package discovery occurs only below `src/`:

- `src/` itself is not a package.
- `src/` must exist before compilation.
- an application project has the fixed direct surface file `src/appl.fol`;
- a standalone packaged-library project has the fixed direct surface file `src/library.fol`;
- a standalone export project has the fixed direct surface file `src/export.fol`;
- exactly one of `src/appl.fol`, `src/library.fol`, or `src/export.fol` must exist;
- if more than one exists, project classification is ambiguous and compilation fails;
- if none exists, no valid project surface exists and compilation fails;
- no other file may occur directly in `src/`;
- every other entry directly under `src/` must be a non-empty package directory containing valid FoLang source;
- package dot paths are relative to `src/`;
- nested package directories may form an arbitrarily deep package hierarchy;
- `src/lib/`, `lib/`, and `build/` never participate in direct `src/` package discovery. Selected `src/lib/exports/` packages and export-kind `.folenc` package contexts are added later to the applicable open package index by their dedicated preparation rules.

Examples:

```
/appl/src/hr/      -> package "hr"
/appl/src/hr/employee/ -> package "hr.employee"
/appl/src/auth/    -> package "auth"
```

Neither the application root nor `src/` is a package.

Multi-File Packages

Multiple `.fol` files in the same package directory automatically belong to that package:

```
src/hr/employee/
├─ Employee.fol      -> hr.employee
├─ EmpService.fol    -> hr.employee
└─ EmpValidator.fol  -> hr.employee
```

Project Layout

FoLang uses one project-root architecture for applications, standalone packaged-library projects, and standalone export projects. The compiler is invoked with the **project root**, not with an individual source surface file. From that root it deterministically classifies the project and discovers every relevant domain.

```
folang compiler <project-root>
|
+-- src/appl.fol      -> application
|
+-- src/library.fol   -> standalone packaged library
|
+-- src/export.fol    -> standalone export project
|
+-- src/lib/          -> optional project-local source libraries/exports/operators
|
+-- lib/              -> optional packaged .folenc dependencies
|
+-- build/            -> generated output; create when absent
```

The project kind changes the meaning/grammar of the fixed direct surface file under `src/`; it does not create a different filesystem architecture.

Project kind	Fixed <code>src/</code> surface	Final product
application	<code>appl.fol</code>	platform/backend executable or application binary
library	<code>library.fol</code>	<code><project-name>.folenc</code> containing only the projected library-surface symbol/API view plus compiled implementation
export	<code>export.fol</code>	<code><project-name>.folenc</code> containing the package contexts and symbols selected by <code>@co.dap.export</code> plus compiled implementation

The build command therefore needs only the project root. The developer does not identify an entry/surface file separately, and the compiler does not search for an arbitrary candidate filename: it checks the reserved structural paths `src/appl.fol`, `src/library.fol`, and `src/export.fol`, then discovers `src/`, optional `src/lib/`, optional `lib/`, and compiler-managed `build/` relative to that same root.

Project Identity and Output Basename

The logical project name is derived from the basename of the canonical project root directory. The fixed structural source filenames `appl.fol`, `library.fol`, and `export.fol` do **not** contribute the application, library, or export-artifact name.

```
/projects/payroll/
└─ src/
   └─ appl.fol

project kind = application
project name = payroll
output basename = payroll
```

```
/libraries/hrlib/
└─ src/
   └─ library.fol

project kind = standalone packaged library
library name = hrlib
compiled library artifact = hrlib.folenc
```

```
/packages/folcore/
└─ src/
   └─ export.fol

project kind = standalone export
export artifact name = folcore
compiled export artifact = folcore.folenc
```

Therefore:

```
project-root basename
    -> project/artifact identity and default output basename

src/appl.fol
    -> application project kind only

src/library.fol
    -> standalone packaged-library project kind only

src/export.fol
    -> standalone export-project kind only
```

`appl.fol` never implies an application named `appl`, `library.fol` never implies a library named `library`, and `export.fol` never implies an export artifact or package named `export`.

For an application, the executable uses the project name as its default basename; the platform/backend determines the final executable file convention or suffix. For example, a project rooted at `payroll/` has executable basename `payroll`, not `appl`.

For a standalone packaged library or standalone export project, the standard compiled FoLang distributable artifact uses the derived project name:

```
<hrlib project root>/    -> hrlib.folenc
<folcore project root>/  -> folcore.folenc
```

The `.folenc` artifact records its artifact kind. A library-kind `.folenc` contains the projected `library.fol` surface API; an export-kind `.folenc` contains the package contexts and symbols selected by `export.fol`.

This identity rule avoids redundant source-level naming and therefore avoids directory-name versus declared-name mismatch states.

The following application tree illustrates all available domains. `src/lib/` and `lib/` are shown to document where those domains live; in a real project they must be omitted entirely when unused and cannot exist as empty directories. `build/` is compiler-managed and may be absent or empty before compilation.

```

/appl/
├── src/                                <- REQUIRED project source domain; NOT a package
│   ├── appl.fol                      <- fixed application entry / single-source application surface
│   ├── hr/                           <- package hr
│   │   ├── employee/                 <- package hr.employee
│   │   │   ├── Employee.fol
│   │   │   └── EmpService.fol
│   │   └── payroll/                  <- package hr.payroll
│   │       └── Payroll.fol
│   └── auth/                         <- package auth
│       └── Auth.fol
├── srclib/                           <- OPTIONAL; omit when unused; if present, non-empty
│   ├── application/
│   │   ├── library.fol
│   │   └── <internal packages>/
│   ├── ffi/
│   │   ├── library.fol
│   │   └── <internal packages>/
│   ├── system/
│   │   ├── library.fol
│   │   └── <internal packages>/
│   ├── advanced/
│   │   ├── library.fol
│   │   └── <internal packages>/
│   ├── dynamicvmrt/
│   │   ├── library.fol
│   │   └── <internal packages>/
│   ├── exports/
│   │   ├── export.fol
│   │   └── <export-source packages>/
│   └── operators/
│       └── library.fol                <- exactly one file; no subdirectories
├── lib/                              <- OPTIONAL; omit when unused; if present, one or more *.folenc
│   ├── postgres.folenc
│   └── crypto.folenc
└── build/                           <- OPTIONAL before compile; compiler creates/manages it
    └── <frontend-artifact>          <- encoding/version selected by backend contract

```

A standalone packaged-library project uses the same root domains:

```

/hrlib/
├── src/                                <- REQUIRED
│   ├── library.fol                  <- fixed standalone-library surface; @co.dap.library(type=...)
│   ├── emp/                         <- internal package emp
│   │   ├── Employee.fol
│   │   └── EmpService.fol
│   └── auth/                        <- internal package auth
│       └── AuthService.fol
├── srclib/                           <- OPTIONAL; omit when unused; if present, non-empty
├── lib/                              <- OPTIONAL; omit when unused; if present, non-empty
└── build/                           <- compiler-managed; may be absent/empty before compilation

```

The empty `srclib/` and `lib/` entries in the second diagram are structural illustration only. An actual project must not contain either directory empty.

A standalone export project uses exactly the same root domains:

```

/fo\core/
├── src/                                <- REQUIRED
│   ├── export.fol                  <- fixed standalone export selector; @co.dap.export(...)
│   ├── co/
│   │   ├── lang/
│   │   ├── meta/
│   │   └── out/
│   └── internal/
├── srclib/                           <- OPTIONAL; omit when unused; if present, non-empty
├── lib/                              <- OPTIONAL; omit when unused; if present, non-empty
└── build/                           <- compiler-managed; may be absent/empty before compilation

```

The package directories below `src/` are ordinary package contexts. `export.fol` selects which of those package contexts are emitted into the export-kind `.folenc`; it does not wrap them in a library surface.

Reserved-Domain Rules

`src/`, `src/lib/`, `lib/`, and `build/` are standardized project-root directory names. They are filesystem/compiler domains only; none is a package. Only `src/` must exist before compilation.

`src/` rules:

- `src/` is the only project directory that must exist before compilation and it cannot be empty.
- `src/app1.fo1` is the fixed application project surface.
- `src/library.fo1` is the fixed standalone packaged-library project surface.
- `src/export.fo1` is the fixed standalone export-project surface.
- exactly one of those three structural files must exist.
- `src/app1.fo1` classifies the project as an application; `src/library.fo1` classifies it as a standalone library; `src/export.fo1` classifies it as a standalone export project.
- more than one structural project surface present is a compile-time project-layout error.
- none of the structural project surfaces present is a compile-time project-layout error.
- no other file may occur directly in `src/`.
- all reusable project source must reside in package subdirectories below `src/`.
- a package directory physically present below `src/` must contain valid FoLang source; empty package directories are invalid.
- package names begin below `src/`; therefore `src/hr/employee/` is `hr.employee`, never `src.hr.employee`.

`src/lib/` rules:

- `src/lib/` is optional and must be omitted when the project uses no project-local source library, source export, or operator bootstrap.
- if `src/lib/` is physically present, it must contain at least one valid standardized child domain and cannot be empty.
- `src/lib/` itself is not a package or a library.
- only the standardized immediate child directories `application/`, `ffi/`, `system/`, `advanced/`, `dynamicvmrt/`, `exports/`, and `operators/` are permitted.
- because every standardized slot has one fixed directory name, a project can contain at most one project-local source library of each library kind and at most one project-local export domain.
- `application/`, `ffi/`, `system/`, `advanced/`, `dynamicvmrt/`, `exports/`, and `operators/` are source-domain roots, not packages, and their names do not become package namespace components.
- `src/lib/application/`, `src/lib/ffi/`, `src/lib/system/`, `src/lib/advanced/`, and `src/lib/dynamicvmrt/`, when present, must each contain exactly one direct file named `library.fo1`; every other entry must be an internal package directory.
- `library.fo1` is a reserved structural filename. It does not create a library named `library`; the fixed enclosing directory determines both the project-local source-library identity and its kind.
- no nested `library.fo1` is permitted. Nested source-library boundaries are forbidden.
- implementation package paths begin below the source-library root and may be arbitrarily deep. For example, `src/lib/ffi/native/marshal/` is internal package `native.marshal` of the FFI source library.
- source-library implementation packages never enter the owning project's ordinary open package index and cannot be bypass-imported as ordinary packages.
- `src/lib/exports/`, when present, must contain exactly one direct file named `export.fo1`; every other entry must be a package directory belonging to that project-local export domain.
- `src/lib/exports/export.fo1` contains only `@co.dap.export(packages={...})`. It selects package contexts from below `src/lib/exports/` for ordinary package access by the owning project's source graph. The `exports/` directory name is structural and never becomes part of a package dot path.
- packages selected by `src/lib/exports/export.fo1` participate as ordinary package contexts, not through a library API gate. They retain normal declaration visibility, type identity, extension, subtype, and member-resolution semantics.
- `src/lib/exports/` does not generate an independent `.fo1enc`; it is compiled as project-owned source. Its selected packages are available through ordinary `@co.ddap.import(package="...")` resolution.
- `src/lib/operators/`, when present, must contain exactly one file named `library.fo1`, with no additional files or subdirectories.
- `src/lib/operators/library.fo1` defines the project operator table for the owning project's open package compilation domain. Library source domains remain isolated from it. Operator metadata is never imported from a `.fo1enc`.
- for the library slots, `library.fo1` is the fixed structural surface filename; in the `operators/` slot it is parsed as the operator bootstrap surface rather than as an importable API surface; in the `exports/` slot the structural filename is `export.fo1`.
- `src/lib/`, `application/`, `ffi/`, `system/`, `advanced/`, `dynamicvmrt/`, `exports/`, `operators/`, `library.fo1`, and `export.fo1` create no package namespace.
- project-wide closed-project reachability is defined in [Unused Symbols, Liveness, and Reachability](#).

`lib/` rules:

- `lib/` is optional and must be omitted when the project has no packaged `.fo1enc` dependencies.

- if `lib/` is physically present, it must contain one or more valid compiled FoLang distributable artifacts named `*.folenc`; an empty `lib/` directory is a project-layout error.
- `.folenc` is the standard Protocol Buffers binary distributable format for standalone library and export projects. Its metadata records whether the artifact is a `library` or an `export`.
- a library-kind `.folenc` exposes only the projected `library.fol` surface API; an export-kind `.folenc` exposes the package contexts and symbols selected by `export.fol`.
- no alternate non-binary compiled distributable format is defined.
- `.fol` source files are invalid in `lib/`, and `lib/` never participates in source discovery.
- multiple separately built library and export artifacts may coexist in `lib/`; the one-per-slot rule applies only to standardized project-local source domains under `srcLib/`.
- packaged `.folenc` artifact liveness and partial-consumption rules are defined in [Unused Symbols, Liveness, and Reachability](#).

`build/` rules:

- `build/` is compiler-managed generated output and never participates in source discovery.
- `build/` need not exist before compilation; the frontend creates it when absent.
- an existing `build/` directory may be empty or non-empty without causing a project-layout error.
- the frontend emits its validated external representation under project-root `build/` using the wire format/version selected by the backend interchange contract.
- `.fol` source files placed under `build/` are invalid project layout.

Project-Domain Presence Invariant

FoLang distinguishes **absence** from an **empty declared domain**:

```
src/ absent          -> error
src/appl.fol only    -> application project
src/library.fol only  -> standalone library project
src/export.fol only   -> standalone export project
more than one project surface -> error
no project surface    -> error
src/ otherwise empty/invalid -> error

srcLib/ absent       -> valid
srcLib/ present but empty -> error

lib/ absent          -> valid
lib/ present but empty -> error

build/ absent        -> frontend creates it
build/ present and empty -> valid
build/ present and non-empty -> valid generated-output state
```

For developer-owned input/dependency domains, physical presence expresses intent to participate in the project. Therefore `srcLib/` and `lib/` must contain valid domain content when present. `build/` is exempt because it is compiler-owned generated state.

Structural Surface and Metadata Filenames

FoLang standardizes structural source positions so the project context is determined from the source domain rather than from arbitrary package placement:

Structural source	Valid location	Meaning	Does the containing directory define a package?
<code>appl.fol</code>	application project's <code>src/</code>	fixed application entry / single-source application surface	no; <code>src/</code> is not a package
<code>library.fol</code>	standalone packaged-library project's <code>src/</code>	fixed packaged-library API surface; kind declared explicitly by <code>@co.dap.library(type=...)</code>	no; <code>src/</code> is not a package
<code>export.fol</code>	standalone export project's <code>src/</code>	selects ordinary package contexts for the export-kind <code>.folenc</code>	no; <code>src/</code> is not a package
<code>library.fol</code>	<code>srcLib/application/</code> , <code>srcLib/ffi/</code> , <code>srcLib/system/</code> , <code>srcLib/advanced/</code> , <code>srcLib/dynamicvmrmt/</code>	project-local source-library API surface	no; each fixed <code><kind>/</code> directory is a library root, not a package
<code>export.fol</code>	<code>srcLib/exports/</code>	selects project-local open package contexts	no; <code>exports/</code> is a source-domain root, not a package

<code>library.fol</code>	<code>src/lib/operators/</code>	project-local operator bootstrap surface	no; <code>operators/</code> is not a package and permits no subdirectories
<code>package.fol</code>	inside an ordinary package directory	package metadata/aliasing	yes; the directory is already the package

The direct source surface under `src/` is both position- and filename-structural. `src/app1.fol` classifies the root as an application project, `src/library.fol` classifies it as a standalone packaged-library project, and `src/export.fol` classifies it as a standalone export project. None of these structural filenames creates a package or an `App1 / Library / Export` declaration.

`library.fol` is consistently used as the library-surface filename: at `src/library.fol` for a standalone library project and at the standardized project-local library slots under `src/lib/`. `export.fol` is consistently used as the package-export selector: at `src/export.fol` for a standalone export project and at `src/lib/exports/export.fol` for the project-local export domain. `package.fol` remains the fixed package metadata filename and does not create its enclosing package.

Package Aliasing

If there is a folder `/appl/src/hr/empl` and under that there is a `fol` file called `Employee.fol` then the import statement as we know will be

```
@co.ddap.import(package="hr.empl", as="emp")
```

where `as` is not a mandatory attribute

An import names a **package**, never a declaration inside it. The package is the folder, so `Employee.fol` under `/appl/src/hr/empl` belongs to package `hr.empl`; the file name is not part of the path. Once the package is imported, the declaration is reached as `emp.Employee`.

Now we want to change `empl` to `emp`, simple way is `change the folder name`, but we want to keep the `physical folder name` as is.

For example, `/appl/src/hr/empl` may be given the logical package name `hr.emp` while the physical folder remains `/appl/src/hr/empl`.

```
// /appl/src/hr/empl/package.fol
_ co.lang.package = {
  name: "emp"
};
```

`package.fol` is a reserved source filename. At most one may exist in a package folder, and its only top-level declaration must be `_ co.lang.package`. In this special source form, `_` occupies the declaration-name position but does not derive a package named `package` from the filename. The required `name` member supplies the logical name of the current package segment. Parent package segments are unchanged.

The import will be as below:

```
@co.ddap.import(package="hr.emp", as="emp")
```

UDT (User defined Data types)

FoLang provides the following user-defined data declaration kinds:

1. cstructs
2. structs
3. unions
4. enums
5. classes
6. modules
7. interfaces
8. signatures

Each ordinary file-backed primary declaration is placed in its own `<Name>.fol` file. The declaration name is never repeated in source; `_` occupies the declaration-name position and the compiler derives the name from the filename.

For canonical declaration-usage rules for structs, cstructs, unions, enums, classes, interfaces, and related primary declarations, see [Unused Symbols, Liveness, and Reachability](#).

For more information about UDTs, see [Built In Kinds](#).

Struct Declaration

```
// Employee.fol
_ co.lang.struct = {
  id   co.lang.int;
  name co.lang.string;
}
```

More about structs: [Structs in detail](#).

C-Struct Declaration

`co.lang.cstruct` is a C-like value type: it is passed by value, has a simple memory layout, and is safe to cross supported ABI boundaries.

```
// Point.fol
_ co.lang.cstruct = {
  x co.lang.int;
  y co.lang.int;
}
```

```
// Rect.fol
_ co.lang.cstruct = {
  origin Point;
  width  co.lang.int;
  height co.lang.int;
}
```

Enum Declaration

`co.lang.enum` declares an enumerated UDT whose body lists the permitted named variants.

```
// Status.fol
_ co.lang.enum = {
  Active,
  Inactive
}
```

Union Declaration

```
// NumberOrText.fol
_ co.lang.union = {
  intValue co.lang.int;
  strValue co.lang.string;
}
```

Class Declaration

```
// Employee.fol
_ co.lang.class = {
  getEmployeeDetails()->(Employee) = empmodule.getEmployeeDetails;

  getEmployeeInfo()->(Employee) =>> empmodule.getEmployeeDetails();
  // delegating – internally redirecting the call to module function
}

// $1, $2, $3 ... are return components of the immediately previous chained function
//Emp.fol
_ co.lang.class = {
  dosomething(a co.lang.int, b co.lang.int)->(co.lang.int)=>>somePack.someMethod(a)=>>someOthPack.someOtherMeth($1, b);
}
```

More about classes: [Classes in detail](#).

For the canonical class usage rule, including the distinction between developer-authored methods and interface-required methods, see [Unused Symbols, Liveness, and Reachability](#).

Interface vs Signature

```
// EmployeeContract.fol
_ co.lang.signature = {
  ...
}
```

More about signatures: [signatures in detail](#).

```
// EmployeeApi.fol
_ co.lang.interface = {
  ...
}
```

More about interfaces: [interfaces in detail](#).

Module Declaration

```
// Employee.fol
_ co.lang.struct = {
  ...
}
```

```
// EmployeeModule.fol
_ co.lang.signature = {
  ...
}
```

```
// EmployeeModImpl.fol
@co.dap.module(signature=EmployeeModule)
_ co.lang.module->(
  signature=EmployeeModule,
  matches=EmployeeModule
) = {
  ...
}
```

More about modules: [Modules in detail](#).

Extension Declarations

// EmployeeExtension.fol

```
_ co.lang.extension={
  someFun(a this)->()={
  }
}
```

An extension is a reusable collection of fully implemented functions. A class participating in the `@co.dap.oops` model may opt into the extension through its `uses` relationship.

Traits

// EmployeeTrait.fol

```

_ co.lang.trait={

    @co.dap.abstract
    someFunction()->();

    somerealFun()->()={
        co.out.println("Doing real work" );
    }

}

```

A trait is interface-like but may provide default function implementations. A trait carries no instance state.

A class participating in the `@co.dap.oops` model may include the trait through the applicable composition relationship.

A consuming class must implement every abstract function that remains unsatisfied.

Virtual methods are not permitted in a trait.

Mixins

//EmployeeMixin.fol

```

_ co.lang.mixin={

    someNum co.lang.int;

    someFun1()->()={
        co.out.println(someNum);
    }

    @co.dap.abstract
    someotherFun()->();

    @co.dap.virtual
    someVirtFun()-()={
        ...
    }

}

```

A mixin is the dedicated abstract-class-like composition form; it avoids declaring an ordinary class merely to mark it abstract.

A mixin may contain state, abstract methods, fully implemented methods, and virtual methods where the mixin rules permit them.

A class may incorporate a mixin through the applicable `extends=true` relationship, implement its abstract methods, and override its virtual methods.

Units

A unit is a stateless source container. A package may contain any number of ordinary unit files, and all their members are consolidated directly into the package namespace.

```

arithmetic.unit.fol
conversion.unit.fol
optional.unit.fol

```

Each ordinary unit file contains:

```

_ co.lang.unit = {
    ...
}

```

The filename is organizational only. It does not create a unit symbol or another qualification level. For package `math`, a function `abs` declared in any ordinary unit is accessed as `math.abs(...)`, not `math.Arithmetic.abs(...)`.

A companion unit uses the reserved filename form `<StructName>.comp.unit.fol`. Its members attach to the matching same-package struct:

```
Employee.fol
Employee.comp.unit.fol
```

Only `co.lang.struct` supports a companion unit. See [Units in detail](#) and [Struct Companion Units](#).

Matchers

Custom Matcher

```
// PositiveEvenMatcher.fol
@co.dap.matcher
_ co.lang.matcher->(type=co.lang.int) = {
  matchCase(
    value co.lang.int,
    pattern co.lang.untyped
  )->(co.lang.int, co.lang.MatchBindings) = {
    // user logic
    // 0 = no match, >0 = match
  }
}
```

A matcher declaration supports exactly one matched-subject type, specified by `type=`. It must declare exactly one protocol function named `matchCase`; overloading `matchCase` inside a matcher is not permitted. A matcher for another subject type must be declared as a separate `<Name>.fol` primary declaration.

At compile time, the compiler resolves the type named by `type=` and the type of the first `matchCase` parameter. The two resolved types must be equivalent. Type aliases are compared after alias resolution. The second parameter must have type `co.lang.untyped`, and the result must be exactly `(co.lang.int, co.lang.MatchBindings)`.

```
// InvalidMatcher.fol
_ co.lang.matcher->(type=co.lang.int) = {
  matchCase(
    value co.lang.string,
    pattern co.lang.untyped
  )->(co.lang.int, co.lang.MatchBindings) = {
    ...
  }
  // compiler error: declared matcher type and first parameter type differ
}
```

The parameter names themselves are not significant; their order and resolved types define the protocol shape.

Matcher liveness is defined in [Unused Symbols, Liveness, and Reachability](#).

Comprehensions

//comprehensionseg1.unit.fol

```
_ co.lang.unit = {
  someFun()->>() = {
    k := (1 .. 10).filter(|x| => x % 2 == 0).map(|x| => x * x);

    result := for (x <- List[1,2,3]).yield(x * 2);           // List[2, 4, 6]
    result := for (x <- Set(1,2,3)).yield(x * 2);           // Set(2, 4, 6)
    result := for (x <- Some(5)).yield(x * 2);              // Some(10)
    result := for (x <- fetchData()).yield(x.process());   // Future

    ages := Map{"A":30,"B":40,"C":66,"E":88};
    upper := for ((name, age) <- ages).yield(name.toUpperCase, age);
  }
}
```

Comprehension Semantics

A FoLang comprehension is a source-driven transformation. Its canonical form is:

```
result := for (pattern <- source).yield(resultExpression);
```

The core comprehension syntax defines the binding and transformation structure; the source type defines the source-specific comprehension behaviour. The syntax does not implicitly convert every source to a `List`, and an arbitrary value does not become a valid comprehension source merely because it appears to the right of `<-`. A source is valid only when it is a FoLang iterable, `Some(T)`, or `Future(T)`. No other non-iterable source category can acquire comprehension capability through extensions, ordinary package APIs, operators, or similarly named methods.

Permitted Comprehension Sources

FoLang comprehensions intentionally accept only the following source categories:

1. **Any FoLang iterable**, including standard iterable forms such as arrays, lists, sets, maps/dictionaries, and ranges.
2. `Some(T)`.
3. `Future(T)`.

This set is closed. Comprehension support is **not** an open extension mechanism. No other non-iterable source category participates in `for ... yield`.

For iterable sources, the comprehension consumes the values exposed by the source's ordinary iteration semantics. The precise traversal order, entry shape, duplicate behaviour, and result-container behaviour remain those defined for that iterable type.

`Some(T)` and `Future(T)` are the only permitted non-iterable comprehension sources. They do not become iterable; instead, `for ... yield` applies their language-defined transformation semantics.

For example:

```
for (x <- List[1,2,3]).yield(x * 2); // valid: iterable
for (x <- 1 .. 10).yield(x * 2);    // valid: iterable range
for ((k, v) <- valuesMap).yield(k, v); // valid: iterable map/dictionary
for (x <- Some(5)).yield(x * 2);    // valid: permitted non-iterable source
for (x <- someFuture).yield(f(x)); // valid: permitted non-iterable source
```

A struct, class, module, or other UDT that is **not iterable** is not a valid comprehension source. Defining `map`, `each`, extensions, operators, or similarly named functions does not make such a type comprehension-capable.

```
emp Employee = ...;

for (x <- emp).yield(x.salary + 1000); // compiler error if Employee is not iterable
```

A user-defined type may therefore appear as a comprehension source only when it participates in FoLang's ordinary iterable model. It cannot define a new non-iterable comprehension meaning comparable to `Some(T)` or `Future(T)`.

The parts of the form have the following meanings:

- `source` is the value being traversed or transformed;
- `pattern` introduces the comprehension-local binding or destructuring pattern for each value produced by the source;
- `<-` associates that binding pattern with the source;
- `.yield(...)` supplies the result expression evaluated for each participating source value according to that source type's comprehension semantics.

Bindings introduced by `pattern` are local to that comprehension. They are visible in the corresponding `yield` expression and do not introduce names into the enclosing scope. A destructuring pattern, such as `(name, age)`, introduces each successful component binding in the same comprehension-local scope.

The result shape is source-defined rather than selected by a universal core-language conversion rule. The examples above establish the following current forms:

```
List(A)  --yield B--> List(B)
Set(A)   --yield B--> Set(B)
Some(A)  --yield B--> Some(B)
Future(A) --yield B--> Future(B)
```

For example:

```

result := for (x <- List[1,2,3]).yield(x * 2);
// List[2, 4, 6]

result := for (x <- Set(1,2,3)).yield(x * 2);
// Set(2, 4, 6)

result := for (x <- Some(5)).yield(x * 2);
// Some(10)

result := for (x <- fetchData()).yield(x.process());
// Future

```

The `Map` form demonstrates source destructuring and pair production:

```

ages := Map{"A":30,"B":40,"C":66,"E":88};
upper := for ((name, age) <- ages).yield(name.toUpperCase, age);

```

Here `((name, age))` destructures the source entry for the current comprehension step. The result-container rules, duplicate-key behaviour, ordering, and other map-specific properties are those defined by the applicable `Map` API rather than by the core `for ... yield` grammar.

Source-specific semantics also govern cardinality, emptiness, deferred execution, failure propagation, ordering, duplicate handling, and similar behaviour. Thus the `Some` and `Future` forms preserve the source abstraction shown by their examples, while the precise empty/failure/scheduling behaviour belongs to the corresponding type's defined API semantics. The comprehension syntax itself does not introduce implicit blocking, concurrency, retries, or error suppression.

The current core `for (pattern <- source).yield(...)` form does not define an inline filter clause. Filtering is expressed separately through the source's supported operations, for example:

```

k := (1 .. 10)
  .filter(|x| => x % 2 == 0)
  .map(|x| => x * x);

```

If a later source-specific comprehension form defines filtering, its filter conditions are evaluated before its result expression as required by the general evaluation-order rules. No filter is implied by the core `for ... yield` syntax shown above.

A comprehension does not imply parallel or concurrent traversal. Execution is sequential unless the selected source semantics or an explicitly requested FoLang execution model states otherwise.

Extensions

Extension functions may be declared in an ordinary package unit:

```

// string_extension.unit.fol
_ co.lang.unit = {

  @co.dap.extension(fortype=co.lang.string, what=extends)
  upperCase()->(co.lang.string) = {
    this.return this.upper();
  }

  @co.dap.extension(fortype=[co.lang.string], what=overrides)
  equals(str co.lang.string)->(co.lang.bool) = {
    this.return this == str;
  }
}

```

Because ordinary unit filenames create no symbol, extension activation identifies the contributing package, not the unit file.

For the current package, omit `from`:

```

@co.ddap.use(methods=[equals, upperCase])
k.upperCase();

```

For another package, use its alias or complete package path:

```

@co.ddap.import(package="text.util", as="tu")

@co.ddap.use(from="tu", methods=[upperCase])
@co.ddap.use(from="text.util", methods=[upperCase])

```

See [Activating Instance Methods](#) for activation of typeclass instances.

Reflections

```
@co.dap.reflection(enable=True, package="co.meta")

x co.lang.int = 10;
x.reflect().getType();    //co.lang.int
x.reflect().getValue();   //10;
x.reflect().getKind();    // value
```

Type Classes

Monads, Applicatives, Functors, Monoids and Transformers

`@co.dap.typeclass(kind=...)` is the single annotation for all typeclass definitions. `kind` specifies the algebraic structure — `Functor`, `Applicative`, `Monad`, `Monoid`, `Transformer`, or any user-defined kind. Instances of any typeclass always use `co.lang.instance`.

Typeclass and instance liveness is defined in [Unused Symbols, Liveness, and Reachability](#).

In a file-backed typeclass declaration, `_` is the filename-derived declaration-name placeholder and the following parenthesized clause declares the typeclass parameters. They are separate grammar components, so the canonical spelling includes a space: `_ (F(_))`, not `_(F(_))`. A parameter such as `T` denotes an ordinary type, while `F(_)` denotes a unary type constructor and `G(_, _)` denotes a binary type constructor. Otherwise-unbound type variables introduced in an operation signature, such as `A` and `B`, are implicitly universally quantified within that operation.

Functor

```
//Functor.fol
@co.dap.typeclass(kind=Functor)
_ (F(_)) co.lang.typeclass = {
  map(value F(A), f (A)->B) -> (F(B));
}

// ListFunctor.fol
_ co.lang.instance->(for=Functor, type=List) = {
  map(value List(A), f (A)->B)->(List(B)) = {
    result = List(B){};
    value.each(_, item, { result.append(f(item)) });
    this.return result;
  }
}
```

Applicative

```
//Applicative.fol
@co.dap.typeclass(kind=Applicative)
_ (F(_)) co.lang.typeclass = {
  pure(x A) -> (F(A));
  apply(fab F(A->B), fa F(A)) -> (F(B));
}

// OptionApplicative.fol
_ co.lang.instance->(for=Applicative, type=Option) = {
  pure(x A)->(Option(A)) = { this.return Some(x); }
  apply(fab Option(A->B), fa Option(A))->(Option(B)) = {
    this.return (fab, fa)
      .match
      .case((Some(f), Some(x)) => Some(f(x)))
      .default(None());
  }
}
```

Monad

```
//Monad.fol
@co.dap.typeclass(kind=Monad)
_ (F(_)) co.lang.typeclass = {
  pure(x A) -> (F(A));
  flatMap(fa F(A), f (A)->F(B)) -> (F(B));
}

// OptionMonad.fol
_ co.lang.instance->(for=Monad, type=Option) = {
  pure(x A)->(Option(A)) = { this.return Some(x); }
  flatMap(fa Option(A), f (A)->Option(B))->(Option(B)) = {
    this.return fa.match().case(Some(x) => f(x)).default(None);
  }
}
```

Monoid

```
//Monoid.fol
@co.dap.typeclass(kind=Monoid)
_ (T) co.lang.typeclass = {
  empty() -> (T);
  combine(a T, b T) -> (T);
}

// IntMonoid.fol
_ co.lang.instance->(for=Monoid, type=co.lang.int) = {
  empty()->(co.lang.int) = { this.return 0; }
  combine(a co.lang.int, b co.lang.int)->(co.lang.int) = { this.return a + b; }
}
```

Transformer

```
//Transformer.fol
@co.dap.typeclass(kind=Transformer)
_ (F(_), G(_)) co.lang.typeclass = {
  map(value F(A), f (A)->B) -> (G(B));
}

// ListToSetTransformer.fol
_ co.lang.instance->(for=Transformer, types=[List, Set]) = {
  map(value List(A), f (A)->B)->(Set(B)) = {
    result = Set(B){};
    value.each(_, item, { result.insert(f(item)) });
    this.return result;
  }
}
```

Using an Instance

An instance is selected **by name**. There is no implicit search.

```
@co.ddap.import(package="abc.tc", as="tc")

xs List(co.lang.int) = [1, 2, 3];
double(x co.lang.int)->(co.lang.int) = { this.return x * 2; }

ys := tc.ListFunctor.map(xs, double);
```

`map` takes the container as its first argument, exactly as the typeclass declares it. The call names `ListFunctor`, so the compiler resolves it the same way it resolves any other imported declaration.

FoLang does **not** search visible packages for an instance that happens to match a type. A typeclass is a contract, an instance is a named implementation of that contract, and the caller names the one it means. Nothing is inferred, so an unresolved call reports a name the developer actually wrote rather than a failed search they never saw.

Method syntax such as `xs.map(f)` is unrelated to typeclasses. It calls an associated function declared in a companion unit for the type. Companion units live in the type's own package, so method calls are unambiguous by construction. A companion function and a typeclass instance may both be named `map`; they are different declarations and are reached differently.

Activating Instance Methods

An instance may also be **activated**, which makes its functions callable as methods on the receiver. Activation uses the same directive that activates extensions.

```
@co.ddap.import(package="abc.tc", as="tc")
@co.ddap.use(from="tc.ListFunctor", methods=[map, reduce])

ys := xs.map(double);           // resolves to tc.ListFunctor.map(xs, double)
```

Activation is explicit and block-scoped. Importing a package does **not** activate anything in it, so adding an import can never change how an existing call resolves.

methods and from

`methods` selects the functions to activate.

For extension functions contributed by ordinary package units:

- omit `from` to select the current package
- use an imported package alias to select another package
- use the complete package path when no alias exists

```
@co.ddap.use(methods=[upperCase])           // current package
@co.ddap.use(from="tu", methods=[upperCase]) // imported package alias
@co.ddap.use(from="text.util", methods=[upperCase]) // complete package path
```

Ordinary unit filenames are not accepted by `from`, because they do not create symbols.

For a typeclass instance, `from` continues to name the instance declaration:

```
@co.ddap.use(from="tc.ListFunctor", methods=[map, reduce])
```

Listing names is optional. Omit `methods` to activate every eligible method from the selected package or instance; provide it to activate a subset. Conflict detection remains receiver-aware and block-scoped.

How a method call resolves

For `xs.map(f)`, where `xs` has type `List(A)`:

1. a class method or companion-unit function on `List`
2. an activated extension for `List`
3. an activated instance function whose typeclass declares `map` with the receiver as its first parameter
4. otherwise, an error

The first match wins. A type's own declarations therefore always take precedence over anything activated into scope, and no activation can silently replace behaviour the type already defines.

Within one scope a given method name may be activated at most once for a given receiver type. Activating `map` for `List` from two sources is an error at the second `@co.ddap.use`, which names both. The conflict is reported where the activation is written, never at a distant call site.

The parser's built-in-method classification is only a lookup candidate. If the frontend recognizes a control-chain shape before receiver-aware lookup, it must retain the complete original member/call chain. A receiver-owned, companion, extension, or activated-instance declaration that wins the order above restores/remains the ordinary method call represented by that chain. An early dedicated control-flow node is therefore only a lowering candidate and becomes final only when the built-in meaning wins.

Activation never affects generic code. A function polymorphic over a typeclass takes the instance as an ordinary parameter and calls it by name.

A Typeclass Is a Type

A typeclass may be used wherever a type is expected. Its values are its instances. This is the same relationship a signature has to a module:

```
mm EmployeeModule = EmployeeModImpl;    // signature as type, module as value
f Functor(List) = tc.ListFunctor;      // typeclass as type, instance as value
```

An instance is therefore an ordinary first-class value. It can be held in a variable, stored in a collection, returned from a function, and chosen at run time.

```
inst := (useCache).then(cachedFunctor).default(plainFunctor);
result := inst.map(xs, transform);
```

Nothing about an instance is special to the compiler, which is why FoLang needs no instance resolution algorithm and no coherence rules.

Writing Code Generic Over a Typeclass

Activation makes `xs.map(f)` work for a **known** container. Code that must work for **any** container cannot activate anything, because the container type is not known until the caller supplies it. Such code takes the instance as an ordinary parameter.

```
@co.dap.generic(type={F:{typename}, A:{typename}, B:{typename}})
mapAll(inst Functor(F), value F(A), fn (A)->B)->(F(B)) = {
  this.return inst.map(value, fn);
}
```

Parameter	What it is
<code>F</code>	the container kind — <code>List</code> , <code>Option</code> , <code>Tree</code>
<code>inst</code>	the instance, an implementation of <code>Functor</code> for <code>F</code>
<code>value</code>	the container itself, an ordinary value

The caller supplies the instance:

```
ys := mapAll(tc.ListFunctor, xs, double);
opt2 := mapAll(tc.OptionFunctor, opt, double);
```

The function never learns what `F` is. It knows only that `inst` provides `map`, which is the whole contract. One definition therefore serves every container that has a `Functor` instance.

A type is never “a Functor” in FoLang. `List` does not become a Functor; an instance implements Functor operations for `List`, and the list stays a plain list. The Functor-ness lives entirely in `inst`.

When a wrapper is worth writing

`mapAll` above forwards directly to `inst.map`, so it adds nothing a caller could not write themselves. A wrapper earns its place when it does more than forward — when it combines several typeclass operations, adds logic, or fixes some parameters and leaves others open.

```
@co.dap.generic(type={F:{typename}})
doubleAll(inst Functor(F), value F(co.lang.int))->(F(co.lang.int)) = {
  this.return inst.map(value, (x co.lang.int)->(co.lang.int){
    this.return x * 2;
  });
}
```

This one fixes the element type and the operation, so it says something a bare `inst.map` call does not.

Where an Instance Is Declared

An instance is declared in **the package that defines the typeclass**, or in **the package that defines the type**. That exact package, not a sub-package.

```
abc.tc.ListFunctor    for=Functor, type=List    // OK  typeclass's package
myapp.ab.TreeFunctor  for=Functor, type=myapp.ab.Tree // OK  type's package
other.util.ListFunctor for=Functor, type=List    // ERR neither is theirs
```

A typeclass is an ordinary declaration and may live in any package; `abc.tc` above is a user package, not a built-in one. Sub-packages are distinct packages, so an instance for `myapp.ab.Tree` belongs in `myapp.ab`, not in `myapp` or `myapp.ab.instances`. This matches the rule for companion units, which also sit in their type's own package. Because a package spans every `.fo1` file in its folder, each instance still gets its own `<InstanceName>.fo1` file and uses `_ co.lang.instance` in source.

The rule is permissive on both sides on purpose. Requiring the typeclass's package alone would make a typeclass usable only by its own author, since nobody could implement it for their own types. Requiring the type's package alone would stop a library from shipping instances for common built-in types.

For library authors. If you define a type and want it usable with someone else's typeclass, declare the instance in your package. If you define a typeclass, you may ship instances for types you do not own, including built-in ones — `IntMonoid` above sits beside `Monoid`, which is exactly this case. If you need an instance for a typeclass and a type you both do not own, wrap the type in a `co.lang.newtype` you do own and declare the instance for the wrapper.

This placement rule is semantic, not syntactic. A misplaced instance parses correctly and is reported during name resolution, so the diagnostic can name the typeclass, the type, and the two packages in which the instance would have been legal.

Labels and Named Blocks

```
// Label
outer:{
    // statements
}

//Blocks Anonymous block
{
    // statements
}

// Named Block
labelBlock co.lang.block={

}

labelBlock.expand();
```

Blocks have their own variable scope and context. A variable declared outside a block remains accessible inside the block unless it is shadowed. A block may declare a variable with the same name and either the same or a different type; that declaration shadows the outer binding and is scoped only to the block. This model is similar to block scoping in C and C++.

Blocks cannot appear outside functions or methods.

Inner executable blocks are prohibited directly inside classes, structs, typeclasses, modules, and other non-function/non-method declarations; such usage is a compile-time error.

```
somefun (a co.lang.int, b co.lang.int)->(co.lang.int)={

    some_other co.lang.float = 20.1f;
    {
        some_other co.lang.char = 'c';
        co.out.println( some_other);    //prints c
    }

    co.out.println(some_other); //prints 20.1;

    {
        some_other co.lang.float = 11.1f;
        co.out.println(some_other); // prints 11.1f
    }

    co.out.println(some_other); // still prints 20.1f;

    {
        some_other = some_other + 1.1f;
        co.out.println(some_other); // prints 21.2
    }

    co.out.println(some_other); // prints 21.2; as this was changed in third block
}
```

imports

FoLang supports package imports and library imports. User packages, exported package artifacts, and libraries must be imported before use. When an import declares `as=`, symbols are accessed through that alias. When `as=` is omitted, the complete imported package path or library identity is used.

An import declaration makes the target context available for name resolution; import liveness and zero-use validation are defined in [Unused Symbols, Liveness, and Reachability](#).

1. Package Import

Use this for ordinary package contexts. A package may come from the current project's `src/` tree, from the project-local `src/lib/exports/` domain, or from an export-kind `.fo1enc` artifact in `lib/`.

```
@co.ddap.import(package="hr.employee", as="emp")

e := emp.getEmployee(1);
co.out.println(e.name);
```

Conceptual resolution:

```
package="hr.employee"
-> exactly one matching package context from:
    <project-root>/src/
    <project-root>/src/lib/exports/ selected packages
    export-kind lib/*.folenc package projections
```

The package path must resolve uniquely in the applicable package index. Multiple providers of the same imported package path are an ambiguity error. The distributable artifact name never becomes an implicit package prefix.

2. Source Library Import

Use this for same-owner project-local libraries whose source is available under the reserved `src/lib/` root.

```
@co.ddap.import(library="ffi", src-library=true, as="ffilib")
```

Resolution:

```
library="ffi", src-library=true
-> <project-root>/src/lib/ffi/library.fol
```

For `src-library=true`, the `library` value selects one standardized project-local source-library slot under `src/lib/`. The permitted values are `application`, `ffi`, `system`, `advanced`, and `dynamicvmrt`. These names are fixed by FoLang and identify both the source library and its library kind. `src/lib/exports/` is not a library slot and is therefore never selected with `library=`; its selected packages use ordinary `package=` imports.

The resolved file must be the fixed source-library surface file:

```
// <project-root>/src/lib/ffi/library.fol
_ co.lang.library={

}
```

Meaning:

- `library` selects one of the standardized project-local source-library identities (`application`, `ffi`, `system`, `advanced`, or `dynamicvmrt`) under `src/lib/`
- `src-library=true` switches `library=` resolution from the packaged-library domain `lib/` to the project-local source-library domain `src/lib/`
- `library.fol` is the mandatory and only direct source file at that source-library root
- the enclosing fixed directory determines the library kind; a `@co.dap.library(type=...)` annotation is neither required nor permitted on a project-local source-library surface
- only the projected surface API is visible to the owning project's `src/` domain
- source in the owning project's `src/` domain can import the source library only through this `library.fol` surface; it cannot directly import any implementation package beneath the source-library root
- implementation packages beneath the source-library root are private to that source-library compilation domain and cannot be imported through ordinary project-package imports
- inside the source-library compilation domain, `library.fol` and other internal packages may import internal packages according to ordinary private package rules, but those imports are still provisional and become live only when at least one symbol from the imported package is actually used

3. Packaged Library Import

Use this for third-party or prebuilt `library-kind .folenc` artifacts.

```
@co.ddap.import(library="hrlib", as="hr")
@co.ddap.import(library="paylib", as="pay")
```

Resolution:

```
library="hrlib" -> <project-root>/lib/hrlib.folenc
```

Only the packaged library's projected surface API is visible to the consumer. An export-kind `.fo1enc` is not imported with `library=`; its exported package contexts are selected with ordinary `package=` imports.

Import Directive Fields

Field	Required	Default	Meaning
<code>package</code>	one of <code>package</code> or <code>library</code>	—	logical package dot path resolved from the applicable package index, including selected project-local exports and export-kind <code>.fo1enc</code> package contexts
<code>library</code>	one of <code>package</code> or <code>library</code>	—	logical library identity; resolves a library-kind <code>.fo1enc</code> from <code>lib/</code> by default or a standardized project-local source library when <code>src-library=true</code>
<code>src-library</code>	✗	<code>false</code>	valid only with <code>library=</code> ; when <code>true</code> , selects a project-local source library (<code>application</code> , <code>ffi</code> , <code>system</code> , <code>advanced</code> , or <code>dynamicvmrt</code>) and resolves its fixed <code>library.fo1</code> surface under <code>srclib/</code>
<code>as</code>	✗	<code>none</code>	local alias; valid FoLang identifier; when omitted, the complete imported package path or library identity is used

Notes:

- `package` and `library` are mutually exclusive; exactly one must be supplied
- `src-library=true` is valid only when `library` is supplied
- `as` is optional — when omitted, no short alias is created and the complete imported package path or library identity is used to access symbols
- dots are not allowed in `as`
- for a project-local source library, the fixed `srclib/<kind>/` path is the source of truth for library identity and kind
- `@co.dap.library(type=...)` is not used on project-local `library.fo1` surfaces
- `srclib/exports/` contributes selected open package contexts and is never addressed through `library=` or `src-library=true`
- packaged `.fo1enc` artifacts carry an artifact kind; library-kind artifacts retain library-kind metadata, while export-kind artifacts retain their selected package/context projection

Examples:

```
@co.ddap.import(package="hr.employee", as="emp")
em emp.Employee;

@co.ddap.import(package="hr.employee")
em hr.employee.Employee;
```

Valid `as` Values

<code>as="hr"</code>	✓
<code>as="v1_hr"</code>	✓
<code>as="v1.hr"</code>	✗
<code>as="123hr"</code>	✗

Cycles

Compiler error if any cycle exists through:

- package imports

Examples:

- `packageA` imports `packageB`, and `packageB` imports `packageA`

Symbol Resolution

Resolution Order

For symbols without a prefix:

1. current package symbol table
2. parent package chain upward
3. error if not found

For symbols with an alias prefix:

1. current package imported-context map
2. parent package chain imported-context maps
3. error if not found

For imports declared **without** `as`:

1. use the full imported package path directly
2. resolve it through imported-context lookup using that full path
3. error if not found

Example Context Graph

```
pp1 (grandparent package)
├─ SymbolTable
├─ ImportedContexts
└─ ChildCtxs

p1 (parent package)
├─ Parent -> &pp1
├─ SymbolTable
├─ ImportedContexts
└─ ChildCtxs

c1 (current package)
├─ Parent -> &p1
├─ SymbolTable
├─ ImportedContexts
└─ ChildCtxs
```

Resolution Examples

```
lookup "OwnType"
  -> current symbol table

lookup "ParentType"
  -> parent chain

lookup "hr.Employee"
  -> imported-context lookup for alias "hr"

lookup "hr.employee.Employee"
  -> imported-context lookup for full imported package path when no alias exists

lookup "unknown.Type"
  -> imported-context lookup fails -> compiler error
```

Short Summary

- FoLang projects standardize the root domains `src/`, `src/lib/`, `lib/`, and `build/`; none is a package; only `src/` is mandatory before compilation
- project kind is determined structurally from `src/`: `src/appl.fo1` means application, `src/library.fo1` means standalone packaged library, and `src/export.fo1` means standalone export; exactly one must exist and no other direct file is permitted
- project identity is derived from the canonical project-root directory basename; `appl.fo1`, `library.fo1`, and `export.fo1` are structural filenames only and never determine the output name
- package dot paths begin below `src/`; export artifacts preserve selected package paths rather than prefixing them with the artifact name
- `src/lib/` permits only `application/`, `ffi/`, `system/`, `advanced/`, `dynamicvmrt/`, `exports/`, and `operators/` as immediate children
- the project-local library roots are not packages and each contains exactly one direct file named `library.fo1`; their remaining source lives in arbitrarily deep internal package subdirectories

- the fixed `src/lib/<kind>/` directory determines the project-local source-library identity and kind; no `@co.dap.library(type=...)` annotation is used there
- `src/lib/exports/` contains exactly one direct `export.fol`; selected descendants participate as ordinary package contexts for the owning project
- `src/lib/operators/` contains exactly one `library.fol`, has no subdirectories, and creates no package namespace
- project-local source-library internals never enter the owning project's open package index; project-local export packages enter it only when selected by `src/lib/exports/export.fol`
- at most one project-local source domain exists for each standardized `src/lib/` slot; independently built library and export artifacts are consumed from `lib/`
- when present, `lib/` contains one or more Protocol Buffers binary `.folenc` distributable artifacts; each artifact records whether it is a library or export
- `src/lib/` and `lib/` are optional by absence but invalid when present and empty; `build/` is compiler-managed, may be absent or empty before compilation, and carries the backend-contract-selected frontend artifact
- each ordinary package source file has exactly one primary top-level declaration
- package functions, templates, macros, and non-UDT type declarations must be enclosed in ordinary `*.unit.fol` files
- all ordinary unit members are consolidated directly into the package namespace
- struct companion behavior must be declared in `<StructName>.comp.unit.fol`
- `src/appl.fol`, `src/library.fol`, and `src/export.fol` are the fixed non-package project surfaces for application, library, and export projects respectively
- `co.*` is always available and never imported
- `@co.ddap.import(package="...")` imports an ordinary package context from project source, project-local exports, or an export-kind `.folenc`
- `@co.ddap.import(library="application|ffi|system|advanced|dynamicvmrt", src-library=true, ...)` imports the corresponding project-local source library through its fixed `library.fol` surface
- `@co.ddap.import(library="...")` imports a library-kind packaged `.folenc` from `lib/`
- imports reference canonical prepared contexts/symbol tables through `ImportedContexts`; imported symbol tables are never copied
- an import is provisional until at least one symbol is actually resolved through it; zero-use imports are pruned from the effective dependency graph and reported as errors
- source in the owning project's `src/` domain can access a project-local source library only through its `library.fol` projected surface; source-library internal packages cannot be directly imported from the owning project's open package domain
- selected `src/lib/exports/` packages are ordinary package contexts and therefore do not use a library gate
- canonical unused-symbol, liveness, reachability, project-local `src/lib`, and packaged `.folenc` rules are defined in [Unused Symbols, Liveness, and Reachability](#)
- every ordinary library surface exports only permitted boundary contracts and public function signatures; implementation packages remain hidden, while export projects expose selected ordinary package contexts

Let and Function Patterns

Bare Function-Pattern Group

A bare function-pattern group does not capture surrounding runtime bindings:

```
classify(0) => { this.return "zero"; }
classify(n).where(n > 0) => { this.return "positive"; }
classify(_) => { this.return "negative"; }
```

The formal clause shape is:

```
name(patterns) [.where(guard)] => result
```

`=>` is mandatory. `=>>` belongs to ordinary function delegation, `==>>` belongs to closure/curry expression and are not accepted here.

A bare group may use:

- its parameters and names introduced by its patterns
- its own name for recursion
- entry-local type names
- `co.*` APIs
- imported package and library APIs
- compile-time symbols that do not require runtime capture

It may not reference an entry-file runtime variable from the surrounding context:

```
offset := 100;

adjust(n) => { this.return n + offset; }
// compiler error: bare function-pattern groups cannot capture `offset`
```

Capturing `let` Function-Pattern Group

The `let` form exists only to declare an entry-local function-pattern group that captures one or more surrounding entry-file runtime bindings:

```
offset := 100;

let adjust(θ) = offset;
let adjust(n) = n + offset;

result := adjust(10);
```

The captured names must resolve to surrounding runtime bindings that are already declared and definitely initialized before the first clause of the group. Built-in names, imported declarations, type names, parameters, and the function's own recursive name are not captures.

A `let` function-pattern group must capture at least one surrounding runtime binding. When no capture is required, the bare form must be used:

```
let fib(θ) = 1;
let fib(1) = 1;
let fib(n) = fib(n - 1) + fib(n - 2);
// compiler error: this group captures nothing; remove `let`

fib(θ) => { this.return 1; }
fib(1) => { this.return 1; }
fib(n) => { this.return fib(n - 1) + fib(n - 2); }
```

The `let` marker is therefore not an optional spelling of a bare function-pattern group. It explicitly requests restricted lexical capture.

Similarities

Bare and capturing `let` function-pattern groups both:

- group compatible clauses by function name and arity
- dispatch by parameter patterns and optional `.where(...)` guards
- evaluate clauses in normal pattern-selection order
- may call themselves recursively
- must maintain compatible parameter and result types across all clauses
- follow the normal pattern ordering, reachability, overlap, and exhaustiveness rules
- are private to the entry file
- cannot be imported, exported, or annotated with package visibility
- cannot be converted to, assigned as, passed as, or returned as function values
- cannot be partially applied or curried

Clause Syntax and Termination

Both forms accept the same pattern list and optional guard:

```
classify(θ) => "zero";
classify(n).where(n > 0) => "positive";
classify(_) => { this.return "negative"; }

offset := 10;
let adjust(θ) = offset;
let adjust(n).where(n > 0) = n + offset;
let adjust(_) = { this.return offset; }
```

Annotations may precede either clause form and are retained on the clause for later checking. Because every function-pattern group is private to the entry file, package visibility/export annotations are invalid even though other entry-local metadata annotations may be used.

An expression-bodied clause is a simple statement and must end with `;`. A block-bodied clause ends at its closing `}` and must not be followed by `;`. Newlines never terminate clauses.

Supported parameter patterns are:

- `_` wildcard patterns;
- literals, including signed numeric literals;
- binding names;

- qualified identity names;
- constructor patterns such as `Some(value)`;
- record patterns such as `Employee{id: value, name: _}`;
- tuple patterns with at least two elements, such as `(left, right)`.

`.where(expression)` is evaluated only after the clause's parameter patterns match. Names bound by those patterns are available to the guard and result. The guard expression must have type `co.lang.bool`.

Clauses are considered in source order. A clause is eligible when its arity matches, all parameter patterns match, and its optional guard evaluates to `co.const.true`. The compiler rejects incompatible arities or result types, unreachable clauses, invalid overlaps, and non-exhaustive groups where exhaustiveness is required by the declared result contract.

Differences

Form	Surrounding runtime capture	Intended use
<code>name(pattern) => result</code>	No	Entry-local pattern dispatch that depends only on its arguments, recursion, built-ins, imports, and compile-time names
<code>let name(pattern) = result</code>	Yes, at least one capture required	Entry-local pattern dispatch that also depends on existing entry-file runtime bindings

Clauses of the same name and arity must all use the same form. Mixing bare and `let` clauses in one function-pattern group is a compiler error.

A capturing `let` function-pattern group is still not a general closure facility. It is named, entry-local, non-first-class, and non-escaping. The compiler may internally retain a capture environment, but FoLang does not expose that environment as a closure object.

`let` cannot be used directly in the entry file for value bindings, anonymous functions, general closure expressions, or curried functions:

```
let base = 10;           // compiler error: entry-file let value binding
let result = base + 1;   // compiler error: entry-file let value binding
let add = (a, b) => a + b; // compiler error: anonymous function
let counter = closure { ... }; // compiler error: general closure expression
let add = a => b => a + b; // compiler error: curried function
```

The compiler may lower either function-pattern form to a private entry helper, but neither source construct is a general-purpose function declaration.

Package in detail

Package Identity

Ordinary project packages exist in subdirectories below the reserved `src/` domain. `src/` itself is not a package and contains exactly one direct project surface file—`appl.fol`, `library.fol`, or `export.fol` according to project kind. No other direct file is permitted there.

- package dot paths start below `src/`
- the project root and `src/` are not packages
- `src/appl.fol`, `src/library.fol`, and `src/export.fol` are the fixed structural surfaces for application, standalone library, and standalone export projects
- `src/lib/`, `lib/`, and `build/` are separate reserved domains and are excluded from direct `src/` package discovery
- source-library roots under `src/lib/` are not packages; their descendant implementation directories define private packages within that source-library compilation domain
- `src/lib/exports/` is also not a package; package paths begin below it, and only the package contexts selected by its `export.fol` enter the owning project's applicable open package index

Examples:

```
/appl/src/hr/           -> package "hr"
/appl/src/hr/employee/ -> package "hr.employee"
/appl/src/auth/         -> package "auth"
```

Multi-File Packages

Multiple `.fol` files in the same package folder below `src/` automatically belong to the same package:

```
src/hr/employee/
├─ Employee.fol      -> hr.employee
├─ EmpService.fol    -> hr.employee
└─ EmpValidator.fol  -> hr.employee
```

Application Project Layout

See the canonical [Application Project Layout](#). The key ownership boundaries are:

```
/appl/
├─ src/                -> appl.fol + application package directories
├─ srclib/             -> standardized project-local special source libraries
│   └─ ffi/            -> library.fol + internal package directories
│   └─ system/         -> library.fol + internal package directories
│   └─ advanced/       -> library.fol + internal package directories
│   └─ dynamicvmrt/    -> library.fol + internal package directories
│   └─ operators/      -> library.fol only
├─ lib/                -> optional; when present, one or more *.folenc Protocol Buffers binary libraries
└─ build/              -> generated compiler output
```

Only paths below `src/` enter the direct project `src/` package index. Package discovery inside a source library begins only below its fixed `srclib/<kind>/` root and remains private to that library. `srclib/exports/` is the deliberate open-package exception: its selected package contexts are added to the owning project's applicable package index without becoming `src/` packages.

Package Access Rules

Four access levels control visibility across package boundaries:

Annotation	Same Package	Parent	Sub-Package	Unrelated	Entry / Surface
<code>@co.dap.public</code>	✓	✓	✓	✓	✓
<code>@co.dap.package</code>	✓	✓	✓	✗	✗
<code>@co.dap.protected</code>	✓	✗	✓	✗	✗
<code>@co.dap.private</code>	✓	✗	✗	✗	✗

Meaning:

- `@co.dap.public` -> visible everywhere
- `@co.dap.package` -> visible across the package family
- `@co.dap.protected` -> visible downward into subpackages only
- `@co.dap.private` -> visible only inside the declaring package

Example:

```
// hr/employee-access.unit.fol - package "hr"

@co.dap.public
_ co.lang.unit = {

    @co.dap.public
    getEmployee(id co.lang.int)->(Employee) = { ... }    // anyone can call

    @co.dap.package
    validateId(id co.lang.int)->(co.lang.bool) = { ... } // hr package family

    @co.dap.protected
    baseQuery()->(co.lang.string) = { ... }              // visible to subpackages

    @co.dap.private
    normalizeId(id co.lang.int)->(co.lang.int) = { ... } // private declaration
}
```

co.* Paths and Aliases

co.* Is Always Available

The FoLang distribution provides the standard `co.*` package tree through its language-supplied export-kind `.fo1enc` artifact. The compiler/toolchain loads that artifact automatically and makes the `co` root implicitly available, so source code never imports standard `co.*` packages through `@co.ddap.import`.

Implicit availability changes package discovery only. Once loaded, `co.*` package declarations use ordinary package, symbol, type, member, visibility, extension, and overload resolution in the same way as declarations reconstructed from another export-kind `.fo1enc` artifact.

```
co.out.println("hello");
co.in.readLine();
x co.lang.int = 42;
```

Being **in scope** does not automatically mean **permitted in every context**.

A package rule, library kind, or other compiler restriction may still forbid use of a particular `co.*` facility. In such cases the name is still known to the compiler, but its use is rejected at compile time.

@co.ddap.alias

Use aliases only to shorten `co.*` paths.

```
@co.ddap.alias(co.out, as="out")
@co.ddap.alias(co.core.list, as="list")
@co.ddap.alias(co.encoding, as="enc")

out.println("hello");
list.of(1, 2, 3);
enc.json.serialize(data);

// full form still works alongside
co.out.println("world");
```

Rules:

- target must be a `co.*` path
- alias must be a valid identifier
- alias scope is file-local
- aliasing user packages is not allowed; use `as=` in `@co.ddap.import`
- duplicate aliases in the same file are compiler errors
- when no alias is declared, the complete `co.*` path is used
- declaring an alias does not disable or hide the complete `co.*` path ***

Package Source Files

A package folder may contain three ordinary source-file categories plus the reserved `package.fo1` metadata form. The fixed project surfaces `src/appl.fo1`, `src/library.fo1`, and `src/export.fo1` are not ordinary package-source files and are invalid inside a package. Likewise, `src/lib/<kind>/library.fo1` and `src/lib/exports/export.fo1` are structural source-domain files and are invalid inside an ordinary package. The compiler classifies ordinary package files from filenames before parsing, using the longest recognized suffix first:

```
<Name>.comp.unit.fo1  -> companion unit
<Fragment>.unit.fo1   -> ordinary package unit
package.fo1           -> reserved package metadata/alias form
<Name>.fo1            -> file-backed primary declaration
```

Reserved structural filenames are recognized before the generic `<Name>.fo1` rule, and a filename ending in `.comp.unit.fo1` is never classified as an ordinary `.unit.fo1` file.

File-Backed Primary Declarations

A primary declaration file has the form `<Name>.fo1` and contains exactly one primary top-level declaration. The source must use `_` in the declaration-name position:

```
// Employee.fo1
_ co.lang.struct = {
  id   co.lang.int;
  name co.lang.string;
}
```

The compiler derives the declaration name `Employee` from the filename. An explicit primary name is invalid:

```
// Employee.fo1
Employee co.lang.struct = { ... }
// compiler error: file-backed primary declarations must use `_
```

This rule applies to ordinary file-backed declarations that use a `<name> co.lang.<kind>` primary form, including classes, structs, cstructs, enums, unions, interfaces, signatures, modules, instances, matchers, objects, and other ordinary `co.lang.*` primary kinds. A `co.lang.library` declaration is not an ordinary `<Name>.fo1` primary: it is valid only in the reserved `library.fo1` surface form described in [Libraries](#). Declaration families with a different surface grammar define their filename binding in their own sections.

The following declaration forms are stated exceptions and keep an explicit name in the head, because filename derivation cannot express what they need:

Form	Why
struct and cstruct declared inside <code>library.fo1</code>	one library surface may carry several boundary declarations
<code>co.lang.data</code> algebraic data type	the head names the variants
parameterized <code>co.lang.type</code>	a filename cannot carry <code>(T)</code>
type declarations in <code>src/app1.fo1</code>	the entry file is not filename-backed as an ordinary package declaration

File-level directives, imports, aliases, annotations, and decorators may appear before the primary declaration. They do not count as additional primary declarations.

ForLang permits the following top-level declaration kinds across ordinary package source files and their reserved special source forms:

1. struct
2. cstruct
3. class
4. module
5. signature
6. interface
7. enum
8. union
9. typeclass
10. instance
11. matcher
12. annotation or object declaration
13. package declaration
14. library declaration
15. unit declaration

Their source-file placement is part of the grammar:

- ordinary filename-backed primary declarations use `<Name>.fo1`
- a package declaration uses the reserved `package.fo1` source form
- an ordinary unit uses `<Fragment>.unit.fo1`
- a companion unit uses `<StructName>.comp.unit.fo1`
- a project-local source-library declaration uses one of the fixed `src/lib/application|ffi|system|advanced|dynamicvmrt/library.fo1` surface forms
- a standalone packaged-library declaration uses the fixed direct `src/library.fo1` project-surface form and must declare its kind explicitly with `@co.dap.library(type=...)`
- a standalone export project uses the fixed direct `src/export.fo1` selector form, and the project-local export domain uses `src/lib/exports/export.fo1`; neither form contains a primary declaration

A source file is invalid when it contains multiple unrelated primary declarations, places a package declaration outside `package.fo1`, places a unit declaration outside a recognized unit filename, places a library declaration outside one of its context-valid surface positions, or places project/library/export metadata outside its dedicated structural source form.

Package-level functions and non-UDT type declarations belong inside ordinary unit files.

Filename Canonicalization

Filename-derived declaration identity is independent of filesystem case sensitivity. The compiler normalizes and case-folds the filename stem to construct the duplicate-detection and lookup key:

```
canonical file key = caseFold(normalize(filename stem))
```

The language-level declaration is then represented using FoLang's canonical declaration spelling. Therefore:

```
employee.fol
Employee.fol
EMPLOYEE.fol
```

all denote the canonical declaration name `Employee`. They conflict rather than declaring separate types, even on a case-sensitive filesystem.

The same rule applies to companion owners:

```
employee.comp.unit.fol -> companion of Employee
```

This rule is specific to filename-derived primary declarations and companion-owner identity; it does not make every FoLang identifier case-insensitive. The compiler must not rely on operating-system filename comparison.

The filename stem must form a valid FoLang declaration identifier after normalization. Case variants and canonically equivalent spellings produce the same package-index key.

Ordinary Package Units

An ordinary package-unit file has the form `<Fragment>.unit.fol`:

```
// arithmetic.unit.fol
_ co.lang.unit = {
  abs(value co.lang.int)->(co.lang.int) = {
    ...
  }
}
```

The fragment name is organizational only. It does not create a public `Arithmetic` symbol. All ordinary unit members in the package are merged into one package namespace:

```
math/
├─ arithmetic.unit.fol
├─ comparison.unit.fol
└─ optional.unit.fol
```

```
math.abs(-10);
math.max(10, 20);
value math.Option(co.lang.int);
```

The following qualification is invalid:

```
math.Arithmetic.abs(-10); // compiler error
```

Any number of ordinary unit files may exist in one package. During package indexing, the compiler shallow-parses their declarations, merges their members with filename-derived primary declarations, and reports duplicate names or duplicate function signatures according to FoLang overload rules.

Companion Unit Files

A companion unit has the form `<StructName>.comp.unit.fol` and must contain `_ co.lang.unit`. The canonical owner name comes from the filename, not from the source body.

```
Employee.fol
Employee.comp.unit.fol
```

The compiler can detect the following before parsing either body:

- duplicate primary declarations after filename canonicalization
- duplicate companion files for the same canonical owner
- an orphan companion file for which no same-package primary declaration exists

After the owner header is known, the compiler also verifies that the owner is a `co.lang.struct`. Other primary declaration kinds cannot own companion units.

Package Indexing Levels

```
Level 1: filename and package indexing
  classify primary, ordinary-unit, and companion-unit files
  canonicalize primary names and companion owners
  detect duplicate primary files
  detect duplicate or orphan companion files
  shallow-parse ordinary units
  merge ordinary-unit members into the package namespace
  report package-namespace conflicts

Level 2: companion validation
  parse companion declarations
  validate explicit receivers against the filename-derived owner
  merge companion members into the owner's companion namespace
  report duplicate companion members and receiver mismatches

Level 3: full semantic resolution
  resolve declaration bodies, expressions, calls, patterns, and types
```

This organization lets an imported source package build most of its symbol index cheaply before full parsing. Compiled packages may store the same index as metadata so imports need no source parsing.

Application Entry File

The application entry file is the fixed **special executable source form** `src/appl.fo1`. It is not a package, does not create an importable namespace, and is not subject to the ordinary package-file rule requiring exactly one primary declaration.

The compiler creates a dedicated **entry-file context** for it:

```
ApplicationEntryContext
├─ file directives, imports, and aliases
├─ entry-local non-structural type declarations
├─ non-capturing function-pattern groups
├─ capturing `let` function-pattern groups
└─ executable statements and expressions
```

Everything declared directly in this context is private to the entry file. Entry-local symbols cannot be imported, exported, or referenced by package or library source files.

Allowed Entry-File Constructs

The application entry file uses exactly the same grammar and allowed-construct rules described in [Single Source Application File](#). The earlier section provides the developer-facing overview; this section defines the entry file's formal context, privacy, and dependency direction.

Entry-Local Function Patterns

Function-pattern groups are allowed as a special entry-file construct even though ordinary function declarations are forbidden. FoLang provides two entry-file forms with the same pattern-dispatch model but different capture semantics.

Entry-File Dependency Direction

The entry file may depend on packages and libraries, but packages and libraries may never depend on the entry context:

```
application entry file
  ↓ uses
packages and libraries
```

This allows the entry file to coordinate application startup while preserving package and library independence.

Export Projects

Purpose

An **Export Project** packages ordinary FoLang package contexts for reuse without introducing a library API gate. It is the open counterpart to a Library Project.

```

Application Project
-> executable

Library Project
-> .folenc
-> only the `library.fol` surface projection is externally visible
-> internal package/type/extension/operator state is isolated

Export Project
-> .folenc
-> selected ordinary package contexts and symbols are externally visible
-> exported declarations keep their ordinary FoLang package/type semantics

```

A Library Project is chosen when the producer wants a closed component whose surface is the only entry point. An Export Project is chosen when the producer wants consumers to work with the actual exported packages and their public or otherwise normally accessible declarations: consumers may subtype extensible types, define legal extensions, provide typeclass instances, and otherwise use the exported declarations according to the ordinary language rules. Export packaging does not create a boundary adapter, snapshot boundary, or separate surface type identity.

Export Project Layout

An Export Project uses exactly the same project-root architecture as every other FoLang project:

```

/mypack/
├── src/
│   ├── export.fol          <- fixed project classifier and package selector
│   ├── hr/
│   │   ├── employee/
│   │   │   ├── Employee.fol
│   │   │   └── EmployeeService.fol
│   │   └── payroll/
│   │       └── Payroll.fol
│   └── internal/
│       └── BuildHelper.fol
├── src/lib/                <- optional; omit when unused
├── lib/                    <- optional; omit when unused
└── build/                  <- compiler-managed

```

`src/export.fol` is structural and is not a package. Every other direct entry under `src/` must be an ordinary non-empty package directory, exactly as for the other project kinds.

`export.fol`

`export.fol` contains no `_ co.lang.* = { ... }` declaration. It contains exactly one project-level `@co.dap.export(packages={...})` annotation and no functions, variables, imports, types, executable statements, or other declarations.

```

// src/export.fol
@co.dap.export(
  packages={
    hr.employee: {recurse=true},
    hr.payroll:  {recurse=false}
  }
)

```

The `packages` map names package contexts relative to the export project's package root:

- `recurse=false` exports exactly the named package;
- `recurse=true` exports the named package and every descendant package;
- recursion selects package contexts only; it never changes declaration visibility;
- an exported package path must exist;
- duplicate or overlapping selections normalize to one exported package context;
- package names are preserved exactly; the export project's filesystem/artifact name is never inserted as a package prefix.

For example:

```

@co.dap.export(
  packages={
    co:      {recurse=false},
    co.lang: {recurse=true},
    co.meta: {recurse=true}
  }
)

```

may export `co`, all packages below `co.lang`, and all packages below `co.meta`. A private declaration remains private, a protected declaration remains protected, and every extensibility/subtyping rule continues to be determined by the declaration itself.

Export Artifact Projection

A standalone Export Project produces the standard `<project-name>.folenc` artifact. The artifact is tagged as kind `export`.

Unlike a Library Project, the compiler does not synthesize a surface API projection. Instead it serializes the selected package contexts and the symbol metadata required to preserve those packages as ordinary FoLang packages for a consumer:

```
mypack.folenc                                kind=export
├─ exported package contexts
│   └─ hr.employee
│       └─ hr.payroll
├─ exported symbol/type metadata
└─ compiled implementation/linkage
```

The consumer sees the actual exported package/type identities. No `co.lang.library` wrapper, boundary contract copy, or adapter function is introduced.

Package imports from an export artifact use the ordinary package form:

```
@co.ddap.import(package="hr.employee", as="emp")

e emp.Employee;
```

They do **not** use `@co.ddap.import(library=...)`.

An export-kind `.folenc` may contain implementation data needed by the backend, but only packages selected by `export.fol` enter the consumer's applicable package index. Unselected package paths are not externally resolvable.

Open Package Semantics

An export boundary is a packaging boundary, not an isolation boundary. Once an exported package context is available to a consumer, qualified-name and member resolution use the same rules as for a source package:

```
source package
project-local exported package
export-kind .folenc package
language-provided co.* export-kind .folenc package
├─
│   v
ordinary package / symbol / member resolution
```

The origin changes how the package is discovered, not how its declarations behave after resolution. Subject to the normal rules of the target declaration, consumer code may therefore:

- reference exported declarations where their ordinary visibility permits;
- subtype exported extensible types;
- define and activate legal extensions for exported types;
- provide legal typeclass instances;
- provide operator overload implementations for exported types when ordinary operator ownership/extension rules permit;
- pass exported types through ordinary application code without library-surface snapshot/ABI adaptation.

The consumer cannot rewrite the already compiled implementation body contained in an export artifact, and export packaging does not override `final`, sealed, private, capability, ownership, or other declaration-level restrictions.

Project-Local Export Domain

`src/lib/exports/` is the project-local source form of the same open-package idea:

```
<project-root>/src/lib/exports/
├─ export.fol
├─ hr/
│   └─ employee/
│       └─ Employee.fol
└─ text/
    └─ format/
        └─ Formatter.fol
```

Its structural file has the same syntax:


```
// src/lib/exports/export.fol
@co.dap.export(
  packages={
    hr.employee: {recurse=true},
    text.format: {recurse=true}
  }
)
```

Rules:

- `src/lib/exports/` is a standardized source-domain root, not a package;
- `export.fol` is the only direct file at that root;
- package paths begin below `src/lib/exports/`; `exports` never contributes a package-name component;
- selected package contexts are available to the owning project through ordinary `@co.ddap.import(package="...")` resolution;
- the domain has no `co.lang.library` surface and no library boundary-adapter restrictions;
- project-local export source does not produce a separate `.folenc`; it is compiled as project-owned source;
- usage, liveness, and cycle checks follow the ordinary package rules for the combined project graph.

Export Projects and `co.*`

The language rule that `co.*` is always available remains unchanged. A compiler may prepare the standard `co.*` package tree from compiler-owned/exported package metadata, but after the `co` root is made implicitly visible, package/member resolution follows the same ordinary context and symbol-table rules as any other exported package tree. The special property is implicit root availability, not a separate qualified-name algorithm.

Export Projects and Operators

Exporting packages does not export a parser operator table. Operator parse metadata remains compilation-domain-local. A standalone Export Project may use its own permitted project operator bootstrap while compiling, but the resulting export-kind `.folenc` does not cause new operator spellings, fixities, precedences, or associativities to appear automatically in the consumer.

Operator implementations that are ordinary exported/visible declarations remain subject to the normal operator ownership, extension, activation, and resolution rules. A consumer that needs a custom operator spelling must have that spelling registered in its own compilation domain before parsing source that uses it.

Libraries

Library Project Layout

A standalone packaged-library project uses the same project-root architecture as an application. Only `src/` must exist before compilation. `src/lib/` and `lib/` are optional and must be omitted when unused; when present they cannot be empty. `build/` is compiler-managed and may be absent or empty before compilation.

```
/hrlib/
├── src/
│   ├── library.fol          <- fixed standalone-library surface; @co.dap.library(type=...)
│   ├── emp/                <- internal package emp
│   │   ├── Employee.fol
│   │   └── EmpService.fol
│   └── auth/                <- internal package auth
│       ├── Auth.fol
│       └── AuthService.fol
├── src/lib/                 <- optional; omit when unused
├── lib/                     <- optional; omit when unused
└── build/                   <- compiler-generated output; created when absent
    └── <frontend-artifact>
```

The empty `src/lib/` and `lib/` entries above are diagrammatic placeholders only; an actual project cannot contain those directories empty.

Consumers see only what the direct library surface under `src/` projects. All package subfolders below the library project's `src/` are internal. The root-level `build/` directory is generated output and is not an internal package.

Library Surface file

FoLang uses surface files in two situations:

1. Packaged library project surfaces
2. Project-local source library surfaces

A library surface contains `_ co.lang.library` and defines the public boundary data contracts and boundary-adaptor functions through which consumers call the library. Its valid structural location depends on the compilation domain. For a project-local source library, the fixed `src/lib/<kind>/library.fol` path supplies identity and kind. For a standalone packaged-library project, the fixed `src/library.fol` file is the library surface and must declare its kind explicitly with `@co.dap.library(type=...)`.

<code>src/appl.fol</code>	-> application entry
<code>src/library.fol</code>	-> standalone packaged-library surface
<code>src/export.fol</code>	-> standalone export selector
<code>src/lib/application/library.fol</code>	-> project-local application-library surface
<code>src/lib/ffi/library.fol</code>	-> project-local FFI-library surface
<code>src/lib/exports/export.fol</code>	-> project-local open-package selector

A library surface is not an ordinary package file. It may contain multiple boundary data declarations and public functions inside one `co.lang.library` declaration.

Packaged Library Project Surface

A packaged-library project has no application entry file. Its mandatory `src/` domain contains exactly one direct FoLang source file named `library.fol`. That fixed file is the packaged-library surface. All other source below `src/` resides in internal package directories.

```
/hrlib/
├─ src/
│   ├── library.fol          <- fixed direct library surface; explicit project-level kind declaration
│   └─ emp/
│       ├── Employee.fol
│       └─ EmployeeService.fol
├─ auth/
│   └─ AuthService.fol
├─ src/lib/                  <- optional; omit when unused
├─ lib/                      <- optional; omit when unused
└─ build/                   <- compiler-generated output
    └─ <frontend-artifact>
```

Rules:

- exactly one direct FoLang source file named `library.fol` exists under the packaged-library project's `src/` directory
- that direct surface file must contain the project-level `@co.dap.library(type=...)` declaration followed by `_ co.lang.library = { ... }`
- the explicit library kind is required because the fixed standalone surface path `src/library.fol` is kind-neutral; unlike `src/lib/<kind>/library.fol`, its filesystem location does not encode a library kind
- the compiler records the declared kind on the library surface context/symbol table before validating the surface, then applies the declaration, boundary-type, public-signature, adapter-body, capability, and other restricted-construct rules for that library kind
- a missing, unsupported, or conflicting standalone library-kind declaration is a compile-time error
- the supported packaged-library kinds are `application`, `dynamicvmrt`, `advanced`, `system`, and `ffi`; `operator` is not a packaged-library kind and operator metadata is never exported through `.folenc`
- the project is compiled into the standard packaged library artifact `.folenc`, which is Protocol Buffers binary
- consumers import it with `@co.ddap.import(library= "...")`
- internal package folders are compiled into the library but are not directly visible to consumers
- the frontend output is emitted under the library project's root-level `build/` directory using the wire format/version selected by the backend interchange contract
- `build/` is generated output and never an internal package or source-discovery root

Project-Local Source Library Surface

Project-local source libraries live only under their fixed `src/lib/` slots. `src/lib/` itself is neither a package nor a library. FoLang also reserves `src/lib/exports/` for project-local open package source and `src/lib/operators/` for the owning project's operator bootstrap:

```

<project-root>/src/lib/
├─ application/
│   └─ library.fol
│       └─ <internal packages>/
├─ ffi/
│   └─ library.fol
│       └─ <internal packages>/
├─ system/
│   └─ library.fol
│       └─ <internal packages>/
├─ advanced/
│   └─ library.fol
│       └─ <internal packages>/
├─ dynamicvmrt/
│   └─ library.fol
│       └─ <internal packages>/
├─ exports/
│   └─ export.fol
│       └─ <export-source packages>/
└─ operators/
    └─ library.fol

```

For `application`, `ffi`, `system`, `advanced`, and `dynamicvmrt`:

- the immediate directory name is fixed by FoLang and determines both library identity and library kind;
- the directory itself is not a package;
- exactly one direct file is permitted: `library.fol`;
- `library.fol` contains `_ co.lang.library = { ... }` and does not use `@co.dap.library(type=...)`;
- all implementation source must reside in descendant package directories;
- descendant package names are relative to the source-library root and may be arbitrarily deep;
- no nested `library.fol` or nested source-library boundary is permitted;
- internal packages remain private to that source-library compilation domain and never enter the ordinary project `src/` package index;
- source in the owning project's `src/` domain accesses the library only through the projected `library.fol` API surface;
- a project-local source library is project-owned source, not an independently distributed library dependency; `src/lib/` exists to segregate code with special capability or implementation intent into a separate source domain instead of mixing that code into the ordinary `src/` package tree;
- project-local source libraries do not generate independent `.folenc` artifacts; they are compiled as part of the owning project's frontend output;
- project-local source-library API usage and reachability follow the strict project-owned rules in [Unused Symbols, Liveness, and Reachability](#).

Example:

```

<project-root>/src/lib/ffi/
├─ library.fol           <- FFI public surface; ffi/ is not a package
├─ native/              <- internal package native
│   └─ marshal/         <- internal package native.marshal
│       └─ Marshal.fol
└─ database/            <- internal package database
    └─ Postgres.fol

```

```

// <project-root>/src/lib/ffi/library.fol
_ co.lang.library = {
  ...
}

```

Import:

```
@co.ddap.import(library="ffi", src-library=true, as="ffilib")
```

Resolution:

```

library="ffi", src-library=true
-> <project-root>/src/lib/ffi/library.fol

```

The same library-surface rule applies to `application`, `system`, `advanced`, and `dynamicvmrt`.

The one-per-kind restriction applies only to **project-local source libraries within one project**. Projects may consume any number of separately built library-kind `.folenc` artifacts from `lib/`.

`srcLib/exports/` is not a library kind and has no `_ co.lang.library` surface. Its fixed `export.fo1` selects open package contexts as defined in [Export Projects](#).

`srcLib/operators/` is the special bootstrap slot described in [Operators](#). It contains exactly `library.fo1` and no package subdirectories. Its filesystem position selects the dedicated operator-source grammar; it is not an ordinary importable source-library API surface.

Unified Surface Model

Every ordinary library kind uses the same conceptual surface shape. The `srcLib/operators/` bootstrap slot is deliberately excluded because it is parser metadata rather than an importable library API:

```
library surface
├─ boundary data contracts
└─ public boundary-adapter functions
```

Library kind changes the permitted boundary representation and transfer semantics, not the conceptual form of the API.

Library kind	Boundary data	Transfer semantics
<code>application</code>	<code>co.lang.struct</code>	automatic deep snapshot
<code>dynamicvmrt</code>	<code>co.lang.struct</code>	automatic deep snapshot
<code>advanced</code>	<code>co.lang.struct</code>	automatic deep snapshot
<code>system</code>	<code>co.lang.cstruct</code>	system ABI value
<code>ffi</code>	<code>co.lang.cstruct</code>	C ABI value

`co.lang.struct` and `co.lang.cstruct` solve different problems:

```
struct -> FoLang semantic data contract
cstruct -> physical ABI-compatible value contract
```

Value transfer must not be confused with C-compatible layout. Application-family libraries therefore use expressive `struct` contracts transferred as deep snapshots, while system and FFI libraries use restricted `cstruct` contracts.

Allowed Surface Declarations

A library surface may contain only:

- file- or library-level imports needed by its adapter implementations
- `co.lang.struct` boundary declarations for `application`, `dynamicvmrt`, and `advanced`
- `co.lang.cstruct` boundary declarations for `system` and `ffi`
- public free-function API declarations with boundary-adapter definitions

The following declaration kinds are forbidden directly in every library surface:

- classes
- modules
- interfaces
- signatures
- units and companion units
- associated functions
- operator functions
- enums, unions, and other ADTs
- type aliases, newtypes, and opaque types
- objects, instances, type classes, and dependent types
- macros, templates, annotations, and decorators
- global variables, pointers, references, addresses, or mutable surface state

Surface `struct` and `cstruct` declarations are data contracts only. They cannot have companion units, associated functions, operators, methods, or other behavior on the surface.

Declarations directly inside the `co.lang.library` body are exported by default. Imports and implementation details are never exported.

Public Signature Type Closure

Every public surface function signature must be closed over the library's exported boundary type set.

For `application`, `dynamicvmrt`, and `advanced` surfaces, parameters and results may use only:

- approved built-in types

- `co.lang.struct` types declared in the same library surface

For `system` and `ffi` surfaces, parameters and results may use only:

- ABI-safe built-in types
- `co.lang.cstruct` types declared in the same library surface

The same closure rule applies recursively to fields of surface boundary types:

- a surface `struct` field may use an approved built-in type or another surface-declared `struct`
- a surface `cstruct` field may use an ABI-safe built-in type or another surface-declared `cstruct`
- an internal package type may never appear in a public function signature or surface boundary-type field
- pointers, references, and addresses may never cross any public library surface

A built-in type is not automatically surface-safe merely because it belongs to `co.lang`. An approved surface built-in must be concrete, fully resolved, and transferable under the library kind's boundary semantics.

The following categories are forbidden in public surface fields and signatures:

- inference-only types such as `co.lang.auto` and `co.lang.infer`
- dynamically typed or unconstrained carriers such as `co.lang.dynamic`, `co.lang.any`, `co.lang.typed`, and `co.lang.untyped`
- function, closure, delegate, loader, AST, reflection, or runtime implementation values
- pointer, reference, address, thunk, and implementation-handle types
- any type whose reachable representation contains a forbidden type

For application-family surfaces, managed built-ins such as `co.lang.string` are permitted when the compiler defines deep-snapshot reconstruction for them. For system and FFI surfaces, only built-ins with a defined ABI representation are permitted; for example, `co.lang.string` is not directly `cstruct`-compatible.

Valid:

```
// Employee.fol
_ co.lang.struct = {
  id      co.lang.int;
  name    co.lang.string;
  address Address;
}

// Address.fol
_ co.lang.struct = {
  city co.lang.string;
}
```

Invalid:

```
getEmployee(id co.lang.int)->(emp.internal.Employee);
// compiler error: an internal type escapes through the public signature
```

Boundary-Adapter Functions

A public surface function may contain a definition, but its definition is restricted to boundary adaptation.

A boundary-adapter function may:

- read its input parameters and their fields
- declare temporary local values used only for conversion
- construct internal input values
- perform deterministic field-to-field or representation conversion
- invoke exactly one internal implementation operation
- receive that operation's result
- construct and return an allowed surface `struct`, surface `cstruct`, or built-in value
- use a direct delegation expression when no conversion is required

A boundary-adapter function may not:

- implement business rules or domain calculations
- perform business validation or authorization decisions
- orchestrate multiple implementation operations
- call another surface function
- perform persistence, networking, file I/O, retries, caching, or transactions

- start concurrency, scheduling, asynchronous work, processes, or continuations
- contain loops, recursion, workflow branching, or externally observable state mutation
- expose an internal value directly without converting it to an allowed boundary type

The compiler enforces the restricted adapter statement set. The restriction is structural: arbitrary executable logic is not accepted merely because it is placed in a surface file.

A direct delegate is the simplest adapter form:

```
health()->(co.lang.bool)
=>> health.internal.Service.health();
```

A converting adapter may map between the public contract and an internal model.

Standalone Packaged Application-Library Surface Example

```
// library.fol – standalone packaged application library
@co.dap.library(type="application")
_ co.lang.library = {

    @co.ddap.import(package="emp", as="emp")

    Employee co.lang.struct = {
        name co.lang.string;
        id   co.lang.int;
    }

    getEmployee(empId co.lang.int)->(Employee) = {
        internalEmployee := emp.EmployeeService.getEmployee(empId);

        this.return Employee{
            name: internalEmployee.name,
            id: internalEmployee.id
        };
    }
}
```

The surface owns conversion from the internal `emp.Employee` representation to the public `hrLib.Employee` contract. The `emp` package owns validation, business rules, persistence, and workflow.

The consumer sees the equivalent API contract:

```
// Employee.fol
_ co.lang.struct = {
    name co.lang.string;
    id   co.lang.int;
}

getEmployee(empId co.lang.int)->(Employee);
```

The consumer does not see the body of `getEmployee` or the `emp` package.

Standalone Packaged System-Library Surface Example

The following example is a separately authored packaged-library project. Project-local `srcLib/system/library.fol` does not repeat the kind annotation because its fixed path already establishes `system`.

```
// library.fol – standalone packaged system library
@co.dap.library(type="system")
_ co.lang.library = {

    @co.ddap.import(package="driver.internal", as="impl")

    Point co.lang.cstruct = {
        x co.lang.int;
        y co.lang.int;
    }

    getOrigin()->(Point) = {
        internalPoint := impl.DriverService.getOrigin();

        this.return Point{
            x: internalPoint.x,
            y: internalPoint.y
        };
    }
}
```

Pointers, addresses, native bindings, and hardware state belong in `driver.internal`. They are forbidden in the public surface and cannot appear in the exported signature.

Boundary Transfer Semantics

For `application`, `dynamicvmrt`, and `advanced` libraries:

```
consumer boundary value
  ↓ automatic deep snapshot of the complete reachable value graph
library-local built-in or surface struct
  ↓ surface conversion
internal implementation value
  ↓ surface conversion
library-local built-in or surface struct
  ↓ automatic deep snapshot
consumer-local boundary value
```

The snapshot rule applies to both surface structs and approved built-in arguments/results. The library never receives the consumer's live object identity, and the consumer never receives a live internal-library object.

For `system` libraries, `cstruct` values cross according to the selected backend/platform system ABI.

For `ffi` libraries, `cstruct` values cross according to the declared C ABI, including its layout, alignment, and calling-convention requirements.

Surface-to-Internal Dependency Direction

The source-level dependency is one-way:

```
library surface
  ↓ imports and invokes
internal packages
```

Internal packages:

- do not import the library surface
- do not use surface-declared `struct` or `cstruct` types
- define their own implementation and domain types
- return internal values to the surface adapter
- contain all business logic, validation, workflow, I/O, and state management

This prevents a surface/internal compilation cycle and keeps the public contract independent from the internal domain model.

An internal implementation symbol may be visible to its own library surface without becoming part of the consumer API. Library encapsulation takes precedence over ordinary package visibility: consumers cannot resolve internal package symbols even when those symbols are callable from the surface during library compilation.

Consumer API Projection

The compiler does not expose the complete surface compilation context to a consumer. It creates a projected imported symbol table.

The projected API contains:

- the library identity and kind

- complete public surface `struct` or `cstruct` definitions, including field names and permitted field types
- public function names
- public function parameter names and types
- public function result types
- calling-convention, effect, error, and linkage metadata where applicable

The projected API does not contain:

- function bodies or implementation AST/HIR
- imports used by the surface
- local conversion variables
- delegate or implementation targets
- internal package paths or symbols
- private compiler-generated helpers
- business classes, modules, units, or internal data types

Therefore, a function definition may exist in the surface source and compiled artifact while the consumer's symbol table contains only its signature.

The backend may retain implementation bodies for code generation, linking, or whole-program optimization. Backend possession of implementation code does not make that code semantically visible to the consumer.

Source libraries follow the same rule. Availability of source code does not weaken the projected API boundary.

Library Compilation Order

A library uses the same import-usage and symbol-usage validator as an application. It is processed in these logical stages:

1. parse the surface header, boundary data declarations, and public function signatures;
2. parse/prepare internal packages without depending on surface types;
3. register internal imports as provisional context references rather than copying symbol tables;
4. type-check and link surface boundary-adapter bodies against internal package symbols, activating an import only when a symbol from it is actually used;
5. complete semantic resolution and mark exact usage-checkable symbols live;
6. prune and report zero-use imports;
7. reject unused usage-checkable symbols and unreachable internal source/packages;
8. emit the compiled implementation and projected consumer API.

The declarations intentionally exported by `library.fol` are semantic roots of the library's internal compilation graph. This order permits the surface to call internal packages while preventing internal packages from depending on the surface. The different producer, project-local source-library, and `.folenc` consumer usage rules are defined in [Unused Symbols, Liveness, and Reachability](#).

Library Kinds

For project-local source libraries, the library kind is structural and comes from the fixed path:

```
src/lib/application/ -> application
src/lib/ffi/         -> ffi
src/lib/system/      -> system
src/lib/advanced/    -> advanced
src/lib/dynamicvmrt/ -> dynamicvmrt
src/lib/exports/     -> open package-export source domain (not a library kind)
src/lib/operators/   -> operator bootstrap context (not importable; not packaged)
```

All project-local library surfaces under `src/lib/application|ffi|system|advanced|dynamicvmrt/library.fol` therefore do **not** repeat the kind with `@co.dap.library(type=...)`. The `exports/` slot is not a library and uses `export.fol`; the `operators/` slot uses `library.fol` only as its dedicated operator-bootstrap structural file.

A separately authored packaged-library project is not classified by an enclosing `src/lib/<kind>/` slot. Its fixed `src/library.fol` surface must therefore contain the project-level `@co.dap.library(type=...)` declaration. The compiler uses that explicit kind to create/tag the library surface context and symbol table, then verifies that the surface contains only constructs allowed for that kind before producing the `.folenc` projection. The `application` kind is used for an ordinary safe packaged library.

Shared Meaning

- `application` -> ordinary safe library implementation
- `dynamicvmrt` -> application features plus full `co.meta` dynamic-runtime support
- `advanced` -> macro, template, concurrency, transformation, and runtime-machinery implementation

- `system` -> unsafe low-level runtime, pointer, process, native, and platform-control implementation
- `ffi` -> foreign-interface implementation

Library kind controls internal capabilities. All kinds still expose only boundary data contracts and public boundary-adaptor functions.

`application` Libraries

Internally allowed:

- ordinary safe language features
- classes, modules, interfaces, signatures, units, and ordinary structs
- normal application-level packages and libraries

Internally forbidden:

- pointers, references, and addresses
- native functions annotated with `@co.dap.native`
- macros and templates
- low-level concurrency machinery
- advanced transformation or dynamic-runtime machinery

Public boundary:

- surface `struct` contracts
- approved built-in types
- deep-snapshot transfer

`dynamicvmrt` Libraries

Internally allowed:

- all `application` capabilities
- full `co.meta` support
- runtime reflection, instrumentation, dynamic loading, patching, and eval-based facilities

Internally forbidden:

- system and FFI capabilities unless reached through an imported public library surface
- raw pointers, addresses, and native functions in its own implementation

Public boundary:

- surface `struct` contracts
- approved built-in types
- deep-snapshot transfer

`advanced` Libraries

Internally allowed:

- all ordinary safe language features
- macros and templates
- async, parallel, continuation, scheduling, and transformation machinery
- compiler/runtime behavior definitions permitted to advanced libraries

Internally forbidden:

- raw pointers
- references and addresses
- native functions

Public boundary:

- surface `struct` contracts
- approved built-in types
- deep-snapshot transfer

`system` Libraries

Internally allowed:

- pointers, references, addresses, and word-level types
- native functions, including `@co.dap.native` declarations and `co.native` facilities

- structs and cstructs
- units containing free functions
- basic value-typed functions
- type aliases
- basic threading and process primitives

Internally forbidden:

- classes
- modules
- interfaces and signatures
- overloading and overriding
- new operator definitions
- generics and templates
- macros
- reflection and dynamic runtime
- dependent types and type classes
- function patterns
- closures and currying
- advanced concurrency machinery such as continuations, defer, lazy evaluation, and thunks
- variadic, named, optional, and default parameters

Public boundary:

- surface `cstruct` contracts
- ABI-safe built-in types
- system ABI value transfer

`ffi` Libraries

Internally allowed:

- extern declarations and foreign symbol bindings
- pointers, references, addresses, and C ABI types
- cstructs and restricted implementation structs
- units containing free functions
- foreign calling-convention and symbol-linkage metadata

Internally forbidden:

- classes
- modules
- interfaces and signatures
- overloading and overriding
- new operator definitions
- generics and templates
- macros
- reflection and dynamic runtime
- dependent types and type classes
- closures, currying, and advanced concurrency machinery

Public boundary:

- surface `cstruct` contracts
- C-ABI-safe built-in types
- C ABI value transfer

Dependency Direction

Allowed dependency flow is one-way:

```

application
  ↓ through public surface APIs only
dynamicvmrt
  ↓ through public surface APIs only
advanced
  ↓ through public surface APIs only
system
  ↓ through public surface APIs only
ffi

```

A library may depend on a library at the same or a lower level only when the dependency does not create a cycle. Reverse dependencies are compiler errors.

Cross-Library Communication

From	To	Allowed?	Notes
application	dynamicvmrt	✓	projected public surface only
application	advanced	✓	projected public surface only
application	system	✓	projected public surface only
application	ffi	✓	projected public surface only
dynamicvmrt	advanced	✓	projected public surface only
dynamicvmrt	system	✓	projected public surface only
dynamicvmrt	ffi	✓	projected public surface only
dynamicvmrt	application	✗	reverse dependency
advanced	system	✓	projected public surface only
advanced	ffi	✓	projected public surface only
advanced	application	✗	reverse dependency
advanced	dynamicvmrt	✗	reverse dependency
system	ffi	✓	projected public surface only
system	application	✗	reverse dependency
system	dynamicvmrt	✗	reverse dependency
system	advanced	✗	reverse dependency
ffi	application	✗	reverse dependency
ffi	dynamicvmrt	✗	reverse dependency
ffi	advanced	✗	reverse dependency
ffi	system	✗	reverse dependency

Units in detail

A `co.lang.unit` is a stateless file-level declaration container. It is not instantiable and does not create an object, runtime scope, or public unit namespace.

FoLang has two unit-file forms:

1. **ordinary package unit** — `<Fragment>.unit.fol`
2. **struct companion unit** — `<StructName>.comp.unit.fol`

Both forms use `_ co.lang.unit`; explicit unit names are invalid.

Ordinary Package Units

Example package layout:

```

math/
├─ arithmetic.unit.fol
├─ comparison.unit.fol
└─ optional.unit.fol

```

```
// arithmetic.unit.fol
_ co.lang.unit = {
  abs(value co.lang.int)->(co.lang.int) = { ... }

  max(
    a co.lang.int,
    b co.lang.int
  )->(co.lang.int) = { ... }
}
```

```
// optional.unit.fol
_ co.lang.unit = {
  Option(T) co.lang.type =
    Some(T) | None();

  isSome(value Option(co.lang.int))->(co.lang.bool) = {
    ...
  }
}
```

All declarations are contributed directly to the `math` package namespace:

```
math.abs(-10);
math.max(10, 20);
value math.Option(co.lang.int);
```

The unit filenames never appear in qualified names:

```
math.Arithmetic.abs(-10); // compiler error
math.Optional.Option(co.lang.int); // compiler error
```

Within the same package, the functions and types may be referenced without the package prefix according to ordinary package name-resolution rules.

An ordinary unit may contain:

- receiverless functions
- `co.lang.type` aliases and ADT/type-constructor declarations
- newtype and opaque-type declarations
- subtype and supertype declarations
- macros and template declarations
- other non-instantiable type declarations explicitly permitted by their sections
- public, package, protected, or private members

An ordinary unit may not contain:

- fields or mutable unit state
- executable loose statements
- classes, structs, cstructs, enums, unions, modules, interfaces, or signatures
- lifecycle declarations for the unit
- a declaration that creates an independent nested primary namespace

A unit file may group closely related declarations or merely organize a convenient set of package-level functions and types. Semantic cohesion is recommended but not enforced.

For canonical unit and per-symbol usage rules, see [Unused Symbols, Liveness, and Reachability](#).

Package-Namespace Consolidation

All ordinary unit files in one package are shallow-parsed and consolidated into one package symbol table. Their order and filenames have no semantic effect.

```
package members =
  filename-derived primary declarations
  + declarations from every ordinary unit file
```

A duplicate is reported as soon as the package index is built. Conflicts include:

- the same type or value name contributed by two ordinary units
- an ordinary-unit member colliding with a filename-derived primary declaration
- duplicate function signatures not permitted by overload rules

- declarations that normalize to the same canonical package symbol

Diagnostics should identify both source files.

Companion Units

A companion unit is governed by the filename and owner-validation rules in [Struct Companion Units](#). Its members do not merge directly into the package namespace; they attach to the owner struct's companion namespace.

Companion-unit functions follow the canonical strict per-symbol usage rule in [Unused Symbols, Liveness, and Reachability](#).

Units have no fields, identity, instances, inheritance, polymorphic dispatch, or independent construction.

CStructs

`co.lang.cstruct` is a C-like value type — passed by value, simple memory layout, safe to cross zone boundaries. Unlike `co.lang.struct` which is passed by reference, `co.lang.cstruct` is always copied on pass.

```
// Point.fol
_ co.lang.cstruct = {
  x co.lang.int;
  y co.lang.int;
}

// Rect.fol
_ co.lang.cstruct = {
  origin Point;
  width co.lang.int;
  height co.lang.int;
}
```

cstruct Rules

```
always passed by value – never by reference
simple memory layout – no metadata
can contain only simple types and other cstructs
cannot contain co.lang.struct           ✗ has metadata
cannot contain co.lang.string           ✗ heap allocated
cannot contain co.lang.dynamic          ✗ runtime type info
cannot contain classes                  ✗ vtable, metadata
cannot contain modules                  ✗
cannot contain any heap allocated type ✗
cannot have methods
cannot have associated functions
cannot embed co.lang.struct
safe to cross direct ABI and zone boundaries

allowed field types:
  co.lang.int, co.lang.uint, co.lang.float  ✓ primitives
  co.lang.bool, co.lang.char, co.lang.byte  ✓ primitives
  co.lang.int->([N])                        ✓ fixed size arrays
  co.lang.cstruct                          ✓ other cstructs
```

Packed cstruct — no padding, exact memory layout

Used for hardware registers, binary protocols, exact memory mapped formats:

```
// Register.fol
@co.dap.packed
_ co.lang.cstruct = {
  flags co.lang.uint8;
  status co.lang.uint8;
  data co.lang.uint16;
}
```

SIMD cstruct — aligned for vector operations

Used for math, graphics, signal processing:

```
// Vec4.fol
@co.dap.simd(align=16)
_ co.lang.cstruct = {
  x co.lang.float;
  y co.lang.float;
  z co.lang.float;
  w co.lang.float;
}
```

Both together

```
// AVXVec.fol
@co.dap.packed
@co.dap.simd(align=32)
_ co.lang.cstruct = {
  data co.lang.float;
}
```

`@co.dap.packed` and `@co.dap.simd` are specialisations of `co.lang.cstruct` — same rules, same zone boundary safety. They are not separate types.

Structs

```
// myStruct.fol
_ co.lang.struct={
  field1 co.lang.int;
  field2 co.lang.string;
  field3 co.lang.bool;
}
```

Struct Rules

```
structs cannot extend other structs
structs cannot contain methods directly
structs can compose other structs
structs cannot have default values to fields/members
structs can embed other structs (Go lang like)
structs can have a same-package companion file named `<StructName>.comp.unit.fol`
```

The struct declaration remains pure data. Associated behaviour is declared separately in `<StructName>.comp.unit.fol`.

Struct Embedding

Embedding promotes fields of an embedded struct directly into the outer struct — they act as the outer struct's own fields at construction and access sites. This is distinct from composition where the embedded struct is a named field.

```
// E.fol
_ co.lang.struct = {
  id co.lang.int;
  name co.lang.string;
}

// ✅ No conflict – id and name promoted as B's own fields
// B.fol
_ co.lang.struct = {
  age co.lang.float;
  E; // embedded – id and name promoted
}

b := B{ age: 25.0, id: 1, name: "Rao" }; // all fields at same level
b.id // direct – no b.E.id needed
b.name // direct – no b.E.name needed
b.age // direct
```

```
// ❌ Compiler error — name conflict between B.name and E.name
// B.fol
_ co.lang.struct = {
  name co.lang.string; // conflicts with E.name
  E;
  age co.lang.float;
}
// Fix 1 — rename B's conflicting field
// Fix 2 — use explicit composition instead: e E;
```

```
// Explicit composition — no promotion, always qualified access
// B.fol
_ co.lang.struct = {
  name co.lang.string;
  e E; // named field — no conflict, no promotion
  age co.lang.float;
}

b.name ; // B's own name
b.e.id ; // E's id — always explicit
b.e.name; // E's name — always explicit
```

Embedding Rules

Situation	Behavior
Embedded field, no conflict	Promoted — acts as the outer struct's own field
Embedded field, name conflict with outer	❌ Compiler error — rename or use composition
Multiple embeds, no conflict between them	All fields promoted
Multiple embeds, conflict between embedded structs	❌ Compiler error
Explicit composition (<code>e E</code>)	No promotion — always accessed via <code>b.e.field</code>

FoLang does **not** silently shadow conflicting fields. Any name conflict is a compiler error — the programmer must make a conscious decision to rename or switch to explicit composition.

Struct Declaration Relationships

A struct cannot physically declare another struct, enum, class, module, function, signature, interface, or other named declaration inside its body. A struct body remains a pure data declaration containing fields and permitted embeddings only.

Use one of the following instead:

- **composition** through a named field
- **embedding** where the declaration kind's embedding rules permit it
- a separately declared target-local declaration restricted with `@co.dap.local`

```
// EmployeeAddress.fol
@co.dap.local(for=hr.employee.Employee)
_ co.lang.struct = {
  street co.lang.string;
  city co.lang.string;
}
```

```
// Employee.fol
_ co.lang.struct = {
  id co.lang.int;
  name co.lang.string;
  address EmployeeAddress; // composition
}
```

The following physical nesting is invalid:

```
// Employee.fol
_ co.lang.struct = {
  Address co.lang.struct = { // ❌ nested declaration
    city co.lang.string;
  }

  address Address;
}
```

structs cannot declare inner structs	❌	compiler error – only through @co.dap.local
structs can declare inner enums/ADTs	❌	compiler error – only through @co.dap.local or @co.dap.nested
structs cannot declare inner classes	❌	compiler error – struct is pure data
structs cannot declare inner modules	❌	compiler error – struct is pure data

@co.dap.local controls declaration visibility only. It does not automatically compose or embed the annotated declaration into the target struct.

Struct Companion Units

A struct companion unit is declared in a file named <StructName>.comp.unit.fol. The matching struct remains in <StructName>.fol, and both files must belong to the same package.

```
// Vector.fol
_ co.lang.struct = {
  x co.lang.float;
  y co.lang.float;
}
```

```
// Vector.comp.unit.fol
_ co.lang.unit = {
  distance(
    left Vector,
    right Vector
  )->(co.lang.float) = {
    ...
  }

  zero()->(Vector) = {
    this.return Vector{x: 0.0, y: 0.0};
  }

  (value Vector) magnitude()->(co.lang.float) = {
    this.return co.math.sqrt(
      value.x * value.x + value.y * value.y
    );
  }

  (Vector) create(
    x co.lang.float,
    y co.lang.float
  )->(Vector) = {
    this.return Vector{x: x, y: y};
  }
}
```

The filename establishes ownership. No companion member needs an ordinary parameter merely to prove association.

Call forms:

```
d      := Vector.distance(first, second);
origin := Vector.zero();
length := value.magnitude();
point  := Vector.create(10.0, 20.0);
```

A companion unit may declare:

- receiverless companion functions, including zero-parameter functions and factories
- instance-associated functions with an explicit value receiver
- type-associated functions with an explicit type receiver
- operator functions permitted for the owner struct
- non-UDT type declarations associated with the owner when their own sections permit them


```

receiverless non-operator companion function
  -> has no explicit receiver
  -> ordinary parameters may have any legal types
  -> no owner-typed parameter is required merely for association

instance-associated function
  -> explicit receiver is a value of the matching struct

type-associated function
  -> explicit receiver is the matching struct type itself

```

These are the general companion rules. Operator functions have the additional operand requirement defined in the Operator Functions subsection.

Companion declarations do not make the struct a class and do not add object identity, inheritance, virtual dispatch, lifecycle methods, or unit-level state.

Filename and Owner Validation

The compiler classifies companion files before parsing:

```
Vector.comp.unit.fol -> canonical owner Vector
```

At package-indexing level, the compiler verifies:

- a canonical primary declaration named `Vector` exists in the same package
- no second companion file resolves to the same canonical owner
- the companion is not orphaned

After the owner declaration header is available, the compiler verifies that the owner is a `co.lang.struct`. A class, cstruct, enum, union, module, interface, signature, or imported declaration with the same short name cannot become the owner.

Receiver Validation

The companion filename determines the required receiver root. Any explicit receiver must match that owner.

```

// Employee.comp.unit.fol
_ co.lang.unit = {
  create(id co.lang.int)->(Employee) = { ... } // valid: receiverless

  (emp Employee) isValid()->(co.lang.bool) = { ... } // valid

  (Employee) empty()->(Employee) = { ... } // valid

  (dept Department) isValid()->(co.lang.bool) = { ... }
  // compiler error: receiver Department does not match companion owner Employee

  (Department) create()->(Department) = { ... }
  // compiler error: type receiver Department does not match companion owner Employee
}

```

Ordinary parameters may have any legal types. They are not used to establish companion ownership.

For a generic struct owner, receiver validation compares the canonical root declaration and generic arity. For example, a receiver based on `Box(T)` may belong to `Box.comp.unit.fol`; `Box(T, E)` or an unrelated root does not.

Companion Namespace and Conflicts

Companion members are resolved through the owner:

```

employee := Employee.create(1);
valid    := employee.isValid();

```

The complete companion namespace is formed from all legal owner-associated declarations and the single companion file. Duplicate names and normalized duplicate signatures are compile-time errors.

The following are invalid:

```

Employee.comp.unit.fol without Employee.fol
Employee.comp.unit.fol and employee.comp.unit.fol in the same package
Employee.comp.unit.fol when Employee.fol declares a class or cstruct
an explicit receiver whose canonical root is not Employee

```

Operator Functions

Operator functions associated with a struct belong in its companion unit. Ownership comes from the companion filename; receiver validation follows the same rule as other companion functions.

```
// Employee.comp.unit.fol
_ co.lang.unit = {
  @co.dap.operator(symbol="+")
  (emp Employee) add(other Employee)->(Employee) = {
    ...
  }

  @co.dap.operator(symbol=="=")
  equals(
    left Employee,
    right Employee
  )->(co.lang.bool) = {
    ...
  }

  @co.dap.operator(symbol">")
  (Employee) greater(
    left Employee,
    right Employee
  )->(co.lang.bool) = {
    ...
  }
}
```

Companion ownership always comes from `<StructName>.comp.unit.fol`. An operator function has an additional operand rule: a value receiver already supplies the owner instance, but a receiverless operator function or a type-receiver operator function must declare the companion owner type as its first ordinary parameter. This allows operator lowering to pass the actual owner instance to the function. The compiler compares the resolved first-parameter type with the companion owner type at compile time. This requirement establishes the operator operand, not companion ownership. An operator declaration that does not operate on the owner struct is invalid.

Unions

`co.lang.union` declares an untagged ADT. Its body lists the alternative members defined by the union.

```
// myUnion.fol
_ co.lang.union={
  intValue co.lang.int;
  strValue co.lang.string;
}
```

Enums

```
// myEnum.fol
_ co.lang.enum={
  Variant1,
  Variant2,
  Variant3
}
```

An enum value's constant expression may use a registered custom operator at any declared precedence. Runtime assignment is forbidden everywhere in that constant-expression subtree, including inside grouping, arguments, and collection or object elements. Whether the resulting expression can actually be evaluated at compile time—including whether a custom operator is foldable—is checked after parsing.

classes

```
// Employee.fol
- co.lang.class = {
  getEmployeeDetails()->(Employee) = empmodule.getEmployeeDetails;
  // assigning module function to class's method

  getEmployeeInfo()->(Employee) =>> empmodule.getEmployeeDetails();
  // delegating - internally redirecting the call to module function
}

// $1, $2, $3 ... are return components of the immediately previous chained function
// Emp.fol
- co.lang.class = {
  dosomething(a co.lang.int, b co.lang.int)->(co.lang.int)=>>somePack.someMethod(a)=>>someOthPack.someOtherMeth($1, b);
}
```

Classes with Operator methods

```
// Employee.fol
- co.lang.class = {
  @co.dap.operator(symbol="+")
  add(other Employee)->(Employee) = {
    // implicit instance method: operands are this and other
  }

  @co.dap.operator(symbol=="")
  @co.dap.static
  equals(
    left Employee,
    right Employee
  )->(co.lang.bool) = {
    // static operator implementation
  }

  @co.dap.operator(symbol=">")
  @co.dap.class
  greater(
    left Employee,
    right Employee
  )->(co.lang.bool) = {
    // class-associated operator implementation
  }
}
```

An unannotated operator function in a class is an implicit instance method; the hidden `this` value contributes the first operand.

`@co.dap.instance` is the explicit spelling of the same category. `@co.dap.static` and `@co.dap.class` do not contribute an implicit operand, so their declared parameters are the complete operator operand list. Operator signature normalization includes these method categories so equivalent declarations are diagnosed as duplicates.

Class Declaration Relationships

A class cannot physically contain named class, struct, enum, module, function, signature, interface, or other declaration definitions as nested declarations. Class bodies contain fields, lifecycle declarations, and methods permitted by the class model.

Types and helper declarations used only by one class are declared in their ordinary legal source locations and restricted to that class with `@co.dap.local`:

```
// EmployeeAddress.fol
@co.dap.local(for=hr.employee.Employee)
- co.lang.struct = {
  street co.lang.string;
  city co.lang.string;
}
```

```
// EmployeeStatus.fol
@co.dap.local(for=hr.employee.Employee)
_ co.lang.enum = {
  Active,
  Inactive,
  Pending
}
```

```
// Employee.fol
_ co.lang.class = {
  address EmployeeAddress;
  status EmployeeStatus;

  getAddress()->(EmployeeAddress) = {
    this.return this.address;
  }
}
```

The local declarations are visible while compiling `hr.employee.Employee`, but they do not become nested names such as `Employee.EmployeeAddress` or `Employee.EmployeeStatus`.

The following is invalid:

```
// Employee.fol
_ co.lang.class = {
  Address co.lang.struct = { // ❌ physical nested declaration
    city co.lang.string;
  }
}
```

Ordinary visibility annotations do not widen a target-local declaration beyond the declaration or closed target set named by `@co.dap.local` or `@co.dap.nested`.

Method Types

```
// Employee.fol
_ co.lang.class = {

  @co.dap.static
  getEmployee()->(Employee) ={}

  @co.dap.instance
  getEmployee()->(Employee)={}

  @co.dap.class
  getEmployee()->(Employee) ={}

  @co.dap.object
  getEmployee()->(Employee)={}
}

@co.dap.oops(
  A: { inherit:true, virtual:true },
  B: { implements:true }, //interfaces
  C: { inherits:true, abstract=true },
  D: { inherits:true }, // inherits are classes abstract classes
  E: { uses:true }, // are extensions
  F: { composes:true }, //concreate classes
  G: { extends:true }, //mixins
  H: { with:true }, //traits
  I: { associate:true }, //concreate classes
)
// test.fol
_ co.lang.class = {
  getTest(id co.lang.int)->(test) ={}
}
```

The @@new and @@init Methods

`@@new` and `@@init` are lifecycle methods — compiler-owned, not user-definable outside the class. `@@` signals they are restricted lifecycle symbols, not regular methods.

Lifecycle names are valid only as class members. A unit (including a struct companion unit), module, interface, signature, local block, or package declaration cannot declare a lifecycle-named function.

```
// Employee.fol
@co.dap.generic(type={T:{typename}, R:{typename}})
_ co.lang.class = {

    id T;
    name R;

    // @@new is provided by default even if not overridden.
    // Override only when you genuinely need to change allocation behavior.

    @co.dap.class
    @co.dap.private
    @@new()->(co.lang.uninit) = { self.return co.const.none; }

    @co.dap.class
    @co.dap.public
    @@new(a co.lang.typevalue, b co.lang.typevalue)->(co.lang.uninit) = {
        // Manual type aliasing - @co.dap.generic handles this automatically
        // Override @@new only when you need custom allocation logic
        T co.lang.type = a;
        R co.lang.type = b;

        // self keyword is allowed only in class methods
        self.parent.new();

        // uninit instance method internally calls new and init
        self.return co.lang.uninit.instance(Employee, self);
    }

    @co.dap.override
    @co.dap.constructor(access=private)
    @@init() = {}

    @co.dap.override
    @co.dap.constructor(access=public)
    @@init(id T, name R) = {
        this.parent.init();
        this.id = id;
        this.name = name;
    }

    getEmployee(id T)->(Employee) = {}
}
```

Anonymous Classes/Types

```
emp := co.lang.class{};

empObj := emp.init();

empobj1 := co.lang.class{
    name co.lang.string;
}.init();
```

Interfaces

```
// IEmployee.fol
_ co.lang.interface = {
    storeEmployee(emp Employee)->(Employee);
}
```

Signatures

```
// Employee is an ordinary package-level declaration.
// MEmployee.fol
_ co.lang.signature = {
    storeEmployee(emp Employee)->(Employee);
}
```

Structurally they look similar — both are lists of contracts. The difference is **who implements them and how**.

	<code>co.lang.signature</code>	<code>co.lang.interface</code>
Implemented by	module via <code>matches=</code>	class via <code>implements=[]</code>
Number of implementations	Any number of distinct modules may match one signature	Any number of classes may implement one interface
Runtime cardinality of one implementation declaration	One module object per loaded module identity	Any number of independently constructed class objects
State model	One shared state for all references to the same module	Separate per-instance state for each class object
May specify required values	✓	✗ — methods only
May reference package-level types	✓	✓
May declare associated or fixed/manifest type components	✓	✗
May require generic type constructors	✓	✗
May declare physical nested/local types	✗	✗
May use <code>@co.dap.local</code>	✗	✗
Instantiation involved	Module is declared once, not constructed	Class objects are constructed
Reference use	Compatible module references may be used through the signature type without creating another module	Interface references may refer to any implementing object instance
OOP dispatch	✗	✓ virtual/dynamic
Contract style	module values, functions, and associated/fixed type components	behavioral methods on object instances
Practical analogy	singleton component contract	object-instance behavioral contract
Origin	ML/OCaml-inspired modules	Java/C#/Go interfaces

- A `signature` is a **module contract** over values, functions, existing package-level types, associated-type requirements, and fixed/manifest type components. A type-component specification is a contract slot, not a physical nested type definition. Multiple modules may match one signature, but each module declaration denotes one module object with shared module state.
- An `interface` is a **behavioral contract** tied to class dispatch and polymorphism. It cannot declare associated or fixed module type components or own nested type definitions. A class implementing an interface may create any number of independent runtime objects.
- The approximation `module + signature ≈ singleton object + interface` is useful for understanding cardinality and shared state, but a module is a language-level component rather than a class-based singleton pattern.

Modules

A module is an ML/OCaml-style abstraction governed by an optional signature. A module may use package-level types, satisfy associated-type requirements declared by its signature, and use fixed/manifest type components established by that signature. It does not physically own or nest arbitrary type declarations. A module should not be introduced merely to prevent functions from appearing loose in a file; use `co.lang.unit` for that simpler structural purpose.

```
// Employee.fol – ordinary package-level type
_ co.lang.struct = {
  Id    co.lang.int;
  Name  co.lang.string;
}

// EmployeeModule.fol
_ co.lang.signature = {
  getEmployee(id co.lang.int)->(Employee);
}

// EmployeeModImpl.fol
@co.dap.module(signature=EmployeeModule)
_ co.lang.module->(signature=EmployeeModule, matches=EmployeeModule) = {

  getEmployee(id co.lang.int)->(Employee) = {
    this.return Employee{
      Id: 10,
      Name: "Rao"
    };
  }
}

mm EmployeeModule = EmployeeModImpl;
v Employee = mm.getEmployee(10);
```

Usage and Liveness

Module, signature, interface, typeclass, instance, matcher, data-shape, and annotation/object liveness is defined centrally in [Unused Symbols, Liveness, and Reachability](#).

Module Cardinality and Singleton Analogy

A module declaration defines exactly one named module object for that loaded module identity. It is not an instantiable blueprint and does not create a new module object each time its name is referenced.

```
first EmployeeModule = EmployeeModImpl;
second EmployeeModule = EmployeeModImpl;
```

`first`, `second`, and `EmployeeModImpl` refer to the same module object. These bindings copy or retain the module reference; they do not clone or instantiate the module. Consequently, values declared directly by the module represent one shared module state for all references to that module.

A signature does not itself create a module object. It is a contract, and any number of separately declared modules may conform to the same signature:

```
DatabaseBackend signature
├ PostgreSQLBackend module -> one named module object
├ MySQLBackend module     -> one named module object
└ TestDatabaseBackend module -> one named module object
```

Each conforming module declaration contributes its own single module object and its own module state. The signature does not restrict the number of distinct conforming module declarations.

A useful analogy is:

```
signature      ≈ interface contract
conforming module ≈ singleton object implementing that contract
class          ≈ instantiable object type
```

The analogy is intentionally limited. A FoLang module is a compiler-recognized named component, not a class made singleton through a private constructor, static field, or runtime pattern. It cannot be repeatedly constructed. Because module references are first-class in FoLang, they may be bound and used through a compatible signature type, but every reference to the same module declaration still denotes the same module object.

Modules are also broader than ordinary interface implementations. A matching module may provide module values and functions, bind associated types required by its signature, and supply generic associated-type constructors where required. Fixed/manifest type components are established by the signature itself and are used by the module without being rebound. An interface constrains object behaviour; it does not provide the same module type-component abstraction.

By contrast, a class declaration defines an instantiable type. Every class construction creates a distinct object with independent identity, state, and lifetime:

PostgreSQLBackend module

└─ one shared module object and module state

PostgreSQLConnection class

└─ connection1 -> independent object and state

└─ connection2 -> independent object and state

└─ connection3 -> independent object and state

Formal mental model: A FoLang module is a single named implementation component that may conform to a signature. It is comparable to a singleton object implementing an interface, but it is not instantiated from a class. Multiple distinct modules may conform to the same signature, while each module declaration denotes one module object. Unlike an ordinary singleton-interface implementation, a module may also satisfy associated-type requirements, including generic associated-type constructors, and use fixed/manifest type components declared by its signature.

Module instantiation A FoLang class or struct declaration introduces an instantiable type but does not create an instance. A FoLang module declaration introduces one named module component directly into its package. The module name acts as the binding for that component, so no separate construction expression is required. The module's runtime state is initialized once according to the language's module-initialization rules.

A module declaration is a singleton component declaration and binding, rather than merely an object definition.

Module Signature Contents

A `co.lang.signature` is a declarative contract for a module. It may specify required module values and functions, associated-type requirements, and fixed/manifest type components. A signature does not allocate storage, initialize variables, execute statements, or provide function bodies.

A signature may contain:

- value specifications
- function signatures
- references to already existing accessible types
- associated-type specifications
- fixed or manifest type-component specifications
- generic associated-type-constructor specifications

A signature may not contain:

- value initializers
- executable statements
- function bodies
- concrete class, struct, enum, module, interface, or signature definitions
- arbitrary nested or target-local declarations

Associated and fixed/manifest type-component specifications are part of module conformance. They are contract slots, not Java-, C++-, or C#-style inner types, and they do not participate in `@co.dap.local`.

Value Specifications

A declaration such as:

```
// Counter.fo1
_ co.lang.signature = {
  count co.lang.int;
  increment(amount co.lang.int)->();
}
```

requires a matching module to provide a value named `count` of type `co.lang.int` and a compatible `increment` function. The signature does not initialize `count` and does not define `co.lang.int`; the built-in type already exists.


```
// CounterImpl.fol
_ co.lang.module->(
  signature=Counter,
  matches=Counter
) = {
  count co.lang.int = 0;

  increment(amount co.lang.int)->() = {
    count.value = count + amount;
  }
}
```

The same rule applies when a value or function specification uses an existing accessible package type:

```
// EmployeeRepository.fol
_ co.lang.signature = {
  current hr.employee.Employee;
  find(id co.lang.int)->(hr.employee.Employee);
}
```

The matching module must provide `current` and `find`. It does not redefine `hr.employee.Employee`.

Associated Type Components

An `associatedType` component declares that every matching module must supply a concrete type binding for that component. When the signature declaration has no initializer, it is an abstract associated-type requirement:

```
// Repository.fol
_ co.lang.signature = {
  Entity co.lang.associatedType;

  current Entity;
  find(id co.lang.int)->(Entity);
}
```

`Entity co.lang.associatedType;` does not define the representation of `Entity`. It declares an associated-type requirement named `Entity`. Every matching module must bind that requirement to a compatible existing type:

```
// EmployeeRepositoryImpl.fol
_ co.lang.module->(
  signature=Repository,
  matches=Repository
) = {
  Entity co.lang.associatedType = hr.employee.Employee;

  current Entity = ...;
  find(id co.lang.int)->(Entity) = { ... }
}
```

Within a matching module, `Entity co.lang.associatedType = ...` is an **associated-type binding**, not an arbitrary nested type declaration. Its name must correspond to an `associatedType` component declared by the matched signature. A module cannot use this form to introduce unrelated module-local types.

An associated-type requirement differs from `forward` and `extern` declarations:

```
associated type requirement
  -> each matching module supplies its own compatible type binding

forward declaration
  -> one specific declaration is completed later

extern declaration
  -> one specific declaration is defined in another linkage or compilation unit
```

Fixed or Manifest Type Components

A signature may fix a type component to an already known type:

```
// IntegerRepository.fol
_ co.lang.signature = {
  Id co.lang.type = co.lang.int;

  find(id Id)->(co.lang.bool);
}
```

Here `Id` is predetermined as `co.lang.int`. This is a fixed/manifest type component rather than an associated-type requirement: the signature itself establishes the type equality, so a matching module uses `Id` but does not bind or redefine it.

```
Entity co.lang.associatedType;      -> associated; matching module supplies the binding
Id co.lang.type = co.lang.int;     -> fixed/manifest; signature supplies the type equality
```

Abstract Generic Type Constructors

A signature may require a generic type constructor without defining its representation:

```
// StackSignature.fol
_ co.lang.signature = {
  Stack(T) co.lang.associatedType;

  empty(T)->(Stack(T));
  push(value T, stack Stack(T))->(Stack(T));
  pop(stack Stack(T))->(T, Stack(T));
}
```

`Stack(T) co.lang.associatedType;` declares a **generic associated-type component**, also described as an abstract type constructor of arity one. The signature specifies that `Stack` accepts one type argument, but it does not define the concrete type constructor represented by `Stack(T)`.

A matching module must provide a compatible type-constructor binding with the same name, arity, and declared constraints:

```
// ListStackModule.fol
_ co.lang.module->(
  signature=StackSignature,
  matches=StackSignature
) = {
  Stack(T) co.lang.associatedType = co.core.list(T);

  empty(T)->(Stack(T)) = { ... }
  push(value T, stack Stack(T))->(Stack(T)) = { ... }
  pop(stack Stack(T))->(T, Stack(T)) = { ... }
}
```

Another matching module may choose another representation:

```
// ArrayStackModule.fol
_ co.lang.module->(
  signature=StackSignature,
  matches=StackSignature
) = {
  Stack(T) co.lang.associatedType = collections.ArrayStack(T);
  ...
}
```

Therefore:

```
StackSignature
  -> requires a generic type constructor Stack(T)

ListStackModule
  -> binds Stack(T) to co.core.list(T)

ArrayStackModule
  -> binds Stack(T) to collections.ArrayStack(T)
```

An associated-type binding does not permit physical type nesting. When an implementation needs a new concrete struct, class, or enum representation, that declaration remains an ordinary package declaration and may be restricted to the implementing module with `@co.dap.local`; the module then binds the associated-type component to that declaration.

Signature Conformance Rules for Type Components

For every type component in a matched signature:

- each unbound `co.lang.associatedType` component must receive exactly one compatible binding from the matching module
- a fixed/manifest component must retain the type equality declared by the signature and must not be rebound by the module
- a generic associated-type binding must preserve generic arity, parameter kinds, bounds, variance, and other declared constraints
- component names must be unique within the signature
- associated-type bindings cannot contain executable code
- an associated-type binding that does not correspond to a component declared by the matched signature is a compiler error
- types referenced by value and function specifications must resolve after applying the module's associated-type bindings and fixed/manifest type equalities

Module Declaration Relationships

A module cannot physically declare nested structs, enums, classes, modules, signatures, interfaces, or other arbitrary named declarations. It references ordinary package-level declarations through its functions and signature. The only type-like declarations permitted directly in a matching module are `co.lang.associatedType` bindings that satisfy associated-type components declared by its matched signature; such bindings do not create independent nested declarations.

A declaration intended only for one module may be restricted with `@co.dap.local`:

```
// EmployeeModuleConfig.fol
@co.dap.local(for=hr.employee.EmployeeModImpl)
_ co.lang.struct = {
  timeout co.lang.int;
  retries co.lang.int;
}
```

```
// EmployeeModImpl.fol
_ co.lang.module = {
  connect(cfg EmployeeModuleConfig)->(co.lang.bool) = {
    ...
  }
}
```

The following remains invalid:

```
// EmployeeModImpl.fol
_ co.lang.module = {
  Config co.lang.struct = { // ❌ physical nested declaration
    timeout co.lang.int;
  }
}
```

A target-local declaration does not automatically become a module member name and is not projected through the module's signature. It becomes part of the signature view only when an associated-type component is explicitly bound to it or when a signature value/function specification references it through an allowed type component.

Structs vs Classes vs Modules vs Units vs Packages

	Struct	CStruct	Class	Module	Unit	Package
Purpose	Pure data shape	C-like value type	Behaviour + data	Signature-backed ML-style abstraction	Stateless package-fragment or struct-companion container	Folder-based grouping
Fields	✅	✅ simple only	✅ per instance	❌	❌	❌
Module-level values	❌	❌	❌	✅ when declared directly or required by a signature	❌	❌
Functions / methods	Companion functions through <code><StructName>.comp.unit.fol</code> ; explicit receivers must match the struct	❌	✅ methods	✅ module functions	✅ package functions in ordinary units; companion functions in companion units	❌

Lifecycle (<code>@@new</code> / <code>@@init</code>)	✗	✗	✓	✗	✗	✗
<code>this</code> / <code>self</code>	✗	✗	✓	✗	✗	✗
Value/literal construction	✓	✓	✗	✗	✗	✗
Lifecycle instantiation (<code>new</code> / <code>init</code>)	✗	✗	✓ multiple objects	✗ — one module object per declaration	✗	✗
Runtime state cardinality	Per bound struct object	Per value	Per class object	One shared state for the module declaration	—	—
First class	✓	✓	✓	✓ module reference; referencing does not instantiate	✗	✗
Pass by	Reference	Value	Reference	Reference to the same module object	—	—
Contract	—	—	<code>interface</code> via <code>implements=[]</code>	<code>signature</code> via <code>matches=</code>	none	—
OOP / inheritance	✗	✗	✓	✗	✗	✗
Physically nested independent named type/container declarations	✗	✗	✗	✗	✗	N/A — packages contain separate source declarations
May be an <code>@co.dap.local</code> target	✓	✗	✓	✓	✗	✗
Pattern matching	✓	✓	✓	✗	✗	✗
Direct ABI / zone boundary safe	✗ — library boundaries require snapshots	✓	✗	✗	✗	✗
Associated functions	✓ through <code><StructName>.comp.unit.fol</code>	✗	—	—	✓ only in a struct companion unit	✗
Embedding	✓	✗	—	—	✗	✗
Declared with	<code>co.lang.struct</code>	<code>co.lang.cstruct</code>	<code>co.lang.class</code>	<code>co.lang.module</code>	<code>co.lang.unit</code>	folder path
C++ backend analogy	struct without methods	plain C struct	class	struct/class abstraction	package namespace fragment or static companion scope	namespace
Closest mental model	Rust struct	C struct	Java/C# class	singleton implementation component with ML-style type members	source fragment merged into a package, or a filename-bound struct companion	filesystem namespace

Mental model:

```
reach for struct    → pure data; use `<StructName>.comp.unit.fol` for associated behaviour
reach for cstruct  → physical ABI-compatible value data crossing direct zone or native boundaries
reach for class    → behaviour, lifecycle, multiple instances
reach for module   → one named implementation component with shared state, governed by an optional signature and capable of
satisfying associated-type requirements
reach for unit     → package fragment (`*.unit.fol`) or struct companion (`*.comp.unit.fol`)
reach for package  → folder-based grouping only, not a value
```

Declaration scoping rule: FoLang does not permit physical nesting of independent file-backed primary declarations. Classes, structs, cstructs, enums, unions, modules, interfaces, signatures, instances, matchers, and other package-owned primary declarations remain in their own `<Name>.fol` files. Ordinary and companion unit files are explicit package containers: they may contain functions and the non-UDT type declarations permitted by the unit rules, but they may not contain independent primary declarations such as classes, structs, enums, modules, interfaces, or signatures. Ordinary local functions and anonymous expressions remain the other explicit nesting exceptions. Supported package declarations may restrict visibility to exact same-package targets with `@co.dap.local`; the annotation changes visibility, not physical ownership. Signature type components and matching-module `co.lang.associatedType` bindings are contract slots rather than arbitrary nested package declarations. ***

Local and/or Nested types and functions

FoLang does not provide Java-, C++-, or C#-style physical nesting of independent named type and container declarations. Such declarations remain in their ordinary legal source locations. Ordinary local functions and anonymous expressions are explicit exceptions governed by the rules below.

Physical Nesting Rules

Prohibited Independent Named Declarations

Independent file-backed primary declarations cannot be physically declared inside another class, struct, cstruct, enum, union, module, unit, interface, signature, function, or executable block. This includes:

- classes, structs, cstructs, enums, and unions;
- modules, interfaces, signatures, and additional units;
- instances, matchers and other file-backed primary declarations.

Non-UDT type declarations are the deliberate unit exception. Type aliases, `co.lang.type` ADTs and type constructors, newtypes, opaque types, refinement types, subtypes, and supertypes may be declared directly inside an ordinary unit, and inside a companion unit where their own rules permit association with the owner. Macros, templates, and decorators follow their own declaration rules. These declarations are not permitted loose at package-file scope or physically inside classes, structs, modules, functions, or executable blocks unless another section explicitly grants that context.

File-backed primary declarations retain package-owned identity and follow their normal `<Name>.fol` placement rules. An association or visibility annotation such as `@co.dap.local`, `@co.dap.nested`, or `@co.dap.inner` does not physically move a separately declared declaration inside its target.

Local-Function Exception

A function body may contain an ordinary named local or inner function where the grammar permits a local-function declaration:

```
outer()->() = {
  value co.lang.int = 10;

  inner()->() = {
    co.out.println(value);
  }

  inner();
}
```

The local function has block-local declaration identity. Its free runtime names are resolved from its lexical declaration context, and its lifetime and escape behavior follow the ordinary inner-function rules. It is not a package member and cannot be independently imported or exported.

This exception permits local functions only. It does not permit a named class, struct, enum, module, unit, interface, signature, or another named type/container declaration inside a function body.

Anonymous-Expression Exception

The physical-nesting restriction does not apply to anonymous constructs that are expressions or type expressions rather than independent named declarations. They may be nested wherever their specific grammar category is permitted.

These include:

- anonymous function expressions;
- lambdas and callback blocks;
- anonymous class or anonymous type expressions;
- permitted anonymous polymorphic type expressions introduced by `forall`;
- ordinary nested block, object-construction, map, collection, and other value-producing expressions.

```
process()->() = {
  operation := (value co.lang.int)->(co.lang.int) {
    this.return value * 2;
  };

  worker := co.lang.class {
    run(value co.lang.int)->(co.lang.int)={
      this.return operation(value);
    }
  }.init();
}

transformer co.lang.type = forall(T).(T)->(T);
```

An anonymous construct has no independently addressable package declaration identity. Its scope, capture, lifetime, type, and escape behavior are determined by the rules for that specific construct. Syntactic containment of an anonymous expression does not create a Java-, C++-, or C#-style named nested declaration and does not violate the one-primary-declaration-per-package-file rule.

Relationship to Association Annotations

The exceptions above are distinct from FoLang's association annotations:

- an ordinary local function is physically declared inside an enclosing function and is lexically scoped;
- an anonymous construct is an expression without independent declaration identity;
- `@co.dap.local`, `@co.dap.nested`, and `@co.dap.inner` apply to separately declared declarations that retain their normal source identity.

The remainder of this section defines those separately declared association forms.

```
@co.dap.local(for=<declaration-reference>)
```

or:

```
@co.dap.local(
  for=[
    <declaration-reference>,
    <declaration-reference>
  ]
)
```

The annotated declaration is called a **target-local declaration**. The declarations named by `for` form its **local target set**.

Supported Declaration and Target Kinds

`@co.dap.local` may be applied only to these declaration kinds:

- `co.lang.class`
- `co.lang.struct`
- `co.lang.enum`
- `co.lang.module`
- a named function declared in a context that normally permits functions

Every entry in the local target set must resolve to exactly one:

- class
- struct
- enum
- module
- function overload

A target list may contain any combination of supported target kinds.

Signatures and interfaces do not support target-local or physically nested declarations. A `co.lang.signature` or `co.lang.interface` cannot be annotated with `@co.dap.local` and cannot appear in a local target set. A signature may declare abstract, fixed, or generic **type-component specifications** as part of its module contract; these are contract requirements rather than local or nested type definitions. Interfaces do not declare signature type components.

Other declaration kinds are not target-local unless a later section explicitly permits them.

Target-Set Syntax

The `for` argument accepts either:

1. one declaration reference; or
2. a non-empty list of declaration references.

The scalar form is equivalent to a singleton target list:

```
@co.dap.local(for=hr.employee.Employee)
```

```
@co.dap.local(for=[hr.employee.Employee])
```

The scalar form is canonical when only one target is required. Use the list form when two or more declarations require access:

```
@co.dap.local(  
  for=[  
    hr.employee.Employee,  
    hr.employee.EmployeeService,  
    hr.employee.EmployeeValidator  
  ]  
)  
// EmployeeState.fol  
_ co.lang.enum = {  
  Active,  
  Inactive  
}
```

Rules:

- the target list must contain at least one declaration reference
- each reference must resolve independently and unambiguously
- duplicate references to the same resolved declaration are a compiler error
- list order does not affect visibility or declaration identity
- the target list is a closed set; access is not granted to declarations omitted from it
- repeated `@co.dap.local` annotations are not used to accumulate targets; use one annotation with one complete target list

Invalid:

```
// State.fol  
@co.dap.local(for=[])  
_ co.lang.struct = { ... }  
// ❌ empty target list
```

```
@co.dap.local(  
  for=[  
    hr.employee.Employee,  
    hr.employee.Employee  
  ]  
)  
// State.fol  
_ co.lang.struct = { ... }  
// ❌ duplicate resolved target
```

Declaration-Reference Syntax

Each `for` entry is a compiler-resolved declaration reference, not a string, runtime expression, function call, or `co.meta.symbol` value.

For a non-function target, use only its complete qualified declaration name:

```
@co.dap.local(for=hr.employee.Employee)  
@co.dap.local(for=hr.employee.EmployeeStatus)  
@co.dap.local(for=hr.employee.EmployeeRules)
```

For a function target, use its complete qualified function signature because FoLang permits overloads:

```
@co.dap.local(  
  for=hr.employee.Employee.calculate(co.lang.float)->()  
)
```

Function references in a target list follow the same rule:

```
@co.dap.local(  
  for=[  
    hr.employee.Employee.calculate(co.lang.float)->(),  
    hr.employee.Employee.calculate(co.lang.int)->(),  
    hr.employee.Employee.validate()->(co.lang.bool)  
  ]  
)  
// CalculationState.fol  
_ co.lang.struct = { ... }
```

Parameter names are not part of the reference:

```
// ✓ canonical  
hr.employee.Employee.calculate(co.lang.float)->()  
  
// ✗ parameter names are not declaration identity  
hr.employee.Employee.calculate(amount co.lang.float)->()
```

The parameter and return types must match one exact overload. An abbreviated function name, unresolved target, or ambiguous overload is a compiler error.

```
@co.dap.local(for=hr.employee.Employee.calculate)  
// ✗ ambiguous when calculate is overloaded
```

Each listed overload is an independent target. Listing one overload does not grant access to sibling overloads with the same function name.

Same-Package Requirement

The target-local declaration and **every declaration in its local target set** must belong to the same exact package. The compiler compares their complete folder-derived package identities; matching only a parent package, package family, import alias, library, or source root is not sufficient.

For a function target, the target package is the package that owns the function's enclosing class, struct companion unit, module, unit, or other legal function container. A target-local function is checked in the same way: the package containing its legal function-owning declaration must match the package of every listed target.

The invariant is:

```
package(target-local declaration)  
== package(target 1)  
== package(target 2)  
== ...
```

```
// hr/employee/Employee.fol  
_ co.lang.class = { ... }  
  
// hr/employee/EmployeeService.fol  
_ co.lang.class = { ... }  
  
// hr/employee/EmployeeState.fol  
@co.dap.local(  
  for=[  
    hr.employee.Employee,  
    hr.employee.EmployeeService  
  ]  
)  
// EmployeeState.fol  
_ co.lang.enum = { Active, Inactive }  
// ✓ all declarations belong to package hr.employee
```

A declaration in another package cannot participate in the local target set, even when ordinary package visibility, imports, aliases, parent/subpackage relationships, or library membership would otherwise make the declaration resolvable.


```
// EmployeeState is declared in package hr.internal
@co.dap.local(
  for=[
    hr.employee.Employee,
    hr.internal.EmployeeService
  ]
)
// EmployeeState.fo1
_ co.lang.enum = { Active, Inactive }
// ✗ local declaration and every target must have the same package
```

```
// CalculationState is declared in package hr.employee.internal
@co.dap.local(
  for=hr.employee.Employee.calculate(co.lang.float)->()
)
// CalculationState.fo1
_ co.lang.struct = { ... }
// ✗ subpackages are distinct packages; both sides must be exactly hr.employee
```

The same-package rule prevents `@co.dap.local` from becoming a cross-package friendship or visibility-bypass mechanism. A local declaration is a selectively shared implementation detail inside one package, not an imported helper attached from elsewhere.

Visibility Domain

A target-local declaration may be resolved only from:

1. each exact declaration in its local target set; and
2. lexical scopes contained within each listed target's implementation.

For targets `A`, `B`, and `C`, the visibility domain is the union of their implementation scopes:

```
visibility(local declaration)
= scope(A) ∪ scope(B) ∪ scope(C)
```

It is not an intersection, and code does not need to be simultaneously inside every target.

Consequences:

- a declaration local to a class is available to that class body and its methods
- a declaration local to a module is available to that module body and its functions
- a declaration local to a struct is available while resolving that struct declaration
- a declaration local to an enum is available while resolving that enum declaration
- a declaration local to a function is available only in the body of the exact targeted overload and its nested statement scopes
- when several targets are listed, each listed target receives access independently
- sibling declarations and sibling overloads not listed cannot resolve it
- subclasses, companion units, extensions, callees, callers, and related declarations do not gain access automatically
- visibility is not inherited, transitive, or propagated through calls, composition, embedding, imports, or other local declarations
- it cannot be imported, exported, or projected through a library surface

```

@co.dap.local(
  for=[
    hr.employee.Employee,
    hr.employee.EmployeeService
  ]
)
// EmployeeState.fol
_ co.lang.enum = {
  Active,
  Inactive
}

// Employee.fol
_ co.lang.class = {
  state EmployeeState; // ✅ listed target
}

// EmployeeService.fol
_ co.lang.class = {
  state EmployeeState; // ✅ listed target
}

// Payroll.fol
_ co.lang.class = {
  state EmployeeState; // ❌ not listed
}

```

Interaction with Ordinary Visibility

`@co.dap.local` establishes the final visibility boundary. Ordinary visibility annotations such as `@co.dap.public`, `@co.dap.package`, `@co.dap.protected`, or `@co.dap.private` cannot widen the declaration beyond its closed local target set.

```

@co.dap.public
@co.dap.local(
  for=[
    hr.employee.Employee,
    hr.employee.EmployeeService
  ]
)
// EmployeeState.fol
_ co.lang.enum = { Active, Inactive }

```

`EmployeeState` is still visible only to the listed targets. The ordinary visibility annotation is redundant for external resolution and may produce a compiler warning, but it never overrides `@co.dap.local`.

When most or all declarations in a package require access, ordinary package visibility should be used instead of maintaining a large target list.

Source Placement and Identity

A target-local declaration remains in the source position normally required for its declaration kind:

- a class, struct, enum, or module remains a normal primary declaration and follows the one-primary-declaration-per-package-file rule
- a function remains inside a unit, class, module, or another declaration context that normally permits functions
- `@co.dap.local` does not permit a free function to appear loose at package-file scope

The declaration retains a stable package-owned compiler identity. Its package identity must be exactly the same as the package identity of every target. It does not become physically nested and is not addressed as `Target.LocalName`.

```

hr.employee.EmployeeState
  local-for [
    hr.employee.Employee,
    hr.employee.EmployeeService
  ]

```

The normalized local target set is declaration metadata. Target-list order does not create a different declaration identity.

No Implicit Relationship or Privilege

`@co.dap.local` changes visibility only. It does not imply:

- composition
- embedding
- inheritance
- lifecycle ownership

- memory ownership
- friendship or privileged private-member access
- automatic membership in any listed target
- shared runtime state among the targets

Composition must still be written explicitly through a field or another supported relation. Embedding remains a separate facility: only struct and enum declarations are eligible for embedding, and each use must follow the embedding rules defined for that declaration kind.

Escape Restrictions

A target-local type must not leak outside its complete visibility domain.

It cannot appear in:

- a public, package, or protected API visible outside the local target set
- a library surface
- the parameter or return types of a function to which it is local
- a field or method signature that exposes it beyond the local target set
- an exported generic specialization or type alias

A value whose static type is target-local must be converted to an externally visible type before leaving the union of the listed targets' visibility domains.

Usage

```
@co.dap.public
@co.dap.local(
  for=[
    hr.employee.Employee,
    hr.employee.EmployeeService
  ]
)
// EmployeeState.fol
_ co.lang.enum = { Active, Inactive }

// Employee.fol
_ co.lang.struct={
  state EmployeeState;
}
```

Invalid Physical Nesting

```
// Employee.fol
_ co.lang.class = {
  Address co.lang.struct = { ... } // ✗
}

// EmployeeModule.fol
_ co.lang.module = {
  Config co.lang.struct = { ... } // ✗
}

process()->() = {
  State co.lang.enum = { Ready, Done } // ✗ named type declaration
}

// EmployeeContract.fol
_ co.lang.signature = {
  Employee co.lang.struct; // ✗
}

// EmployeeApi.fol
_ co.lang.interface = {
  Result co.lang.struct; // ✗
}
```

The `process` example rejects the named enum declaration; it does not reject ordinary local functions or anonymous expressions. Those are permitted only under the explicit exceptions in [Physical Nesting Rules](#).

Use separately declared package declarations, composition, embedding where allowed, and `@co.dap.local` when a closed set of declarations requires selective access.

There is another annotation `@co.dap.nested` which is similar to local but captures target state.

@co.dap.nested(target=hr.emp.Employee)

Instead of `for` we use `target` and `target` is always a single fully qualified type/function.

`@co.dap.nested` provides nested/inner-style target-state capture while the declaration itself remains physically separate, as defined above.

Comparison Table

Annotation	attribute	multiple targets	capture the target state
@co.dap.local	for	✓ as list for single target can mention without list syntax	✗
@co.dap.nested	target	✗	✓

@co.dap.inner Declarations

`@co.dap.inner` is an association annotation. It does **not** create a physically nested type, method, or function. The annotated declaration remains in its ordinary legal source location, retains its package-owned identity, and must satisfy the same exact-package rule that applies to `@co.dap.local` and `@co.dap.nested`.

Unlike `@co.dap.local` and `@co.dap.nested`, `@co.dap.inner` does not contain a `for` or `target` attribute. Its target relationship is established at the use site by explicitly embedding or attaching the declaration to a legal target. An `@co.dap.inner` declaration cannot be used as a standalone declaration.

Usage

```
@co.dap.public
// EmployeeState.fol
@co.dap.inner
_ co.lang.enum = {
    Active,
    Inactive
}

// Employee.fol
_ co.lang.struct = {
    EmployeeState;
    state EmployeeState;
}
```

The use of `EmployeeState;` establishes the inner association for this target. The annotation itself does not name the target.

Scope of @co.dap.inner Executable Declarations

For an `@co.dap.inner` function or method, parameters and local declarations resolve normally. A free runtime name is resolved from the **lexical context of the active call or attachment site**, not from the declaration site of the separately declared `@co.dap.inner` function.

The lookup order is:

1. parameters and local declarations of the `@co.dap.inner` function or method
2. the innermost lexical scope at the active call or attachment site
3. enclosing lexical scopes of that call or attachment site
4. statically resolved types, annotations, imports, and built-in `co.*` names
5. compiler error

This model is called **call-site lexical-context resolution**.

It is distinct from an ordinary inner function. An ordinary inner function is physically declared inside another function and resolves free runtime names from its enclosing lexical **declaration** context. An `@co.dap.inner` function is declared separately and therefore obtains its runtime context from the lexical scope in which it is actively attached or called.

It is also distinct from `@co.dap.dynamicscope`. Call-site lexical-context resolution does not search arbitrary runtime caller frames beyond the lexical scope chain of the active call site, and it does not use mixed-scope fallback. At every statically known use or call site, the compiler validates that each required runtime binding exists, is definitely initialized, has a compatible type, and permits the requested access or mutation.

For `@co.dap.inner` types and other non-executable declarations, ordinary static name and type resolution continues to apply. Call-site lexical-context resolution concerns only free runtime names used by executable function or method bodies.

An implementation may lower an `@co.dap.inner` declaration by copying, inlining, specialization, or another equivalent mechanism. Such lowering is not observable language semantics; the association and scope rules above are normative.

Exception

`@co.dap.inner`, `@co.dap.local`, and `@co.dap.nested` cannot be used with `co.lang.cstruct`.

Statements

A statement is a complete executable or declarative instruction. It may contain one or more expressions and may change program state, control execution, introduce declarations, or produce observable effects.

Common statement categories in `Folang`

1. Declaration Statement
2. Initialization Statement
3. Expression Statement
4. Conditional Statement
5. Loop Statement etc.,

Expressions

An expression is a construct that evaluates to a value, produces an effect, or both, and may be contained within a larger expression or statement.

Expression Evaluation Order

FoLang defines a deterministic evaluation order for expressions.

Expression structure is determined first by:

1. explicit grouping with parentheses;
2. operator precedence;
3. operator associativity.

After the expression structure has been determined, operands and subexpressions are evaluated from left to right in their source-code order, except where a construct explicitly defines conditional, lazy, asynchronous, concurrent, or otherwise specialized evaluation behavior.

Ordinary Expressions

For an ordinary expression:

```
result = first() + second();
```

the evaluation order is:

1. evaluate `first()`;
2. evaluate `second()`;
3. apply the `+` operator;
4. assign the resulting value to `result`.

Precedence and Evaluation Order

Left-to-right evaluation does not override operator precedence.

For example:

```
result = first() + second() * third();
```

operator precedence determines the expression structure as:

```
result = first() + (second() * third());
```

The function calls are evaluated in source order:

1. evaluate `first()`;
2. evaluate `second()`;
3. evaluate `third()`;
4. multiply the results of `second()` and `third()`;
5. add the multiplication result to the result of `first()`;
6. assign the final result to `result`.

Therefore, precedence determines how values are combined, while evaluation order determines when each operand or subexpression is evaluated.

Function Calls

For a function call, the callable expression is evaluated first, followed by the arguments from left to right.

```
result = calculate(first(), second(), third());
```

The evaluation order is:

1. evaluate the callable expression `calculate`;
2. evaluate `first()`;
3. evaluate `second()`;
4. evaluate `third()`;
5. invoke `calculate` with the resulting argument values;
6. assign the returned value to `result`.

Name resolution performed at compile time does not constitute runtime evaluation.

Method Calls

For a method call, the receiver expression is evaluated first, followed by the arguments from left to right.

```
result = getEmployee().calculate(first(), second());
```

The evaluation order is:

1. evaluate `getEmployee()`;
2. resolve the resulting object as the method receiver;
3. evaluate `first()`;
4. evaluate `second()`;
5. invoke `calculate`;
6. assign the returned value to `result`.

The receiver expression must be evaluated exactly once.

Member Access

For member access, the expression producing the containing object is evaluated before the member is accessed.

```
value = getEmployee().address.city;
```

The evaluation order is:

1. evaluate `getEmployee()`;
2. access `address` on the resulting object;
3. access `city` on the resulting address;
4. assign the resulting value to `value`.

Each intermediate receiver is evaluated only once.

Indexing

For an indexing expression, the indexed object is evaluated first, followed by index expressions from left to right.

```
value = getMatrix()[row()][column()];
```

The evaluation order is:

1. evaluate `getMatrix()`;
2. evaluate `row()`;
3. perform the first indexing operation;
4. evaluate `column()`;
5. perform the second indexing operation;
6. assign the resulting value to `value`.

Assignment

For assignment, FoLang evaluates the assignment target first, then evaluates the right-hand-side expression, and finally performs the assignment.

```
array[index()] = calculate();
```

The evaluation order is:

1. evaluate `array`;
2. evaluate `index()`;
3. determine the destination location;
4. evaluate `calculate()`;
5. assign the resulting value to the destination.

Determining the assignment target does not read the previous value stored at that location unless the assignment operation explicitly requires it.

Simple Variable Assignment

For assignment to a simple variable:

```
result = first() + second();
```

the evaluation order is:

1. determine the binding represented by `result`;
2. evaluate `first()`;
3. evaluate `second()`;
4. apply the `+` operator;
5. assign the resulting value to `result`.

Compound Assignment

A compound assignment evaluates its target exactly once.

```
array[index()] += calculate();
```

The evaluation order is:

1. evaluate `array`;
2. evaluate `index()`;
3. determine the destination location;
4. read the current value from that location;
5. evaluate `calculate()`;
6. apply the `+` operator;
7. assign the resulting value to the same destination.

The expression above must not behave as though it were expanded into a form that evaluates `array` or `index()` more than once.

Multiple Assignment

When an assignment contains multiple right-hand-side expressions, those expressions are evaluated completely from left to right before values are assigned to their targets.

```
a, b = first(), second();
```

The evaluation order is:

1. determine the target for `a`;
2. determine the target for `b`;
3. evaluate `first()`;
4. evaluate `second()`;
5. assign the first resulting value to `a`;
6. assign the second resulting value to `b`.

Right-hand-side evaluation must complete before any target receives its new value.

This permits value exchange without an explicit temporary variable:

```
a, b = b, a;
```

Operator Expressions

Built-in and user-defined operators follow the same operand-evaluation rules.

For a binary operator:

```
left() + right();
```

FoLang evaluates `left()` before `right()`.

For a prefix unary operator:

```
-operand();
```

FoLang evaluates `operand()` before applying the operator.

For a postfix unary operator:

```
operand()!;
```

FoLang evaluates `operand()` before applying the operator.

Operator overloading must not change the specified evaluation order of operands.

Short-Circuit Boolean Operations

Short-circuit Boolean operators evaluate the left operand first.

The right operand is evaluated only when it is required to determine the result.

For logical AND:

```
left() && right();
```

the evaluation order is:

1. evaluate `left()`;
2. when the result is false, return false without evaluating `right()`;
3. otherwise, evaluate `right()` and use its Boolean result.

For logical OR:

```
left() || right();
```

the evaluation order is:

1. evaluate `left()`;
2. when the result is true, return true without evaluating `right()`;
3. otherwise, evaluate `right()` and use its Boolean result.

A skipped operand produces no value, mutation, side effect, or error.

Conditional Expressions

A conditional expression evaluates its condition first and then evaluates exactly one selected result expression.

```
result = condition()  
    .then(whenTrue())  
    .default(whenFalse());
```

The evaluation order is:

1. evaluate `condition()`;
2. when the condition is true, evaluate `whenTrue()` only;
3. otherwise, evaluate `whenFalse()` only;
4. assign the selected result to `result`.

The unselected expression must not be evaluated.

Conditions and Branches

For a condition-and-branch construct, conditions are evaluated sequentially from left to right.

```
firstCondition().then({
  firstBranch();
}).otherwise(secondCondition()).then({
  secondBranch();
}).default({
  finalBranch();
});
```

The evaluation order is:

1. evaluate `firstCondition()`;
2. when true, execute `firstBranch()` and skip the remaining conditions and branches;
3. otherwise, evaluate `secondCondition()`;
4. when true, execute `secondBranch()` and skip the final branch;
5. otherwise, execute `finalBranch()`.

Only the selected branch is evaluated. `otherwise(condition)` is considered only after every earlier condition fails. If all conditional branches fail, `default(...)`, when present, supplies the terminal branch.

Pattern Matching

A pattern-matching subject is evaluated exactly once.

Cases are examined in source order unless a particular matcher explicitly defines another ordering rule.

```
getValue().match
  .case(firstPattern => firstResult())
  .case(secondPattern => secondResult())
  .default(defaultResult());
```

The evaluation order is:

1. evaluate `getValue()` exactly once;
2. test `firstPattern`;
3. when it matches, evaluate `firstResult()` and stop;
4. otherwise, test `secondPattern`;
5. when it matches, evaluate `secondResult()` and stop;
6. otherwise, evaluate `defaultResult()`.

Results belonging to unmatched cases are not evaluated.

Pattern guards are evaluated only after their corresponding structural pattern has matched.

Collection Literals

Elements of an array, list, tuple, set, map, or other collection literal are evaluated from left to right as they appear in the source.

```
values = [first(), second(), third()];
```

The evaluation order is:

1. evaluate `first()`;
2. evaluate `second()`;
3. evaluate `third()`;
4. construct the collection from the resulting values.

For a map entry, the key expression is evaluated before its corresponding value expression.

```
values = {
  firstKey(): firstValue(),
  secondKey(): secondValue()
};
```

The evaluation order is:

1. evaluate `firstKey()`;
2. evaluate `firstValue()`;

3. evaluate `secondKey()`;
4. evaluate `secondValue()`;
5. construct the map.

Range Expressions

For a bounded range, the lower-bound expression is evaluated before the upper-bound expression.

```
range = lower() .. upper();
```

The evaluation order is:

1. evaluate `lower()`;
2. evaluate `upper()`;
3. construct the range.

An omitted bound is not evaluated and does not produce an implicit function call or side effect.

Comprehensions and Iteration

Comprehension binding, result-shape, and source-specific behaviour are defined in [Comprehension Semantics](#). This section defines only their interaction with the general evaluation-order rules.

Within each iteration:

- source expressions are evaluated in the order required by the applicable comprehension or iteration form;
- when an explicitly defined filter is present, its condition is evaluated before the corresponding result expression;
- a result expression is evaluated only for a source value that participates according to the source's defined semantics and any explicit filter;
- individual operand evaluation remains left to right.

The core `for (pattern <- source).yield(...)` form does not itself introduce an inline filter. The language does not implicitly evaluate comprehension iterations concurrently unless concurrency is explicitly requested or the selected source semantics explicitly define another execution model.

Lazy Expressions

An expression declared lazy is not evaluated when the lazy binding is created.

Its evaluation occurs only when demanded according to the rules of the corresponding lazy construct.

Once evaluation begins, the expression itself follows the normal FoLang evaluation-order rules unless the lazy construct specifies otherwise.

The specification for each lazy construct must also state whether its result is:

- evaluated once and cached;
- evaluated again for every demand;
- safe for concurrent demand;
- permitted to propagate errors more than once.

Asynchronous and Concurrent Expressions

Left-to-right evaluation determines the order in which asynchronous or concurrent operations are created, invoked, or submitted.

It does not necessarily determine the order in which independently executing operations complete.

```
submit(first());  
submit(second());
```

FoLang evaluates and submits `first()` before `second()`. However, completion order depends on the execution-model rules unless explicit ordering is requested.

Concurrent execution never arises implicitly from an ordinary expression. It must be introduced by an explicit FoLang concurrency, parallelism, asynchronous-execution, scheduling, or execution-model construct.

Errors During Evaluation

When evaluation of a subexpression produces an error that prevents further evaluation, later subexpressions are not evaluated.

```
result = first() + failing() + third();
```

When `failing()` terminates evaluation with an error:

1. `first()` has already been evaluated;

2. `failing()` produces the error;
3. `third()` is not evaluated;
4. no final value is assigned to `result`.

The applicable error-handling rules determine whether the error is propagated, matched, converted, recovered from, or terminates execution.

Side Effects

All observable side effects occur according to the defined evaluation order.

Observable effects include:

- object mutation;
- variable rebinding;
- input and output;
- file-system operations;
- network operations;
- synchronization;
- exception or error generation;
- interaction with foreign code;
- interaction with the runtime environment.

An implementation must not reorder operations when doing so could change any observable effect.

Compiler Optimizations

A compiler may optimize, combine, eliminate, or internally reorder operations only when the transformation preserves all behavior required by this specification.

In particular, an optimization must not change:

- the resulting value;
- the order or presence of observable side effects;
- the identity or aliasing behavior of objects;
- the timing or presence of specified errors;
- the selected conditional or pattern-matching branch;
- the number of times an effectful expression is evaluated;
- synchronization or concurrency guarantees.

Pure internal operations may be reordered when no conforming FoLang program can observe the difference.

Implementation Requirement

Every conforming FoLang implementation must preserve the evaluation order defined in this section.

An implementation may use any parser, intermediate representation, optimizer, runtime, or backend, but those implementation choices must not alter the externally observable behavior required by these rules.

Operators

FoLang distinguishes a **symbolic spelling** from an **operator**. A sequence of symbol characters is not automatically an operator merely because it resembles one. Its syntactic role is determined by the grammar context in which the whole spelling occurs.

For example:

```
a co.lang.int->(**);
```

In this declaration, `->` is the structural type-derivation marker and `**` is pointer-degree metadata inside the derivation. Neither occurrence is parsed as an expression operator. The same `**` spelling in `left ** right` is parsed as the registered power operator because it occurs in an operator-expression position.

FoLang uses one uniform implementation model for every expression operator. The difference between built-in, pre-declared, and project-local custom operators is only who registers the symbol and its parse properties.

```

built-in operator
  symbol and parse properties registered by the language
  one or more implementations may already exist

pre-declared operator glyph
  symbol and parse properties registered by the language
  no implementation is required to exist

project-local custom operator
  symbol and parse properties registered once in src/lib/operators/library.fol
  no implementation is contained in the operator declaration

all three categories
  implementations are ordinary mode=override operator functions
  duplicate normalized operand signatures are errors

```

Symbolic Runs, Classification, and Boundaries

After comments, literals, and closed scanner-known composite spellings such as `@@new` and `@@init` are recognized, the lexer consumes each remaining complete contiguous run of symbol characters as one symbolic token candidate. It does not backtrack or split an unrecognized run into shorter valid tokens. Whitespace, comments, and delimiters such as `(`, `)`, `[`, `]`, `{`, and `}` end a symbolic run. Comment openers are recognized before ordinary symbolic-run scanning.

The complete run is then classified by its grammar context as one of the following:

1. a fixed structural or declaration spelling such as `->`, `=>`, `:=`, or `?=`;
2. a contextual metadata spelling, such as a contiguous run of one or more `*` characters inside `->(…)` to express pointer degree;
3. a registered expression operator; or
4. an unrecognized symbolic token, which is a compilation error.

There is no fallback splitting. Therefore, when `++` is not registered:

```

++ a;      // error: unrecognized symbolic token `++`
a ++ b;    // error: unrecognized symbolic token `++`
a++b;      // error: the contiguous run is still the single candidate `++`

+ +a;      // valid: two `+` operators separated by whitespace
+(+a);     // valid: `(` creates the token and operand boundary
a + +b;    // valid: binary `+` followed by unary `+`

```

A registered expression operator whose spelling contains more than one symbol character must have an explicit boundary on every operand-facing side. A boundary may be whitespace, a comment, or an applicable delimiter. Boundary presence is checked from the original source before whitespace and comments are discarded. The rule is based on the operator's fixity:

- an infix multi-symbol operator requires a boundary before and after it;
- a prefix multi-symbol operator requires a boundary after it;
- a postfix multi-symbol operator requires a boundary before it.

Suppose `+-` is registered as an infix operator:

```

a +- b;      // valid
(a)+-(b);    // valid: the parentheses provide both operand boundaries
a+-b;        // error: missing boundaries
a +-b;       // error: missing the right boundary
a+- b;       // error: missing the left boundary
a + -b;      // valid: separate `+` and unary `-` operators
a + (-b);    // valid: the parenthesis separates the operators

```

This boundary requirement applies uniformly to built-in, pre-declared, and custom multi-symbol expression operators. It does not apply merely because a structural token or metadata spelling contains multiple symbols. Thus `co.lang.int->(**)` remains valid without spaces around `->` or inside the pointer metadata.

The statement-level definition spellings `:=` and `?=` are not expression operators, but they deliberately use the analogous two-sided boundary rule. Write `name := value` and `name ?= value`; compact forms such as `name:=value` and `name?=value` are invalid.

Operator `mode=override` and `mode=extends` are unsupported. Ordinary class method overriding through `@co.dap.override` remains a separate class feature.

Operator Implementations

An operator implementation is a normal function with operator metadata. It cannot be declared loose at package scope. Every implementation must be in one of these legal function-owning locations:

Operand owner	Required implementation location
built-in type	an <code>@co.dap.extension</code> function inside a unit
<code>co.lang.struct</code>	the struct's same-package companion unit
<code>co.lang.class</code>	an operator method declared by the class
module, enum, union, interface, signature, <code>co.lang.cstruct</code>	unsupported

An implementation uses `mode=overload`, or omits `mode` because omission is shorthand for `mode=overload`. A one-character symbol may use a character literal; a multi-character symbol must use a string literal:

```
@co.dap.operator(symbol='+', mode=overload)
add(left Employee, right Employee)->(Employee) = {
    ...
}

@co.dap.operator(symbol="+-", mode=overload)
combine(left Employee, right Employee)->(Employee) = {
    ...
}
```

FoLang does not attempt to prove that an overload follows the conventional meaning of its symbol. A developer may overload `+` with any implementation. Using one symbol consistently for one recognizable concept is recommended for readability, but it is not a compiler-enforced semantic restriction.

For a receiverless struct-companion operator function, the first declared operand must have the matching struct type. A matching struct type in a later operand is insufficient. For a unary operator, the sole operand must have the matching struct type.

A matching instance receiver contributes the first operand. A matching type receiver establishes ownership but contributes no operand; its ordinary parameter list is the complete operator signature. In a class, an ordinary or `@co.dap.instance` operator method has an implicit `this` first operand, while `@co.dap.static` and `@co.dap.class` operator methods use only their declared parameters and require the first declared operand to have the enclosing class type.

After receiver normalization, the operand count must match the registered arity of the operator. Each distinct normalized operand signature is an overload. A second implementation with the same normalized signature is a duplicate and is rejected. This rule applies equally to built-in operators, pre-declared glyphs, and project-local custom operators.

Operator Type Resolution, Conversion, and Overflow

FoLang does **not** perform C-, C++-, or Java-style automatic numeric promotion or implicit casting as part of operator overload resolution. Operator resolution uses the operand types actually present in the expression. An overflow condition does not cause the compiler or runtime to silently widen an operand or select a different numeric operator overload.

An operator implementation has its own declared result type. After that operator has been selected, an enclosing declaration, assignment, argument, or return position may require the produced value to satisfy another target type. That is handled by FoLang's ordinary conversion model, not by operator promotion.

The language-provided simple types, such as `co.lang.int` and `co.lang.float`, provide standard `to` and `from` conversion methods through their ordinary package APIs. When the required target type has an applicable conversion and the value is representable by that target, FoLang uses that conversion to satisfy the target context. The operator overload that produced the value is not changed by that conversion.

User-defined types may participate in the same model by supplying supported `to` and/or `from` conversions through the extension mechanisms defined by FoLang. A conversion provided by an extension is an ordinary type conversion; it does not create automatic promotion between the operands of an operator expression.

Consequently:

```
operand types
-> exact operator-overload resolution
-> operator result type
-> optional target-context conversion through to/from
```

Overflow never implies promotion. If a declared result/target context provides an applicable conversion to a type capable of representing the produced value, that conversion handles the value at the target boundary. Otherwise the selected operator and result type retain responsibility for their own overflow, division-by-zero, range, and other value-level arithmetic behavior.

For the standard language-provided simple types, arithmetic, ordering, equality, logical, and bitwise operators follow their conventional meanings. `!=` is the logical negation of `==`; `&&` and `||` are short-circuit Boolean operators; and `&`, `|`, and `^` are the ordinary bitwise operators for operand types that support those operations. Exact exceptional-value behavior belongs to the selected type/operator implementation rather than to parser precedence or operator-token classification.

Language-Owned Operators

Language-owned operators already have a registered symbol, fixity, precedence, associativity, and arity. They must not be redeclared in the project operator source. They receive additional implementations through `mode=overload`.

This category includes:

1. ordinary built-in operators such as `+`, `-`, `*`, and `==`;
2. pre-declared operator glyphs whose parse properties are fixed by the language but which may initially have no implementation.

Pre-Declared Operator Glyphs

The current alpha profile pre-declares exactly two mathematical operator glyphs: `⊔` and `⊗`. Their parse properties are fixed by the language and are listed in C.9: both are binary infix operators with precedence `500` and left associativity. These glyphs are language-owned and therefore cannot be redeclared with `co.lang.operator`.

Unlike a reserved-but-unimplemented operator spelling, `⊔` and `⊗` are enabled expression operators in the current alpha profile. The lexer recognizes them as registered operator tokens and the parser applies their language-defined fixity, precedence, associativity, and arity. Their concrete behavior is supplied through ordinary `mode=overload` operator implementations using the same ownership, signature-normalization, and overload-resolution rules as other language-owned operators.

A use of `⊔` or `⊗` therefore parses as an operator expression even when no applicable implementation exists. If overload resolution finds no matching implementation for the operand types, compilation fails during operator resolution rather than during lexing or parsing.

See C.10 for the normative current-alpha rule.

Hard-reserved spellings such as `:=`, `->>`, `<->`, backtick, backslash, `#`, and comment openers are different: they are not overloadable or declarable unless a later language revision explicitly assigns them operator semantics.

Project-Local Custom Operator Source

A custom operator is a symbol that is neither language-owned nor hard-reserved. Its symbol and parse properties are registered only in the fixed project-local operator bootstrap surface:

```
<project-root>/src/lib/operators/library.fo1
```

`src/lib/` and `src/lib/operators/` are not packages. `operators/` is the one standardized operator slot under `src/lib/`. If it is absent, the project introduces no project-local custom operator symbols. If it is present, it must contain exactly one file named `library.fo1` and no additional files or subdirectories. The fixed `library.fo1` name is shared by all `src/lib/` library slots; the enclosing `operators/` directory selects the dedicated operator bootstrap semantics.

The fixed file is parsed by the dedicated operator-source lexer and parser before the ordinary FoLang lexer and parser run. Its filesystem position already establishes the operator bootstrap context, so no library-kind annotation is used:

```
// src/lib/operators/library.fo1
_ co.lang.library = {

  ⊗ co.lang.operator = {
    fixity: co.operator.fixity.infix,
    precedence: 60,
    associativity: co.operator.associativity.left,
    arity: co.operator.arity.binary,
    commutative: co.const.false,
    idempotent: co.const.false,
    identity: co.const.none,
    foldable: co.const.false,
    vectorizable: co.const.false,
    distributes_over: [],
    desugar: "intrinsic:tensor_product"
  };

  +- co.lang.operator = {
    fixity: co.operator.fixity.infix,
    precedence: 60,
    associativity: co.operator.associativity.left,
    arity: co.operator.arity.binary
  };
}
```

`_ co.lang.library` identifies the structural operator bootstrap surface. It is not an ordinary importable library surface and does not produce a `.fo1enc` artifact. Its body may contain only `co.lang.operator` declarations. Imports, functions, types, variables, expressions, implementation packages, and nested libraries are forbidden.

`co.lang.operator` is valid only in this dedicated source grammar. It is not an ordinary FoLang declaration kind and cannot appear in package source, `src/app1.fol`, or an ordinary library surface.

Operator declaration attributes

Every custom declaration must contain each required parse attribute exactly once. Optional metadata may appear at most once. Duplicate and unknown attributes are errors.

Attribute	Required	Accepted value	Meaning
<code>fixity</code>	yes	<code>infix</code> , <code>prefix</code> , <code>postfix</code> , or a reserved future fixity name (<code>circumfix</code> , <code>postcircumfix</code> , <code>precircumfix</code> , etc)	parser placement
<code>precedence</code>	yes	decimal integer from <code>0</code> through <code>100</code>	binding strength
<code>associativity</code>	yes	<code>left</code> , <code>right</code> , or <code>none</code>	grouping for equal precedence
<code>arity</code>	yes	<code>unary</code> , <code>binary</code> , <code>ternary</code> , or a positive decimal integer	normalized operand count
<code>commutative</code>	no	<code>co.const.true</code> or <code>co.const.false</code>	semantic/optimization metadata
<code>idempotent</code>	no	<code>co.const.true</code> or <code>co.const.false</code>	semantic/optimization metadata
<code>identity</code>	no	a FoLang literal, including <code>co.const.none</code>	identity-value metadata
<code>foldable</code>	no	<code>co.const.true</code> or <code>co.const.false</code>	compile-time folding metadata
<code>vectorizable</code>	no	<code>co.const.true</code> or <code>co.const.false</code>	vectorization metadata
<code>distributes_over</code>	no	a list of quoted operator symbols	distributivity metadata
<code>desugar</code>	no	a string literal	intrinsic or lowering hook metadata

In the alpha profile, only `infix`, `prefix`, and `postfix` are implemented. Consequently, an alpha `prefix` or `postfix` declaration must be unary and an alpha `infix` declaration must be binary. Other fixity names remain reserved for future delimiter and slot grammars and are rejected by the alpha conformance profile.

The four required parse attributes construct the tokenizer and precedence table. Optional semantic/optimization attributes do not change tokenization or parsing.

Declaration and implementation are separate

The operator library registers only the symbol and metadata. It contains no implementation. Implementations use the same `mode=overload` form as built-in and pre-declared operators:

```
// vector/Vector.comp.unit.fol
_ co.lang.unit = {

  @co.dap.operator(symbol='⊗', mode=overload)
  tensorProduct(left Vector, right Vector)->(Vector) = {
    ...
  }
}
```

A custom symbol has exactly one declaration in the operator library, but it may have any number of distinct implementation overloads in legal owners:

```
one ⊗ declaration
Vector ⊗ Vector -> Vector
Matrix ⊗ Matrix -> Matrix
Tensor ⊗ Tensor -> Tensor
```

The declaration cannot be duplicated, aliased, merged, selected, or remapped. The implementations participate in ordinary operator overload resolution.

Symbol registration and exact recognition

The dedicated operator-source lexer reads the declaration name as one maximal contiguous symbol run. After the operator table is built, the ordinary lexer uses the same whole-run rule. A registered custom symbol is recognized only when the complete run matches that symbol; an unknown run is rejected without being split into shorter operators.

Multi-symbol operator uses must also satisfy the operand-boundary rule defined in [Symbolic Runs, Classification, and Boundaries](#). Consequently, registering `++` permits `a ++ b`, not `a++b`. Unicode is not required for custom operators.

A custom symbol:

- must not be language-owned;
- must not be a hard-reserved spelling;
- must not contain `//` or `/*`, because a comment opener terminates the preceding symbolic run and is recognized before further operator matching;
- must have exactly one declaration in the operator library.

Symbols are global; implementations are resolved by scope and type

Once a custom symbol is registered—or a language-owned symbol is loaded—the ordinary tokenizer recognizes it throughout that compilation. The parser can therefore build the same operator expression in every package.

Whether an implementation is available is determined later through ordinary operator resolution. Receiver-owned class and companion implementations follow their normal lookup rules. Extension implementations require normal `@co.ddap.use` activation. A symbol with no matching visible implementation produces a resolution error, not a syntax error.

Using one symbol for one recognizable concept across its overloads is strongly recommended for readability, but the compiler does not and cannot enforce that recommendation. The type/signature rules, not the conventional meaning of the symbol, determine whether an overload is legal.

Operators Do Not Cross a Library Boundary

Operator notation is compilation-domain-local syntax, not a library API element. An ordinary library surface exports named function signatures and boundary data contracts. It exports no operator table.

```
geometry.folenc
├─ projected symbol table
└─ compiled implementation/linkage      (no operator table)
```

The project-local table produced from `src/lib/operators/library.fol` belongs to the owning project's open package compilation domain. That domain includes ordinary package source under `src/` and, when present, the selected open package source under `src/lib/exports/`. This is true for application, standalone packaged-library, and standalone export projects. The table is never reused while parsing or resolving isolated source-library domains under `src/lib/application/`, `src/lib/ffi/`, `src/lib/system/`, `src/lib/advanced/`, or `src/lib/dynamicvmrt/`.

Project-local `ffi`, `system`, and `advanced` source libraries cannot declare custom operators or operator overload implementations. They are compiled against the language-owned operator set.

An application project and a standalone export project may have their own optional project-local `src/lib/operators/library.fol`. A standalone packaged library may do so when its declared library kind permits project-local custom operators. For a standalone library project, the frontend shallow-reads the fixed `src/lib.fol` surface header early enough to obtain `@co.dap.library(type=...)` before validating whether the operator bootstrap is permitted for that library kind.

That operator table belongs only to the producer's open package compilation domain. Its operator expressions are fully resolved and lowered before a library- or export-kind `.folenc` is emitted. The resulting `.folenc` does not contain an importable parser operator table, and loading it cannot add operator spellings or parse properties to the consumer.

A project whose `src/` code wants operator notation for an imported boundary type must register the symbol in its own `src/lib/operators/library.fol` when necessary and provide a legal local implementation that delegates to the library's named public function.

Parallel Frontend Preparation, Imports, and Reachability

FoLang projects use eager preparation for performance and explicit import edges for semantics. Eagerly parsing or deserializing a project component does **not** make its symbols visible in another context.

Eager parallel preparation

After fixed-layout validation, the frontend may prepare independent domains in parallel:


```

project discovery
|
+-- src/lib/operators/library.fol
|   -> dedicated operator parser
|   -> ProjectOperatorTable
|
+-- src/lib/application/**
|   -> ASTs + isolated library contexts/symbol tables
|
+-- src/lib/ffi/**
|   -> ASTs + isolated library contexts/symbol tables
|
+-- src/lib/system/**
|   -> ASTs + isolated library contexts/symbol tables
|
+-- src/lib/advanced/**
|   -> ASTs + isolated library contexts/symbol tables
|
+-- src/lib/dynamicvmrt/**
|   -> ASTs + isolated library contexts/symbol tables
|
+-- src/lib/exports/**
|   -> open package contexts; parsed with owning ProjectOperatorTable
|
+-- lib/*.folenc
|   -> parallel protobuf deserialization
|   -> library-surface or exported-package context/symbol metadata

```

The isolated source-library domains do not depend on the owning project's custom operator table, so they may be parsed concurrently with the operator bootstrap. `src/lib/exports/` is different: it is open package source and uses the owning project's active operator table. Once that table is complete when applicable, `src/lib/exports/`, the fixed project surface under `src/`, and package source below `src/` may be parsed, including in parallel where implementation scheduling permits.

Source parsing produces ASTs and canonical context/symbol-table entries. Deserializing a `.folenc` does not reparse source. A library-kind artifact reconstructs its projected surface context; an export-kind artifact reconstructs its selected package contexts and exported symbol/type metadata.

One canonical symbol table; imports store references

The frontend keeps one canonical context/symbol-table representation for each prepared package, exported package, or library surface. An import does not copy or merge the target symbol table into the importer.

Conceptually:

```

ContextMap
├─ application-entry
├─ hr.employee
├─ ffi-surface
├─ thirdparty.text
├─ crypto-folenc-surface
└─ ...

application-entry.ImportedContexts
├─ "ffi"          -> ffi-surface ContextId
├─ "thirdparty.text" -> exported package ContextId
└─ "crypto"       -> crypto-folenc-surface ContextId

```

When the compiler resolves an import declaration it:

```

resolve import identity
  -> locate the already-prepared ContextId / SymbolTableId
  -> add a provisional reference to the importing context's ImportedContexts

```

The target symbol table remains canonical in the global maps. Import visibility is therefore represented by context references, not by symbol copying.

The import becomes semantically live only when name/type resolution actually consumes at least one symbol through that imported context:

```

resolve symbol through imported context
  -> mark exact SymbolId used
  -> mark the import used
  -> activate the dependency edge
  -> make the target context reachable

```

Using one imported symbol does not mark sibling symbols in that target context used. At finalization, every provisional import with zero successfully used symbols is removed from the effective `ImportedContexts` /dependency graph and reported as an unused-import compile-time error.

Unused Symbols, Liveness, and Reachability

This section is the **single canonical definition** of FoLang unused-symbol, declaration liveness, import liveness, project reachability, project-local source library/export usage, standalone library/export producer usage, and packaged `.fo1enc` consumer usage. Other sections define syntax and mechanics and link here rather than restating these rules.

Core Meaning

FoLang distinguishes availability from semantic participation:

```
parsed / loaded / deserialized
  != used
  != live
  != reachable
```

- **used** — semantic resolution actually consumes the declaration or operation, or a language-defined contract/protocol rule establishes its required liveness;
- **live** — the declaration or project entity satisfies the usage condition in the matrix below;
- **reachable** — the package, source-library domain, internal package, or packaged artifact is connected to the applicable project root through effective dependencies;
- **unused** — a usage-checkable project-owned declaration remains non-live after semantic resolution reaches a fixed point;
- **orphan** — physically participating project content remains unreachable after the effective dependency graph is constructed.

Parsing a declaration, importing a package/library, loading a `.fo1enc`, preparing a symbol table, naming a declaration without exercising the required semantic use, or merely making something visible does not by itself establish liveness.

Governing Principle

FoLang checks **behavior, not shape**.

A function or method has to be written, understood, tested, and maintained. An orphan one is cost with no return, so callables are usage-checked. A type or declaration is lighter, but someone must still know it exists, so it is checked as a whole. A field, variant, or member is a slot in a data shape: it carries no logic to test or reason about, so it is not checked at all.

Everything below follows from that. It is also why a contract-required member is exempt: it is not orphan behavior, because the contract needs it and the developer could not remove it. And it is why a function-local variable is exempt: a local is not behavior anyone can call, test, or depend on — it is scratch inside behavior that is already checked.

FoLang applies three member-level rules:

```
developer-chosen behavior
  -> independently usage-checked

contract/protocol-required behavior
  -> live by conformance once the implementation itself is live

data/object shape
  -> declaration/type is usage-checked;
      fields, members, or variants are not checked independently
```

A callable the developer could remove without breaking a declared contract is independently usage-checked. A callable that must exist because a contract or protocol requires it does not need an artificial separate call merely to avoid an unused-symbol error.

Complete Usage and Liveness Matrix

Entity / declaration	What makes it live / valid	What is independently checked	Unused result
class	class/type participates in the live semantic graph	every developer-authored method not required by an implemented interface	unused class or unused freely-authored method = compile-time error; fields are not checked
module	module participates as a live module value/declaration	every developer-authored function/method not required by its signature	unused module or unused freely-authored member = compile-time error; state values are not checked
struct	struct type is semantically used	declaration/type only	unused struct = error; fields are not independently checked

cstruct	cstruct type is semantically used	declaration/type only	unused cstruct = error; fields are not independently checked
union	union type is semantically used	declaration/type only	unused union = error; individual union members are not checked
enum	enum type or any enum variant is semantically used	declaration/type only	unused enum = error; individual variants are not checked
user-defined object	object declaration is semantically referenced	declaration as a whole	unused object = error; object members are not independently checked
user-defined annotation object	annotation is actually applied in source	declaration as a whole	never-applied annotation object = error; annotation fields are not independently checked
matcher	matcher is actually selected in a live <code>.match(...)</code> chain	declaration as a whole	unused matcher = error; required <code>matchCase</code> is live by matcher protocol
interface	at least one live class implements it	declaration/contract as a whole	interface with no live implementing class = error; interface members are contract shape
signature	at least one live module conforms to/matches it	declaration/contract as a whole	signature with no live conforming module = error; signature members are contract shape
typeclass	at least one live instance implements it	declaration/contract as a whole	typeclass with no live instance = error; typeclass methods are contract shape
typeclass instance	at least one typeclass operation is actually exercised through that specific instance	instance as a whole	instance with no exercised operation = error; remaining required methods are live by conformance
ordinary unit	the unit filename itself creates no usage identity	every usage-checkable package symbol contributed by the unit	every unused contributed symbol is an independent compile-time error
struct companion unit	the companion filename itself creates no usage exemption	every user-declared companion function	every unused companion function is an independent compile-time error; owner-struct use does not use companions
ordinary import	at least one symbol is successfully consumed through the imported context	import edge	zero-use import = compile-time error and no effective dependency edge
ordinary <code>src/</code> package	reachable from the applicable project surface through effective symbol-use edges, or intentionally selected by <code>src/export.f01</code> in an export producer	all usage-checkable project-owned declarations according to their declaration rows above	unreachable package or unused declaration = compile-time error unless the declaration is an intentional standalone-export artifact root
project-local <code>src/lib/<kind>/library.f01</code> surface	every exported surface API is consumed by the owning project's <code>src/</code> graph	every exported surface API	any unconsumed exported API = compile-time error
project-local <code>src/lib/<kind>/</code> internal package	reachable from that source library's <code>library.f01</code> through actual symbol use	all project-owned declarations according to this matrix	unreachable internal package or unused declaration = compile-time error
project-local <code>src/lib/exports/export.f01</code>	structurally active when the <code>exports/</code> domain is present and its package selections are valid	selector validity; selected packages themselves follow ordinary package/declaration rules	invalid/missing selections = compile-time error; selection alone does not exempt project-local declarations from unused checks
project-local <code>src/lib/exports/</code> package	participates in the owning project's ordinary package graph when selected by <code>export.f01</code>	all usage-checkable declarations according to their ordinary declaration rows	same unused/reachability rules as project-owned ordinary packages
<code>src/lib/operators/library.f01</code>	structurally active when present and permitted	operator-specific validity/implementation rules	handled by the operator rules; it is the structural exception to ordinary <code>src/lib</code> API consumption

standalone library producer <code>src/library.fol</code> export	intentional export from the standalone library surface	producer surface declarations are roots; their internal implementation dependencies remain strict	export does not require an internal consumer; unused implementation source remains an error
standalone library producer internal <code>src/</code> package	reachable from <code>src/library.fol</code> through effective symbol-use dependencies	all project-owned declarations according to this matrix	unreachable package or unused declaration = compile-time error
standalone export producer <code>src/export.fol</code>	valid package selections intentionally define the artifact's exported package roots	selector validity and producer implementation reachability	selected export roots do not require an internal consumer; unselected/unreachable producer source remains an error
standalone export producer selected <code>src/</code> package	selected directly or recursively by <code>src/export.fol</code>	exported declarations are artifact roots; internal dependencies remain strict	external export intent keeps selected exported declarations in the artifact even when no producer-internal call uses them
packaged <code>lib/<name>.folenc</code>	at least one projected library symbol or exported-package symbol is actually used by the consumer	artifact/import liveness only; sibling surface/package exports are not consumer-unused checked	zero used exports = compile-time error; unused sibling exports are valid
private implementation inside <code>.folenc</code>	producer already validated it when creating the artifact	not revalidated by consumer	no consumer-side unused-symbol analysis
application project	dependency graph rooted at <code>src/appl.fol</code> reaches all physically participating project-owned source/entities	all applicable rows above	any orphan project-owned source/entity = compile-time error
standalone library project	dependency graph rooted at <code>src/library.fol</code> reaches all physically participating producer implementation source/entities	all applicable rows above	any orphan producer source/entity = compile-time error
standalone export project	dependency graph rooted at the <code>src/export.fol</code> selections reaches the exported package roots and all required producer implementation dependencies	all applicable rows above, with selected exports treated as intentional producer roots	invalid selection or orphan unselected producer source/entity = compile-time error

Contract and Protocol Relationships

Interfaces, signatures, typeclasses, and matchers are contracts/protocols rather than collections of unrelated callable utilities:

```
interface  -> live implementing class
signature  -> live conforming/matching module
typeclass  -> live implementing instance
matcher    -> actual matcher selection in .match(...)
```

Their required members are not independently usage-checked.

For a class or module implementing a contract:

```
contract-required member
  -> live by conformance

additional developer-authored member
  -> independently usage-checked
```

Conformance therefore does not excuse helper/debug APIs that the developer freely added.

Typeclass instance

A typeclass instance becomes live only when at least one operation is actually exercised through that specific instance:

```

OptionMonad instance
├─ pure()          required by Monad
├─ flatMap()       required by Monad

no operation exercised through OptionMonad
-> OptionMonad UNUSED

OptionMonad.flatMap(...) resolves successfully
-> OptionMonad LIVE
-> pure() valid by conformance
-> flatMap() valid by conformance

```

An exercised operation includes:

- a qualified operation call through the instance;
- an activated method call that resolves to that instance through `@co.ddap.use`;
- live generic code that is supplied that instance and actually invokes a typeclass operation through it.

Merely importing, naming, binding, storing, passing around, or activating the instance without an operation being exercised does not by itself satisfy the instance's usage requirement.

Direct Usage Declarations

The following are live by their own semantic use rather than by implementation of a separate contract:

```

struct / cstruct / union / enum
-> type use

object
-> object declaration/reference use

annotation object
-> annotation application

matcher
-> matcher selection in .match(...)

```

For data-shape declarations, fields/members/variants are not independently checked. For matchers, `matchCase` is protocol-required and therefore not independently checked.

Language-owned annotations, directives, and pragmas such as `@co.dap.*`, `@co.ddap.*`, and `@co.pdap.*` are compiler/language facilities rather than project-owned declaration symbols and are outside project unused-symbol validation.

Strict Per-Symbol Declarations

Ordinary units, struct companion units, freely-authored class methods, and freely-authored module functions/methods are strict because each symbol was independently chosen by the developer.

```

utility.unit.fol
├─ Id alias          USED
├─ Result type      USED
├─ parse()          USED
└─ format()         UNUSED -> compile-time error

```

```

Employee struct          USED
Employee.comp.unit.fol
├─ calculateTax()       USED
├─ serialize()          USED
└─ debugDump()         UNUSED -> compile-time error

```

Using one unit symbol does not use its siblings. Using a struct does not use its companion functions. Using one companion function does not use its siblings.

Imports, Source Libraries, and Packaged Libraries

All dependency-bearing entities follow the same general rule:

```

written / discovered / loaded
!= live dependency

actual semantic contribution
-> effective dependency edge

```

For an ordinary import:

```
resolve symbol through imported context
-> mark exact SymbolId used
-> mark import used
-> activate dependency edge
-> make target context reachable
```

A zero-use import is pruned and reported as unused.

Project-local source libraries and independently produced `.folenc` artifacts differ because their ownership differs:

```
project-local source library
-> same project owns producer surface and consumer
-> every exported library-surface API must be consumed

independent library-kind .folenc
-> separately produced closed artifact
-> consumer may legitimately use only part of its surface API

independent export-kind .folenc
-> separately produced open package artifact
-> consumer may legitimately use only part of its exported package graph

either independent .folenc kind
-> at least one projected/exported symbol must be consumed for the physical dependency to be live
```

A project-local source library is therefore strict across both its surface and its implementation:

```
src/lib/<kind>/library.fol
-> every exported API consumed by owning src graph

src/lib/<kind>/<internal packages>
-> every package reachable from library.fol
-> every usage-checkable declaration follows this matrix
```

A packaged `.folenc` is intentionally relaxed at the consumer boundary:

```
lib/foo.folenc

0 projected/exported symbols used
-> UNUSED artifact -> compile-time error

1+ projected/exported symbols used
-> LIVE artifact
-> unused sibling library-surface symbols or exported-package symbols valid
-> private implementation not revalidated by consumer
```

For an export-kind artifact, package imports activate the artifact by consuming symbols through one of its exported package contexts. Merely placing the artifact in `lib/` does not make all of its packages live.

A standalone library **producer** remains strict about its own implementation. Its exported `src/library.fol` declarations are producer surface roots because they are intentionally built for external consumers; internal packages and usage-checkable source remain subject to this matrix.

A standalone export **producer** follows the same producer-side principle. Packages selected by `src/export.fol` are intentional artifact roots and remain in the output even when no producer-internal call uses them. Unselected helper packages and implementation dependencies must still be reachable from those roots and satisfy the ordinary declaration-usage rules.

Closed-Project Reachability

Physical project-owned source expresses participation in the project.

For an application:

```
root = src/appl.fol
```

For a standalone library producer:

```
root = src/library.fol
```

For a standalone export producer:

```
roots = package contexts selected by src/export.fol
```

After semantic/name/type/operator resolution reaches a fixed point:

```
CandidateImport -- no used SymbolId --> prune + unused-import error  
CandidateImport -- used SymbolId ----> effective dependency edge  
  
DiscoveredProjectEntities - ReachableProjectEntities = Orphans
```

The compiler validates, using the single matrix above:

1. zero-use imports;
2. declaration and member usage;
3. reachability of every applicable `src/` package;
4. complete exported-API consumption and internal reachability for project-owned library `src/lib` domains, plus ordinary package usage/reachability for `src/lib/exports/`;
5. at least one used projected/exported symbol for every physical `lib/*.folenc` artifact;
6. intentional producer roots from `src/library.fol` or `src/export.fol` where applicable;
7. orphan/unreachable project-owned entities.

All accumulated unused-symbol and reachability failures are compile-time errors.

Bootstrap Order

1. The compiler receives the project root as the build input.
2. Validate the project-root domains:
 - src/ is mandatory;
 - src/lib/ and lib/ are optional but cannot be empty when present;
 - build/ is compiler-managed and is created when absent.
3. Derive the project identity from the canonical project-root directory basename, then classify the project from the fixed src surface:
 - project-root basename -> project/artifact name and default output basename;
 - src/appl.fol -> application project;
 - src/library.fol -> standalone packaged-library project;
 - src/export.fol -> standalone export project;
 - more than one present -> error;
 - none present -> error.
4. For a standalone packaged-library project, shallow-read src/library.fol sufficiently to obtain and validate the project-level @co.dap.library(type=...) kind before kind-dependent bootstrap decisions. For a standalone export project, recognize src/export.fol as the package selector surface; its package selections are validated after package indexing has established the available src package contexts.
5. From the same project root:
 - discover all package source below src/;
 - discover standardized src/lib slots when src/lib/ is present, including application, ffi, system, advanced, dynamicvmrt, exports, and operators;
 - discover every lib/*.folenc artifact when lib/ is present;
 - prepare build/.
6. Start independent preparation work permitted by the classified project/kind:
 - a. parse src/lib/operators/library.fol when present and permitted;
 - b. parse each existing source-library src/lib domain in its own isolated library context;
 - c. deserialize all lib/*.folenc artifacts, reconstructing either a library-surface projection or export-package contexts according to the artifact kind.The src/lib/exports/ domain is open package source and is parsed after the owning ProjectOperatorTable is available.
7. For src/lib/operators/library.fol:
 - validate the operator-only surface and build ProjectOperatorTable.
 - Its table belongs only to the owning project's open package compilation domain ('src/' plus selected 'src/lib/exports/' source when present).
8. After ProjectOperatorTable is ready when applicable, parse the fixed project surface (src/appl.fol, src/library.fol, or src/export.fol), all package source below src/, and selected/open package source under src/lib/exports/, producing ASTs and canonical context/symbol-table entries. For an export project or project-local export domain, validate @co.dap.export package selections against the package contexts discovered for that export root.
9. Resolve import declarations by reference:
 - locate the prepared ContextId / SymbolTableId and add a provisional reference to the importing context's ImportedContexts.
 - Never copy imported symbol tables and do not mark the dependency live merely because the import was written.
10. During semantic/name/type resolution, when a symbol is successfully resolved through a provisional imported context:
 - mark that exact SymbolId used;
 - mark the import used;
 - activate the dependency edge;
 - make the target context reachable.Using one symbol never marks sibling symbols used.
11. Resolve each project-local source library's private internal dependencies independently from its src/lib/<kind>/library.fol surface. Ordinary/open package source can enter that domain only through the projected surface, never through a direct import of an internal library package.

Resolve src/lib/exports/ package dependencies as ordinary package edges in the owning project's open package graph. No library-surface gate is inserted for selected export packages.

12. Continue semantic/name/type/operator resolution until dependency, usage, and resolution state reaches a fixed point.
13. Run final unused-symbol, liveness, and reachability validation exactly as defined in [Unused Symbols, Liveness, and Reachability](#unused-symbols-liveness-and-reachability), including import pruning, declaration-kind usage rules, project-owned `srcLib` strictness, standalone library/export producer roots, `.folenc` consumer relaxation, and closed-project orphan detection. Report accumulated unused and unreachable entities as compile-time errors.
14. Complete final semantic validation.
15. Serialize the validated frontend result under <project-root>/build/ using the wire format/version selected by the backend interchange contract.

Imports contribute no parser operator metadata. Project-local operator bootstrap, source-library preparation, project-local export preparation, and packaged `.folenc` metadata loading therefore do not depend on import order. Visibility is established later by explicit `ImportedContexts` references. Those references are provisional until semantic resolution consumes at least one symbol through them; zero-use imports are pruned and rejected during final validation.

Forward / Extern Declarations

Variables extern declaration

// someUnit.unit.fol

```
_ co.lang.unit = {
  @co.dap.declare(extern)
  someBool co.lang.bool;
}
```

Functions forward declaration

// someOtherUnit.unit.fol

```
_ co.lang.unit = {

  @co.dap.declare(forward)
  getEmployee(id co.lang.int)->(somepack.Employee);

  // or - @co.dap.declare is optional for functions
  getEmployee(id co.lang.int)->(somepack.Employee);

}
```

Types external declaration

```
// Employee.fol
_ co.lang.class = {

  @co.dap.declare(extern)
  Dept co.lang.struct;

}
```

For functions and types `@co.dap.declare` is optional. For variables it is required.

Functions

FoLang does not allow free-flowing package functions. Package functions must be declared inside an ordinary `<Fragment>.unit.fol` file. Their public identity is the package member name, not the unit filename.

Normal

```
// general.unit.fol
_ co.lang.unit = {

    fun1(k co.lang.int, b co.lang.char)->(co.lang.int, co.lang.char) = {
        // function body
    }
}
```

A FoLang function may return multiple values.

Default Parameters

// somefununit.unit.fol

```
_ co.lang.unit = {
    fun1(k co.lang.int, b co.lang.char = 'A')->(co.lang.int, co.lang.char)={
    }
}

usage:
fun1(10, 'B'); // k = 10 and b = 'B'
fun1(10);      // k = 10 and b = 'A', the declared default value
```

When an argument for a default parameter is not supplied, the parameter assumes its declared default value.

Variadic Functions

Curried functions are not allowed to be variadic, and vice versa.

//someCurried.unit.fol

```
_ co.lang.unit = {
    fun1 (k co.lang.int, ...b co.lang.char)->(co.lang.int, co.lang.char)={
    }
}
```

Optional Parameters

//someOptional.unit.fol

```
_ co.lang.unit = {
    fun1(k? co.lang.int)->()={
        k.omitted.then({

        }).default({

        });
    }
}
```

When an optional argument is not supplied, the parameter is in the language-defined omitted/unprovided state and its `omitted` flag is `true`. This state does **not** assign `co.const.none` (or any other sentinel value) to the parameter. `k.omitted` tests whether a value was supplied; developers must not assume that reading or printing an omitted `k` produces `co.const.none`.

Named Parameters

// someNamedParam.unit.fol

```

_ co.lang.unit = {
  fun1(~k co.lang.int, ~v co.lang.int)->()={

  }
}

```

Usage:

```

fun1(v=10,k=20); // valid
fun1(10,20); //valid here k =10 and v=20

```

Named-parameter declaration is all-or-none within a function: a function cannot mix `~`-named parameters with ordinary parameters. When a call uses argument names, their order does not matter. A function whose parameters are all declared with the named form may still be called positionally where otherwise valid, as shown above.

Named Returns

//someNamedResults.unit.fo1

```

_ co.lang.unit = {

  doManythings(a co.lang.int, b co.lang.int->(&, meta={type=out}))->(r co.lang.int, e co.lang.exception)={}
  doSomething(input co.lang.int)->(a co.lang.int, b co.lang.bool) = {
    this.return 20, co.const.true;
  }
}

```

Usage:

// Before the call, a and b do not yet exist; using either name here is a compiler error.

```
doSomething(10);
```

// Immediately after the call, the named return values are available at the call site.

// If a or b already exists with the matching type, that existing binding is used.

// If either name already exists with an incompatible type, it is a compiler error.

```
co.out.println(a); // 20
```

```
co.out.println(b); // prints True (boolean)
```

> Named returns create call-site bindings using the declared return names when compatible bindings do not already exist.

Function Delegates

// someFunDelegate.unit.fo1

```

_ co.lang.unit = {
  myFunc(s co.lang.int, t co.lang.int)->(co.lang.int, co.lang.int) = {
    this.return 10, 10;
  }

  mySecondFun(s co.lang.int, t co.lang.int)->(co.lang.int, co.lang.int) = {
    this.return 20, 20;
  }
}

```

Usage inside a legal executable function or method body:

```

@co.dap.delegate someDelegate co.lang.delegate =
  (a co.lang.int, b co.lang.int)->(co.lang.int, co.lang.int);

someDelegate = myFunc;
someDelegate(10, 20); // invokes the currently registered function

someDelegate = myFunc;
someDelegate += mySecondFun;
someDelegate(10, 20); // invokes the registered delegate functions

```

Delegates are primarily used when a developer needs to register multiple compatible functions. When a call should simply redirect or chain into another function, use function chaining as shown below.

Multicast Delegate Invocation

When a delegate has multiple registered compatible functions, invocation proceeds sequentially through the registered functions in delegate invocation order. A function added with `+=` is invoked after the functions already registered before it.

The value returned by the delegate invocation is always the result of the **last invoked function**. Results produced by earlier registered functions do not become the result of the delegate invocation. For a multi-return delegate, the complete set of return components produced by the last invoked function is returned.

For example, if `myFunc` returns `(10, 10)` and `mySecondFun` returns `(20, 20)`, then a call made after:

```
someDelegate = myFunc;
someDelegate += mySecondFun;
```

invokes `myFunc` and then `mySecondFun`, and the delegate call returns `(20, 20)`.

Function Chaining
// somefunctionChaining.unit.fol ```folang _ co.lang.unit = { fetchEmployee(empld co.lang.string)-> (Employee)=>>empMod.getEmployee(this, empld);
dosomething(a co.lang.int, b co.lang.int)->(co.lang.int) =>> somePack.someMethod(a) =>> someOthPack.someOtherMeth(\$1, b); } ```
For a chained invocation, bind variables refer only to the return values of the immediately preceding function invocation. Return components are numbered from one in declared return order: <code>\$1</code> is the first return value, <code>\$2</code> the second, and so on. Therefore a preceding multi-return function exposes all of its return components to the next chaining step as <code>\$1 ... \$N</code> . A single-return function exposes only <code>\$1</code> . ### Anonymous Functions // someAnonymousFun.unit.fol ```folang _ co.lang.unit = { add co.lang.function = (a co.lang.int, b co.lang.int) -> (co.lang.int){ this.return a + b; };
res co.lang.function = (a co.lang.int, b co.lang.int) -> (co.lang.int){ this.return a * b; }(10, 20); } ``` > Why is there no equals sign after the function signature? It is deliberately omitted because the function signature acts as the type and the body acts as the literal value. The declaration is therefore a function-object initialization, analogous to initializing any other UDT from an object literal. > For more information about functions, see Functions in Detail .


Functions in detail
Inline
folang // math_functions.unit.fol _ co.lang.unit = { @co.dap.inline add(a co.lang.int, b co.lang.int)-> (co.lang.int)={ this.return a + b; } }
Anonymous Functions
// someOtherAnonymousfun.unit.fol
```folang _ co.lang.unit = { add co.lang.function = (a co.lang.int, b co.lang.int) -> (co.lang.int) { this.return a + b; };
res co.lang.function = (a co.lang.int, b co.lang.int) -> (co.lang.int) { this.return a * b; }(10, 20); } ```
### Lambda
Only allowed as an inline callback argument to receiver-qualified collection operations (e.g. <code>each</code> , <code>map</code> , <code>filter</code> , <code>reduce</code> , <code>forEach</code> , <code>sortBy</code> , <code>groupBy</code> ). Transparent grouping around the member callee is permitted, so <code>(nums.map)( x  =&gt; x*x)</code> is equivalent to the ordinary call spelling. A bare function call such as <code>map( x  =&gt; x)</code> is not a collection-method context. Using <code> ... </code> anywhere else is a syntax/lint error.
// somelambda.unit.fol ```folang // Callback literal syntax, shown only in its valid collection-call context _ co.lang.unit = { nums.map( x  => xx); words.filter( s  => s.len() > 3); pairs.reduce( acc, e  => acc + e, 0); dict.map( k, v  => v 10); list.sortBy( a, b  => a.score - b.score); } ```
The lambda must be a direct argument of the allowed collection call. That call may itself be nested, for example <code>consume(nums.map( x  =&gt; x*x))</code> ; the enclosing call does not make the lambda an argument of <code>consume</code> .
### Inner Function // someInnerFun.unit.fol
<code>@co.dap.local</code> , <code>@co.dap.nested</code> , and <code>@co.dap.inner</code> apply to separately declared functions and do not change the lexical-scope rule for an ordinary inner function. They provide visibility, target-state capture, or call-site association without requiring the function body to be physically placed inside the target declaration.
The permitted local-function form above is useful when the inner function is small and belongs naturally to one enclosing function. This is the named local-function exception described in <a href="#">Physical Nesting Rules</a> .
### Curried // someOtherCurried.unit.fol ```folang _ co.lang.unit = { add(first co.lang.int)(second co.lang.int)->(co.lang.int)={ this.return first + second;

}} ``
### Closure //someClosure.unit.fol
### Functions Taking and Returning Functions
#### Syntax 1 — Inline signature //someInlineSignature.unit.fol ``folang
_ co.lang.unit = { someFunction (r (co.lang.int, co.lang.int)->(co.lang.int))->((co.lang.int)->(co.lang.int))={ } } ``
#### Syntax 2 — Named type alias //someNamedTypeAliases.unit.fol ``folang _ co.lang.unit = { someFArg co.lang.type = (co.lang.int, co.lang.int)->(co.lang.int); someFRet co.lang.type = (co.lang.int)->(co.lang.int);
someFunction (r someFArg)->(someFRet)={ } } ``
#### Syntax 3 — Function objects //someFunObject.unit.fol ``folang _ co.lang.unit = { someFArg co.lang.function = (a co.lang.int, b co.lang.int) -> (co.lang.int){ this.return a + b; };
someFRet co.lang.function = (a co.lang.int) -> (co.lang.int){ this.return a * 2; }; } ``
### Other Ways to Declare Closures, Function Objects, Function Types, and Curried Functions //someAdditionalLeg.unit.fol ``folang _ co.lang.unit = { myobj co.lang.function = (a co.lang.int, b co.lang.int)->(co.lang.int){ this.return a + b; };
add (a co.lang.int, b co.lang.int)->(co.lang.int)={ this.return a + b; } oObj co.lang.function = add;
funtype co.lang.type = (a co.lang.int, b co.lang.int)->(co.lang.int);
closure=(factor co.lang.int, val co.lang.int) ==>> factory * val;
curry = (factor co.lang.int) (x co.lang.int) ==>> x * factor; } `` ***
### Associated Functions
For a user-defined struct, associated functions must be declared inside the same-package companion unit whose name matches the struct. For more information, see <a href="#">Associated Functions in a Companion Unit</a> .
### Some Restrictions on Special Functions
#### Non-overloadable Function Forms
The following function forms cannot be overloaded:
1. functions with named parameters; 2. functions with optional parameters; 3. functions with default parameters; 4. variadic functions; 5. curried functions; 6. functions that use inline function-signature syntax directly in a parameter or return position; 7. functions that use a function type as a parameter or return type; 8. multi-return functions; 9. named-return functions; 10. functions having a pointer, address, reference, thunk, or slice in any parameter or return position.
These categories are signature-level restrictions. A declaration that falls into more than one category remains non-overloadable for the same purpose; the categories do not create separate overload families.
#### Existing Callback and Execution-Model Restrictions
Curried functions, functions with named, optional, variadic, or default parameters, functions that take or return functions or function types, and dynamically scoped or mixed-scoped functions retain the existing restrictions that they cannot be used as ordinary callbacks and cannot participate in <a href="#">Execution Models and Control Abstractions</a> .
#### Scoping Rules for Functions
All functions in FoLang have a defined scope — the set of variables a function can access.
***
##### Default — Lexical / Static Scope
All of the following are <b>always</b> lexically scoped — this cannot be changed:
Lexical scope means a function resolves names from its <b>declaration site</b> , not from the scope of its caller. Unit-level variables are forbidden, so a unit-scoped free function receives runtime values through parameters or introduces them locally. An ordinary inner function captures from the enclosing lexical declaration context. Calling that inner function does not replace its captured context with the caller's runtime scope.
``folang // scope_example.unit.fol _ co.lang.unit = {
foo()->() = { x co.lang.int = 10;
printX()->() = { co.out.println(x); //  x from the enclosing lexical scope }
printX(); } } ``
***
##### @co.dap.inner — Call-Site Lexical Context
@co.dap.inner is not the syntax for an ordinary inner function. It annotates a separately declared function or method that becomes associated with a target at an explicit use or attachment site.

Its free runtime names use <b>call-site lexical-context resolution</b>:
text 1. <code>@co.dap.inner</code> parameters and local declarations 2. the innermost lexical scope at the active call or attachment site 3. enclosing lexical scopes of that call or attachment site 4. statically resolved types, annotations, imports, and <code>co.*</code> names 5. compiler error
The compiler validates these requirements at every statically known use or call site. This is a fixed property of <code>@co.dap.inner</code>; it is not selected through <code>@co.dap.lexicalscope</code>, <code>@co.dap.dynamicscope</code>, or <code>@co.dap.mixedscope</code>.
This differs from ordinary lexical inner functions, whose free runtime names are fixed by their declaration site, and from dynamically scoped associated functions, which may search outward through the runtime caller activation chain.
***
##### Associated Functions — Additional Scope Options
Only associated functions support non-lexical scoping via annotations:
The examples below are members of the same-package <code>Employee</code> companion unit.
<code>@co.dap.lexicalscope</code> — default, explicit declaration
<code>@co.dap.dynamicscope</code> — accesses caller's scope
<code>@co.dap.mixedscope</code> — accesses both scopes
***
##### Callback Scope Inside Dynamically Scoped Associated Functions
A callback block or lambda does not independently select lexical, dynamic, or mixed scope. The associated function that executes the callback determines how the callback's free runtime names are resolved.
Callback parameters and declarations made inside the callback always belong to the callback's local scope. Names that are not callback parameters or callback-local declarations follow the executing associated function's scope policy. <code>//someOperation.unit.fol</code>
For this example:
text <code>acc</code>, <code>e</code> -&gt; callback parameters temporary locals -&gt; callback-local context <code>total</code> -&gt; <code>reduce</code>'s dynamic caller context <code>co.*</code> and imports -&gt; normal built-in/import resolution
The callback syntax remains uniform. <code>reduce</code> is dynamically scoped, so unresolved runtime names in the callback are resolved through the active caller-context chain. A lexically scoped associated function would instead resolve free runtime names through its declaration context.
***
##### Why Dynamic Scope Exists — <code>.then</code>, <code>.loop</code>, <code>.each</code>, and Collections
FoLang's control-flow model is built on dynamically scoped associated functions. The executing function supplies the scope policy, so blocks do not require separate capturing and non-capturing forms.
<pre>//someScopeEg1.unit.fol ``folang _ co.lang.unit = { someFun()-&gt;()={ x co.lang.int = 10; total co.lang.int = 0; arr co.lang.int-&gt;([5]) = [1, 2, 3, 4, 5];  // .then reads and modifies the caller's x (x &gt; 5).then({ x.value = 20; co.out.println(x); });  // .loop modifies the caller's x (x &gt; 0).loop({ x.value -= 1; });  // .each(..., action) performs the iteration and modifies the caller's total arr.each(_ val, { total.value += val; });  // .filter, .map, and .reduce use the same dynamic caller context nums.filter( x  =&gt; x &gt; 10); nums.reduce([acc, e] =&gt; { total.value += e; acc + e }, 0); } } ``</pre>
The compiler does not create a separate capture description for each control block. It resolves names according to the scope mode of the associated function executing that block.
***
##### FoLang Control Flow Uses Dynamic Scope
``text no if/else keywords -&gt; <code>.then</code> / <code>.otherwise</code> / <code>.default</code> — dynamic scope no for/while keywords -&gt; <code>.loop</code> — dynamic scope no foreach keywords -&gt; <code>.each(..., action)</code> — dynamic scope no in keyword -&gt; <code>.contains</code> — dynamic scope no map/filter keywords -&gt; <code>.map</code> / <code>.filter</code> — dynamic scope
all control flow is implemented as associated functions control associated functions are dynamically scoped caller variables are accessible and mutable under the dynamic lookup rules block syntax requires no separate capture mode ``
***
##### Dynamic and Mixed Lookup Order
For <code>@co.dap.dynamicscope</code>, runtime variable lookup follows this order:



## Type System Dimensions



- $\lambda \rightarrow$ : Simply Typed  $\lambda$ -Calculus  
Terms depend on Terms (Functions)
- $\lambda 2$ : System F  
Terms depend on Types (Polymorphism)
- $\lambda_\omega$ :  $\lambda$ -Calculus with Operators  
Types depend on Types (Type Operators)
- $\lambda P$ :  $\lambda P$  (LF)  
Types depend on Terms (Dependent Types)
- $F_\omega$ : System  $F_\omega$   
Generics + Type Operators
- $\lambda P2$ :  $\lambda P2$   
Dependent Types + Polymorphism
- $\lambda P_\omega$ :  $\lambda P_\omega$   
Dependent Types + Kinds
- $\lambda C$ : Calculus of Constructions  
The Ultimate Unification (All dependencies)

***

### ## Refinement Types

A refinement type restricts the admissible values of an existing base type by a Boolean predicate. It does not introduce a separate binder name merely so the predicate can refer to the value being tested. The contextual token `_` denotes that **candidate value** inside the refinement predicate.

The canonical declaration form is:

```
folang // types.unit.fol _ co.lang.unit = { positiveInt co.lang.refinementType = (co.lang.int).where(_ > 0);
}
```

The declaration above means that `positiveInt` admits values of `co.lang.int` that satisfy `_ > 0`. The `_` placeholder has the base type (`co.lang.int` in this example), and every occurrence of `_` in the same refinement predicate refers to the same candidate value.

```
``folang // types.unit.fol _ co.lang.unit = { percentage co.lang.refinementType = (co.lang.int).where(_ >= 0 && _ <= 100);
```

```
evenInt co.lang.refinementType = (co.lang.int).where(_ % 2 == 0);
```

```
nonEmptyString co.lang.refinementType = (co.lang.string).where(_.length > 0); } ``
```

The predicate supplied to `.where(...)` must resolve to `co.lang.bool`. The candidate placeholder is available only while resolving that refinement predicate. It does not introduce a variable into the enclosing unit, does not escape the declaration, and cannot be rebound or assigned. Member access and ordinary expressions may use it according to the operations supported by the base type.

`_` remains contextual. In pattern and discard positions it retains its wildcard/discard meaning, and in filename-derived primary declarations it retains its declaration-name-placeholder meaning. Only the predicate belonging to a `co.lang.refinementType` declaration gives `_` the candidate-value meaning.

For a statically known candidate, the compiler can evaluate the refinement predicate directly. A statically known value that makes the predicate false cannot inhabit the refinement type:

```
``folang // types.unit.fol _ co.lang.unit = { positiveInt co.lang.refinementType = (co.lang.int).where(_ > 0);
```

```
good positiveInt = 10; // valid: 10 > 0 bad positiveInt = -5; // compile-time error: -5 does not satisfy the predicate } ``
```

Refinement validation is normatively a **run-time operation**. Whenever a value enters a `co.lang.refinementType` through assignment, argument passing, return, conversion, initialization, or another admissible value-transfer operation, the refinement predicate is validated at run time. If the predicate evaluates to `co.const.true`, the value is admitted; if it evaluates to `co.const.false`, refinement validation fails.



The compiler may additionally evaluate the predicate when the candidate value is statically known. This permits an invalid statically known candidate to be rejected earlier, as in the <code>bad</code> declaration above. Such compile-time analysis is an early diagnostic and does not change the language's validation model: refinement validity is defined by the run-time predicate check rather than by a requirement that the predicate be statically provable.
Refinement, Dependent, and Associated Types
These three mechanisms solve different problems:
Form   What determines or constrains the type   FoLang role   <code>co.lang.refinementType</code> = <code>(Base).where(predicate)</code>   a predicate restricts which values of <code>Base</code> are admitted   value-set restriction     <code>Vector(n)</code> / a function returning <code>co.lang.dependentType</code>   a value participates in the resulting type identity or shape   value-indexed type     <code>Entity co.lang.associatedType;</code>   a signature requires a matching module to supply a compatible type binding   module contract type component
A refinement predicate therefore does not make its candidate value an index of the type in the dependent-type sense. Likewise, an associated type is not a predicate-restricted value set; it is a type component selected by a matching module.
A <code>co.lang.refinementType</code> declaration is also distinct from <code>co.lang.subtype</code> and <code>co.lang.supertype</code>. Refinement adds a value predicate to a base type. It does not by itself define the inheritance, variance, or assignability rules of the separate subtype/supertype declaration kinds.
***
Dependent Types
Type-Level Functions — Functions That Return Types
A function that accepts a type or value and returns a type is a <b>type-level function</b>. When its result depends on a value argument, it defines or selects a dependent type. This is distinct from a parameterized <code>co.lang.type</code> constructor such as <code>Option(T)</code>, whose declaration directly defines a family of types.
<pre>// sometypes4.unit.fol ``folang _ co.lang.unit = { // Vector — value-indexed type-level function // takes → co.lang.int (size) // returns → co.lang.dependentType (a type) Vector(n co.lang.int)-&gt;(co.lang.dependentType) = co.lang.int-&gt;([n]);</pre>
<pre>// calling Vector(3) returns a TYPE at compile time // that type is co.lang.int-&gt;([3]) v3 Vector(3) = [1, 2, 3]; // type is Vector(3) v4 Vector(4) = [1, 2, 3, 4]; // type is Vector(4) — different type!</pre>
<pre>// Vector(3) ≠ Vector(4) — completely different types // size is part of the type — compiler knows at compile time }</pre>
***
More About Type
Name(T) co.lang.data = variants; → concrete ADT type-constructor definition → right-hand-side definition is mandatory
Name(T) co.lang.associatedType; → generic associated-type requirement → permitted only inside a signature
Name(T) co.lang.associatedType = ExistingType(T); → associated type-constructor binding → permitted directly in a matching module for the corresponding signature component
Name(T) co.lang.type = ExistingType(T); → concrete type alias in an ordinary type context, or a fixed/manifest type component in a signature
***
A Type-Level Function Returns a Type
<pre>Vector → type-level function Vector(3) → function call → returns type co.lang.int-&gt;([3]) Vector(4) → function call → returns type co.lang.int-&gt;([4])</pre>
<pre>just like: add(1, 2) → returns a value (3) Vector(3) → returns a type (int[3])</pre>
***
Compiler Enforced Size Safety
<pre>//sometypes5.unit.fol ``folang _ co.lang.unit = { // dot product — only valid for same size vectors // compiler enforces this via dependent types dotProduct(a Vector(n), b Vector(n))-&gt;(co.lang.int) = { // n is same for both — compiler verified }</pre>
<pre>v3 Vector(3) = [1, 2, 3]; v4 Vector(4) = [1, 2, 3, 4];</pre>
<pre>dotProduct(v3, v3); // ✓ same type Vector(3) dotProduct(v3, v4); // ✗ compiler error — Vector(3) ≠ Vector(4) }</pre>
***
Matrix — Two-Parameter Type-Level Function
<pre>somematrix.unit.fol ``folang _ co.lang.unit = { // Matrix — takes rows and cols, returns dependent type Matrix(r co.lang.int, c co.lang.int)-&gt;(co.lang.dependentType) = co.lang.int-&gt;([r, c]);</pre>
<pre>m34 Matrix(3, 4) = [[1,2,3,4],[5,6,7,8],[9,10,11,12]]; m45 Matrix(4, 5) = ...;</pre>
<pre>// matrix multiply — cols of A must equal rows of B // n must match — compiler verified multiply(a Matrix(r, n), b Matrix(n, c))-&gt;(Matrix(r, c)) = { // compiler ensures dimensions are compatible }</pre>
<pre>multiply(m34, m45); // ✓ Matrix(3,4) × Matrix(4,5) = Matrix(3,5) multiply(m34, m34); // ✗ compiler error — 4 ≠ 3 }</pre>
***

### Stack — Value and Type Parameter //somestack.unit.fol ```folang _ co.lang.unit = { // Stack — takes size and element type Stack(n co.lang.int, T co.lang.type)-&gt;(co.lang.dependentType) = T-&gt;([n]);
s Stack(10, co.lang.int) = ...; // stack of max 10 ints t Stack(5, co.lang.string) = ...; // stack of max 5 strings } ```
***
### Type Is Value + Kind Combined ``` Vector(3): kind = Vector (what it is) value = 3 (how many) type = Vector(3) (both together — the dependent type)
Vector(3) ≠ Vector(4) ← different types entirely Vector(3) = Vector(3) ← same type ```
***
### Type Constructors and Type-Level Functions
```folang // option.unit.fol _ co.lang.unit = { // Parameterized type declaration: Option is a type constructor. Option(T) co.lang.type = Some(T)   None();
// Value-indexed type-level function: Vector computes a dependent type. Vector(n co.lang.int)->(co.lang.dependentType) = co.lang.int->([n]); } ```
Option and Vector both operate at the type level, but they are different declaration categories:
```text Option(T) co.lang.type -&gt; type-constructor declaration -&gt; substitution produces Option(T)
Vector(n)-&gt;(co.lang.dependentType) -&gt; type-level function -&gt; computation produces a type ```
***
### Simple Dependent Type //someiden.unit.fol
### Compile-Time and Runtime Values in Type-Related Computation
FoLang distinguishes value-indexed dependent types, compile-time type computation, runtime type descriptors, and runtime values whose concrete types may differ. These mechanisms are related, but they are not interchangeable.
#### 1. Value-Indexed Dependent Types
A dependent type may contain a value as part of its type identity. The value index may be a compile-time constant, a symbolic type-level value, or a runtime function parameter whose relationship to the result is tracked by the compiler. //someIndex.unit.fol
The compiler does not need to know the concrete value of n while compiling readVector. It records the relationship:
text input index = n result type = Vector(n)
Likewise: // sometypes6.unit.fol
requires both vectors to have the same value index. Dependent typing means that a type contains or is constrained by a value; it does not mean that an arbitrary runtime branch can silently change the static type of an already compiled variable.
***
#### 2. Compile-Time Type Computation
A function may compute and return a type when it is guaranteed to execute during compilation. // someFuncomp.unit.fol
The arguments must be compile-time evaluable when the result is used in a static type position: //someTest1.unit.fol
Invalid: //someTest2.unit.fol
Conceptually:
text compile-time value -&gt; compiler executes the type function -&gt; one concrete static type is available before ordinary type checking completes *** #### 3. Built-in compile-time type computation
A function may compute and return a type when it is guaranteed to execute during compilation.
&gt; decltype is a built-in method.
The arguments must be compile-time evaluable when the result is used in a static type position: // someFun5.unit.fol
#### 4. Runtime Type Descriptors
An ordinary function returning co.lang.type produces a runtime type descriptor when it is not executed at compile time. //someruntime2.unit.fol ```folang _ co.lang.unit = { selectType(value co.lang.int)-&gt;(co.lang.type) = { (value &lt; 100) .then(co.lang.string) .default(co.lang.bool); }
selectedType co.lang.type = selectType(input); } ```
Here, selectedType is a runtime object that describes a type. It may be used for reflection, runtime validation, dynamic loading, metadata inspection, or dynamic object creation where those capabilities are permitted.
A runtime type descriptor cannot ordinarily be used as the static type of a declaration: //someruntime3.unit.fol ```folang _ co.lang.unit = { someFun()-&gt;()={ value selectType(input);

// compiler error: runtime type descriptor used in a static type position } } ``
Conceptually:
text runtime value -> runtime type-selection function -> co.lang.type descriptor object
A value that represents a type is not automatically a compile-time-resolved static type. ***
#### 5. @co.dap.typefromvalue
@co.dap.typefromvalue derives a type from the type of a returned compile-time value. In the initial FoLang specification, it is permitted only for compile-time evaluation. //sometypeval.unit.fol
The compiler derives:
text "Hello" -> co.lang.string co.const.true -> co.lang.bool
Every argument must be compile-time evaluable when the result is used in a type position. Runtime value-based type selection must instead use an explicit runtime representation.
***
#### 6. Runtime Values with Different Concrete Types
When a runtime branch may produce unrelated concrete value types, the function must return one stable outer type. FoLang may use a tagged value, an ADT, co.lang.dynamic where permitted, or another explicitly packaged representation.
Using co.lang.tag : //someruntime1.unit.fol
The outer static type is always co.lang.tag :
text co.lang.tag  — runtime type descriptor  — value compatible with that descriptor
An ADT provides a more strongly typed alternative:
``folang SelectedValue co.lang.data = StringValue(co.lang.string)   BoolValue(co.lang.bool);
selectValue(value co.lang.int)->(SelectedValue) = { (value < 100) .then(StringValue("Hello")) .default(BoolValue(co.const.true)); } `` ***
#### Summary
``text Vector(n) -> value-indexed dependent type
@co.dap.comptime function returning co.lang.type -> compile-time type computation
ordinary runtime function returning co.lang.type -> runtime type descriptor
runtime branch returning unrelated value types -> tag, ADT, dynamic value, or another explicit package ``
A runtime type descriptor is a value that represents a type. It must not be confused with a statically resolved dependent type.
***
### Parameterized Type Declarations and Type-Level Functions
Two declaration families produce types from parameters. The spelling depends on whether the declaration directly defines a type family or computes a type through a function body. //someEg7.unit.fol ``folang _ co.lang.unit = { // all parameters are types -> parameterized co.lang.type declaration Option(T) co.lang.type = Some(T)   None(); someAlias(F) co.lang.type = Functor(F);
// a value parameter is present -> type-level function syntax Vector(n co.lang.int)->(co.lang.dependentType) = co.lang.int->([n]); Stack(n co.lang.int, T co.lang.type)->(co.lang.dependentType) = T->([n]); } ``
A parameterized co.lang.type declaration is a type constructor. Its type parameters appear directly in the declaration head and it does not use @co.dap.generic .
A function that accepts values or type values and returns co.lang.dependentType is a type-level function. Stack demonstrates why the function form exists: it can mix value parameters and type-valued parameters and compute the resulting type.
co.lang.dependentType is both a type-producing return kind and a direct type-declaration kind. A type-level function uses it when a value parameter determines the produced type. A direct declaration may use it when no parameter list is required:
folang LengthBound co.lang.dependentType = co.lang.int;
The kind is also usable in a declarator. If a function returns co.lang.dependentType , a binding receiving that result may therefore be declared co.lang.dependentType .
A type-level function has exactly one unnamed type-producing result. That result may be a union using   , but comma-separated multiple results are invalid: //somebadeg1.unit.fol
#### Parameterized aliases are transparent
An alias declared with co.lang.type names the same type, not a new one. //someParameg1.unit.fol ``folang _ co.lang.unit = { someAlias(F) co.lang.type = Functor(F);
someFun()->()={ someAlias(List); // the same type as Functor(List) someAlias(Option); // the same type as Functor(Option) } } ``

Because the alias adds no identity, it creates no separate instance slot: an instance cannot be declared for <code>someAlias</code> as distinct from one for <code>Functor</code> . Declaring one would be a duplicate.
Named parameters also allow reordering and partial application, which a positional placeholder could not express. <code>//someTypeClz1.unit.fol</code>
Use <code>co.lang.newType</code> instead when a <b>distinct</b> identity is wanted, such as the wrapper described in <a href="#">Where an Instance Is Declared</a> .
### Dependent Type Index Rules
An <b>index</b> is an argument to a dependent type, such as the <code>n</code> in <code>Vector(n)</code> , or a dimension in an array derivation, such as the <code>n</code> in <code>co.lang.int-&gt;([n])</code> . Both positions obey the same rules.
#### An index is a literal or a name
An index is an integer literal or a name. Arithmetic, function calls, indexing and every other operator are rejected. <code>// someldxEG1.unit.fol ``folang _ co.lang.unit = { someFun()-&gt;()={ @co.dap.const SIZE co.lang.int = 1024;</code>
<code>v Vector(3); // ✓ literal v Vector(SIZE); // ✓ @co.dap.const name buf co.lang.int-&gt;([SIZE]); // ✓ same rule for array sizes</code>
<code>v Vector(n + 1); // ✗ arithmetic is not permitted in an index v Vector(computeSize()); // ✗ a call is not permitted in an index buf co.lang.int-&gt;([n * 2]); // ✗ same rule for array sizes }</code> ``
This restriction applies only to the <b>size</b> of an array, never to element access. Indexing an array is an ordinary expression and arithmetic is fine.
<code>folang buf co.lang.int-&gt;([SIZE]); // size – restricted buf[i + 1] = 42; // access – unrestricted buf[compute(x)] = 7; // access – unrestricted</code>
#### What a named index may resolve to
A name used as an index resolves in exactly one of two ways.
<b>A parameter bound by the enclosing signature.</b> A type constructor or a function signature introduces the name, and every use of it inside that signature and its body refers to the bound parameter. The name is not a constant; it stands for whatever value the caller supplies. <code>//someEG2.unit.fol ``folang _ co.lang.unit = { // n is introduced here, and bound for the whole declaration Vector(n co.lang.int)-&gt;(co.lang.dependentType) = co.lang.int-&gt;([n]);</code>
<code>// n is introduced by this signature, and both parameters must share it dotProduct(a Vector(n), b Vector(n))-&gt;(co.lang.int) = { // ... }</code>
<code>// n is introduced as a value parameter and reused in the return type readVector(n co.lang.int)-&gt;(Vector(n)) = { // ... } ``</code>
<b>A <code>@co.dap.const</code> compile-time constant.</b> Outside a signature that binds it, a name has nothing to bind to, so it must be a constant the compiler can substitute.
<code>folang @co.dap.const SIZE co.lang.int = 1024; buf co.lang.int-&gt;([SIZE]); // ✓ SIZE substitutes to 1024 v Vector(SIZE); // ✓ same rule for dependent types</code>
Nothing else qualifies. <code>@co.dap.final</code> marks an immutable binding, and an immutable value need not be known while compiling, so it cannot be substituted.
<code>``folang @co.dap.final n co.lang.int = readInput(); bad Vector(n); // ✗ immutable, but not known at compile time</code>
<code>m co.lang.int = 10; alsoBad Vector(m); // ✗ an ordinary variable is not an index ``</code>
So in a plain variable declaration, where no signature is binding anything, the only legal names are <code>@co.dap.const</code> constants.
#### An index is non-negative
Zero is permitted; a negative index is not.
<code>``folang empty co.lang.int-&gt;([0]); // ✓ zero-length array</code>
<code>buf co.lang.int-&gt;([-1]); // ✗ rejected while parsing v Vector(-1); // ✗ rejected while parsing</code>
<code>@co.dap.const OFFSET co.lang.int = -1; buf co.lang.int-&gt;([OFFSET]); // ✗ rejected after substitution ``</code>
A negative literal cannot be written at all, because no prefix operator is reachable in an index position. A negative constant is rejected when the compiler substitutes it. Both are compile-time errors.
#### When two dependent types are equal
Two dependent types are equal when their constructors are the same and their indices are pairwise equal. An index comparison has exactly three cases.
Index form   Compared by   <code> — </code>   integer literal   value   <code>@co.dap.const</code> name   substituted literal value   parameter   name identity
<code>folang Vector(3) vs Vector(3) // equal Vector(n) vs Vector(n) // equal Vector(n) vs Vector(m) // NOT equal – rejected Vector(SIZE) vs Vector(1024) // equal when @co.dap.const SIZE = 1024</code>
Rejecting <code>Vector(n)</code> against <code>Vector(m)</code> is the point of the feature. It is what lets the compiler catch a size mismatch without the developer writing a single check.

#### What FoLang deliberately does not do
FoLang does not decide index equality up to arithmetic. <code>Vector(n+1)</code> and <code>Vector(1+n)</code> are not merely unequal — they cannot be written.
Accepting them would require symbolic reasoning, and there is no partial version of it. Once <code>n+1 == 1+n</code> is accepted, the next reasonable request is <code>2*n == n+n</code> , and the type checker becomes a theorem prover by accretion. That is the complexity FoLang is built to avoid.
The cost is narrow. Length-arithmetic signatures such as <code>concat(Vector(n), Vector(m)) -&gt; Vector(n+m)</code> are out of scope; return a dynamically sized type and check at run time instead. Everything that needs only same-parameter identity still works, and that covers the common cases.
<code>foLang multiply(a Matrix(r, n), b Matrix(n, c)) -&gt; (Matrix(r, c)) dotProduct(a Vector(n), b Vector(n)) -&gt; (co.lang.int) zip(a Vector(n), b Vector(n)) -&gt; (Vector(n))</code>
Matrix multiplication, the usual demonstration of dependent types, needs only that the shared <code>n</code> matches.
#### Dependent types are checked, never inferred
Every dependent type appears in a written signature. FoLang never infers one.
This is what keeps checking decidable without a constraint solver, and it is why FoLang does not adopt Hindley-Milner style whole-program inference. Inferring a dependent type would mean inferring the index <b>value</b> , not merely the type, which is the step that makes checking undecidable in general.
***
## Indexer
Indexer functions for a struct are associated functions and must be declared inside <code>&lt;StructName&gt;.comp.unit.fo1</code> .
<code>``foLang // MyList.fo1 _ co.lang.struct = { eles co.lang.int-&gt;([...]); }</code>
<code>// MyList.comp.unit.fo1 _ co.lang.unit = {</code>
<code>@co.dap.indexer(symbol="[]") (g MyList) get(index co.lang.int)-&gt;(co.lang.int) = { this.return g.eles[index]; }</code>
<code>@co.dap.indexer(symbol="[]=") (g MyList) set(index co.lang.int, value co.lang.int)-&gt;() = { g.eles[index] = value; }</code>
<code>!st MyList; co.out.println(!st[0]); !st[1] = 22; ``</code>
*** ## Generics
<code>foLang @co.dap.generic( at=callsite, types=[ {name= T, variance=invariant, bound=Number, inference=param}, {name= R, variance=invariant, bound=Number, inference=result} ], impredicative=false, resolution=compiletime ) add(a T, b T)-&gt;(R) = { this.return a + b; }</code>
<b>Generic annotation fields:</b>
Attribute   Values    — —    types   map with type and other details    requires   resolution  <code>runtime</code> , <code>compiletime</code>    reified  <code>true</code> or <code>false</code>    at  <code>usesite</code> or <code>callsite</code>    specializable  <code>true</code> or <code>false</code>    impredicative  <code>true</code> or <code>false</code>
<b>types</b> attributes   Attribute   Values    — —   name   constraints   upper-bound   lower-bound    bound   not needed when <code>upper-bound</code> and/or <code>lower-bound</code> is specified    default   variance  <code>covariant</code> , <code>invariant</code> , <code>contravariant</code>    nullable   inference  <code>param</code> , <code>result</code> , <code>arg</code> , <code>var</code>    capabilities    isAKind   an interface, class, or trait; e.g. <code>SomeInterface</code>    typekind  <code>type</code> , <code>class</code> , <code>function</code> , <code>struct</code> , <code>typeconstructor</code>    inclusive
> The entries above are sub-attributes of the <code>types</code> attribute and are represented in map form.
> Example: <code>@co.dap.generic(types={ T: {variance:..., bound:..., ...} })</code>
> These are not independent attributes; they describe each declared generic type entry.
#### Generic Functions — Parameters and Return Values
#### Rank-1: Outer function is generic; parameter uses the same type variable
<code>T</code> is fixed at the call site before the function parameter is used. The passed function is already monomorphic inside the body.
<b>Syntax 1 — Inline signature</b> <code>//somGen1.unit.fo1</code>
<b>Syntax 2 — Named type alias</b> <code>//somGen2.unit.fo1 ``foLang _ co.lang.unit = { @co.dap.generic(types=[{ name=T}]) someFArg co.lang.type = (T, T)-&gt;(T);</code>
<code>@co.dap.generic(types=[{ name=T}]) someFunction(f someFArg, a T, b T)-&gt;(T) = { } ``</code>
***
#### Rank-2: The function parameter is itself polymorphic (higher-rank)
The passed function stays generic <b>inside the callee</b> . The callee decides what <code>T</code> is. Uses existing <code>forall</code> .
<b>Syntax 1 — Inline signature</b> <code>//someGen3.unit.fo1</code>

<b>Syntax 2 — Named type alias</b> <code>//someGen4.unit.fol ``folang _ co.lang.unit = { someFArg co.lang.type = forall(T).(T, T)-&gt;(T);</code>
<code>someFunction(f someFArg)-&gt;(co.lang.int) = {} } //someGen5.unit.fol1 folang _ co.lang.unit = { // Correct — Syntax 2 with co.lang.type someFArg co.lang.type = forall(T).(T, T)-&gt;(T);</code>
<code>someFunction(f someFArg)-&gt;(co.lang.int) = {} } ``</code>
***
#### Returning Generic Functions
<b>Rank-1 return</b> <code>//someGen6.unit.fol</code>
<b>Rank-2 return — returning a polymorphic function</b> <code>//somGen7.unit.fol</code>
***
#### Rank-3: A Parameter is Itself a Rank-2 Function
Rank-3 works naturally in FoLang via <code>forall</code> nesting. No new constructs needed.
<b>Syntax 1 — Inline</b> <code>//someGen8.unit.fol</code>
<b>Syntax 2 — Named type aliases (cleaner)</b> <code>//someGen9.unit.fol ``folang _ co.lang.unit = { rank2FnType co.lang.type = forall(T).(T, T)-&gt;(T); rank3ArgType co.lang.type = (rank2FnType) -&gt; (co.lang.int);</code>
<code>applyRank2(f rank3ArgType) -&gt; (co.lang.int) = { this.return f(1, 1); } } ``</code>
<b>Rank-3 return</b> <code>//somGen10.unit.fol</code>
***
#### Impredicativity — Instantiating <code>T</code> with a <code>forall</code> Type
Impredicativity is when a type variable <code>T</code> in a generic is itself instantiated with a <code>forall</code> type. Example of what this means: <code>//somGen11.unit.fol ``folang _ co.lang.unit = {</code>
<code>@co.dap.generic(types=[{name=T}]) box(x T) -&gt; (Box(T)) = {</code>
<code>someFun()-&gt;()={ // Impredicative call — T being set to forall(U).(U)-&gt;(U) result := box(forall(U).(U)-&gt;(U)); // ❌ not legal without explicit opt-in } } ``</code>
Most type systems reject this by default. FoLang takes an opt-in approach.
<b>Initial alpha release Workaround — Option C: Wrapping with <code>co.lang.type</code></b>
Not true impredicativity but solves 90% of practical cases: <code>//somGen12.unit.fol ``folang _ co.lang.unit = { polyId co.lang.type = forall(U).(U)-&gt;(U);</code>
<code>// box takes co.lang.type — no impredicative unification needed box(x co.lang.type) -&gt; (Box(co.lang.type)) = {</code>
<code>someFun()-&gt;()={ result := box(polyId); // ✅ works — x is co.lang.type, not a forall type } } ``</code>
<b>1.0 release — Option A: <code>impredicative:true</code> in <code>@co.dap.generic</code></b>
Explicit opt-in via the existing annotation. The compiler only permits <code>forall</code> instantiation where declared: <code>//somGen13.unit.fol ``folang _ co.lang.unit = {</code>
<code>@co.dap.generic( types=[{name=T,variance=invariant}], impredicative=true ) box(x T) -&gt; (Box(T)) = {</code>
<code>polyId co.lang.type = forall(U).(U)-&gt;(U);</code>
<code>someFun()-&gt;()={ result := box(polyId); // ✅ legal — impredicative:true explicitly opts in } } ``</code>
***
#### Generic Function Rank Support Matrix
Scenario   Allow?   Notes    — — —    Rank-1 generic param (Syntax 1, 2, 3)   ✅ Yes   Natural extension, no new concepts     Rank-1 generic return (Syntax 1, 2, 3)   ✅ Yes   Same as above     Rank-2 param via <code>forall</code> (Syntax 1, 2)   ✅ Yes   <code>co.lang.type</code> naturally holds polymorphic types; no new kind needed     Rank-2 param via Syntax 3 <code>co.lang.function</code>   ❌ Compiler error   Function objects are concrete values; use <code>co.lang.type = forall(T).(T)-&gt;(T)</code> instead     Rank-2 return via <code>forall</code> (Syntax 1, 2)   ✅ Yes   Same reasoning as param     Rank-3 via <code>forall</code> nesting (Syntax 1, 2)   ✅ Yes   No new constructs — <code>forall</code> nesting works naturally     Rank-3 return   ✅ Yes   Same reasoning as Rank-3 param     Rank-3 via Syntax 3 <code>co.lang.function</code>   ❌ Compiler error   Same rule as Rank-2; function objects are concrete     Impredicative — workaround (Option C)   initial alpha release   ✅ Yes   Wrap <code>forall</code> type in <code>co.lang.type</code> ; solves 90% of real cases     Impredicative — true opt-in (Option A)    1.0   <code>impredicative:true</code> in <code>@co.dap.generic</code> ; explicit opt-in
<code>@co.dap.generic</code> remains the declaration annotation for constraints, variance, reification, and the other generic metadata described below. A direct generic clause supplies the named parameters and their arity; the two forms may be used together when that metadata is required.
<code>``folang // LinkedList.fol @co.dap.generic(types=[{name=T}]) _ co.lang.struct={ value T; next LinkedList; prev LinkedList; }</code>

k := LinkedList.new(co.lang.int); // when we call new it returns an object of type co.lang.uninit
or
k LinkedList->(T=co.lang.int);
actualList := k.init(); // this is what create a fully formed object of type class
// Employee.fol @co.dap.generic(types=[{name=T},{name= R}]) _ co.lang.class = { id T; name R;
@co.dap.override @co.dap.constructor(access=private) @@init() = {}
@co.dap.override @co.dap.constructor(access=public) @@init(id T, name R) = { this.parent.init(); this.id = id; this.name = name; }
getEmployee(id T)->(Employee)={} }
a := Employee.new(co.lang.int, co.lang.string); b := a.init(1, "Rao");
Ordinary generic construction does not require overriding @@new or @@init. These lifecycle methods are used only when generic construction requires behavior different from the default construction model.
Normal conditions to create/instantiate object of class we just call init which internally call new In specific cases as above we need to do two calls or use call chain like below
c := Employee.new(co.lang.int,co.lang.string).init(1,"Rao");
This new and init methods will be available without you overriding/overloading new and init
// Employee.fol @co.dap.generic(types=[{name=T},{name= R}]) _ co.lang.class = { id T; name R;
}
c := Employee.new(co.lang.int,co.lang.string).init(1,"Rao");
This works folang automaticall provides all type parameters new and all field init implementations.
or
c Employee->(T=co.lang.int, R=co.lang.string);
'''
### Generics Inheritances and Types
''' This is in conceptual stage not supported.
A) Abstract vs concrete type members B) Path-dependent types 1. Type-level projection 2. Path-dependent In folang how it would be
'''
### forall
#### What forall Is — and Is Not
forall is <b>not</b> a general-purpose generic declaration keyword. It is a <b>type-level expression only</b> , restricted to contexts where a polymorphic type must appear as an anonymous value inline — specifically Rank-2 and Rank-3 parameter and return positions, and co.lang.type aliases.
Named generic structs, classes, functions, and methods use @co.dap.generic as their sole generic-parameter declaration mechanism. forall is not a declaration mechanism; forall at declaration level is a <b>compiler error</b> .
***
#### Where forall Is Allowed — Type Expression Form Only
forall(T). followed by an anonymous type body. The . is the syntactic signal that what follows is a type body, not a declaration name.
Pattern:
'''folang // co.lang.type alias — naming a polymorphic type for reuse someFArg co.lang.type = forall(T).(T, T)->(T);
// Rank-2 inline parameter — callee decides what T is someFunction(f forall(T).(T)->(T)) -> (co.lang.int) = {}
// Rank-2 return type — returning a polymorphic function makeldentity() -> (forall(T).(T)->(T)) = {}
// Rank-3 inline parameter — f takes a Rank-2 function applyRank2(f (forall(T).(T, T)->(T)) -> (co.lang.int)) -> (co.lang.int) = {} '''
***
#### Where forall Is Banned — Use @co.dap.generic Instead
'''folang // ❌ compiler error — forall at declaration level forall(T) identity(x T)->(T) = {}
// ✅ correct @co.dap.generic(types=[{name=T,variance=invariant}]) identity(x T)->(T) = {} '''
'''folang // ❌ compiler error // LinkedList.fol forall(T) _ (T) co.lang.struct = { value T; next LinkedList; }
// ✅ correct // LinkedList.fol @co.dap.generic(types=[{name=T}]) _ co.lang.struct = { value T; next LinkedList; } '''

<pre> folang // ❌ compiler error — Rank-1 generics belong to @co.dap.generic forall(T) someFunction(f (T,T)-&gt;(T), a T)-&gt;(T) = {}  // ✅ correct @co.dap.generic(types=[{name=T,variance=invariant}]) someFunction(f (T,T)-&gt;(T), a T)-&gt;(T) = {} </pre>
***
#### Quick Reference
Form   Status   Context    — — —    forall(T) name ...   ❌ Compiler error   Declaration level — use @co.dap.generic instead     forall(T).(T)->(T)   ✅ Allowed   Type level only — Rank-2/3 param, return, co.lang.type alias
<b>The rule in one sentence:</b> forall(T). forms an anonymous polymorphic type expression; it is never a declaration keyword or a file-backed declaration-name mechanism.
&gt; Generic declarations are supported only for structs, classes, functions, and methods. Their type parameters are introduced exclusively by @co.dap.generic.
The following declaration-head generic forms are invalid:
<pre> folang // Cache.fol _ (T) co.lang.module = {} // compiler error // operations.unit.fol _ (F(_)) co.lang.unit = {} // compiler error Callback(T) co.lang.delegate = (T)-&gt;(T); // compiler error </pre>
A parameterized co.lang.type declaration is the separate type-constructor form and does not use @co.dap.generic:
<pre> folang // option.unit.fol _ co.lang.unit = { Option(T) co.lang.type = Some(T)   None(); } </pre>
Generic structs and classes remain file-backed primary declarations. Their names come from filenames:
<pre> folang // LinkedList.fol @co.dap.generic(types=[{name=T}]) _ co.lang.struct = { value T; next LinkedList; prev LinkedList; }  myIntList LinkedList = LinkedList.withTypes(co.lang.int); </pre>
<pre> folang // Employee.fol @co.dap.generic(types=[{name=T}, {name=R}]) _ co.lang.class = { id T; name R; }  emp Employee = Employee.new(co.lang.int, co.lang.string).init; </pre>
Generic functions use the same annotation but are declared inside a legal function-owning context such as an ordinary unit, class, or companion unit: //sommGen1.unit.fol  <pre> folang _ co.lang.unit = {   @co.dap.generic(types=[{name=T}, {name=R}]) add(a T, b T)-&gt;(R) = { ... } someFun()-&gt;()={ add_int_int := add.withTypes(co.lang.int,co.lang.int);    // or    add_int_int co.lang.function = add.withTypes(co.lang.int,co.lang.int);    k := add_int_int(12,10); } } </pre>
***
## Specialization
@co.dap.specialize to specialize generics for specific types upfront //sommGen2.unit.fol <pre> folang _ co.lang.unit = {   @co.dap.generic( types=[ {name=T}], requires=[ co.lang.Add(left=T, right=T, result=T)] ) add(a T, b T)-&gt;(T) = { this.return a + b; } } </pre>  for the above generic want to specialize for co.lang.int //sommGen5.unit.fol  <pre> folang _ co.lang.unit = { @co.dap.specialize(   target=add, types=[ {name=T, type=co.lang.int}] ) addInt(a co.lang.int, b co.lang.int)-&gt;(co.lang.int) = { this.return co.intrinsic.intAdd(a, b); } } </pre>
...
folang provides partial specialization below is the example for partial specializationn //sommGen7.unit.fol
*fields of specialize
Attribute Values   — —    target  the generic fully qualified name includes package name if omitted it is current package     types  resoulution types    priority     strategy intrinsic
***
## Generic Declarations and Type Constructors
FoLang distinguishes annotation-based generic declarations from parameterized co.lang.type constructors.
### Generic Structs, Classes, Functions, and Methods
@co.dap.generic is the sole mechanism for declaring generic structs, classes, functions, and methods. Generic parameters for these declaration kinds must not appear in the declaration head.
<pre> folang // Box.fol @co.dap.generic(types=[{name=T}]) _ co.lang.struct = { value T; } </pre>
<pre> folang // conversion.unit.fol _ co.lang.unit = { @co.dap.generic(types=[{name=T}, {name=R}]) convert(value T)-&gt;(R) = { ... } } </pre>
These forms are invalid:



<pre>folang // Box.fol _(T) co.lang.struct = { ... } // compiler error // Container.fol _(T) co.lang.class = { ... } // compiler error</pre>
co.lang.type Constructors
A <code>co.lang.type</code> constructor does not use <code>@co.dap.generic</code>. Its type parameters appear directly in the type declaration head, and the declaration must be inside an ordinary unit or another explicitly legal type-declaration context.
<pre>folang // option.unit.fol _ co.lang.unit = { Option(T) co.lang.type = Some(T)   None(); }</pre>
<code>Option</code> denotes a unary type constructor:
<pre>text Option : Type -&gt; Type</pre>
Applying it produces a type: <code>//applyingEg1.unit.fol</code>
<code>Some(T)</code> and <code>None()</code> are value constructors declared by the <code>Option</code> definition itself. They do not require separate function implementations.
<code>@co.dap.generic</code> is invalid on <code>co.lang.type</code>, and declaration-head type parameters are invalid on structs, classes, functions, methods, signatures, interfaces, modules, enums, unions, cstructs, units, and other declaration kinds unless a later specification version explicitly adds support.
No Type-Constructor or Type-Function Annotation
FoLang requires neither <code>@co.dap.typeconstructor</code> nor <code>@co.dap.typefunction</code>.
<pre>```text Option(T) co.lang.type = ... -&gt; recognized syntactically as a type-constructor declaration</pre>
<pre>ElementType(container co.lang.type)-&gt;(co.lang.type) = ... -&gt; recognized syntactically as a type-level function ```</pre>
The declaration form already determines the category unambiguously.
<pre>***</pre>
Templates
Typed
<pre>```folang // mytypedtemplate.unit.fol</pre>
<pre>_ co.lang.unit = { @co.dap.template add(a co.lang.int, b co.lang.int)-&gt;(co.lang.int)={ this.return a + b; } } ```</pre>
Untyped
<pre>```folang // MyTemplate.unit.fol</pre>
<pre>_ co.lang.unit = {</pre>
<pre>@co.dap.template add(a, b)-&gt;(co.lang.untyped)={ this.return a + b; } } ``` ***</pre>
Annotations and Decorators
<pre>```folang // Annotation — static object, can carry data</pre>
<pre>// myAnnotation.fol</pre>
<pre>_ co.lang.object-&gt;(for=annotation) = { value co.lang.string; enabled co.lang.bool; }</pre>
<pre>// Decorator — function, transforms target, returns // Decorator.unit.fol</pre>
<pre>_ co.lang.unit = {</pre>
<pre>@co.dap.decorator myDecorator(target co.lang.function)-&gt;(co.lang.function) = { }</pre>
User-defined directives and pragmas cannot be declared. They are language-internal metadata categories. FoLang provides declaration constructs for user-defined annotations and decorators only.
<pre>```</pre>
User-defined annotation/object liveness is defined in <a href="#">Unused Symbols, Liveness, and Reachability</a>.
<pre>***</pre>
Macros  <pre>```folang // a. Basic macro //macro1.unit.fol</pre>
<pre>_ co.lang.unit = {</pre>
<pre>@co.dap.macro say()-&gt;()={ this.return co.macro.quote({ println("Line 1") println("Line 2") }); }</pre>
<pre>// b. Escape assign @co.dap.macro yes_esc_assign()-&gt;(co.lang.untyped)={ this.return co.macro.quote({ co.macro.esc(y) = 42; co.out.println("Inside macro: y =", y); }); }</pre>
<pre>}</pre>
<pre>// c. Debug macro with gensym</pre>
<pre>// macro2.unit.fol</pre>

<pre>_ co.lang.unit = { @co.dap.macro debug(expr)-&gt;(co.lang.untyped)={ tmp := co.macro.gensym(co.lang.var, "tmp"); this.return co.macro.quote({ tmp = co.macro.esc(expr); co.out.println("Result:", tmp); tmp; }); }</pre>
<pre>// d. if/else condition macro @co.dap.macro( group = {items:["if","else"], chain:true}, sugarform={forms:["if expr block"]}, bind={vars: ["x"]}, isolate={vars:["temp", "index"]}, gensym={prefix:"tmp_"}, hygienic=true, argtransform={param:"body", wrap:"lambda", whentype:"block"}, desugar={exprs:["if(\$cond) { \$block }" =&gt; "if(\$cond,\$block)"}}, mode="inject" ) if(condition expr, body block)-&gt;()= {} }</pre>
<pre>blockormacro co.lang.kind = block   macro</pre>
<pre>@co.dap.macro( group={items:["if","else"], chain:true}, sugarform={forms:["else block","else if"]}, chainswith={macro:"if", position:"immediate", required:true}, argtransform={param:"body", wrap:"lambda", whentype:"block"}, standalone=false, desugar= {exprs:[ "else if(\$cond) { \$block }" =&gt; "else(if(\$cond, \$block))", "else { \$elseblock }" =&gt; "else(\$elseblock)"}}, ) else(body blockormacro)- &gt;()= {} } ``</pre>
Other macro utilities: 1. <code>@co.dap.compose(using=["base_if", "blockify"])</code> 2. <code>@co.dap.guard(expr="is_bool_expr(expr)")</code> 3. Quasiquote macros use <code>co.macro.quote</code> and <code>co.macro.unquote</code>
<pre>***</pre>
<pre>## Collections</pre>
<pre>``folang</pre>
<pre>x co.core.List-&gt;(co.lang.string) = co.core.List["A","B","C"];</pre>
<pre>y co.core.Set-&gt;(co.lang.int) = co.core.Set(1,2,3);</pre>
<pre>map co.core.Map-&gt;(key=co.lang.string, val=co.lang.int)=co.core.Map{"A":1, "B":2, "C":3};</pre>
<pre>arr co.core.Array-&gt;(dims=2,type=co.lang.int, sizes=[2,4]);</pre>
<pre>matr co.core.Matrix-&gt;(rows=2,cols=4,type=co.lang.float);</pre>
<pre>//variable with type deduction</pre>
<pre>y := co.core.Set-&gt;(co.lang.int)(1,2,3);</pre>
<pre>x := co.core.List-&gt;(co.lang.string)["A","B","C"];</pre>
<pre>map := co.core.Map-&gt;(key=co.lang.string, val=co.lang.int){"A": 1, "B": 2, "C": 3};</pre>
<pre>`` ***</pre>
<pre>## Execution Models and Control Abstractions (library type=advanced)</pre>
FoLang executes ordinary code <b>sequentially by default</b>. A normal function or method declaration therefore requires no execution-model decorator merely to be called sequentially.
The built-in decorator <code>@co.dap.executionmodel(...)</code> is used only when a declaration requires non-default execution semantics that must remain observable across conforming FoLang implementations. The language exposes an execution-model choice only when that choice changes required FoLang behaviour. A distinction that changes only a backend's internal implementation strategy is not a separate FoLang execution model.
This gives the following separation:
<pre>``text source-level execution semantics -&gt; @co.dap.executionmodel(...)</pre>
<pre>concurrent / parallel / asynchronous submission and communication -&gt; co.cpa</pre>
<pre>continuation / CPS control operations -&gt; co.control</pre>
<pre>backend/runtime implementation -&gt; threads, pools, virtual/green threads, event loops, processes, work-stealing schedulers, or another conforming mechanism ``</pre>
A backend may choose any internal mechanism that preserves the semantics required by the declaration. For example, a logical FoLang task may be implemented using an operating-system thread, a thread pool, a virtual thread, a fiber-like runtime mechanism, an event loop, or another conforming scheduler. Such implementation choices do not become FoLang source-level execution kinds merely because a backend uses them.
Execution-model declarations use ordinary function/method declaration syntax and expose a callable signature, but the frontend represents them with a specialized execution-model declaration node rather than reducing them to an ordinary function node. The specialized node preserves the execution semantics explicitly for semantic analysis and backend interchange. This AST representation does not remove the declaration's callable function/method interface at the language level.
The same principle applies to other function-shaped special declarations such as templates, macros, and native functions: shared surface syntax does not require the frontend to represent every declaration with the same AST node.
<pre>### Default Sequential Execution</pre>
Sequential execution is the default. FoLang therefore does <b>not</b> define <code>@co.dap.executionmodel(type=sequential)</code> as a standard execution-model form; an ordinary undecorated function or method already has sequential semantics.
<pre>``folang calculate(a co.lang.int)-&gt;(co.lang.int) = { this.return a + 1; }</pre>

<code>value := calculate(10); ``</code>
The call proceeds as an ordinary FoLang function call according to the normal call, evaluation-order, error, and return-value rules.
### Concurrent Execution
<code>concurrent</code> describes work whose execution may overlap with other work. The standard source-level kinds are deliberately small:
- <code>task</code> — a logical schedulable unit of work. This is the preferred general concurrent abstraction and leaves the backend substantial scheduling freedom. - <code>thread</code> — requests thread-semantic execution when thread identity, affinity, thread-local behaviour, or another specification-defined thread property is intentionally observable.
Backend-specific names such as <code>goroutine</code> or <code>greenlet</code> are not FoLang execution kinds. <code>subroutine</code> is ordinary callable code, while <code>generator</code> and <code>coroutine</code> primarily describe suspension/resumption behaviour rather than an independent concurrent execution carrier. <code>fiber</code> likewise is not exposed merely because a backend may use a fiber-like mechanism internally. Such abstractions may exist as ordinary package/runtime facilities when separately defined, but they are not synonyms for <code>concurrent.kind</code> .
<code>``folang // concurrent.unit.fol _ co.lang.unit = { @co.dap.executionmodel(type=concurrent, kind=task) someConcurrent(a co.lang.int)-&gt;(co.lang.int,co.lang.error) = { ... } }</code>
<code>co.cpcp.submit(someConcurrent, params=[10], results=[val, errors]); (errors.isEmpty).then( co.out.println(val) ).default( co.out.println(errors) ); ``</code>
A declaration may request thread semantics explicitly when those semantics are required:
<code>``folang _ co.lang.unit = { @co.dap.executionmodel(type=concurrent, kind=thread) threadBoundWork(a co.lang.int)-&gt;(co.lang.int,co.lang.error) = { ... } }</code>
<code>co.cpcp.submit(threadBoundWork, params=[10], results=[val, errors]); ``</code>
Scheduling is a separate dimension from execution kind. Where supported by the selected kind, an optional scheduling policy may be declared:
<code>folang @co.dap.executionmodel( type=concurrent, kind=task, scheduling=cooperative )</code>
The standard scheduling values are <code>cooperative</code> and <code>preemptive</code> . When the field is omitted, the backend/runtime may choose any conforming scheduling policy. A structural policy such as <code>fork-join</code> , when requested, is likewise separate from both the execution kind and the scheduling policy; omitting it provides no fork-join guarantee.
### Parallel Execution
<code>parallel</code> states that the declaration has parallel-execution semantics. FoLang does not currently require a <code>parallel.kind</code> merely to expose an implementation mechanism:
<code>folang @co.dap.executionmodel(type=parallel) parallelWork(data SomeData)-&gt;(Result) = { ... }</code>
<code>process</code> , <code>fork</code> , <code>spawn</code> , and <code>exec</code> are not parallel execution kinds. A process is an execution/isolation container, while <code>fork</code> , <code>spawn</code> , and <code>exec</code> are process/runtime operations. They remain ordinary facilities of the applicable <code>co.cpcp</code> , <code>co.os</code> , or system API rather than values of <code>@co.dap.executionmodel(type=parallel, kind=...)</code> .
A conforming backend may implement parallel semantics with threads, worker pools, processes, accelerators, or another mechanism when the program cannot observe a forbidden difference.
### Asynchronous Execution
<code>async</code> describes execution whose completion is not required to block the submitting context. Completion representation is separate from the execution model:
<code>folang @co.dap.executionmodel(type=async, completion=future) loadData(request Request)-&gt;(Data, co.lang.error) = { ... }</code>
The standard completion forms are:
- <code>future</code> — the caller receives or is associated with a future completion value; - <code>callback</code> — completion is delivered through a compatible callback mechanism.
<code>promise</code> is not a separate execution-model value. If a producer-side completion object is required by a particular runtime/API, it is an implementation or package abstraction rather than another async execution kind. <code>await</code> is a consumption or control operation supplied by <code>co.cpcp</code> ; it is not an execution-model kind.
An actor is a concurrency abstraction with state/message semantics, not an async completion kind. Channels are communication abstractions, events are notification abstractions, and distributed scaling is a deployment/runtime policy. These may be provided through <code>co.cpcp</code> , but none is a top-level <code>@co.dap.executionmodel</code> type merely because it participates in concurrent or asynchronous programs.
### Continuations and Control
Continuation semantics are deliberately separated from concurrent, parallel, and asynchronous execution facilities. A declaration may use:

- <code>kind=full</code> — full/undelimited continuation semantics; - <code>kind=delimited</code> — continuation capture is bounded by a delimiter.
Continuation and CPS operations are supplied by <code>co.control</code> , not <code>co.cpcap</code> .
<code>foolang @co.dap.executionmodel(type=continuation, kind=delimited, control="shift-reset")</code> <code>continuationWork(...)-&gt;(…) = { ... }</code>
For delimited continuations the currently documented control families are <code>shift-reset</code> , <code>prompt-control</code> , and <code>spawn-yield</code> . <code>csp</code> is not a continuation kind; Communicating Sequential Processes belong to the concurrency/communication domain.
<code>spawn-yield</code> has a specific continuation meaning in FoLang and must not be confused with an ordinary process/task spawn operation or with a scheduler-only yield. In the <code>spawn-yield</code> control family, <code>spawn</code> establishes an <b>implicit delimited-control boundary</b> . When execution reaches <code>yield</code> , the continuation from that suspension point back to the spawned control boundary is preserved internally by the control runtime. The suspended execution can later resume from the yield point.
Conceptually:
<code>text spawn</code> — establishes implicit continuation delimiter — A — B — <code>yield</code> — preserve/suspend continuation internally up to the spawned control boundary ↓ resume ↓ continue after yield
Unlike <code>shift-reset</code> and <code>prompt-control</code> , <code>spawn-yield</code> does not require the captured continuation to be exposed directly to application code as an explicit continuation value such as <code>k</code> . The continuation may remain an internal control-runtime object. The observable contract is suspension at <code>yield</code> , preservation of the delimited execution state, and later resumption from that suspension point according to the applicable <code>co.control</code> operation.
This distinction is important:
““text ordinary <code>co.cpcap</code> / process spawn -> creates or submits concurrent execution
ordinary scheduler yield -> voluntarily gives execution opportunity to another schedulable unit
<code>co.control spawn-yield</code> -> establishes an implicit delimited continuation boundary and suspends/resumes execution by preserving the continuation internally ““
Accordingly, the word <code>spawn</code> may occur in more than one conceptual domain, but the package and operation determine its semantics. Concurrent/process creation remains under the applicable <code>co.cpcap</code> , <code>co.os</code> , or system API. Continuation-oriented <code>spawn-yield</code> operations belong to <code>co.control</code> .
Continuation kind   Control family   Delimiter model   Continuation exposure   Meaning    — — — —    <code>full</code>   supplied by the applicable <code>co.control</code> full-continuation API   no delimited boundary   according to the full-continuation API   captures the continuation according to the full-continuation contract     <code>delimited</code>   <code>shift-reset</code>   explicit nearest/applicable <code>reset</code> boundary   exposed according to Shift/Reset semantics   captures only the continuation bounded by <code>reset</code>     <code>delimited</code>   <code>prompt-control</code>   explicit/applicable prompt boundary   exposed according to Prompt/Control semantics   captures only the continuation bounded by the applicable prompt     <code>delimited</code>   <code>spawn-yield</code>   implicit spawned control boundary   internal by default   suspends at <code>yield</code> , preserves the bounded continuation internally, and later resumes from the yield point
Parameter table for delimited continuations
Operator Pair   Delimiter Targeting   Captures Delimiter into <code>\$k\$?</code>   Keeps Delimiter on Stack During <code>\$expr\$?</code>   Primary Modern Use Case / Paradigm     <code>—</code>   <code>—</code>   <code>—</code>   <code>—</code>   <code>—</code>     <b>Shift / Reset</b>   Nameless (Closest)   <b>Yes</b> (Implicitly)   <b>Yes</b>   Standard algebraic effects, functional stream processing     <b>Prompt / Control</b>   Nameless (Closest)   <b>No</b>   <b>Yes</b>   Dynamic context modification, legacy Scheme systems     <b>Spawn / Yield</b>   Implicit (Fiber Bound)   <b>Internal</b>   <b>No</b>   Cooperative multitasking, <code>async/await</code> engines, generators
Example declaration:
<code>foolang @co.dap.executionmodel( type=continuation, kind=delimited, control="spawn-yield" ) resumableWork(...)-&gt;(…) = { ... }</code>
The decorator identifies the required continuation family; the actual suspension, yield, resume, and continuation-control operations are supplied by <code>co.control</code> .
### Canonical Submission and Placement
<code>co.cpcap.submit(...)</code> is the canonical submission operation for declarations whose execution is governed by <code>co.cpcap</code> . The execution-model decorator states the required semantics; the submission call requests execution of that declaration.
A specialized operation such as <code>submitToPool</code> , <code>submitThread</code> , or <code>submitToEventLoop</code> is justified only when the standard <code>co.cpcap</code> API explicitly defines an additional observable placement/resource constraint. Such a function must not exist merely as another spelling of <code>submit(...)</code> . If no additional semantic guarantee is defined, <code>submit(...)</code> is used and the backend/runtime is free to select the underlying mechanism.
Application pragmas may configure runtime resources without changing declaration semantics:
<code>foolang @co.pdap.threadpool(min=10,max=40,increment=10) @co.pdap.schedulerpool(min=10,max=40,increment=5)</code>
Pragmas are applicable to applications only; they are not applied by libraries, source libraries, packages, or exports. A library may document resource assumptions or recommendations, but it cannot impose application pragmas. Explicit pool-size configuration is optional; when absent, the FoLang runtime/backend supplies its configured defaults.

### Execution-Model Summary			
Type	Source-level kind/model	Optional semantic dimension	Package responsible for operations           default
sequential   —   —	ordinary function/method call	concurrent   task, thread   scheduling=cooperative preemptive ;	
structural policies such as fork-join	when separately defined	co.cpcap   parallel   —	parallel-specific policy only when
separately specified	co.cpcap     async   —	completion=future callback   co.cpcap   continuation   full ,	
delimited   delimited	control="shift-reset" "prompt-control" "spawn-yield"	co.control	
The following are intentionally <b>not</b> execution-model types or kinds:			
``text backend/runtime mechanisms: goroutine, greenlet, virtual thread, event-loop implementation, work-stealing implementation			
ordinary process/runtime operations: process, fork, spawn, exec			
completion/control operations: await			
communication/notification abstractions: channel, event			
higher-level concurrency abstractions: actor, CSP			
control/suspension abstractions unless separately specified: coroutine, generator, fiber, yield/resume			
deployment/runtime policy: distributed scaling, scale up/down/out/in ``			
The governing rule is:			
> Expose an execution-model choice in <code>@co.dap.executionmodel</code> only when the > choice changes externally observable FoLang semantics. If two choices differ > only in how a backend implements the same required behaviour, they are not > separate FoLang execution-model choices.			
***			
## Native Code (Library type=system)			
<code>@co.dap.native</code> marks a function or method declaration as a <b>native implementation declaration</b> . It does not grant native capability merely because the annotation is present. The annotation is valid only inside a <code>system</code> library/source-library domain when the installation permits the system capability.			
A system library may use the <code>co.native</code> package to express low-level implementation such as assembly or machine-level operations through facilities including <code>co.native.asm</code> and <code>co.native.inline</code> . This is the FoLang facility analogous to embedding an assembly block in a low-level systems language.			
The frontend preserves a native declaration as a specialized declaration node and carries its required semantics through the backend interchange contract. The reference backend demonstrates one implementation of that contract. A conforming third-party backend is not required to reproduce the reference backend's internal native-code lowering, instruction representation, allocation strategy, or code-generation mechanism, but the externally observable behavior required by the specification must be preserved.			
<code>ffi</code> is a separate foreign-function-interface domain. It is used to describe and invoke foreign functions through extern declarations, C-ABI-compatible types, calling conventions, symbol linkage, pointer/address forms, and any permitted FoLang-side marshalling or invocation code. An <code>ffi</code> library does <b>not</b> gain <code>@co.dap.native</code> or direct FoLang assembly/machine-code implementation capability merely because FFI itself permits low-level ABI types.			
### Native Functions // native.unit.fol			
*** ## Dynamic Runtime (library type=dynamicvmrt)			
The <code>@co.ddap.dynamicruntime</code> annotation enables full access to the <code>co.meta</code> package. Through this package, developers can use dynamic class and type loading, monkey patching, runtime reflection, instrumentation, eval-based code execution, and other advanced metaprogramming capabilities.			
A <code>dynamicvmrt</code> library provides these capabilities only within its own library boundary; they do not automatically escape into ordinary application or library code.			
1. runtime type creation from strings, streams, files, and other supported dynamic-runtime inputs 2. complete reflection and introspection 3. Runtime code modification add/remove/update methods etc.			
> When a library is marked <code>dynamicvmrt</code> and enables <code>@co.ddap.dynamicruntime</code> , the final binary includes the runtime support required to create and manage dynamic types, methods, and objects.			
> Dynamically created objects may interact with compiled types according to the dynamic-runtime boundary rules; ordinary compiled code does not gain unrestricted reverse access to dynamic-runtime facilities.			
> Loaders are the dynamic-runtime containers used to manage these runtime-created objects and types.			
> FoLang provides BasicLoader but you can extend this as follows			
//MySpecLoader.fol			
``folang			
@co.dap.extends(co.meta.BaseLoader) _ co.lang.loader={			
} ``			

> The basic loader provides the operations required to create, update, delete, and manage runtime-created objects.
> Loaders form a hierarchy. When a referenced runtime type is not found in the current loader realm, lookup proceeds through the base-loader chain and finally to the compiled-type environment.
***
## Variable Kinds Support
Kind   Where   <code>---</code>   Normal   All     Pointers   Library of type <code>system</code> and <code>ffi</code>   Arrays   All     References Heap, Lvalue, Rvalue   Library of type <code>system</code> and <code>ffi</code>   Addresses   Library of type <code>system</code> and <code>ffi</code>   Thunks   Library <code>application</code> or <code>system</code> or <code>ffi</code> or <code>advanced</code> or <code>dynamicvmrt</code>   Ranges   All     Slices   All
***
## Builtin Data Types
Type   Kind   <code>---</code>   <code>co.lang.string</code>    <code>co.lang.int</code>    <code>co.lang.bit</code>    <code>co.lang.double</code>    <code>co.lang.float</code>    <code>co.lang.long</code>    <code>co.lang.byte</code>    <code>co.lang.char</code>    <code>co.lang.any</code>    <code>co.lang.dynamic</code>    <code>co.lang.auto</code>    <code>co.lang.bool</code>    <code>co.lang.void</code>    <code>co.lang.value</code>   value types stores values when take snapshot   <code>co.lang.typed</code>    <code>co.lang.untyped</code>    <code>co.lang.word</code>    <code>co.lang.MatchBindings</code>    <code>co.lang.tag</code>    <code>co.lang.typevalue</code>    <code>co.lang.uninit</code>    <code>co.lang.error</code>    <code>co.lang.literal</code>   literal representation for simple and compound literal objects   <code>co.lang.operator</code>   operator-source-only declaration kind; invalid in ordinary FoLang source
A name appearing in this registry is not necessarily an enabled source-language feature. A built-in kind is usable only when this specification defines its declaration syntax and semantics. An undocumented or explicitly reserved kind remains unavailable and must produce an unsupported-feature diagnostic when used.
***
## Built-in Directives   Kind   <code>---</code>   <code>PRAGMA</code>   <code>"@co.pdap.threadpool", "@co.pdap.schedulerpool"</code>   <code>DIRECTIVE</code>   <code>"@co.ddap.import", "@co.ddap.dynamicruntime", "@co.ddap.use", "@co.ddap.alias"</code>    <code>ANNOTATION</code>   <code>"@co.dap.template", "@co.dap.macro", "@co.dap.operator", "@co.dap.annotation", "@co.dap.library", "@co.dap.module", "@co.dap.native", "@co.dap.class", "@co.dap.static", "@co.dap.instance", "@co.dap.object", "@co.dap.inline", "@co.dap.ctfe", "@co.dap.friend", "@co.dap.sealed", "@co.dap.extension", "@co.dap.override", "@co.dap.virtual", "@co.dap.abstract", "@co.dap.delegate", "@co.dap.dynamicscope", "@co.dap.lexicalscope", "@co.dap.staticscope", "@co.dap.mixedscope", "@co.dap.typeclass", "@co.dap.matcher", "@co.dap.constructor", "@co.dap.oops", "@co.dap.extends", "@co.dap.hokrlt", "@co.dap.indexer", "@co.dap.generic", "@co.dap.comptime", "@co.dap.typefromvalue", "@co.dap.local", "@co.dap.private", "@co.dap.public", "@co.dap.package", "@co.dap.protected", "@co.dap.internal", "@co.dap.export", "@co.dap.eager", "@co.dap.lazy", "@co.dap.packed", "@co.dap.declare", "@co.dap.simd", "@co.dap.reflection", "@co.dap.mop", "@co.dap.nested", "@co.dap.inner", "@co.dap.final", "@co.dap.const", "@co.dap.decorator", "@co.dap.specialize"</code>   <code>//mop =&gt; meta object programming</code>   <code>DECORATOR</code>   <code>"@co.dap.before", "@co.dap.after", "@co.dap.around", "@co.dap.onErrExcept", "@co.dap.InvokeAlways", "@co.dap.HandleEffect", "@co.dap.defer", "@co.dap.callable", "@co.dap.executionmodel"</code>
***
## Builtin Kinds   Kind   Purpose   <code>---</code>   <code>co.lang.type</code>    <code>co.lang.struct</code>    <code>co.lang.cstruct</code>    <code>co.lang.class</code>    <code>co.lang.interface</code>   all abstract methods   <code>co.lang.union</code>    <code>co.lang.object</code>    <code>co.lang.instance</code>    <code>co.lang.matcher</code>    <code>co.lang.loader</code>    <code>co.lang.trait</code>   interfaces with default implementations   <code>co.lang.mixin</code>   abstract classes alias   <code>co.lang.extension</code>   reusable implemented functions that can be composed with supported classes without inheritance   <code>co.lang.delegate</code>    <code>co.lang.typeclass</code>    <code>co.lang.module</code>    <code>co.lang.unit</code>   stateless file-level container; ordinary units merge into the package namespace and <code>*.comp.unit.fol</code> attaches to a struct   <code>co.lang.block</code>    <code>co.lang.package</code>    <code>co.lang.signature</code>    <code>co.lang.function</code>    <code>co.lang.newtype</code>    <code>co.lang.opaquetype</code>    <code>co.lang.subtype</code>    <code>co.lang.supertype</code>    <code>co.lang.dependentType</code>   result kind of a value-indexed type-level function   <code>co.lang.refinementType</code>   base type restricted by a Boolean predicate over the candidate value   <code>co.lang.associatedType</code>   signature associated-type requirement or matching-module associated-type binding   <code>co.lang.data</code>    <code>co.lang.enum</code>    <code>co.lang.library</code>
## Builtin Collections
Name   Purpose   <code>---</code>   <code>co.core.List</code>    <code>co.core.Set</code>    <code>co.core.Map</code>    <code>co.core.Tree</code>    <code>co.core.Trie</code>    <code>co.core.Array</code>    <code>co.core.Tuple</code>
## Builtin Operators
### Arithmetic operators <code>+</code> , <code>-</code> , <code>*</code> , <code>/</code> , <code>%</code> , <code>**</code>
### Logical operators <code>&amp;&amp;</code> , <code>  </code> , <code>!</code>
### Bitwise operators <code>&amp;</code> , <code> </code> , <code>^</code>
### Comparison operators <code>==</code> , <code>!=</code> , <code>&lt;</code> , <code>&gt;</code> , <code>&lt;=</code> , <code>&gt;=</code>
### Standard operator examples
```folang left co.lang.int = 10; right co.lang.int = 3;
notEqual co.lang.bool = left != right; // true

bitsAnd co.lang.int = 6 & 3; // 2 bitsOr co.lang.int = 6 3; // 7 bitsXor co.lang.int = 6 ^ 3; // 5
mulAssign co.lang.int = 6; mulAssign *= 3; // 18
divAssign co.lang.int = 18; divAssign /= 3; // 6
modAssign co.lang.int = 17; modAssign %= 5; // 2
powAssign co.lang.int = 2; powAssign **= 3; // 8
andAssign co.lang.int = 6; andAssign &= 3; // 2
xorAssign co.lang.int = 6; xorAssign ^= 3; // 5
orAssign co.lang.int = 6; orAssign = 3; // 7 ```
For a compound assignment <code>lhs op= rhs</code> , FoLang resolves the corresponding binary operator <code>op</code> , evaluates the left-hand location only once, and stores the resulting value back through that same location. The ordinary target-type conversion rules apply to the stored result.
Other operator and language-token spellings <code>@</code> , <code>#</code> , <code>!</code> , <code>~</code> , <code>\$</code> , <code>^</code> , <code>(</code> , <code>)</code> , <code>_</code> , <code>`</code> , <code>?</code> , <code>{</code> , <code>[</code> , <code>]</code> , <code>}</code> , <code>\</code> , <code>:</code> , <code>;</code> , <code>"</code> , <code>'</code> , <code>=</code> , <code>.</code> , <code>?=</code> , <code>:=</code> , <code>::=</code> , <code>,</code> , <code>..</code> , <code>...</code> , <code><..</code> , <code>..<code><</code></code> , <code><..<code><</code></code> , <code>==></code> , <code>=></code> , <code>=</code> , <code>-></code> , <code><-</code> , <code>->></code> , <code><-></code> , <code>@@</code> , <code>+=</code> , <code>-=</code> , <code>*=</code> , <code>/=</code> , <code>%=</code> , <code>**=</code> , <code>&=</code> , <code>^=</code> , <code> =</code>
Contiguous symbolic spellings that are absent from this inventory and from the active custom-operator table are unrecognized symbolic tokens; the lexer does not split them into shorter operators.
This inventory includes punctuation and reserved token spellings. Presence in this list does not make a spelling declarable as a <code>co.lang.operator</code> , nor usable with <code>mode=overload</code> , <code>mode=implements</code> , <code>mode=extends</code> , <code>mode=inherits</code> or <code>mode=override</code> .
Pre-Declared Operator Glyphs <code>U</code> , <code>n</code>
Every glyph in this list is language-owned and already has fixed parse properties. It cannot be redeclared with <code>co.lang.operator</code> , but it can receive implementations through <code>mode=overload</code> in a class, struct companion unit, or built-in extension package contribution. Until a matching implementation is visible, use of the glyph fails during resolution.
// union_inter_eg.unit.foI
```foIang
_ co.lang.unit = {
@co.dap.operator(symbol='U', mode=overload) @co.dap.extension(fortype=co.core.Set, what=extends) union(other co.core.Set)-> (co.core.Set) = { .... }
@co.dap.operator(symbol='n', mode=overload) @co.dap.extension(fortype=co.core.Set, what=extends) intersection(other co.core.Set)->(co.core.Set) = { .... } }
// Uses of the pre-declared glyphs after matching overloads are visible. v co.core.Set = co.core.Set(1, 2, 3); p co.core.Set = co.core.Set(4, 5, 2); w co.core.Set = co.core.Set(7, 8);
u := v U p; i := v n p; combined := v n p U w; // same precedence and left associativity: (v n p) U w
z := p + v U w; // parses as p + (v U w) x := p * v U w; // parses as (p * v) U w ```
> for UDT like classes and companion unit there is no @co.dap.extension as a developer can directly implement into the new type upfront
See <a href="#">Pre-Declared Operator Glyphs</a> .
***
### Reserved words <code>co</code> , <code>let</code> , <code>this</code> , <code>for</code> , <code>forall</code> , and <code>fo</code> are hard-reserved words. <code>self</code> is a contextual keyword available in every method declared by a <code>co.lang.class</code> , including the lifecycle methods <code>@@new</code> and <code>@@init</code> .
<code>fo</code> is permanently reserved. The language originally reserved both <code>co</code> and <code>fo</code> during its naming evolution; FoLang retains <code>fo</code> as a language-owned word rather than allowing it to become an ordinary identifier. This preserves the established lexer/parser classification and prevents programs from using <code>fo</code> as a variable or declaration name.
### Difference between <code>this</code> and <code>self</code> - <code>this</code> refers to the applicable instance/object receiver. - <code>self</code> is available throughout class methods, including <code>@@new</code> and <code>@@init</code> . - <code>self</code> is contextual rather than globally hard-reserved; outside a class-method context it has no special class-method meaning. - <code>static</code> has no special shortcut; use the applicable variable or class name explicitly. - Where the class-member rules permit access, both <code>self</code> and <code>this</code> may be used to reach class/member state according to their receiver context.

## Special methods

Method	Responsibility
@@new	Lifecycle method for class
@@init	Lifecycle method for class

## Builtin Packages

### `co` — root (reserved word)

`co` is the root of the standard package tree supplied with FoLang. The tree is distributed as a separate language-provided export-kind `.fo1enc` artifact and is loaded automatically by the compiler/toolchain. The developer therefore receives implicit access to the `co` root without an import, while the declarations inside the artifact retain ordinary exported-package semantics.

The table below describes the current standard package hierarchy and API responsibilities. After version 1.0, the package/subpackage paths in this hierarchy are fixed, but the declarations contained inside an existing package are not frozen. Later standard-package artifact versions may add or update ordinary types, unit-level functions, methods, data structures, algorithms, modules, and other declarations without creating new FoLang grammar or a new package path.

Sub-package	Responsibility
<code>co.lang</code>	All data types and kinds
<code>co.sys</code>	file, concurrent, parallel, goto, invoke, bind, call, apply, setTimeout, setInterval, scheduler, cron, event
<code>co.os</code>	signal, cmd, execute, run, env, getenv, setenv, sleep, exit, cwd, chdir, fork, wait, pipe, dup, dup2, close, readfd, writefd, random
<code>co.meta</code>	ast, instrument, transform, augment, reflect, introspect, patch, inject, create, runtime(eval,proto, prototype,etc), realm
<code>co.core</code>	List, Set, Map, Tree, Trie, Sort, Search, Array, Pointer, Ref, Address, Ptr, Matrix, Word
<code>co.native</code>	load, register, asm, inline, emit, ffi, spawnnon[gpu,cpu, npu,apu, fpga,asic, tpu,mki,mcu], arch[x86,x86-64,risc,arm,vliw]
<code>co.in</code>	read, readln
<code>co.out</code>	println, print
<code>co.regex</code>	stex, pattern, match, search
<code>co.crypto</code>	rsa, aes, hash, md5, rand, uuid, ssl, tls
<code>co.dap</code>	built-in decorators, annotations
<code>co.ddap</code>	built-in directives
<code>co.pdap</code>	built-in pragmas
<code>co.net</code>	tcp, udp, http
<code>co.const</code>	<code>true</code> , <code>false</code> , <code>none</code>
<code>co.encoding</code>	base64Encode, base64Decode, json, yaml, bson
<code>co.utils</code>	makeImmutable, makeShared, copyOnWrite, toSnapshot — object behaviour policies
<code>co.dynamic</code>	dynamic capabilities
<code>co.runtime</code>	
<code>co.compiletime</code>	
<code>co.macro</code>	
<code>co.pattern</code>	
<code>co.control</code>	continuation, CPS, full/delimited continuation control, shift/reset, prompt/control, and continuation-oriented control abstractions
<code>co.cpa</code>	concurrent/parallel/async submission, task/thread execution facilities, future/callback completion, await, pools, channels, events, actors, process/distributed facilities, scheduling, fiber/coroutine facilities, defer, lazy, and related execution APIs
<code>co.hokrlt</code>	
<code>co.operator</code>	



## FoLang Philosophy — Uniform Object Model

---

### Core Principle

Everything in **FoLang** is an object. That is the reason for the name **FO — Functional Objects**.

FoLang gives the programmer one uniform object model across all types instead of forcing them to switch between separate type families for:

- ordinary values
- immutable values
- shared/concurrent values
- copy-on-write values
- literal values

The programmer writes against a single conceptual model and opts into the required object behaviour only when needed.

---

### Uniform Object Principle

In FoLang, **everything is an object**.

Unless explicitly stated otherwise, the reference and mutation rules in this section apply to **managed FoLang objects**. `co.lang.cstruct` remains an object in the language model, but it is an explicitly value-semantic ABI representation: assignment and parameter passing copy its value rather than a managed object reference.

This includes:

- scalar objects
- collection objects
- user-defined objects
- ADT objects
- function objects

Functions are objects in the same sense as all other values.

This is true for both **named functions** and **anonymous functions**.

FoLang does not introduce a separate semantic model for functions, scalars, UDTs, or ADTs.

They all follow the same core object principles:

- assignment of managed objects copies references
- assignment of `co.lang.cstruct` copies its value
- `==` compares values
- mutation is applied through the object
- behaviour policies such as immutability, shared, copy-on-write, and literal conversion apply uniformly where meaningful

So in FoLang, the programmer does not need one mental model for data objects and another for function values.

The language treats them under one consistent object model.

---

### 1. Default Object Model

Every value in FoLang is an object and is **mutable by default**.

This includes:

```
co.lang.int, co.lang.float, co.lang.string → built-in scalar types
user-defined types (structs, classes, ADTs) → user-defined objects
co.core.List, co.core.Map, co.core.Array → built-in collections
```

Example:

```
positive_int co.lang.int = 10;
```

`positive_int` denotes an object.

Its current value is `10`.

By default, it is mutable.

---

### 2. Assignment, Aliasing, and Mutation

FoLang distinguishes three related ideas:

- **assignment**

- aliasing
- mutation

## 2.1 Assignment copies references

When one object variable is assigned to another, FoLang copies the **reference**, not the internal contents.

```
b := a;
```

After this, `a` and `b` refer to the same object.

## 2.2 Mutation changes object contents

If two names refer to the same object, mutating through one name is visible through the other.

```
// Employee.fol
_ co.lang.class = {
 Name co.lang.string
}

a Employee = { Name = "Kamesh" };
b := a;
c Employee = { Name = "Kamesh" };

a == b; // true
a == c; // true
b == c; // true

b.Name = "Ramesh"; // changes a.Name too
c.Name = "Ramesh"; // does not change a.Name
```

Why:

- `a` and `b` are the same object
- `c` is a different object with the same earlier value

## 2.3 `==` means value equality

In FoLang, `==` compares **values**, not object identity.

That rule is uniform across all object kinds, including built-in types and user-defined types.

So:

```
a == b;
```

means:

- compare contents deeply
- return `true` when the values match
- even if `a` and `b` are not the same object

## 2.4 Mutation accessors

The accessor used for mutation depends on the object type:

scalar (int, float, bool, string, char, byte)	→ <code>.value</code>
struct field	→ <code>.fieldname</code>
class member	→ <code>.membername</code>
map entry	→ <code>.key</code>
array / list element	→ <code>.index</code>

Examples:

```
count.value = 30;
emp.name = "Rao";
marks.math = 95;
nums[0] = 42;
```

## 2.5 Function-call behaviour

Function parameters introduce a **new local binding**.

So:

- rebinding a parameter does **not** affect the caller
- mutating the object through the parameter **does** affect the caller if both still refer to the same object

```
checkPositive(a co.lang.int)->(co.lang.bool) = {
 a = 20; // local rebinding only – caller unchanged
 a.value = 30; // object mutation – caller sees 30
}
```

If `positive_int` is passed to `checkPositive(a)`:

```
a = 20 → local rebinding only
a.value = 30 → mutates the passed object
```

### 3. Literal Objects

All literals in FoLang create **Literal Objects**.

```
"Kumar" → Literal Object – string
42 → Literal Object – int
true → Literal Object – bool
```

A Literal Object is an **anonymous object created from a literal expression**.

Literal Objects participate in value equality just like every other object in FoLang:

```
10 == 10; // true
```

In FoLang, `==` compares **values**, not object identity.

So two literal expressions with the same value compare equal, but that does not mean they are the same object.

#### Binding a Literal Object

When a literal is assigned to a named object, the name refers to the object created from that literal expression.

```
a co.lang.int = 10;
```

Here, `a` refers to the anonymous integer object created from literal `10`.

So:

```
a.value = 30;
```

mutates that same object, and `a` now has value `30`.

The important points are:

- literal expressions create anonymous objects
- literal-created objects are still ordinary objects unless another policy is applied
- `==` compares values, not identity
- once bound to a name, a literal-created object behaves like any other normal mutable object
- immutability applies only when explicitly requested, for example through `makeImmutable(...)`

#### Why Literal Objects Cannot Be Mutated Directly

Mutation can occur in two ways:

##### 1. Through rebinding

```
a co.lang.int = 10; a = 20;
```

The above statement is valid because `a` is a valid identifier and named handle to `co.lang.int` type.

```
10 = 20; // ❌ invalid
```

The above statement is invalid because `10` is not a valid identifier and if `10` in literal object type there is no named handle.

The rebinding happens through named handles which are valid identifiers of `foLang`

##### 2. Through a property or method

An object may also be mutated through one of its mutable properties or methods. `a co.lang.int = 10; a.value = 20;`

A literal object cannot be mutated directly because it does not provide properties/methods that can be accessed.

10.value = 20; // ❌ invalid

```
a co.lang.int = 10;
a.value = 30;
```

is valid after binding, a bare literal such as `10` cannot be mutated directly because there is no name or handle through which to perform mutation.

### Not the Same as `toSnapshot(...)`

A literal object created by a literal expression is **not** the same thing as the internal snapshot representation produced by

`co.utils.toSnapshot(...)`.

- literal expression → anonymous object
- `toSnapshot(x)` → internal snapshot representation used to reconstruct a fresh local object later

These are related concepts, but they are not the same mechanism.

### Types of Literal objects:

1. Simple types
2. Complex or compound types (UDTS)

Simple literal forms include values such as ``10``, ``'A'``, and ``"A"``.

```
k co.lang.int=10;
```

Compound types are in Json form

```
// Employee.fol
_ co.lang.class = {
 id co.lang.string;
 name co.lang.string;
}
```

```
k Employee = Employee{ id: "10", name: "ABC" };
```

> Even though compound literals are ended with brace block we need to end with semicolon as it is value not block.

### Summary

- literal objects are real objects
- literal expressions create anonymous objects
- identical literals compare equal by value
- identical literals are not automatically the same object
- literal-created objects are mutable by default once bound to a handle
- only `makeImmutable(...)` makes an object immutable

## 4. Object Behaviour Policies

Any managed object can be given a behaviour policy using `co.utils.*`.

The policy operations are **in-place transformations of the object graph**:

- the object itself changes behaviour kind;
- there is no wrapper object;
- there is no alternate binding to capture;
- aliases continue to refer to the same transformed object graph.

Object policy belongs to the **object graph**, not to a variable name or alias. If two or more bindings refer to the same managed object, applying Immutable, Shared, or CopyOnWrite behaviour through one alias is observed through every alias that still refers to that object. Rebinding one alias later does not change the policy of the original object graph.

All policies are **deep by default**. The reachable object graph is the transitive closure obtained by starting at the policy root and recursively following managed-object references through members, nested objects, collection elements, and other managed references. Repeated references and cycles identify the same reachable object rather than creating additional logical objects. The applicable policy therefore governs

the complete reachable graph, not only the root object.

Immutable, Shared, and CopyOnWrite are mutually exclusive object policies. Once an existing object graph enters one of these policy states, that policy is permanent for the lifetime of that graph; it cannot later be changed into either of the other policy states. `makeValueImmutable(x)` and `makeImmutable(x)` both make the current object graph Immutable; `makeImmutable(x)` additionally prevents rebinding of the binding supplied as `x`. Binding immutability is distinct from object policy and does not make other aliases non-rebindable.

#### 4.1 Immutable

```
co.utils.makeImmutable(positive_int);
```

After this call the object itself is immutable.

This is an in-place transformation, not a wrapper.

Any attempt to mutate the object or any reachable object beneath it is a **compiler error** where detectable statically, and a **runtime error** otherwise.

```
positvie_int co.lang.int = 10;
positive_int.value = 30 // ✗ compiler error
positive_int = 30 // ✗ compiler error
```

For nested objects:

```
// Employee.fol
_ co.lang.struct = { address Address; }
Address co.lang.Address = {city co.lang.string; state co.lang.string; lane co.lang.string; pin co.lang.string;}

emp Employee = Employee{
 address: Address{
 city: "Pune"
 }
};
co.utils.makeImmutable(emp);

emp = Employee{
 address: Address{
 city: "ABC"
 }
}; // ✗ compiler error

emp.address = newAddress // ✗ compiler error
emp.address.city.value = "Mumbai" // ✗ compiler error
```

Immutability is deep and total.

#### 4.2 Value Immutable

```
positive_int co.lang.int = 20;

co.utils.makeValueImmutable(positive_int);

positive_int.value = 30; // ✗ current object is immutable
positive_int = 30; // ✔ binding may reference a new object

co.utils.makeValueImmutable(emp);

emp.address.city = "ABC"; // ✗
emp.address.city.value = "ABC"; // ✗

emp = Employee{
 address: Address{
 city: "ABC"
 }
}; // ✔
```

## Difference between Immutable and Immutable Value

```
makeValueImmutable(x)
└─ current object graph cannot change

makeImmutable(x)
└─ current object graph cannot change
└─ binding cannot be reassigned
```

### Table

Operation	Binding	Current value/object graph
<code>makeValueImmutable(x)</code>	Mutable	Immutable
<code>makeImmutable(x)</code>	Immutable	Immutable

## 4.3 Shared

```
co.utils.makeShared(positive_int);
```

A Shared Object is safe for concurrent and multi-threaded use.

The programmer continues to use the same object model — same accessors, same syntax — while FoLang guarantees the required safety properties internally.

Shared behaviour is deep:

- all reachable objects within the shared object are also shared

### Note on Analogies

Comparisons to things like Java's `AtomicInteger` or `ConcurrentHashMap` are **analogies only**.

They may help explain the concept, but they are **not part of the formal language contract**.

FoLang specifies the behavioural guarantee, not the exact runtime implementation.

A good explanatory statement is:

A shared `co.lang.int` may be thought of similarly to an atomic integer, and a shared map may be thought of similarly to a concurrent map. These comparisons are explanatory only. FoLang does not require any particular internal runtime representation.

## 4.4 CopyOnWrite

```
co.utils.copyOnWrite(positive_int);
```

A CopyOnWrite object passes by reference like a normal managed object until mutation is attempted. Merely passing or reading the object does not change the source state.

### Normative CopyOnWrite Semantics

CopyOnWrite in FoLang has **whole-reachable-object-graph isolation semantics**. If any member or reachable object beneath a CopyOnWrite root is mutated in a context that must not affect the source object, the mutating context must observe an independent modified state for the complete logical graph rooted at that CopyOnWrite object, while external aliases that continue to refer to the source graph must continue to observe the unchanged source state.

For example:

```
a Employee = Employee{
 dept: Dept{ id: 10 },
 address: Address{ city: "Pune" }
};

co.utils.copyOnWrite(a);

process(emp Employee)->() = {
 emp.dept.id = 20;
}

process(a);
```

After `process(a)`, the caller's `a` must still observe the original `Employee` graph with `dept.id == 10`, while the mutating context must observe an `Employee` state in which `dept.id == 20`. The isolation applies to the complete logical graph rooted at `a`, including `dept`, `address`, and all other reachable managed objects.

The specification defines this **observable isolation contract**, not a required physical copying algorithm. A conforming backend is not required to allocate a complete duplicate graph if it can preserve the same observable semantics by another implementation technique.

Internal alias topology and cycles are part of the logical graph semantics. If multiple members logically refer to the same object, that relationship must remain correct in the isolated COW state. Cyclic relationships must likewise remain semantically intact. An implementation must not allow a mutation in the isolated COW state to become visible through an external alias that still denotes the original source state.

### Reference Backend Implementation

The default FoLang backend is the **reference implementation** for these semantics. Its straightforward reference strategy is whole-graph materialization: on the first mutation requiring isolation, it clones the entire reachable managed-object graph rooted at the `CopyOnWrite` object, preserves cycles and repeated internal references, redirects the mutating context to the cloned root, and applies the mutation to the clone. External aliases remain attached to the original graph.

Conceptually, the reference implementation behaves as follows:

```
caller before mutation

a → Employee A
 ├── Dept B (id = 10)
 ├── Address C
 └── ... other reachable objects

reference-backend COW materialization

caller mutating context
a → Employee A emp → Employee A'
 ├── Dept B (id = 10) ├── Dept B' (id = 20)
 ├── Address C
 └── ... ├── Address C'
 └── ...
```

This reference algorithm is deliberately simple and is intended to make the required semantics easy to understand, implement, inspect, and test. It is not a requirement that production backends use the same physical representation or copying strategy.

### Reference Backend Whole-Graph Cost

Because the reference backend materializes the complete reachable graph, the cost is determined by the size of the COW root's reachable graph, not by the size of the individual member being modified.

For example:

```
employees co.core.list(Employee) = ...;
co.utils.copyOnWrite(employees);

process(emps co.core.list(Employee))->() = {
 emps[500].dept.id = 20;
}

process(employees);
```

Although the logical change is only to `emps[500].dept.id`, the COW root is the employee list. The reference backend therefore materializes an isolated copy of the entire reachable `List(Employee)` graph before applying the mutation. This includes the list, all reachable `Employee` objects, their departments, addresses, and other reachable managed objects. The caller's original list graph remains intact.

Conceptually:

```
original caller graph reference-backend COW graph

List<Employee> A List<Employee> A'
├── Employee 0 ├── Employee 0'
├── Employee 1 ├── Employee 1'
├── ... ├── ...
├── Employee 500 ├── Employee 500'
│ ├── Dept (id = 10) │ ├── Dept' (id = 20)
│ └── Address │ ├── Address'
│ └── ... │ └── ...
├── ... ├── ...
└── Employee N └── Employee N'
```

This simple reference strategy may be unsuitable for some production workloads. A tiny nested mutation can require work proportional to a very large reachable graph, may temporarily keep both source and isolated graphs live, and can increase allocation pressure, peak memory use, memory-bandwidth consumption, and allocator fragmentation. These are characteristics of the **reference backend implementation**, not additional FoLang language semantics.

Conforming Alternative Backends

A third-party or specialized backend may implement the same CopyOnWrite semantics using structural sharing, path copying, persistent data structures, lazy materialization, reference counting, arenas, page-based techniques, virtual-memory mechanisms, a garbage-collected runtime, a custom memory manager, or another strategy. Such implementation choices are permitted provided that all externally observable FoLang semantics are preserved.

In particular, a backend may physically share unchanged storage between the source and isolated COW states as long as later mutations cannot leak across those logical states. FoLang source code must not be able to distinguish the optimized backend from the reference semantics except through non-semantic characteristics such as performance, memory consumption, allocation behaviour, or diagnostics that do not alter program meaning.

Backend implementers should use the written specification as the normative contract and may test their implementation against the reference backend and the FoLang conformance suite. Matching the reference backend’s internal whole-graph allocation algorithm is **not** a conformance requirement; matching the required observable behavior is.

4.5 toSnapshot

```
co.utils.toSnapshot(positive_int);
```

`toSnapshot` converts an object into a **snapshot representation** — a value descriptor, not a live object. The snapshot representation itself cannot be mutated.

When passed to a function, the compiler/runtime uses the snapshot representation to construct a fresh independent local variable bound to the parameter name. That local is a normal mutable object with no shared identity with the original. All of this happens automatically — the developer writes only `co.utils.toSnapshot(positive_int)`.

```
process(a co.lang.int)->() = {
 a.value = 99; // mutates the fresh local - positive_int completely unaffected
}


process(co.utils.toSnapshot(positive_int))

or

k co.lang.value = co.utils.toSnapshot(positive_int);

// positive_int unchanged
```

The flow:



`toSnapshot` is deep — the snapshot representation covers the entire reachable object graph.

5. Policy Summary

Object Kind	Mutation allowed	Caller sees mutation	Thread safe	Copy on write	Deep
Literal expression object	✗ directly in source — no handle	—	value-like use	—	✓
Normal (default)	✓ via accessor	✓	✗	✗	—
Immutable	✗	—	✓	—	✓
Shared	✓ via accessor	✓	✓	✗	✓
CopyOnWrite	✓ on own copy	✗	✓	✓	✓
toSnapshot result	✓ on reconstructed local object	✗	independent snapshot	—	✓



## 6. Literal vs Value/snapshot

`co.lang.value` vs `co.lang.literal`

Literal (`co.lang.literal`) is literal representation of objects literals are object.

Literal representations use `to` conversion methods to produce typed objects. When no suitable conversion exists for a custom type, the developer must provide the required overload through the supported extension mechanism. A literal representation by itself does not carry the complete reconstruction metadata stored by a snapshot value.

A `co.lang.value` snapshot carries more information than a literal representation: it records the type information, literal/value representation, and reconstruction information required to create the corresponding object according to its declaration kind.

## 7. No Type Fragmentation

FoLang deliberately avoids special types for mutability or concurrency concerns.

There is no need for separate public type families such as:

```
ImmutableInt
SharedMap
CopyOnWriteList
AtomicInt
ConcurrentMap
```

Instead, FoLang keeps the type model uniform and applies behaviour policies through `co.utils.*`.

In many ecosystems, a programmer must constantly choose between:

normal integer	vs atomic integer
normal map	vs concurrent map
mutable value	vs immutable wrapper
raw structure	vs copy-on-write structure

FoLang instead aims to provide:

- one object model
- one programming style
- one mental model

while still allowing the programmer to opt into immutability, sharing, copy-on-write, or literal conversion when needed.

## 8. Formal Object-Policy and Identity Rules

The following rules define the previously open object-model precision points.

### 8.1 Aliases and Policy Ownership

Managed-object assignment copies a reference. Therefore multiple bindings may refer to the same object. Immutable, Shared, and CopyOnWrite state belongs to that object graph rather than to an individual alias. Applying one of these policies through any alias changes the policy observed through every alias that still refers to the same graph.

```
a Employee = Employee{...};
b := a;

co.utils.makeShared(a);

// a and b still refer to the same object graph.
// The graph is Shared through both aliases.
```

### 8.2 Rebinding Is Independent of Object Policy

Rebinding an alias changes only which object that binding refers to. It does not remove, transfer, or alter the policy of the previously referenced object.

```
a Employee = Employee{...};
b := a;

co.utils.makeShared(a);
b = Employee{...};
```

After the rebinding, `a` still refers to the original Shared graph. `b` refers to a different newly created object graph. The original graph remains Shared for its lifetime.

`makeImmutable(x)` is special only in that it also freezes the supplied binding `x`; the Immutable object policy itself still belongs to the object graph. Another alias to the same immutable graph cannot mutate that graph, although that other alias may itself be rebound unless its own binding has separately been made non-rebindable.

### 8.3 Deep Policy Scope

A deep policy covers the entire managed-object graph reachable from its root. Reachability is transitive: members, nested objects, collection elements, and other managed-object references are followed recursively until no new object identity is encountered. Cycles terminate by identity rather than by traversal depth, and repeated references to the same object remain repeated references to one object.

Thus applying a deep policy to an `Employee` also applies that policy to its reachable `dept`, `address`, nested collections, collection elements, and other managed objects in that graph. The policy invariant is maintained for the graph while that policy is active.

### 8.4 Mutation Visibility Across Calls

Function arguments use the managed reference model. Mutation visibility therefore follows the object's policy:

- **Normal** — mutation through the parameter mutates the same object graph and is visible to the caller and other aliases.
- **Shared** — mutation is visible to aliases of the same graph, with the concurrency-safety guarantees of Shared behaviour.
- **Immutable** — mutation is prohibited.
- **CopyOnWrite** — a mutation requiring isolation gives the mutating context an independent logical state for the complete reachable graph while the caller's original graph remains unchanged. The reference backend realizes this by materializing a complete private graph copy; other conforming backends may use different internal techniques.

The policy of an existing graph is not temporary or call-local. Once that graph becomes Immutable, Shared, or CopyOnWrite, it remains in that state for its lifetime.

### 8.5 Value Equality and Reference Identity

`==` performs deep value equality. It answers whether two values/object graphs are equal by value; it does not answer whether two bindings refer to the same managed object.

Reference identity is exposed through `sameRef()` for managed reference-semantic objects:

```
a Employee = Employee{name: "Rao"};
b := a;
c Employee = Employee{name: "Rao"};

a == b; // true
a == c; // true

a.sameRef(b); // true - same managed object
a.sameRef(c); // false - equal value, different managed object
```

`co.lang.cstruct` remains value-semantic and is outside the managed-object reference-identity model.

### 8.6 Policy Lifetime and Non-Stacking

Immutable, Shared, and CopyOnWrite are mutually exclusive permanent states of an existing managed object graph. A graph may transition from its normal mutable state into one of these states, but it cannot subsequently transition from one policy state into another.

```
Normal ———┬─ Immutable
 ├── Shared
 └─ CopyOnWrite

Immutable -X-> Shared
Immutable -X-> CopyOnWrite
Shared -X-> Immutable
Shared -X-> CopyOnWrite
CopyOnWrite -X-> Immutable
CopyOnWrite -X-> Shared
```

`makeValueImmutable(x)` and `makeImmutable(x)` both place the current object graph into the Immutable state. Their difference concerns the supplied binding: `makeValueImmutable(x)` permits that binding to be rebound later, whereas `makeImmutable(x)` also prevents rebinding of that binding. Rebinding to a different object does not transfer the old object's policy to the new object.

## 9. Formal Philosophy Statement

All managed FoLang objects use reference semantics by default. `co.lang.cstruct` is an explicitly value-semantic ABI representation and is an exception to managed-object reference semantics.

In FoLang, everything is an object and managed objects are mutable by default.

Assignment of managed objects copies references, `co.lang.cstruct` assignment copies values, `==` compares values deeply, and `sameRef()` exposes managed-object reference identity.

Developers may opt into Immutable, Shared, or CopyOnWrite behaviour without changing the public type of the object. These policies belong to the object graph, are deep, are observed through all aliases to that graph, are mutually exclusive, and are permanent for that graph's lifetime.

Rebinding an alias changes the object referenced by that binding but does not modify or transfer the policy of the previously referenced object.

CopyOnWrite has whole-graph isolation semantics: a mutation anywhere in the reachable graph that requires isolation must leave the source logical graph intact while the mutating context observes an independent modified logical graph. The reference backend demonstrates this by cloning the complete reachable graph, but conforming backends may use any internal mechanism that preserves the same observable semantics.

This policy model is uniform across managed types, so programmers do not need separate type families for ordinary, immutable, concurrent, or snapshot-oriented use.

Familiar analogies such as atomic integers or concurrent maps may help explain the design, but they are not part of the formal implementation contract.

---

## FoLang Definition and Documentation License

Unless otherwise stated, the copyrightable material contained in the FoLang language definition and documentation is licensed under the [Creative Commons Attribution 4.0 International License](#).

The CC BY 4.0 licence applies to the copyrightable expression, organization, and presentation of the FoLang language definition and documentation, including:

- the FoLang language specification;
- original FoLang-specific syntax forms;
- grammar productions and grammar notation;
- rules governing how FoLang syntax forms may be combined;
- semantic rules and behavioural descriptions;
- name-resolution and scope-resolution rules;
- type-system definitions and constraints;
- execution-model and control-abstraction descriptions;
- compiler-behaviour descriptions;
- source-code examples demonstrating FoLang syntax and semantics;
- diagrams, tables, illustrations, and formal notation;
- reference documentation;
- explanatory and instructional material.

This includes the documented expression and presentation of FoLang-specific constructs such as:

```
(booleanExpression).then({
 ...
}).default({
 ...
});
```

and other original FoLang syntax, grammar, and semantic-rule descriptions contained in this document.

Under CC BY 4.0, this material may be:

- copied and redistributed in any medium or format;
- adapted, modified, translated, and extended;
- used for commercial or non-commercial purposes.

A person exercising these permissions must:

- give appropriate attribution;
- provide a link to the CC BY 4.0 licence;
- indicate whether modifications were made;
- retain notices supplied with the licensed material where required;
- not imply endorsement by the FoLang project or its contributors.

## Independent Use and Implementation

This licence applies to the copyrightable expression contained in the FoLang language definition and documentation.

It does not restrict the independent use or implementation of programming-language ideas, concepts, features, procedures, systems, or methods of operation.

For example, this licence does not claim exclusive rights over general programming-language concepts such as:

- continuations;
- modules;
- structs and classes;
- dynamic, lexical, or mixed scope;
- functions and closures;
- pattern matching;
- algebraic data types;
- dependent types;
- concurrency and parallelism;
- coroutines and asynchronous execution. etc.,

Another project may independently implement such concepts without copying the copyrightable expression of the FoLang documentation.

The FoLang-specific documentation describing how these concepts are expressed through FoLang syntax, grammar, semantic rules, examples, diagrams, and explanatory text remains subject to this licence when that material is copied or adapted.

## Software Licences

This licence does not change or replace the licences applied to:

- the FoLang frontend source code;
- the default backend source code;
- other FoLang software components;
- generated compiler binaries;
- third-party software;
- third-party documentation or assets.

The FoLang frontend and backend continue to be governed by their separately stated software licences.

## Other Rights

CC BY 4.0 does not grant rights relating to:

- trademarks;
- project names and branding;
- logos;
- patents;
- third-party material;
- privacy, publicity, or personality rights.

Such rights, where applicable, remain governed separately.

## Suggested Attribution

FoLang Language Definition and Documentation, including its syntax, grammar, and semantic-rule descriptions, by [Kemeswara Rao Mithipati](#) and FoLang contributors, licensed under [CC BY 4.0](#).

When modifications are made, the attribution should also indicate that the material was modified.

Example:

Based on the FoLang Language Definition and Documentation by [Kemeswara Rao Mithipati](#) and FoLang contributors, licensed under [CC BY 4.0](#). Modified from the original. ***

## Appendix A - Complete FoLang EBNF Grammar

---

The following grammar is the normative lexical and syntactic grammar for FoLang.

```
(*
FoLang EBNF Grammar – current alpha profile

Generated from:
 language-ref.md
 SHA-256: 18333739295affc4839b7ba3bfeefda16191a40239773adbab8b8d134d7b9cb3

Authority:
 The language reference is normative. This grammar is a syntactic
 consolidation of the source forms defined there, including Appendix C.

Notation:
 = defines a production
 ; ends a production
 | alternative
 [...] optional
 { ... } zero or more
 (...) grouping
 "..." terminal text
 ? ... ? lexical or context-sensitive condition described in prose

Parser/lexer layering:
- Source text is UTF-8.
- U+FEFF is permitted only as the first source code point.
- white-space/comments are discarded between tokens.
- newlines never terminate statements.
- the parser consumes the token stream produced by the lexer.

Filesystem-selected parser roots:
 src/appl.fol
 -> application-entry-file

 src/library.fol
 -> standalone-library-surface-file

 src/lib/ffi/library.fol
 src/lib/system/library.fol
 src/lib/advanced/library.fol
 src/lib/dynamicvmrt/library.fol
 -> source-library-surface-file

 src/lib/operators/library.fol
 -> operator-source-file
 (dedicated lexer/parser; not ordinary compilation-unit syntax)

ordinary package files
 <Name>.fol
 <Fragment>.unit.fol
 <StructName>.comp.unit.fol
 package.fol
 -> the matching package-source-file alternative selected from the
 externally validated filename class

Current-alpha rules that are deliberately explicit here:
- A named function block body requires "=" before the body.
- An anonymous function literal does not use "=" between signature/body.
- Direct declaration/function/block bodies end at "}" and cannot take ";".
- Braced expressions/literals still require the enclosing statement ";".
- Object construction is Type{field: value, ...}.
- A type's arrow tail "T->(...)" carries three unrelated meanings that
 only the tail's own shape and the base's meaning separate: a derivation
 applied to T, the result list of a function type, or a generic
 type-argument binding list. The binding list admits positional arguments
 and two named spellings, both of which the reference writes:
 co.lang.int->([5]) derivation
 co.lang.int->(&, meta={type=out}) derivation, attributes
 (co.lang.int)->(co.lang.int) function type
 co.core.List->(co.lang.string) positional type args
 co.core.Map->(key co.lang.string, val co.lang.int) named, space form
 co.core.Map->(key=co.lang.string, val=co.lang.int) named, "=" form
 co.core.Array->(dims=2, type=co.lang.int, sizes=[2,4]) bound values
 The positional and space-named spellings are return-item-list's shape.
 The "name=value" spelling is derivation-attribute-list's shape, and a
 bound argument's value may be a type, an integer, or a bracketed list,
 which is what the Array and Matrix declarations bind. That makes
 "name=value" the ONE tail shape a derivation and a generic binding list
 share, so those two readings are separated by the base type and not by
```

the tail's first token. See parenthesized-type-list.

- A built-in collection value is a collection type followed directly by its literal body, which is typed-collection-literal:  
co.core.List["A","B","C"]      "[" elements "]"  
co.core.Map{"A": 1, "B": 2}    "{" entries  "}"  
co.core.Set(1,2,3)            "(" arguments ")"  
co.core.Set->(co.lang.int)(1,2,3) generic args, then arguments  
The body form is fixed by the collection, not chosen by the writer.  
builtin-collection-type-name records the reference's Builtin Collections registry, which is the closed set of built-in names admissible in that prefix position.
- There is no untyped object literal, and a braced map body is an object literal representation, so a braced compound value ALWAYS names its type and map-literal is not a primary-expression alternative. An array literal is not an object literal: "[ ... ]" is a simple literal like a string or an integer and stays available untyped. See C.4 and C.4.1.
- .match and .match() mean automatic matcher selection;  
.match(expr) selects an explicit matcher.
- Control-chain vocabulary uses .then(X) for one-shot conditional selection, .loop(block) for single-condition repeated execution, .otherwise(condition) only for another selection condition, and .default(X) only for a terminal selection fallback. A .loop(block) form has exactly one condition and cannot be followed by .otherwise(condition) or .default(X). Element iteration is owned by .each(...) itself: its explicit-binding form takes index/key, value, and the per-element action directly, while its single-argument form takes a callable callback. An .each(...) call cannot be followed by .loop(...). There is no conditionless .otherwise form and no .do control verb in the current alpha profile.
- Built-in precedence follows Appendix C.9.
- predeclared-operator-glyph tokens are recognized but are not accepted as current-alpha expression operators.
- reserved-future-operator-fixity is a diagnostic-recognition production: encountering it in current alpha must produce an unsupported-feature parser error rather than enabling that fixity.
- table-listed/reserved co.* names with no implemented source form remain reserved/unimplemented and must not be treated as ordinary user syntax. This covers a co.* name in declaration-kind, type, callable or operator position. It does NOT cover the @co.* metadata forms – directives, annotations, pragmas and decorators – which all parse through the single annotation production and are collected whether or not anything implements them, then silently ignored. See Appendix C.7.1.

Context-sensitive/semantic requirements intentionally remain outside EBNF:  
filename canonicalization and package ownership; project layout and reachability; import field uniqueness/mutual exclusion; library-kind capability restrictions; interface/signature/typeclass conformance; instance placement/liveness; overload resolution; dependent-index name resolution; generic-argument binding, where a named argument must name a declared type parameter of the type being applied, names and positions must not be mixed in one list, and a `name=value` tail is a derivation unless the base type is a generic declaration; whether a type prefix on a collection body names a collection type; annotation target/type checks; operator metadata uniqueness, ranges and arity/fixity compatibility; and all type/runtime semantics.

*)

(* ===== *)

(* 1. Source roots and externally selected file forms *)

(* ===== *)

```
compilation-unit = package-source-file
 | application-entry-file
 | library-surface-file ;
```

```
source-filename = companion-unit-filename
 | ordinary-unit-filename
 | package-metadata-filename
 | application-entry-filename
 | library-surface-filename
 | ordinary-source-filename ;
```

```
companion-unit-filename = filename-identifier, ".comp.unit.fol" ;
```

```
ordinary-unit-filename = filename-identifier, ".unit.fol" ;
```

```

package-metadata-filename = "package.fol" ;

application-entry-filename = "appl.fol" ;

library-surface-filename = "library.fol" ;

ordinary-source-filename = filename-identifier, ".fol" ;

filename-identifier = identifier-head, { "_", identifier-segment } ;

package-source-file = package-primary-source-file
 | ordinary-unit-source-file
 | companion-unit-source-file
 | package-metadata-source-file ;

package-primary-source-file = file-preamble, primary-declaration ;

ordinary-unit-source-file = file-preamble, unit-declaration ;

companion-unit-source-file = file-preamble, unit-declaration ;

package-metadata-source-file = file-preamble, package-alias-declaration ;

application-entry-file = file-preamble, { entry-item } ;

library-surface-file = standalone-library-surface-file
 | source-library-surface-file ;

standalone-library-surface-file =
 file-preamble, standalone-library-kind-annotation, library-declaration ;

source-library-surface-file =
 file-preamble, library-declaration ;

standalone-library-kind-annotation =
 "@co.dap.library", "(", "type", "=", library-kind-string, ")" ;

library-kind-string =
 ? string-literal whose decoded value is exactly one of:
 application, dynamicvmrt, advanced, system, ffi ? ;

file-preamble = { file-directive } ;

file-directive = import-directive
 | alias-directive
 | use-directive
 | dynamic-runtime-directive ;

entry-item = file-directive
 | entry-type-declaration
 | bare-function-pattern-clause
 | capturing-function-pattern-clause
 | entry-statement ;

entry-statement = variable-declaration
 | inferred-variable-declaration
 | grouped-variable-declaration
 | multiple-assignment-statement
 | expression-statement
 | empty-statement ;

primary-declaration = struct-declaration
 | cstruct-declaration
 | enum-declaration
 | union-declaration
 | class-declaration
 | interface-declaration
 | signature-declaration
 | module-declaration
 | typeclass-declaration
 | object-declaration
 | instance-declaration
 | matcher-instance-declaration ;

(* ===== *)

(* 2. Directives, annotations, and metadata *)

```

```

(* ===== *)

annotation = { annotation } ;

annotation = "@", qualified-name,
 ["(", [annotation-argument-list], ")"] ;

annotation-argument-list = annotation-argument,
 { ",", annotation-argument }, [",",] ;

annotation-argument = [annotation-key, annotation-binder], annotation-value ;

annotation-binder = "=" | ":" ;

annotation-key = annotation-key-segment,
 { "-", annotation-key-segment } ;

annotation-key-segment = identifier | "for" ;

annotation-value = literal
 | type-expression
 | qualified-name
 | declaration-reference
 | annotation-list
 | annotation-map
 | annotation-arrow-pair ;

annotation-list = "[", [annotation-value,
 { ",", annotation-value }, [",",]], "]" ;

annotation-map = "{", [annotation-map-entry,
 { ",", annotation-map-entry }, [",",]], "}" ;

annotation-map-entry = annotation-key, annotation-binder, annotation-value
 | annotation-key ;

annotation-arrow-pair = string-literal, "=>", string-literal ;

import-directive = "@co.ddap.import", "(", import-field,
 { ",", import-field }, ")" ;

import-field = "package", "=", string-literal
 | "library", "=", string-literal
 | "src-library", "=", "true"
 | "as", "=", string-literal ;

alias-directive = "@co.ddap.alias", "(", co-path, ",",
 "as", "=", string-literal, ")" ;

use-directive = "@co.ddap.use", "(", use-field,
 { ",", use-field }, ")" ;

use-field = "from", "=", string-literal
 | "methods", "=", "[", [identifier,
 { ",", identifier }], "]" ;

dynamic-runtime-directive = "@co.ddap.dynamicruntime",
 ["(", [annotation-argument-list], ")"] ;

(* ===== *)

(* 3. Names and references *)

(* ===== *)

filename-derived-name = "_" ;

qualified-name = (identifier | "co"), { ".", identifier } ;

co-path = "co", ".", identifier, { ".", identifier } ;

declaration-reference = qualified-function-reference | qualified-name ;

qualified-function-reference = qualified-name, "(", [type-list], ")",
 return-type-clause ;

```



```

lifecycle-name = special-method ;

special-method = "@@new" | "@@init" ;

special-binding = result-binding | self-binding ;

self-binding = "$" ;

result-binding = "$", digit, { digit } ;

wildcard = "_" ;

(* ===== *)

(* 4. Type syntax *)

(* ===== *)

type-expression = forall-type | union-type-expression ;

forall-type = "forall", "(", type-parameter-list, ")", ".", type-expression ;

type-parameter-list = identifier, { ",", identifier } ;

union-type-expression = arrow-type-expression,
 { "|", arrow-type-expression } ;

arrow-type-expression = type-postfix-expression,
 ["->", arrow-type-tail]
 | "(", [function-type-parameter,
 { ",", function-type-parameter }],
 ")", "->", arrow-type-tail ;

arrow-type-tail = type-derivation
 | parenthesized-type-list
 | type-expression ;

type-postfix-expression = type-atom, { type-argument-list } ;

type-atom = qualified-name
 | "(", type-expression, ")" ;

type-argument-list = "(", [type-or-value-argument,
 { ",", type-or-value-argument }], ")" ;

type-or-value-argument = type-expression | dependent-index ;

dependent-index = integer-literal | qualified-name ;

type-list = type-expression, { ",", type-expression } ;

(* This is one shape with two readings: the result list of a function type, or
a generic type-argument binding list applied to the base type. Both admit an
optional binder name per item, which is exactly return-item-list's shape, so
the two are spelled with one production rather than two indistinguishable
ones. Which reading applies is decided from the base type, not from the
tail.

The third arrow tail, type-derivation, is a separate alternative told apart
by its first token in every case but one. The reference now also writes a
named generic argument as `name=Type` and binds non-type values the same way:

 co.core.Map->(key=co.lang.string, val=co.lang.int)
 co.core.Array->(dims=2, type=co.lang.int, sizes=[2,4])
 LinkedList->(T=co.lang.int)

An `identifier "="` opener is therefore a shape a derivation-attribute-list
and a generic binding list share, and only the base type separates them: a
tail on a generic declaration binds that declaration's type parameters, while
a tail on any other type is the derivation-attribute-list of C.11. The two
spellings coexist in the reference – C.11 and C.4.1 write `key co.lang.string`
and `key=co.lang.string` for the same instantiation – so both are admitted
here and the choice between them carries no meaning. *)
parenthesized-type-list = "(", [return-item-list], ")" ;

type-derivation = "(", derivation-specification, ")" ;

```

```

derivation-specification = pointer-specification
 | array-specification
 | reference-specification
 | range-type-specification
 | slice-type-specification
 | thunk-type-specification
 | address-type-specification
 | derivation-attribute-list ;

pointer-specification = pointer-stars,
 [",", derivation-attribute-list] ;

pointer-stars = ? one contiguous symbolic run consisting only of one or
 more "*" characters; its length is the pointer degree ? ;

reference-specification = ("&" | "&&" | "~"),
 [",", derivation-attribute-list] ;

address-type-specification = "@", [",", derivation-attribute-list] ;

thunk-type-specification = "^", [",", derivation-attribute-list] ;

slice-type-specification = "[:]", [",", derivation-attribute-list] ;

range-type-specification = "..", [",", derivation-attribute-list] ;

array-specification = array-dimension-group, { array-dimension-group },
 [",", derivation-attribute-list] ;

array-dimension-group = "[", array-dimension-content, "]" ;

array-dimension-content = [array-dimension], { ",", [array-dimension] } ;

array-dimension = "... " | ". " | dependent-index ;

(* Attributes carried by a derivation, and – on a generic base type – the
 `name=value` spelling of a generic type-argument binding list. One shape,
 two readings, separated by the base type exactly as parenthesized-type-list
 describes. annotation-value is what admits a bound value that is not a type:
 the integer of `dims=2` and the bracketed list of `sizes=[2,4]`. *)
derivation-attribute-list = derivation-attribute,
 { ",", derivation-attribute } ;

derivation-attribute = annotation-key, "=", annotation-value ;

return-type-clause = "->", "(", [return-item-list], ")" ;

return-item-list = return-item, { ",", return-item } ;

return-item = [identifier], type-expression ;

(* ===== *)

(* 5. Common declaration components *)

(* ===== *)

generic-parameter-clause = "(", generic-parameter,
 { ",", generic-parameter }, ")" ;

generic-parameter = identifier, [generic-arity-clause] ;

generic-arity-clause = "(", generic-arity-slot,
 { ",", generic-arity-slot }, ")" ;

generic-arity-slot = "_" | identifier ;

kind-options = "->", "(", [annotation-argument-list], ")" ;

field-declaration = annotations, identifier, type-expression,
 ["=", expression], statement-end ;

embedded-field-declaration = annotations, type-expression, statement-end ;

value-specification = annotations, identifier, type-expression, statement-end ;

variable-declaration = annotations, typed-variable-declarator,

```

```

 { ",", typed-variable-declarator }, statement-end ;

typed-variable-declarator = identifier, type-expression,
 ["=", expression] ;

inferred-variable-declaration = annotations, inferred-variable-declarator,
 { ",", inferred-variable-declarator },
 statement-end ;

inferred-variable-declarator = identifier, definition-operator, expression ;

definition-operator = ("!=" | "?="),
 multi-symbol-infix-operator-boundary-guard ;

(* ===== *)

(* 6. Data and type declarations *)

(* ===== *)

struct-declaration = annotations, filename-derived-name,
 "co.lang.struct", "=", struct-body ;

surface-struct-declaration = annotations, identifier,
 "co.lang.struct", "=", struct-body ;

struct-body = "{", { struct-member }, body-close ;

pure-field-declaration = annotations, identifier, type-expression,
 statement-end ;

struct-member = pure-field-declaration | embedded-field-declaration ;

cstruct-declaration = annotations, filename-derived-name,
 "co.lang.cstruct", "=", cstruct-body ;

surface-cstruct-declaration = annotations, identifier,
 "co.lang.cstruct", "=", cstruct-body ;

cstruct-body = "{", { pure-field-declaration }, body-close ;

enum-declaration = annotations, filename-derived-name,
 "co.lang.enum", "=", enum-body ;

enum-body = "{", [enum-variant,
 { enum-separator, enum-variant }, [enum-separator]],
 body-close ;

enum-separator = "," ;

enum-variant = annotations, identifier,
 ["(", [type-list], ")"],
 ["=", constant-expression] ;

union-declaration = annotations, filename-derived-name,
 "co.lang.union", "=", union-body ;

union-body = "{", { pure-field-declaration }, body-close ;

data-declaration = annotations, identifier,
 [generic-parameter-clause], "co.lang.data", "=",
 data-variant, { "|", data-variant }, statement-end ;

data-variant = qualified-name,
 ["(", [type-list], ")"] ;

type-declaration = parameterized-type-declaration
 | simple-type-declaration ;

parameterized-type-declaration = annotations, identifier,
 generic-parameter-clause, "co.lang.type",
 [kind-options], ["=", type-expression],
 statement-end ;

simple-type-declaration = annotations, identifier, type-declaration-kind,
 [kind-options], ["=", type-expression],
 statement-end ;

```

```

type-declaration-kind = "co.lang.type"
 | "co.lang.newtype"
 | "co.lang.opaquetype"
 | "co.lang.subtype"
 | "co.lang.supertype"
 | "co.lang.dependentType"
 | "co.lang.kind" ;

entry-type-declaration = entry-parameterized-type-declaration
 | entry-simple-type-declaration ;

entry-parameterized-type-declaration = annotations, identifier,
 generic-parameter-clause,
 "co.lang.type",
 [kind-options],
 ["=", type-expression],
 statement-end ;

entry-simple-type-declaration = annotations, identifier,
 ("co.lang.type"
 | "co.lang.newtype"
 | "co.lang.opaquetype"
 | "co.lang.subtype"
 | "co.lang.supertype"
 | "co.lang.dependentType"),
 [kind-options],
 ["=", type-expression],
 statement-end ;

package-alias-declaration = filename-derived-name, "co.lang.package", "=",
 package-alias-body, statement-end ;

package-alias-body = "{", "name", ":", string-literal, "}" ;

(* ===== *)

(* 7. Containers, contracts, instances, and library surfaces *)

(* ===== *)

unit-declaration = annotations, filename-derived-name,
 "co.lang.unit", "=", unit-body ;

unit-body = "{", { unit-member }, body-close ;

unit-member = function-declaration
 | closure-declaration
 | data-declaration
 | type-declaration
 | type-level-function-declaration
 | function-object-declaration
 | delegate-declaration
 | extern-variable-declaration ;

extern-variable-declaration =
 "@co.dap.declare", "(", "extern", ")", identifier, type-expression,
 statement-end ;

class-declaration = annotations, filename-derived-name,
 "co.lang.class", [kind-options], "=", class-body ;

class-body = "{", { class-member }, body-close ;

class-member = field-declaration
 | function-declaration
 | lifecycle-method-declaration ;

lifecycle-method-declaration = annotations, lifecycle-name,
 parameter-list, [return-type-clause],
 function-definition ;

interface-declaration = annotations, filename-derived-name,
 "co.lang.interface", "=", interface-body ;

interface-body = "{", { function-specification }, body-close ;

```

```

signature-declaration = annotations, filename-derived-name,
 "co.lang.signature", "=", signature-body ;

signature-body = "{", { signature-member }, body-close ;

signature-member = value-specification
 | function-specification
 | signature-type-component ;

signature-type-component = annotations, identifier,
 [generic-parameter-clause], "co.lang.type",
 ["=", type-expression], statement-end ;

module-declaration = annotations, filename-derived-name,
 "co.lang.module", [kind-options], "=", module-body ;

module-body = "{", { module-member }, body-close ;

module-member = variable-declaration
 | inferred-variable-declaration
 | function-declaration
 | signature-type-component ;

library-declaration = filename-derived-name,
 "co.lang.library", "=", library-body ;

library-body = "{", { library-member }, body-close ;

library-member = import-directive
 | surface-struct-declaration
 | surface-cstruct-declaration
 | function-declaration ;

object-declaration = annotations, filename-derived-name,
 "co.lang.object", [kind-options], "=", object-body ;

object-body = "{", { field-declaration | function-declaration }, body-close ;

instance-declaration = annotations, filename-derived-name,
 "co.lang.instance", [kind-options], "=",
 instance-body ;

instance-body = "{", { function-declaration | variable-declaration }, body-close ;

typeclass-declaration = annotations, filename-derived-name,
 typeclass-parameter-clause, "co.lang.typeclass", "=",
 contract-body ;

typeclass-parameter-clause = generic-parameter-clause ;

matcher-instance-declaration = annotations, filename-derived-name,
 "co.lang.matcher", matcher-options, "=",
 matcher-body ;

matcher-options = "->", "(", "type", annotation-binder, type-expression, ")" ;

matcher-body = "{", { function-declaration }, body-close ;

contract-body = "{", { function-specification | value-specification }, body-close ;

named-block-declaration = annotations, identifier,
 "co.lang.block", "=", block, body-closure-guard ;

(* ===== *)

(* 8. Functions and callable declarations *)

(* ===== *)

delegate-declaration = annotations, identifier,
 "co.lang.delegate", "=", function-type, statement-end ;

function-object-declaration = annotations, identifier,
 "co.lang.function", "=",
 function-object-binding ;

function-object-binding = expression, statement-end ;

```

```

type-level-function-declaration = annotations, function-name, parameter-list,
 { parameter-list }, type-level-return-clause,
 type-level-binding ;

type-level-return-clause = "->", "(", type-level-result-kind,
 { "|", type-level-result-kind }, ")" ;

type-level-result-kind = "co.lang.dependentType"
 | "co.lang.type" ;

type-level-binding = function-definition
 | function-delegation
 | "=", type-expression, statement-end
 | "=", non-block-expression, statement-end
 | statement-end ;

function-declaration = annotations, [receiver-clause], function-name,
 parameter-list, { parameter-list },
 [return-type-clause], function-binding ;

function-name = identifier ;

receiver-clause = "(", (type-expression
 | identifier, type-expression), ")" ;

parameter-list = "(", [parameter,
 { ",", parameter }], ")" ;

parameter = ["...", ["~"], identifier, ["?"],
 [type-expression], ["=", expression] ;

function-binding = function-definition
 | function-delegation
 | function-alias-binding
 | statement-end ;

function-definition = "=", block, body-closure-guard ;

function-delegation = "=>", expression,
 { "=>", expression }, statement-end ;

function-alias-binding = "=", non-block-expression, statement-end ;

function-specification = annotations, [receiver-clause], function-name,
 parameter-list, { parameter-list },
 [return-type-clause], statement-end ;

function-type = "(", [function-type-parameter,
 { ",", function-type-parameter }], ")",
 return-type-clause ;

function-type-parameter = type-expression
 | identifier, type-expression ;

anonymous-function-expression =
 ["forall", "(", type-parameter-list, ")", "."],
 parameter-list, return-type-clause, block ;

lambda-expression = "|", [lambda-parameter,
 { ",", lambda-parameter }], "|", "=>",
 (expression | block) ;

lambda-parameter = identifier, [type-expression] ;

closure-declaration = annotations, identifier, "=", parameter-list,
 { parameter-list }, "=>", expression,
 statement-end ;

local-function-declaration = annotations, function-name, parameter-list,
 { parameter-list }, return-type-clause,
 function-definition ;

(* ===== *)

(* 9. Function-pattern clauses and patterns *)

```

```

(* ===== *)

bare-function-pattern-clause = annotations, identifier, pattern-parameter-list,
 [where-clause], ">=", pattern-result ;

capturing-function-pattern-clause = annotations, "let", identifier,
 pattern-parameter-list,
 [where-clause], "=", pattern-result ;

pattern-parameter-list = "(", [pattern,
 { ",", pattern }], ")" ;

where-clause = ".where", "(", expression, ")" ;

pattern-result = block, body-closure-guard
 | non-block-expression, statement-end ;

pattern = wildcard
 | literal-pattern
 | binding-pattern
 | constructor-pattern
 | record-pattern
 | tuple-pattern
 | qualified-name ;

literal-pattern = literal
 | ("+" | "-"),
 (integer-literal | floating-literal) ;

binding-pattern = identifier ;

constructor-pattern = qualified-name, "(", [pattern,
 { ",", pattern }], ")" ;

record-pattern = qualified-name, "{", [record-pattern-field,
 { ",", record-pattern-field }], "}" ;

record-pattern-field = identifier, [":", pattern] ;

tuple-pattern = "(", pattern, ",", pattern,
 { ",", pattern }, ")" ;

match-case = ".case", "(", match-case-body, ")" ;

match-case-body = pattern, [":", expression], ">=",
 (expression | block) ;

match-default = ".default", "(", (expression | block), ")" ;

(* ===== *)

(* 10. Statements, blocks, and termination *)

(* ===== *)

block = "{", { block-item }, [block-tail-expression], "}" ;

block-item = use-directive | statement ;

block-tail-expression = expression ;

statement = named-block-declaration
 | variable-declaration
 | inferred-variable-declaration
 | grouped-variable-declaration
 | let-value-declaration
 | local-function-declaration
 | closure-declaration
 | multiple-assignment-statement
 | return-statement
 | expression-statement
 | labeled-block
 | block-statement
 | empty-statement ;

grouped-variable-declaration = "(", typed-variable-declarator,
 { ",", typed-variable-declarator }, ")",

```

```

 statement-end ;

let-value-declaration = "let", identifier, [type-expression], "=",
 expression, statement-end ;

multiple-assignment-statement = assignment-target, ",", assignment-target,
 { ",", assignment-target }, "=",
 expression-list, statement-end ;

assignment-target = postfix-expression
 | tuple-assignment-target ;

tuple-assignment-target = "(", assignment-target, ",", assignment-target,
 { ",", assignment-target }, ")" ;

return-statement = ("this" | "self"), ".return",
 [expression-list], statement-end ;

expression-statement = annotations, non-block-expression, statement-end ;

labeled-block = identifier, ":", block, body-closure-guard ;

empty-statement = ";" ;

expression-list = expression, { ",", expression } ;

statement-end = ";" ;

body-close = "}", body-closure-guard ;

body-closure-guard =
 ? zero-width condition: the next significant token is not ";", or there is no next token ? ;

block-statement = block, body-closure-guard ;

non-block-expression = expression, non-block-expression-guard ;

non-block-expression-guard =
 ? zero-width condition: the complete source span is not admissible as an
 unparenthesized block production; when both block and another braced
 expression are possible, the block reading has priority ? ;

(* ===== *)

(* 11. Expressions and built-in operator precedence *)

(* ===== *)

expression = assignment-expression
 | extended-operator-expression ;

assignment-expression = logical-or-expression,
 [runtime-assignment-operator,
 assignment-expression] ;

runtime-assignment-operator = "="
 | compound-assignment-operator ;

compound-assignment-operator = ("+=" | "-=" | "*=" | "/=" | "%="
 | "**=" | "&=" | "^=" | "|="),
 multi-symbol-infix-operator-boundary-guard ;

constant-expression = (logical-or-expression
 | extended-operator-expression),
 ? zero-width condition: no runtime-assignment-operator
 occurs anywhere in this constant-expression subtree ? ;

logical-or-expression = logical-and-expression,
 { logical-or-operator, logical-and-expression } ;

logical-or-operator = "||", multi-symbol-infix-operator-boundary-guard ;

logical-and-expression = bitwise-or-expression,
 { logical-and-operator, bitwise-or-expression } ;

logical-and-operator = "&&", multi-symbol-infix-operator-boundary-guard ;

```



```

bitwise-or-expression = bitwise-xor-expression,
 { "|", bitwise-xor-expression } ;

bitwise-xor-expression = bitwise-and-expression,
 { "^", bitwise-and-expression } ;

bitwise-and-expression = equality-expression,
 { "&", equality-expression } ;

equality-expression = relational-expression,
 { equality-operator, relational-expression } ;

equality-operator = ("==" | "!="),
 multi-symbol-infix-operator-boundary-guard ;

relational-expression = range-expression,
 { relational-operator, range-expression } ;

relational-operator = "<" | ">" | multi-symbol-relational-operator ;

multi-symbol-relational-operator = ("<=" | ">="),
 multi-symbol-infix-operator-boundary-guard ;

range-expression = additive-expression,
 [range-operator, [additive-expression]]
 | range-operator, additive-expression ;

range-operator = (".." | "<.." | "..<" | "<..<"),
 multi-symbol-range-operator-boundary-guard ;

additive-expression = multiplicative-expression,
 { additive-operator, multiplicative-expression } ;

additive-operator = "+" | "-" ;

multiplicative-expression = unary-expression,
 { multiplicative-operator, unary-expression } ;

multiplicative-operator = "*" | "/" | "%" ;

unary-expression = { prefix-operator }, power-expression ;

prefix-operator = "+" | "-" | "!" ;

power-expression = postfix-expression,
 [power-operator, unary-expression] ;

power-operator = "**", multi-symbol-infix-operator-boundary-guard ;

postfix-expression = primary-expression,
 { postfix-suffix | postfix-operator } ;

postfix-operator = "!" ;

postfix-suffix = call-suffix
 | index-suffix
 | member-suffix
 | match-suffix ;

call-suffix = "(", [argument-list], ")" ;

argument-list = argument, { ",", argument } ;

argument = ([identifier, "="], expression)
 | block
 | lambda-expression
 | wildcard ;

index-suffix = "[", [expression-list], "]" ;

member-suffix = ".", (member-identifier | "for" | lifecycle-name) ;

member-identifier =
 ? an identifier token other than "match"; "match" is routed through
 match-suffix so matcher-chain syntax has one parser path ? ;

(* map-literal is deliberately absent. C.4/C.4.1 admit no untyped object
literal, and a braced `{ ... }` map body is an object literal representation,

```

so it is a collection BODY and never a value in its own right – it is reachable only through typed-collection-literal, behind a type prefix. array-literal is present because an array literal is not an object literal: C.4.1 makes `[ ... ]` a simple literal in the same sense as a string, character, or integer literal, so it carries no type prefix and needs none. The asymmetry between the two is that rule, not an oversight. A bare braced group in expression position is therefore always the block alternative. *)

```
primary-expression = literal
 | special-binding
 | "this"
 | "self"
 | qualified-name
 | grouped-expression
 | tuple-expression
 | array-literal
 | typed-collection-literal
 | object-construction
 | anonymous-class-expression
 | block
 | anonymous-function-expression
 | let-expression
 | comprehension-expression ;

grouped-expression = "(" , expression , ")" ;

tuple-expression = "(" , expression , "," , expression ,
 { "," , expression }, ")" ;

array-literal = "[" , [expression ,
 { "," , expression }, [","]], "]" ;

map-literal = "{", [map-entry ,
 { "," , map-entry }, [","]], "}" ;

map-entry = expression , ":" , expression ;
```

(* A built-in collection value names its collection type and then supplies the literal body that type takes. The generic arguments may be written on the type prefix, in which case the body follows the completed arrow tail:

```
co.core.List["A","B","C"]
co.core.Map{"A": 1, "B": 2}
co.core.Set(1,2,3)
co.core.List->(co.lang.string)["A","B","C"]
co.core.Map->(key co.lang.string, val co.lang.int){"A": 1}
co.core.Set->(co.lang.int)(1,2,3)
```

Without the arrow tail, `Type[ ... ]` and `Type( ... )` are also admissible as an index-suffix and a call-suffix on the same source span, and `Type{ ... }` is also admissible as object-construction. The overlap is resolved by typed-collection-literal-guard, not by a syntactic distinction.

The no-arrow-tail `Type( ... )` collection value shares the same token span with postfix-expression's ordinary call-suffix. Naming call-suffix in the first alternative and applying the guard makes the contextual collection reading explicit: when the prefix names a supported collection type the body is the collection's arguments, and otherwise it remains an ordinary call.

After an explicit arrow tail the body is always a collection body. *)

```
typed-collection-literal =
 type-postfix-expression,
 (array-literal | map-literal | call-suffix),
 typed-collection-literal-guard
| type-postfix-expression , "->" , parenthesized-type-list ,
 (array-literal | map-literal | call-suffix) ;

typed-collection-literal-guard =
 ? zero-width condition: the prefix names a supported collection type rather
 than a
 value in scope – a builtin-collection-type-name, a file-local alias bound
 to one by an alias-directive, or a user-declared collection type – and a
 following body form is the one fixed for that collection; an empty braced
 body on a recognized supported braced collection is an empty typed
 collection; for
 every other type an empty braced body is object construction; and a
 non-empty braced body is object construction whenever every entry has the
 object-field-initializer shape ? ;
```

(* The reference's Builtin Collections registry: the built-in names that can

```

stand in a typed-collection-literal prefix. It is a name registry consulted
by typed-collection-literal-guard, not a parser entry path – the prefix is
parsed as an ordinary type-postfix-expression, so no production references
this one. Which body form each name takes, and every other collection
semantic, stays outside this grammar. *)
builtin-collection-type-name = "co.core.List"
 | "co.core.Set"
 | "co.core.Map"
 | "co.core.Tree"
 | "co.core.Trie"
 | "co.core.Array"
 | "co.core.Tuple" ;

(* co.core.Tree, co.core.Trie, co.core.Array, and co.core.Tuple are reserved
built-in collection names, but their constructor body forms are unsupported
in the current alpha profile. The registry reserves their names, while
typed-collection-literal-guard refuses a constructor reading until the
language reference defines those forms. Their use as collection constructors
requires an unsupported-feature diagnostic under the reference Disclaimer. *)

object-construction = type-postfix-expression, "{",
 [object-field-initializer,
 { ",", object-field-initializer }, [",", "]], "}" ;

object-field-initializer = identifier, ":", expression ;

anonymous-class-expression = "co.lang.class", "{",
 { class-member }, "}" ;

let-expression = "let", "(", "{", let-binding,
 { ",", let-binding }, "}", ")",
 ".in", "(", "{", expression, "}", "}" ;

let-binding = (identifier | special-binding), "=", expression ;

comprehension-expression = "for", "(", comprehension-binding, ")",
 ".yield", "(", expression-list, ")" ;

comprehension-binding = pattern, "<-", expression ;

match-suffix = ".match", ["(", [expression], ")"],
 { match-case }, [match-default] ;

multi-symbol-infix-operator-boundary-guard =
 ? zero-width condition: a multi-symbol infix operator has an explicit
 boundary on both operand-facing sides; a boundary is whitespace, a
 comment, or an applicable delimiter, checked before separators are
 discarded ? ;

multi-symbol-range-operator-boundary-guard =
 ? zero-width condition: a multi-symbol range operator has an explicit
 boundary on every side for which an operand is present ? ;

extended-operator-expression =
 ? expression containing a registered project-local custom operator,
 parsed by precedence climbing from its declared alpha-supported fixity,
 precedence, associativity, and arity; pre-declared reserved glyphs and
 reserved future fixities are not accepted as expression operators in the
 current alpha profile; every registered multi-symbol operator satisfies
 the operand-facing boundary rule for its fixity ? ;

(* ===== *)

(* 12. Informative control-chain shapes (not parser entry paths) *)

(* ===== *)

(* Control vocabulary in the current alpha profile:
- .then(X) conditionally evaluates X once; X may be a block or value.
- .then(...) may participate in .otherwise(condition) selection chains.
- .loop(X) is a single-condition repetition form and cannot participate in
 .otherwise(condition) chains or take a .default(X) fallback.
- .otherwise(condition) always introduces another selection condition; there
 is no conditionless .otherwise form.
- .default(X) is the optional terminal fallback of a selection chain; it is
 not part of loop syntax.
- .each(...) is itself element iteration and cannot be followed by .loop.

```

The explicit-binding form is `.each(index-or-key, value, action)`, where action may be any expression (including a block or anonymous function) or a lambda-expression. The single-argument form accepts a callable callback; semantic analysis requires it to match the receiver's iteration tuple. These informative productions document recognizable shapes; ordinary parsing still proceeds through the general postfix/member/call grammar above. *)

```
informative-condition-chain =
 "(", expression, ")", ".then", "(", block, ")",
 { ".otherwise", "(", expression, ")", ".then", "(", block, ")" },
 [".default", "(", block, ")"] ;

informative-loop-chain =
 "(", expression, ")", ".loop", "(", block, ")" ;

informative-ternary-chain =
 "(", expression, ")", ".then", "(", expression, ")",
 { ".otherwise", "(", expression, ")", ".then", "(", expression, ")" },
 ".default", "(", expression, ")" ;

informative-each-chain =
 postfix-expression, ".each", "(",
 ((identifier | "_"), ",", identifier, ",",
 (expression | lambda-expression)
 | (expression | lambda-expression)), ")" ;

informative-contains-chain =
 postfix-expression, ".contains", "(", expression, ")",
 ".then", "(", block, ")",
 [".default", "(", block, ")"] ;

informative-pipeline-chain =
 postfix-expression,
 { (".filter" | ".map" | ".reduce" | ".forEach"
 | ".sortBy" | ".groupBy" | ".fold"),
 "(", argument-list, ")" } ;

(* ===== *)

(* 13. Literals and lexical grammar *)

(* ===== *)

literal = builtin-literal ;

builtin-literal = integer-literal
 | floating-literal
 | string-literal
 | character-literal
 | boolean-literal
 | none-literal ;

integer-literal = (binary-integer-literal
 | octal-integer-literal
 | decimal-integer-literal
 | hexadecimal-integer-literal),
 [integer-suffix] ;

binary-integer-literal = ("0b" | "0B"), binary-digit-sequence ;

octal-integer-literal = "0", [octal-digit-sequence] ;

decimal-integer-literal = nonzero-digit, { decimal-digit } ;

hexadecimal-integer-literal = hexadecimal-prefix,
 hexadecimal-digit-sequence ;

hexadecimal-prefix = "0x" | "0X" ;

binary-digit-sequence = binary-digit, { binary-digit } ;

octal-digit-sequence = octal-digit, { octal-digit } ;

decimal-digit-sequence = decimal-digit, { decimal-digit } ;

hexadecimal-digit-sequence = hexadecimal-digit, { hexadecimal-digit } ;
```

```

integer-suffix = unsigned-suffix,
 [long-suffix | long-long-suffix | size-suffix]
 | long-suffix, [unsigned-suffix]
 | long-long-suffix, [unsigned-suffix]
 | size-suffix, [unsigned-suffix] ;

unsigned-suffix = "u" | "U" ;

long-suffix = "l" | "L" ;

long-long-suffix = "ll" | "LL" ;

size-suffix = "z" | "Z" ;

floating-literal = decimal-floating-literal
 | hexadecimal-floating-literal ;

decimal-floating-literal = fractional-constant,
 [exponent-part],
 [floating-point-suffix]
 | decimal-digit-sequence,
 exponent-part,
 [floating-point-suffix] ;

fractional-constant = decimal-digit-sequence, ".",
 decimal-digit-sequence ;

hexadecimal-floating-literal = hexadecimal-prefix,
 (hexadecimal-fractional-constant
 | hexadecimal-digit-sequence),
 binary-exponent-part,
 [floating-point-suffix] ;

hexadecimal-fractional-constant = hexadecimal-digit-sequence, ".",
 hexadecimal-digit-sequence ;

exponent-part = ("e" | "E"), [sign], decimal-digit-sequence ;

binary-exponent-part = ("p" | "P"), [sign], decimal-digit-sequence ;

sign = "+" | "-" ;

floating-point-suffix = "f" | "F" | "l" | "L"
 | "f16" | "F16"
 | "f32" | "F32"
 | "f64" | "F64"
 | "f128" | "F128"
 | "bf16" | "BF16" ;

character-literal = single-quote, alpha-basic-c-character, single-quote ;

alpha-basic-c-character =
 ? any translation character except apostrophe, backslash,
 carriage return, or line feed ? ;

string-literal = double-quote, { alpha-basic-s-character }, double-quote ;

alpha-basic-s-character =
 ? any translation character except double quote, backslash,
 carriage return, or line feed ? ;

double-quote = ? Unicode scalar value U+0022 ? ;

single-quote = ? Unicode scalar value U+0027 ? ;

backslash = ? Unicode scalar value U+005C ? ;

boolean-literal = "co.const.true" | "co.const.false" ;

none-literal = "co.const.none" ;

byte-order-mark =
 ? U+FEFF, permitted only as the first code point of a source file ? ;

identifier = identifier-head, { "_", identifier-segment },
 identifier-trailing-guard ;

identifier-trailing-guard =

```

```

 ? a zero-width assertion that the next character is not "_" ? ;

identifier-head = ascii-letter, { ascii-alphanumeric } ;

identifier-segment = ascii-alphanumeric, { ascii-alphanumeric } ;

ascii-alphanumeric = ascii-letter | decimal-digit ;

ascii-letter = "A" | "B" | "C" | "D" | "E" | "F" | "G" | "H"
 | "I" | "J" | "K" | "L" | "M" | "N" | "O" | "P"
 | "Q" | "R" | "S" | "T" | "U" | "V" | "W" | "X"
 | "Y" | "Z"
 | "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h"
 | "i" | "j" | "k" | "l" | "m" | "n" | "o" | "p"
 | "q" | "r" | "s" | "t" | "u" | "v" | "w" | "x"
 | "y" | "z" ;

binary-digit = "0" | "1" ;

octal-digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" ;

hexadecimal-digit = decimal-digit
 | "a" | "b" | "c" | "d" | "e" | "f"
 | "A" | "B" | "C" | "D" | "E" | "F" ;

digit = decimal-digit ;

decimal-digit = "0" | nonzero-digit ;

nonzero-digit = "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;

hard-reserved-word = "co" | "let" | "this" | "for" | "forall" | "fo" ;

contextual-keyword = "self" ;

reserved-operator = "[:=" | "->>" | "<->" | "`" | backslash | "#" ;

predeclared-operator-glyph =
 "λ" | "(o)" | "â" | "ĩ" | "v" | "ε" | "o" | "ö" | "u"
 | "š" | "š" | "m" | "Π" | "⇒" | "f" | "f" | "v" | "f"
 | "↓" | "ð" | "⊥" | "↓" | "↓" ;

reserved-future-operator =
 ? a complete documented reserved/future symbolic spelling that is not a
 current built-in operator and not a predeclared-operator-glyph; the lexer
 recognizes the complete spelling and the parser reports unsupported ? ;

token = identifier
 | keyword-token
 | literal
 | delimiter-token
 | predeclared-operator-glyph
 | reserved-operator
 | symbolic-token
 | reserved-future-operator ;

keyword-token = hard-reserved-word | contextual-keyword ;

delimiter-token = "(" | ")" | "{" | "}" | "[" | "]"
 | "," | ";" | double-quote | single-quote ;

symbolic-token = ? the complete maximal contiguous run of one or more symbol
 characters after comments, literals, and closed composite
 spellings are recognized; the run is preserved whole for
 contextual classification and is never split as a
 fallback ? ;

token-separator = white-space ;

line-comment = "//", { ? any Unicode scalar value except CR or LF ? } ;

block-comment = "/*", { block-comment-character }, "*/" ;

block-comment-character = ? any Unicode scalar value that does not begin the
 two-character sequence */ ? ;

line-break = "\r\n" | "\n" | "\r" ;

```

```

horizontal-white-space = " " | "\t" | "\f" ;

white-space = horizontal-white-space | line-break | line-comment | block-comment ;

(* ===== *)

(* 14. Dedicated project-local operator-source grammar *)

(* ===== *)

operator-source-file = operator-library-declaration ;

operator-library-declaration =
 "_", "co.lang.library", "=", operator-library-body ;

operator-library-body = "{", { operator-declaration }, body-close ;

operator-declaration = operator-symbol, "co.lang.operator", "=",
 operator-body, statement-end ;

operator-body = "{", operator-property,
 { ",", operator-property }, "}" ;

operator-property = "fixity", ":", operator-fixity
 | "precedence", ":", ("0" | decimal-integer-literal)
 | "associativity", ":", operator-associativity
 | "arity", ":", operator-arity
 | "commutative", ":", boolean-literal
 | "idempotent", ":", boolean-literal
 | "identity", ":", operator-identity-value
 | "foldable", ":", boolean-literal
 | "vectorizable", ":", boolean-literal
 | "distributes_over", ":", operator-symbol-list
 | "desugar", ":", string-literal ;

operator-fixity = "co.operator.fixity.infix"
 | "co.operator.fixity.prefix"
 | "co.operator.fixity.postfix"
 | reserved-future-operator-fixity ;

reserved-future-operator-fixity =
 "co.operator.fixity.circumfix"
 | "co.operator.fixity.postcircumfix"
 | "co.operator.fixity.precircumfix"
 | "co.operator.fixity.mixfix"
 | "co.operator.fixity.ternary"
 | "co.operator.fixity.distfix" ;

operator-associativity = "co.operator.associativity.left"
 | "co.operator.associativity.right"
 | "co.operator.associativity.none" ;

operator-arity = "co.operator.arity.unary"
 | "co.operator.arity.binary"
 | "co.operator.arity.ternary"
 | decimal-integer-literal ;

operator-identity-value = literal ;

operator-symbol-list = "[", [operator-symbol-reference,
 { ",", operator-symbol-reference }], "]" ;

operator-symbol-reference = character-literal | string-literal ;

operator-symbol = ? a maximal run of one or more symbol characters, where a
 symbol character is any character that is not an ASCII
 letter, digit, underscore, whitespace, or one of the
 delimiters () { } [] , ; " ' ; the run must not be a
 language-owned or hard-reserved symbol and must not contain
 // or /* ? ;

```

# Appendix B - Grammar Decisions and Rationale

# Appendix C - Grammar Clarifications and Alpha Parser Policy

This appendix records grammar decisions that are normative for the current FoLang alpha language profile. It synchronizes lexical rules already present in the consolidated FoLang EBNF with language-reference decisions finalized after that grammar draft.

Where an older example or explanatory paragraph elsewhere in this document conflicts with a rule in this appendix, the rule in this appendix governs until the older example is corrected. Appendix C does not enable a feature that the [Disclaimer](#) marks as unimplemented.

## C.1 Source Encoding and Identifier Lexical Grammar

FoLang source text is UTF-8. A U+FEFF byte-order mark is permitted only as the first code point of a source file; U+FEFF anywhere else is an error.

Ordinary FoLang identifiers are intentionally ASCII-only even though the source encoding is UTF-8. No Unicode normalization rule is required for ordinary identifiers because non-ASCII characters are not admitted into an identifier token.

An ordinary identifier:

- begins with an ASCII letter `A-Z` or `a-z` ;
- may continue with ASCII letters or decimal digits;
- may contain `_` only between two non-empty alphanumeric segments;
- cannot begin with `_` ;
- cannot contain consecutive underscores;
- cannot end with `_` ;
- cannot be the single spelling `_` because `_` is a dedicated contextual token used for wildcard/discard positions, filename-derived declaration positions, unnamed type-constructor parameter slots such as `F(_)` , and the candidate-value placeholder inside a `co.lang.refinementType` predicate;
- is checked against the reserved-word table after its character sequence has been recognized. A hard-reserved word is emitted as its reserved token, not as an ordinary identifier.

Normative lexical grammar:

```
identifier = identifier-head, { "_", identifier-segment },
 identifier-trailing-guard ;

identifier-trailing-guard =
 ? a zero-width assertion that the next character is not "_" ? ;

identifier-head = ascii-letter, { ascii-alphanumeric } ;
identifier-segment = ascii-alphanumeric, { ascii-alphanumeric } ;

ascii-alphanumeric = ascii-letter | decimal-digit ;

ascii-letter = "A" | "B" | "C" | "D" | "E" | "F" | "G" | "H"
 | "I" | "J" | "K" | "L" | "M" | "N" | "O" | "P"
 | "Q" | "R" | "S" | "T" | "U" | "V" | "W" | "X"
 | "Y" | "Z"
 | "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h"
 | "i" | "j" | "k" | "l" | "m" | "n" | "o" | "p"
 | "q" | "r" | "s" | "t" | "u" | "v" | "w" | "x"
 | "y" | "z" ;

decimal-digit = "0" | nonzero-digit ;
nonzero-digit = "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
```

Examples:

name	valid
myVar2	valid
v1_hr	valid
a_b_c	valid
123hr	invalid: identifier cannot begin with a digit
_x	invalid: identifier cannot begin with underscore
_1	invalid: identifier cannot begin with underscore
a_	invalid: trailing underscore
a__b	invalid: consecutive underscores
_	contextual token, not an identifier



The hard-reserved words in the current alpha grammar are `co`, `let`, `this`, `for`, `forall`, and `fo`. `fo` is permanently reserved and cannot be used as an ordinary identifier. `self` is contextual rather than globally hard-reserved: it has its class-method meaning in every method declared by a `co.lang.class`, including `@@new` and `@@init`.

## C.2 Numeric Literal Lexical Grammar

FoLang uses the numeric-literal subset already selected by the consolidated grammar. Numeric digit separators are not supported. In particular, forms such as `1'000`, `0x1'a`, and `0b1011'0010` are invalid.

A numeric sign is not part of the ordinary numeric literal token. In an ordinary expression, unary `+` or `-` is parsed as a prefix operator. In a pattern, the grammar explicitly admits a leading `+` or `-` before an integer or floating literal so signed numeric literal patterns remain valid.

### C.2.1 Integer literals

```
integer-literal = (binary-integer-literal
 | octal-integer-literal
 | decimal-integer-literal
 | hexadecimal-integer-literal),
 [integer-suffix] ;

binary-integer-literal = ("0b" | "0B"), binary-digit-sequence ;
octal-integer-literal = "0", [octal-digit-sequence] ;
decimal-integer-literal = nonzero-digit, { decimal-digit } ;
hexadecimal-integer-literal = hexadecimal-prefix,
 hexadecimal-digit-sequence ;

hexadecimal-prefix = "0x" | "0X" ;

binary-digit-sequence = binary-digit, { binary-digit } ;
octal-digit-sequence = octal-digit, { octal-digit } ;
decimal-digit-sequence = decimal-digit, { decimal-digit } ;
hexadecimal-digit-sequence = hexadecimal-digit, { hexadecimal-digit } ;

binary-digit = "0" | "1" ;
octal-digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" ;
hexadecimal-digit = decimal-digit
 | "a" | "b" | "c" | "d" | "e" | "f"
 | "A" | "B" | "C" | "D" | "E" | "F" ;

integer-suffix = unsigned-suffix,
 [long-suffix | long-long-suffix | size-suffix]
 | long-suffix, [unsigned-suffix]
 | long-long-suffix, [unsigned-suffix]
 | size-suffix, [unsigned-suffix] ;

unsigned-suffix = "u" | "U" ;
long-suffix = "l" | "L" ;
long-long-suffix = "ll" | "LL" ;
size-suffix = "z" | "Z" ;
```

The literal `0` is valid through the octal production. A decimal integer with a nonzero value begins with `1-9`.

### C.2.2 Floating literals

A decimal or hexadecimal floating literal that contains a radix point must contain at least one digit on both sides of the point. Therefore `1.0` and `0.10` are valid, while `1.` and `.10` are invalid. Scientific notation without a radix point, such as `1e5`, remains valid.

```

floating-literal = decimal-floating-literal
 | hexadecimal-floating-literal ;

decimal-floating-literal = fractional-constant,
 [exponent-part],
 [floating-point-suffix]
 | decimal-digit-sequence,
 exponent-part,
 [floating-point-suffix] ;

fractional-constant = decimal-digit-sequence, ".",
 decimal-digit-sequence ;

hexadecimal-floating-literal = hexadecimal-prefix,
 (hexadecimal-fractional-constant
 | hexadecimal-digit-sequence),
 binary-exponent-part,
 [floating-point-suffix] ;

hexadecimal-fractional-constant = hexadecimal-digit-sequence, ".",
 hexadecimal-digit-sequence ;

exponent-part = ("e" | "E"), [sign], decimal-digit-sequence ;
binary-exponent-part = ("p" | "P"), [sign], decimal-digit-sequence ;
sign = "+" | "-" ;

floating-point-suffix = "f" | "F" | "l" | "L"
 | "f16" | "F16"
 | "f32" | "F32"
 | "f64" | "F64"
 | "f128" | "F128"
 | "bf16" | "BF16" ;

```

Backend-conditionally-supported floating suffixes are accepted only when the configured backend/compiler contract supports the corresponding representation.

Signed numeric patterns use:

```

literal-pattern = literal
 | ("+" | "-"),
 (integer-literal | floating-literal) ;

```

### C.3 Comments, Whitespace, and Line Breaks

FoLang supports `//` line comments and non-nesting `/* ... */` block comments.

- A line comment starts with `//` and extends up to, but does not consume as part of its body, the next CR or LF line terminator.
- A block comment starts with `/*` and ends at the first following `*/`.
- Block comments do not nest. An inner `/*` is ordinary block-comment content; the first subsequent `*/` closes the block comment.
- Comments are recognized before ordinary symbolic-run scanning.
- Spaces, horizontal tabs, form-feed characters, line breaks, and comments are token separators and are discarded between tokens.
- A line break never terminates a FoLang statement. FoLang has no automatic semicolon insertion.

Normative lexical grammar:

```

line-comment = "//", { ? any Unicode scalar value except CR or LF ? } ;

block-comment = "/*", { block-comment-character }, "*/" ;
block-comment-character = ? any Unicode scalar value that does not begin the
 two-character sequence */ ? ;

line-break = "\r\n" | "\n" | "\r" ;
horizontal-white-space = " " | "\t" | "\f" ;

white-space = horizontal-white-space
 | line-break
 | line-comment
 | block-comment ;

token-separator = white-space ;

```

## C.4 Canonical Compound and Object Construction Syntax

A user-defined object/struct/class value is constructed with an explicit type followed immediately by a braced field initializer. The canonical form is Go-like:

```
b := B{age: 25.0};
emp := Employee{name: "Rao", id: 1};
point := Point{x: 10.0, y: 20.0};
```

Object field initializers use `:` between the field name and value and `,` between fields. `=` is not an object-field initializer binder.

Normative shape:

```
object-construction = type-postfix-expression, "{",
 [object-field-initializer,
 { ",", object-field-initializer }, [",", " "], "]" ;

object-field-initializer = identifier, ":", expression ;
```

Therefore these are invalid object-construction spellings:

```
b := B{age = 25.0}; // invalid
emp := {name: "Rao"}; // invalid as a UDT construction: explicit type required
emp := Employee{name = "Rao"}; // invalid
```

There is no untyped object literal in FoLang. A braced value is a body, never a value in its own right, so every object value names its type. This is a general rule and not a restriction on user-defined types alone: it applies equally to the braced built-in collection value of C.4.1.

A typed map literal remains a distinct expression from object construction. Its entries also use `:`, but the left side of an entry is an expression rather than a field identifier.

A closing `}` belonging to object construction closes the expression only. It does **not** terminate the enclosing statement; the statement still requires `;` as specified in C.6.

### C.4.1 Built-in collection values

There are two current-alpha construction forms. When the declaration itself is type-deduced, the constructor carries its generic arguments explicitly and the literal body follows the completed arrow tail described in C.11:

```
x := co.core.List->(co.lang.string) ["A", "B", "C"];
map := co.core.Map->(key=co.lang.string, val=co.lang.int) {"A": 1, "B": 2};
y := co.core.Set->(co.lang.int) (1,2,3);
```

When the surrounding declaration already supplies the generic collection type, the value constructor has no arrow tail and does not repeat those generic arguments:

```
x co.core.List->(co.lang.string) = co.core.List ["A", "B", "C"];
y co.core.Set->(co.lang.int) = co.core.Set (1,2,3);
map co.core.Map->(key=co.lang.string, val=co.lang.int) = co.core.Map {"A": 1, "B": 2};
```

Normative shape:

```
typed-collection-literal =
 type-postfix-expression,
 (array-literal | map-literal | call-suffix),
 typed-collection-literal-guard
| type-postfix-expression, "->", parenthesized-type-list,
 (array-literal | map-literal | call-suffix) ;
```

The type prefix on a braced collection value is required, for the reason C.4 gives: there is no untyped object literal, and a `{ ... }` map body is an object literal representation. An untyped `{ ... }` is therefore never a value.

An array literal is not an object literal. `[ ... ]` is a simple literal in the same sense as a string, character, or integer literal, so it carries no type prefix and needs none. C.6 governs its termination exactly as it governs any other simple literal:

```
xs := [1, 2, 3]; // simple literal, valid untyped
cfg := co.core.Map->(key=co.lang.string, val=co.lang.any) {"host": "db", "port": 5432};

cfg := {"host": "db", "port": 5432}; // invalid: untyped map literal
```

## C.4.2 Disambiguating a typed collection body from the same source shape

Without an arrow tail, a typed collection body is spelled identically to three existing forms. The readings are separated as follows, and the separation is contextual rather than syntactic:

```
Type{ ... } -> for a recognized supported braced collection constructor,
 an empty body is an empty typed collection and a non-empty
 body is a typed collection unless every entry has the shape
 identifier ":" expression; for every other type, an empty
 body or identifier-field body is object construction
Type[...] -> typed list literal when Type names a supported collection type;
 otherwise an index-suffix applied to the value Type
Type(...) -> typed collection value when Type names a supported collection type;
 otherwise an ordinary call
```

`Employee{name: "Rao", id: 1}` and `Employee{}` are therefore object construction. `co.core.Map{"A": 1}` is a typed map literal because a string key is not an object-field initializer, and `co.core.Map{}` is the empty typed map because the recognized collection name resolves the otherwise empty shape. An explicit arrow tail removes the overlap entirely: after `Type->( ... )` the following body is always the collection's literal body.

Only `co.core.List`, `co.core.Set`, and `co.core.Map` have current-alpha collection-constructor forms. `co.core.Tree`, `co.core.Trie`, `co.core.Array`, and `co.core.Tuple` remain reserved built-in collection names, but their body forms are intentionally unspecified and unsupported. Under the [Disclaimer](#), using one of those names as a collection constructor must produce an unsupported-feature diagnostic until a later language-reference revision supplies an example, body form, and semantics. They must not silently inherit the body syntax of List, Set, or Map.

A closing `}`, `]`, or `)` belonging to a typed collection value closes the expression only, exactly as in C.4 and C.6.3. The enclosing statement still requires its `;`.

## C.5 Match Invocation and Matcher Selection

FoLang distinguishes automatic matcher selection from explicit matcher selection by whether a matcher argument is supplied.

### C.5.1 No matcher argument

When `.match` has no matcher argument, the compiler selects the applicable built-in/default matcher from the subject type and case-pattern forms.

Both existing no-argument spellings are accepted and equivalent:

```
value.match.case(...).default(...);
value.match().case(...).default(...);
```

The empty parentheses do not select a different matcher.

### C.5.2 Explicit matcher argument

When an argument is supplied to `.match(...)`, that expression explicitly identifies the matcher to use. It may identify a named built-in matcher or a user-defined matcher.

```
value.match(co.pattern.Type).case(...);
value.match(co.pattern.Value).case(...);
value.match(PositiveEvenMatcher).case(...).default(...);
```

A user-defined matcher such as `PositiveEvenMatcher` follows the normal matcher declaration and import/name-resolution rules.

The parser-level shape remains:

```
match-suffix = ".match", ["(", [expression], ")"],
 { match-case }, [match-default] ;
```

Semantic interpretation is:

```
.match -> no matcher argument -> automatic built-in/default matcher selection
.match() -> no matcher argument -> automatic built-in/default matcher selection
.match(matcher) -> explicit matcher selection
```

## C.6 Statement and Expression Termination

FoLang uses explicit termination. Newlines never terminate statements, and there is no automatic semicolon insertion.

The governing rule is:

simple statement or expression statement	-> terminated by ;
direct block/body	-> terminated by its closing }
braced expression/literal	-> } closes the expression, then ; closes its statement

### C.6.1 Semicolon termination

A semicolon is mandatory after every simple statement and expression statement, including:

- variable declarations and initializations;
- assignments and compound assignments;
- function/method calls used as statements;
- return statements expressed through `this.return ...`;
- expression-bodied function-pattern clauses;
- object-construction expressions used in declarations, assignments, or standalone expression statements;
- typed collection values such as `co.core.List[...]`, `co.core.Map{...}`, and `co.core.Set(...)`;
- generic instantiations written as a typed declaration, such as `k LinkedList->(T co.lang.int)`;
- array, tuple, or other simple literal expressions when they form a simple statement; a map is written as the typed collection value of the previous item, since there is no untyped map literal;
- forward declarations and other declaration forms whose syntax is a simple declaration rather than a block body.

Examples:

```
x co.lang.int = 10;
x = 20;
co.out.println(x);
emp := Employee{name: "Rao", id: 1};
cfg := co.core.Map->(key=co.lang.string, val=co.lang.any){"host": "db", "port": 5432};
xs := [1, 2, 3];
```

### C.6.2 Block/body termination

A direct block or declaration body terminates at its closing `}` and is not followed by a semicolon.

```
// Employee.fol
_ co.lang.struct = {
 id co.lang.int;
 name co.lang.string;
}

// function body
calculate()->(co.lang.int) = {
 this.return 10;
}

// block-bodied function-pattern clause
classify(n) => {
 this.return "positive";
}
```

Writing `};` after one of these direct block/body forms is invalid.

### C.6.3 A braced expression is not a block terminator

An object construction, typed map literal, typed collection value, anonymous value expression, or other braced **expression** is not a declaration/function/block body merely because it ends with `}`. The enclosing simple statement still requires its semicolon.

```
emp := Employee{id: 1, name: "Rao"};
this.return Employee{id: 1};
cfg := co.core.Map->(key=co.lang.string, val=co.lang.int){"a": 1, "b": 2};
ages := co.core.Map->(key=co.lang.string, val=co.lang.int){"A": 30, "B": 40};
```

Built-in directives and annotations are self-delimiting metadata forms rather than simple statements and therefore do not acquire a trailing semicolon merely because they appear on their own source line.

## C.7 Current Alpha, Reserved, Future, and Unimplemented Syntax

The current alpha parser follows the existing Disclaimer policy with an explicit phase distinction:

```
source contains a reserved/future/unimplemented language spelling
-> lexer recognizes the complete token/spelling
-> lexer does not reject it merely because the feature is not implemented
-> parser recognizes that the spelling belongs to a reserved/unimplemented construct
-> parser reports an unimplemented/unsupported-feature parse error
-> semantic analysis is not used to pretend the feature is implemented
```

This policy applies to already documented future or conceptual forms such as future impredicative generic syntax, generic inheritance/path-dependent-type work, reserved future fixities, and other table-listed but not alpha-enabled constructs.

The diagnostic should identify the feature, for example:

```
parser error: feature `impredicative generic instantiation` is not implemented in the current alpha profile
```

This rule does not turn arbitrary unknown text into reserved syntax. A character sequence or symbolic run that is neither a valid current token nor a documented reserved/future spelling remains a lexical error.

### C.7.1 Exception for `@co.*` metadata forms

The unsupported-feature parse error above applies to reserved or unimplemented **core source syntax**. It does not apply to declared built-in `@co.*` directives, annotations, pragmas, or decorators, which are exempt from the Disclaimer's example requirement.

Every such metadata application uses the generic metadata grammar:

```
annotation = "@", qualified-name,
 ["(" , [annotation-argument-list] , ")"] ;
```

The parser must preserve the complete application independently of whether an example exists:

```
declared @co.* directive/annotation/pragma/decorator
-> parse the qualified metadata name
-> parse and collect every supplied positional/named argument, field,
 attribute, and argument expression
-> attach the complete metadata node to its target
-> do not reject or discard a field merely because no example documents it
-> if no semantic handler implements the metadata form, silently ignore the
 collected form and its arguments
-> if a semantic handler exists, validate its field schema and target rules
 during semantic resolution
```

Collection is therefore not activation. An ignored metadata form has no semantic effect and contributes nothing to liveness, but its syntax and supplied fields were still parsed and preserved. Examples document intended semantics; they are not required for parser acceptance or attribute collection.

An annotation whose name is outside `co.*` is an ordinary user-defined annotation and is resolved by the normal rules in C.8. Nothing here exempts external metadata from name resolution.

### C.7.2 `co.*` package declarations are package API, not reserved syntax

Ordinary declarations provided by the language's standard `co.*` export-kind `.foIenc` artifact are not enabled or disabled by the Disclaimer's example rule. The toolchain loads the standard artifact implicitly, reconstructs its exported package contexts, and then resolves those declarations through the ordinary package/symbol/member model.

Therefore a type or unit-level function newly added to an existing standard package—for example a future `co.core.RBTree`—may be used through ordinary existing FoLang syntax as soon as the applicable standard package artifact provides that declaration. The language reference does not need to introduce a new grammar form merely to permit the new package member.

This exception does not allow an API declaration to manufacture new syntax. A special constructor body, new operator spelling, new declaration form, or other grammar-level facility still requires explicit core-language definition under C.7 and the Disclaimer.

## C.8 User-Defined Annotation Application

A user-defined annotation is applied with the same annotation syntax as a built-in annotation. There is no separate application grammar for user annotations.

Normative annotation shape:

```
annotation = "@", qualified-name,
 ["(" , [annotation-argument-list] , ")"] ;
```

The difference is name resolution: built-in `co.*` annotations are intrinsically available, while a user-defined annotation is an ordinary user declaration and must be resolved through normal package/import rules. When the annotation is declared in another package, that package must be imported before the annotation is used.

Example declaration:

```
// my/annotations/Audited.fol
_ co.lang.object->(for=annotation) = {
 value co.lang.string;
 enabled co.lang.bool;
}
```

Example use from another package:

```
@co.ddap.import(package="my.annotations", as="ann")

@ann.Audited(value="payroll", enabled=co.const.true)
_ co.lang.class = {
 ...
}
```

An import without `as=` may use the complete imported package path according to the normal import rules:

```
@co.ddap.import(package="my.annotations")

@my.annotations.Audited(value="payroll")
_ co.lang.class = {
 ...
}
```

The annotation’s permitted targets, required fields, field types, defaults, and liveness remain semantic checks; application syntax is identical to built-in annotation syntax.

### C.8.1 Refinement-Type Candidate Placeholder

The token `_` has one additional contextual meaning inside the predicate of a `co.lang.refinementType` declaration. In the canonical form:

```
positiveInt co.lang.refinementType = (co.lang.int).where(_ > 0);
```

`_` denotes the candidate value of the immediately preceding base type. It is not an ordinary identifier and does not create a binding in the surrounding scope. Multiple `_` occurrences in the same refinement predicate denote that same candidate value.

This interpretation is restricted to the refinement predicate. It does not change `_` into a general expression identifier and does not alter its wildcard, discard, filename-derived declaration-name, or type-constructor-placeholder meanings in their respective contexts.

The predicate must resolve to `co.lang.bool`. For the remaining validation rule for candidates that are not statically known, see [Refinement Types](#).

### C.9 Built-In Operator Parse Table

FoLang follows the conventional precedence ordering used by mainstream C-family languages for the ordinary arithmetic, relational, equality, bitwise, logical, and assignment families, with FoLang’s exponentiation and range levels inserted explicitly. Larger precedence numbers bind more tightly.

Precedence	Operator/form	Fixity	Associativity	Arity / parse role
700	call <code>(...)</code> , index <code>[...]</code> , member <code>.</code> , postfix <code>!</code>	postfix	left	call/index/member syntax; postfix <code>!</code> unary
650	<code>**</code>	infix	right	binary
600	<code>+</code> , <code>-</code> , <code>!</code>	prefix	right	unary
550	<code>*</code> , <code>/</code> , <code>%</code>	infix	left	binary
500	<code>u</code> , <code>n</code>	infix	left	binary
450	<code>+</code> , <code>-</code>	infix	left	binary

400	<code>..</code> , <code>&lt;..</code> , <code>..<code>&lt;</code></code> , <code>&lt;..<code>&lt;</code></code>	infix/range	none	binary with one bound optionally omitted according to range grammar
350	<code>&lt;</code> , <code>&lt;=</code> , <code>&gt;</code> , <code>&gt;=</code>	infix	left	binary
300	<code>==</code> , <code>!=</code>	infix	left	binary
250	<code>&amp;</code>	infix	left	binary
200	<code>^</code>	infix	left	binary
150	<code> </code>	infix	left	binary
100	<code>&amp;&amp;</code>	infix	left	binary; short-circuit
50	<code>  </code>	infix	left	binary; short-circuit
10	<code>=</code> , <code>+=</code> , <code>-=</code> , <code>*=</code> , <code>/=</code> , <code>%=</code> , <code>**=</code> , <code>&amp;=</code> , <code>^=</code> , <code> =</code>	infix assignment	right	binary assignment

Examples of grouping:

```
a + b * c; // a + (b * c)
a ** b ** c; // a ** (b ** c)
a || b && c; // a || (b && c)
a = b = c; // a = (b = c)
```

FoLang's left-to-right operand **evaluation order** is independent from operator grouping. Precedence and associativity determine the parse tree; the evaluation-order rules determine when the operands of that tree are evaluated.

The definition spellings `:=` and `?=` are statement-level definition operators, not general expression operators. They therefore do not receive a general expression-precedence level and cannot be chained as ordinary binary expressions.

Structural spellings such as `=>`, `=>>`, `==>>`, `->`, `<-`, `:=`, `->>`, and `<->` are not ordinary expression operators merely because they contain symbol characters.

Multi-symbol expression operators continue to obey the explicit operand-facing boundary rule: whitespace, a comment, or an applicable delimiter must delimit each operand-facing side required by the operator's fixity.

## C.10 Pre-Declared Operator Glyphs in the Current Alpha Profile

The current alpha profile defines the following language-owned pre-declared operator glyphs:

`U`, `n`

Both glyphs are enabled expression operators. Their parser properties are fixed by C.9:

```
fixity = infix
precedence = 500
associativity = left
arity = binary
```

A project must not redeclare either glyph with `co.lang.operator`. Implementations are supplied through ordinary `mode=overload` operator functions in the legal operator owner for the operand type. Multiple implementations may coexist when their normalized operand signatures are distinct.

The lexer recognizes `U` and `n` as registered language-owned operator tokens, and the parser constructs the corresponding infix expression according to the properties above. Absence of an applicable overload is a resolution error, not a lexical or parse-time unsupported-feature error.

## C.11 Type Arrow Tail and Generic Type Arguments

`->` is a structural spelling, not an expression operator (C.9). In **type** position, a `->` tail attached to a type carries three unrelated meanings that only the tail's shape and the base type's meaning separate:

<code>co.lang.int-&gt;([5])</code>	derivation applied to <code>co.lang.int</code>
<code>co.lang.int-&gt;(&amp;, meta={type=out})</code>	derivation applied to <code>co.lang.int</code>
<code>(co.lang.int)-&gt;(co.lang.int)</code>	function type from <code>int</code> to <code>int</code>
<code>co.core.List-&gt;(co.lang.string)</code>	generic type arguments applied to <code>List</code>
<code>LinkedList-&gt;(T=co.lang.int)</code>	generic type arguments, named



A tail whose first token starts a derivation specification — `*`, `&`, `&&`, `~`, `@`, `^`, `[:]`, `..`, `[]`, or an `attribute=` pair — is a type derivation. Any other parenthesized tail is a `parenthesized-type-list`, which is read as a function result list when the base is a parameter list and as a generic type-argument list when the base is a generic declaration.

Items in that list carry an optional binder name, so both spellings below are admitted and mean the same instantiation:

```
co.core.List->(co.lang.string)
co.core.Map->(key=co.lang.string, val=co.lang.int)
Employee->(T=co.lang.int, R=co.lang.string)
```

Normative shape:

```
arrow-type-expression = type-postfix-expression, ["->", arrow-type-tail]
 | "(", [function-type-parameter,
 { ",", function-type-parameter }],
 ")", "->", arrow-type-tail ;

arrow-type-tail = type-derivation
 | parenthesized-type-list
 | type-expression ;

parenthesized-type-list = "(", [return-item-list], ")" ;

return-item-list = return-item, { ",", return-item } ;

return-item = [identifier], type-expression ;
```

A binder name in a generic argument list names the type parameter it binds. When names are used, they must be declared parameter names of the type being applied; when they are omitted, arguments bind positionally. Names and positions must not be mixed in one list. These are semantic checks, not parse-level ones.

### C.11.1 Both generic application spellings are admitted

FoLang accepts two spellings for applying a type to type arguments, and they are not in conflict:

<code>Option(co.lang.int)</code>	<code>type-argument-list</code> , positional only
<code>co.core.List-&gt;(co.lang.string)</code>	<code>arrow tail</code> , positional or named

`type-argument-list` is the direct application form used by `co.lang.type` constructors and admits a `dependent-index` argument. The arrow tail additionally admits binder names, which is what a multi-parameter built-in collection such as `co.core.Map` relies on. Where a declaration reads naturally in either spelling, both denote the same applied type.

### C.11.2 Instantiating a generic declaration through a typed declaration

A typed declaration whose type is a generic declaration with its arrow tail supplied is a complete instantiation. It is the declarative equivalent of the explicit `new`/`init` pair:

```
k := LinkedList.new(co.lang.int); // explicit two-step form
k LinkedList->(T co.lang.int); // equivalent instantiation

c := Employee.new(co.lang.int, co.lang.string).init(1, "Rao");
c Employee->(T co.lang.int, R co.lang.string);
```

This introduces no new statement form. It is an ordinary `variable-declaration` whose `type-expression` carries an arrow tail, so C.6.1 applies unchanged and the declaration is terminated by `;`.

`@@new` and `@@init` remain reserved for the cases described in [Generics](#), where the two steps must be separated or overridden explicitly.